



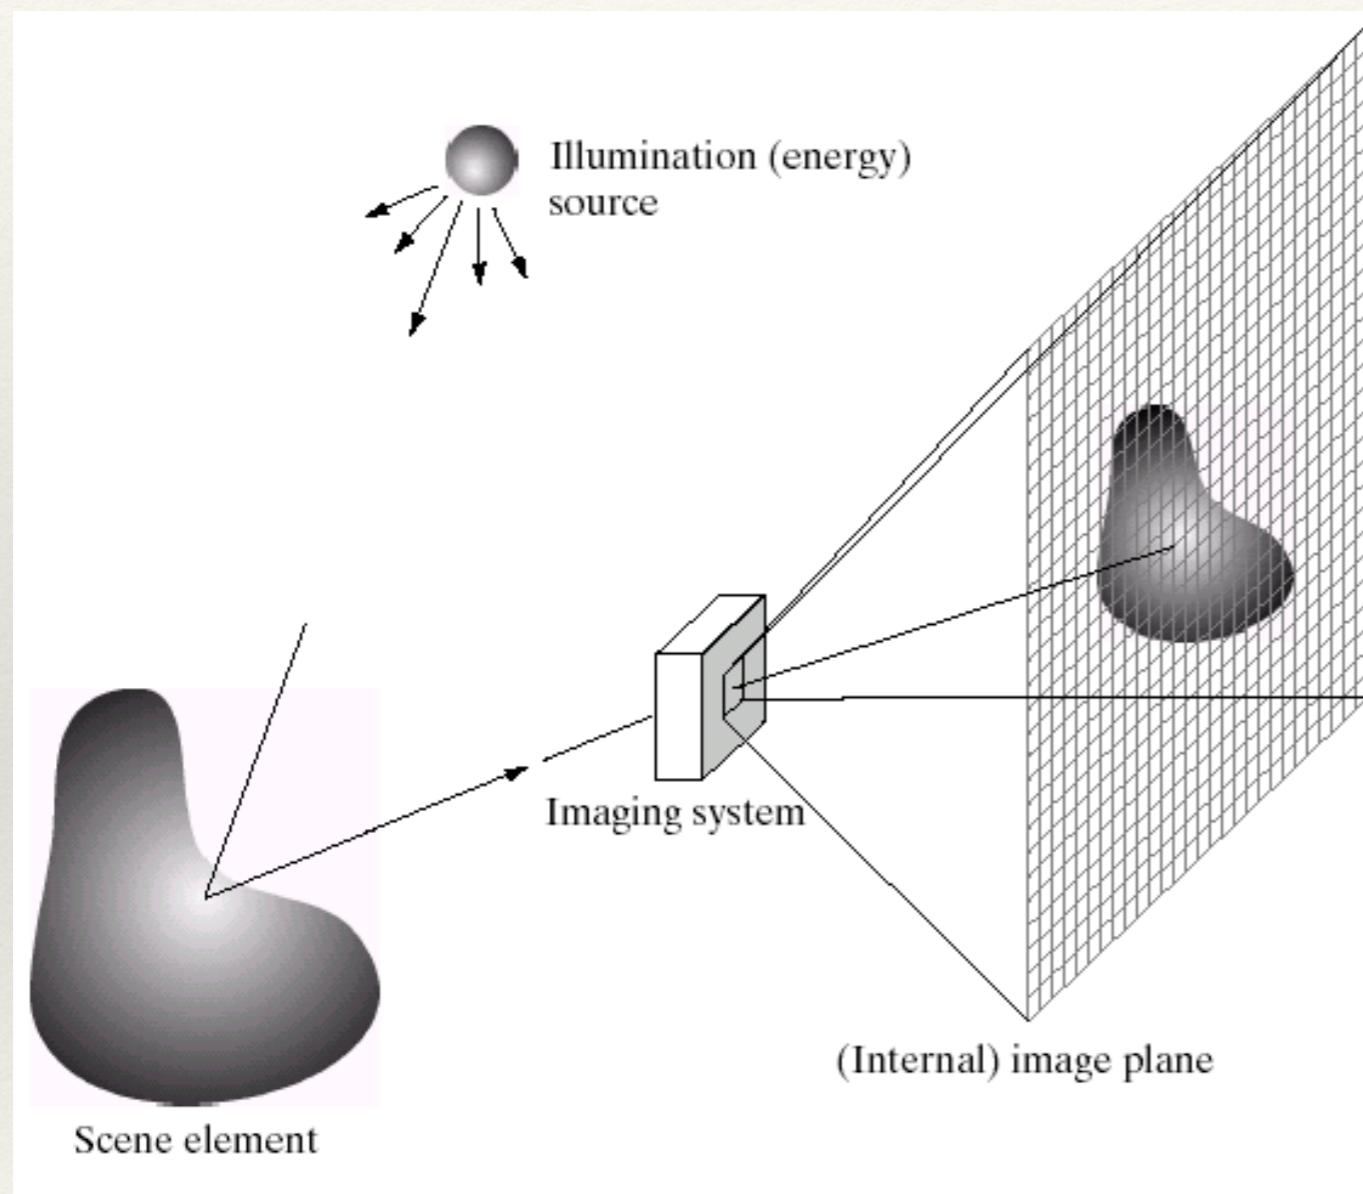
Xu Zhao @ Shanghai Jiao Tong university

Lecture 3-1: Image Filtering

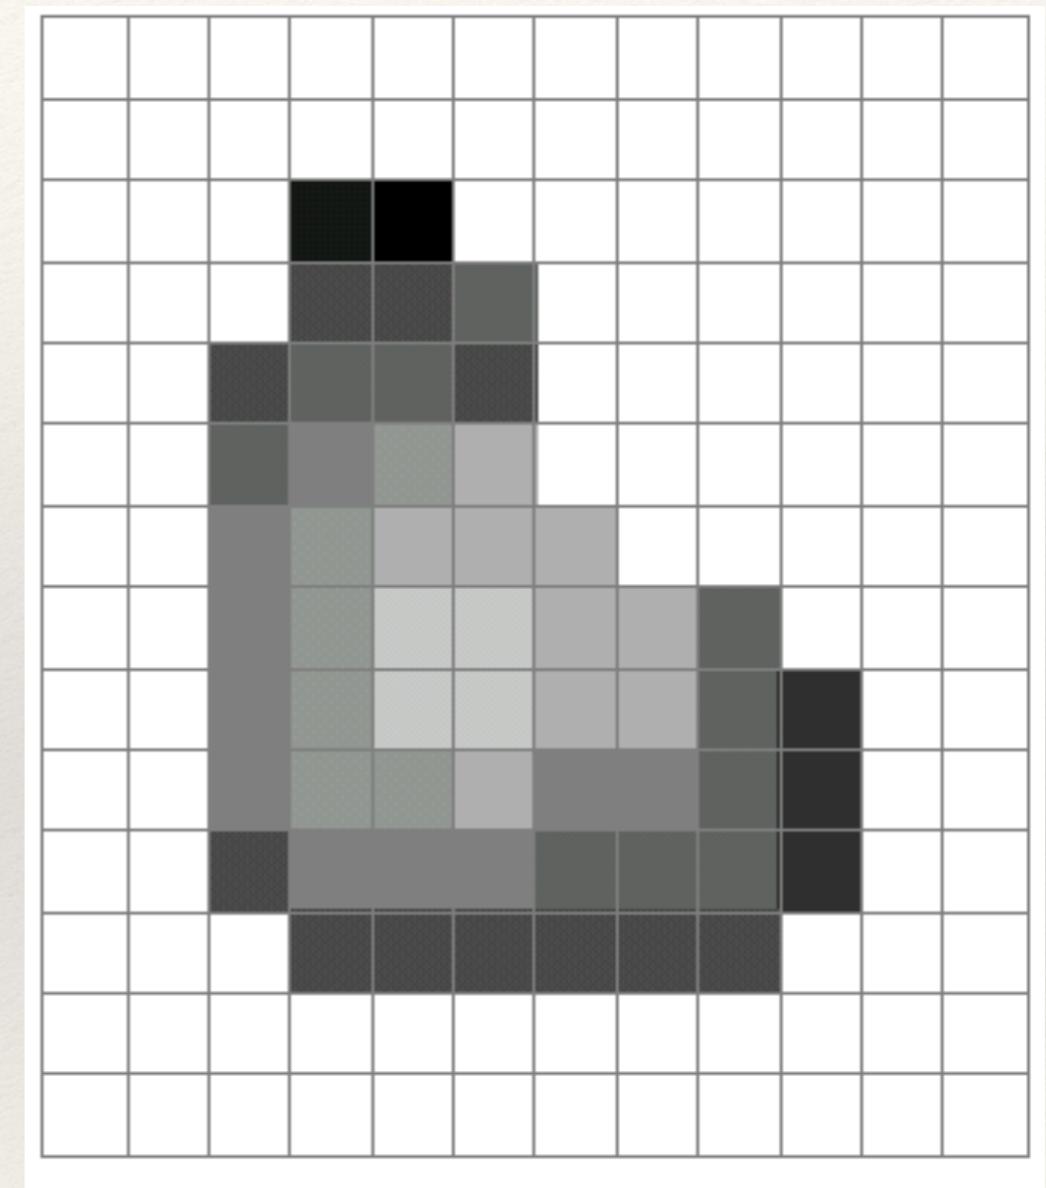
Contents

- ❖ **Spatial filtering**
 - Linear filter
 - Non-linear filter
- ❖ **Frequency perspective**
 - Spatial frequency
 - Fourier transform

Image formation



Camera Sensor

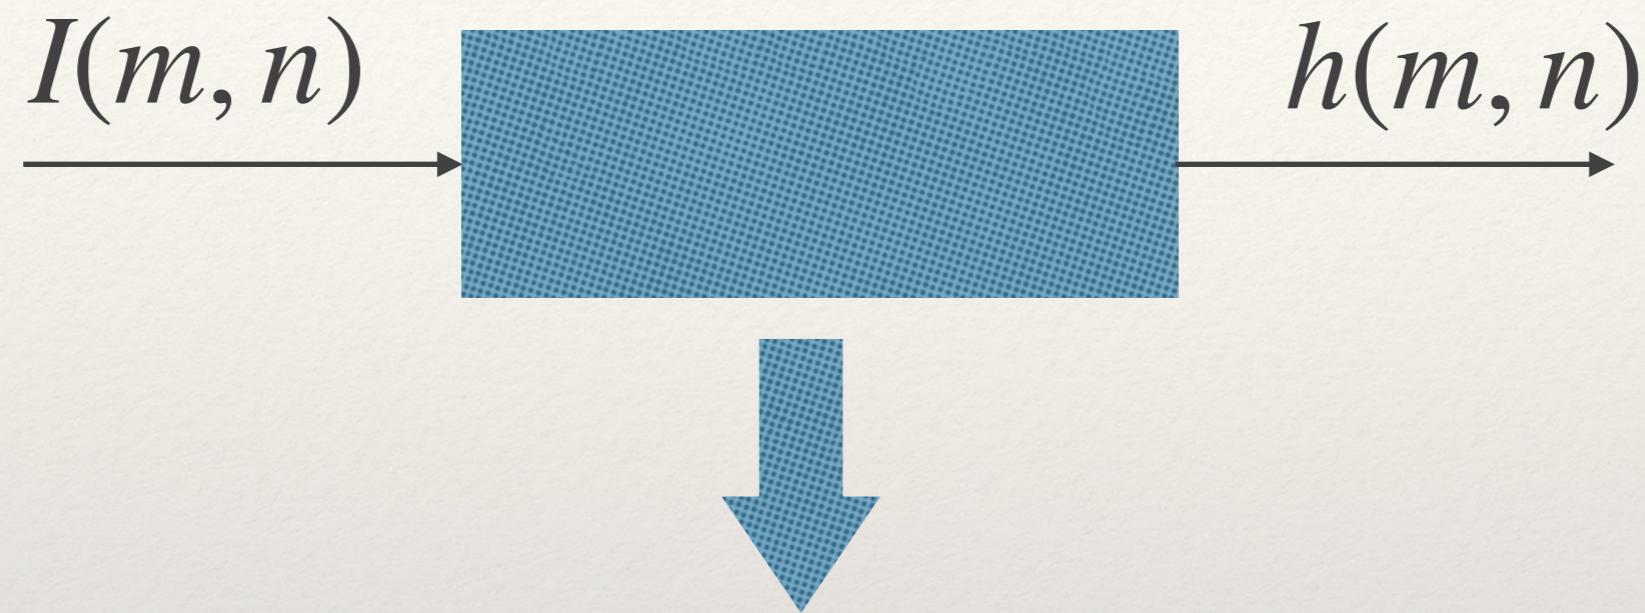


Output Image

Image filtering, why and how?

- ❖ Motivation
 - ❖ Enhance an image (denoise, resize, etc)
 - ❖ Extract information (texture, edges, etc)
 - ❖ Detect patterns (template matching)
- ❖ Compute a function of the local neighborhood at each pixel in the image
 - ❖ Function specified by a “filter” or mask saying how to combine values from neighbors.

Linear filtering



$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

For a linear system, each output is a linear combination of all the input values

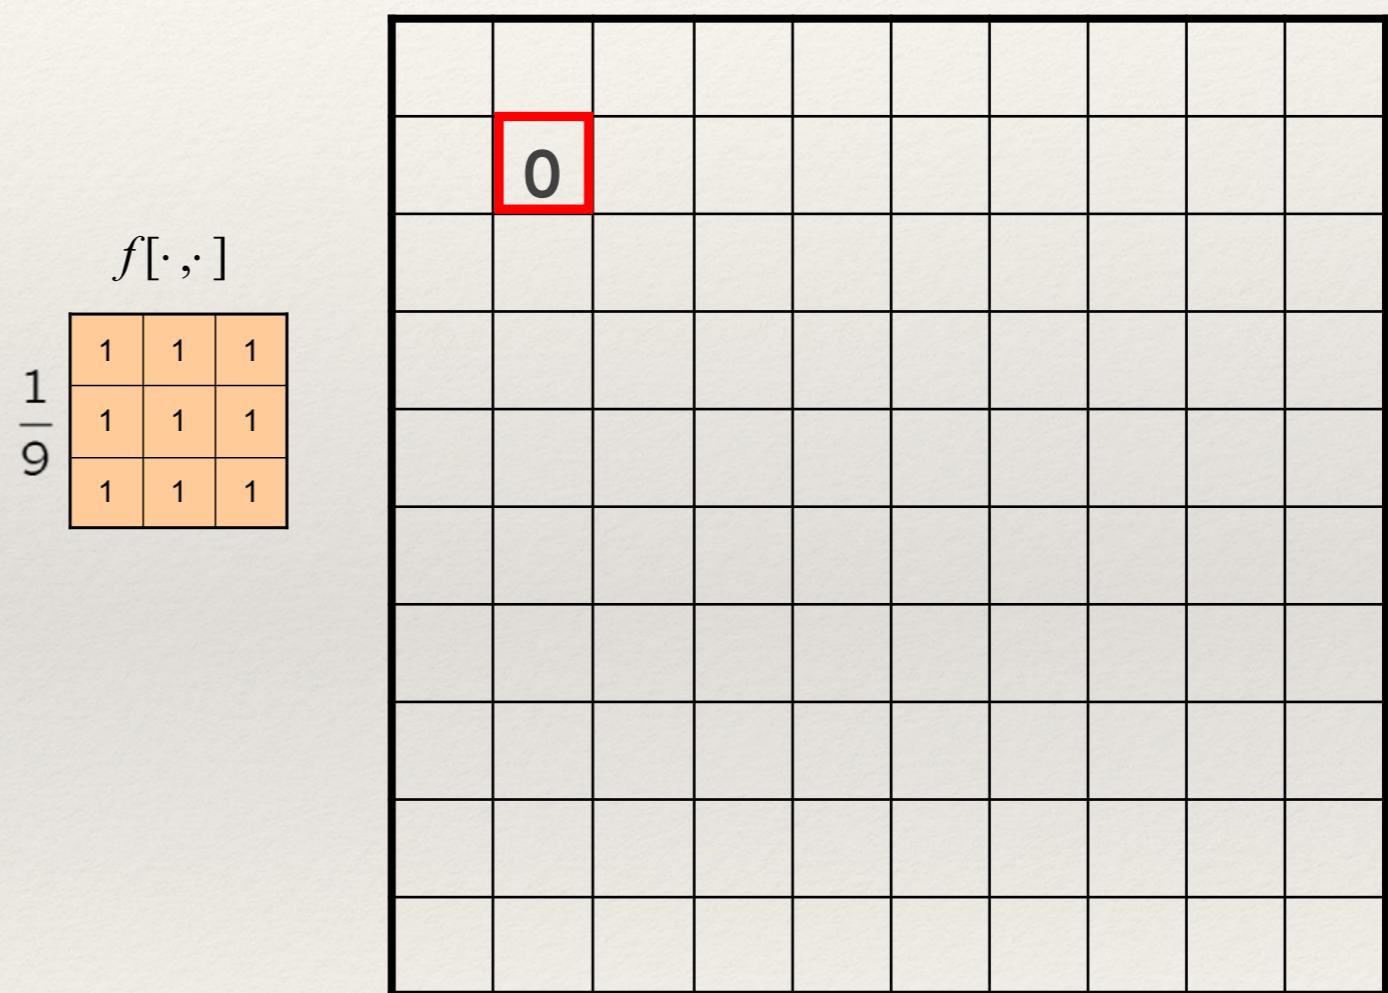
Example - Box filter

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Example - Box filter

$$I[.,.] \quad h[m,n] = \sum_{k,l} f[k,l] I[m+k, n+l] \quad h[.,.]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0



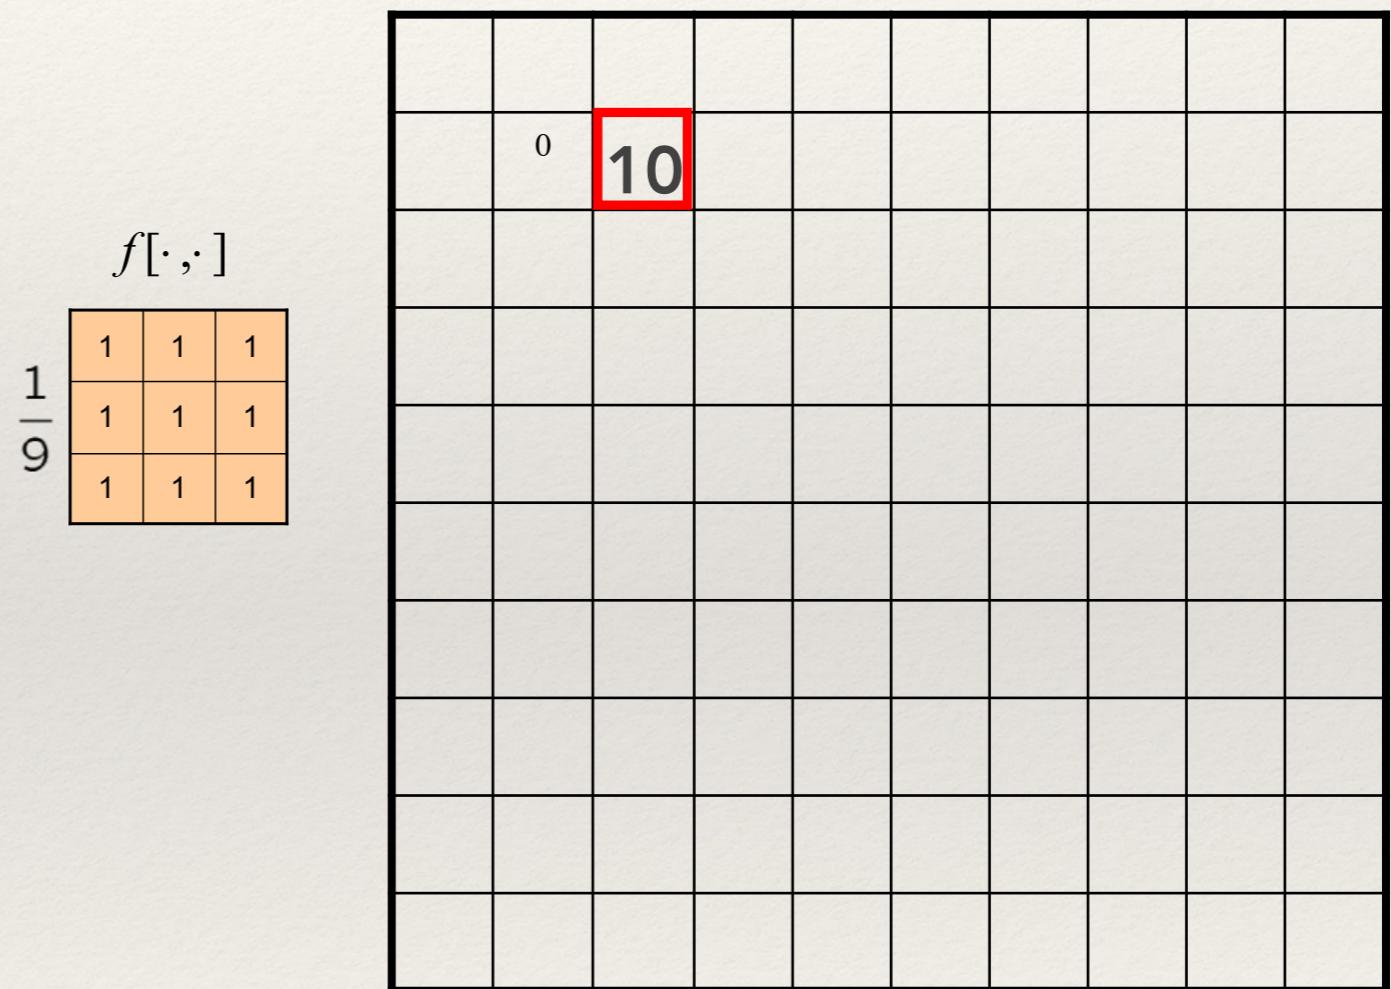
$$m = 1, n = 1$$

$$k, l = [-1, 0, 1]$$

Example - Box filter

$$I[.,.] \quad h[m,n] = \sum_{k,l} f[k,l] I[m+k, n+l] \quad h[.,.]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0



$$\begin{aligned} m &= 2, n = 1 \\ k, l &= [-1, 0, 1] \end{aligned}$$

Example - Box filter

$$I[.,.] \quad h[m,n] = \sum_{k,l} f[k,l] I[m+k, n+l] \quad h[.,.]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$f[\cdot, \cdot]$$

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

		0	10	20					

$$m = 3, n = 1$$

$$k, l = [-1, 0, 1]$$

Example - Box filter

$$I[.,.] \quad h[m,n] = \sum_{k,l} f[k,l] I[m+k, n+l] \quad h[.,.]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$f[\cdot, \cdot] = \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

			0	10	20	30			

$$\begin{aligned} m &= 4, n = 1 \\ k, l &= [-1, 0, 1] \end{aligned}$$

Example - Box filter

$$I[.,.] \quad h[m,n] = \sum_{k,l} f[k,l] I[m+k, n+l] \quad h[.,.]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$f[\cdot, \cdot] = \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

			0	10	20	30	30		

$$m = 5, n = 1$$

$$k, l = [-1, 0, 1]$$

Example - Box filter

$$I[.,.] \quad h[m,n] = \sum_{k,l} f[k,l] I[m+k, n+l] \quad h[.,.]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$f[\cdot, \cdot]$$

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

	0	10	20	30	30				

$$m = 4, n = 6$$

$$k, l = [-1, 0, 1]$$

Example - Box filter

$$I[.,.] \quad h[m,n] = \sum_{k,l} f[k,l] I[m+k, n+l] \quad h[.,.]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$f[\cdot, \cdot]$$

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

	0	10	20	30	30					

$$m = 6, n = 4$$

$$k, l = [-1, 0, 1]$$

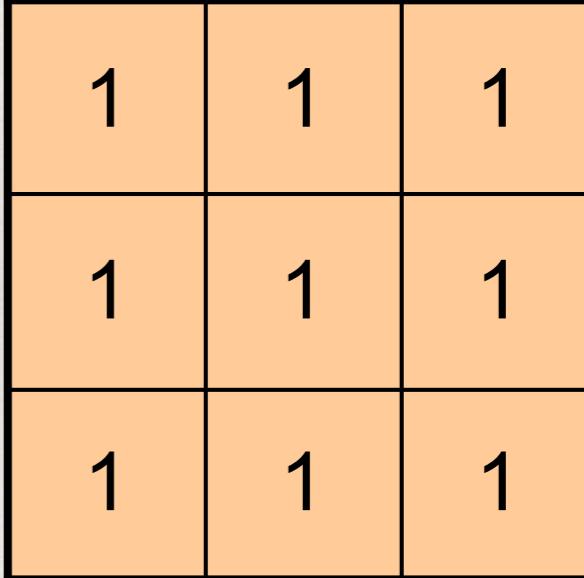
Example - Box filter

$$I[\cdot, \cdot] \quad h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l] \quad h[\cdot, \cdot]$$

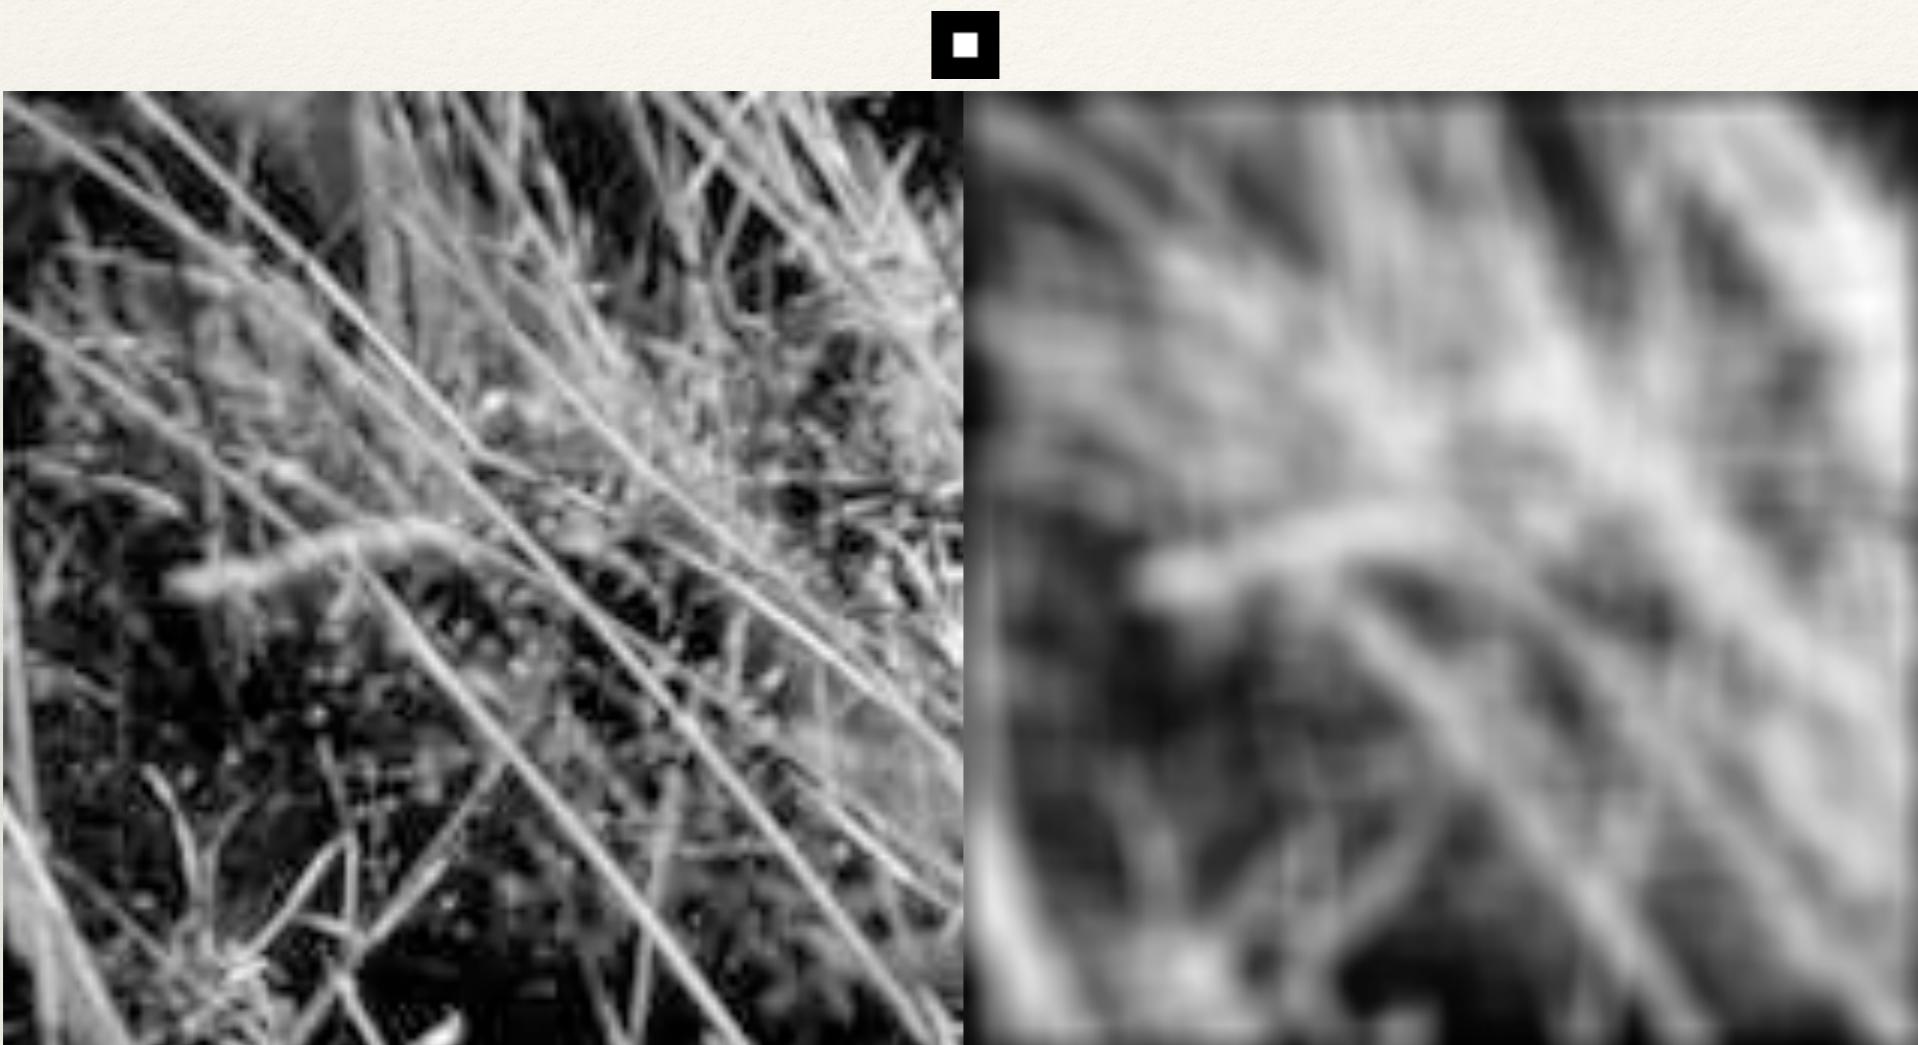
$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Box filter

- ❖ What does it do?
 - ❖ Replaces each pixel with an average of its neighborhood
 - ❖ Achieve smoothing effect
(remove sharp features)

$$\frac{1}{9} \begin{bmatrix} f[\cdot, \cdot] \\ 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$


Smoothing with box filter



$f[\cdot, \cdot]$

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

Integral image

- ❖ Don't forget to use integral image to improve computational efficiency for utilization of bunch of box filters.

3	2	7	2	3
1	5	1	3	4
5	1	3	5	1
4	3	2	1	6
2	4	1	4	8

(a) $S = 24$

3	5	12	14	17
4	11	19	24	31
9	17	28	38	46
13	24	37	48	62
15	30	44	59	81

(b) $s = 28$

3	5	12	14	17
4	11	19	24	31
9	17	28	38	46
13	24	37	48	62
15	30	44	59	81

(c) $S = 24$

Correlation and Convolution

- ❖ 2D correlation:

$$h[m,n] = \sum_{k,l} f[k,l] I[m+k, n+l]$$

- ❖ 2D convolution:

$$h[m,n] = \sum_{k,l} f[k,l] I[m-k, n-l]$$

- ❖ **Convolution is the same as correlation with a 180° rotated filter kernel.**
- ❖ **Correlation and convolution are identical when the filter kernel is symmetric.**

Key properties of linear filters

- ❖ **Linearity:**

$$\text{imfilter}(I, f_1 + f_2) = \text{imfilter}(I, f_1) + \text{imfilter}(I, f_2)$$

- ❖ **Shift invariance:** same behavior given intensities regardless of pixel location m, n

$$\text{imfilter}(I, \text{shift}(f)) = \text{shift}(\text{imfilter}(I, f))$$

- ❖ Any linear, shift-invariant operator can be represented as a convolution.

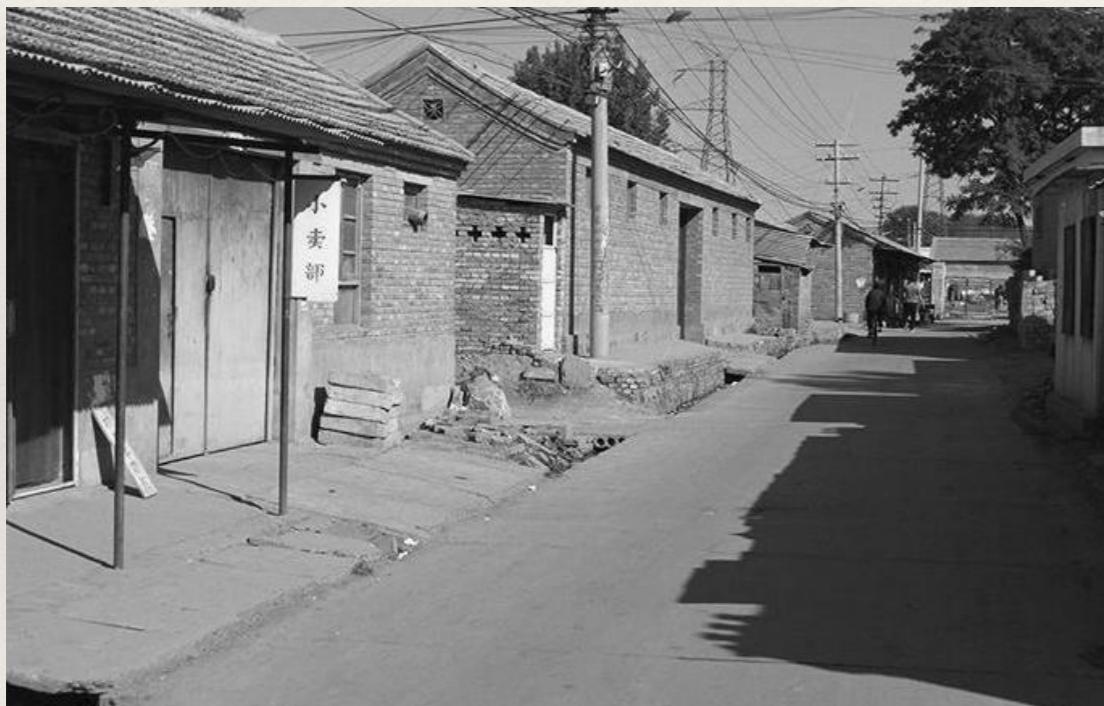
Convolution properties

- ❖ Commutative: $a * b = b * a$
 - ❖ Conceptually no difference between filter and signal
- ❖ Associative: $a * (b * c) = (a * b) * c$
 - ❖ Often apply several filters one after another: $((a * b_1) * b_2) * b_3$
 - ❖ This is equivalent to applying one filter: $a * (b_1 * b_2 * b_3)$
 - ❖ Correlation is not associative (rotation effect)
- ❖ Distributes over addition: $a * (b + c) = (a * b) + (a * c)$
- ❖ Scalars factor out: $ka * b = a * kb = k(a * b)$
- ❖ Identity: unit impulse $e = [0, 0, 1, 0, 0]$, $a * e = a$

Correlation for template matching

Let's see if we can use correlation to 'find' the parts of the image that look like the filter.

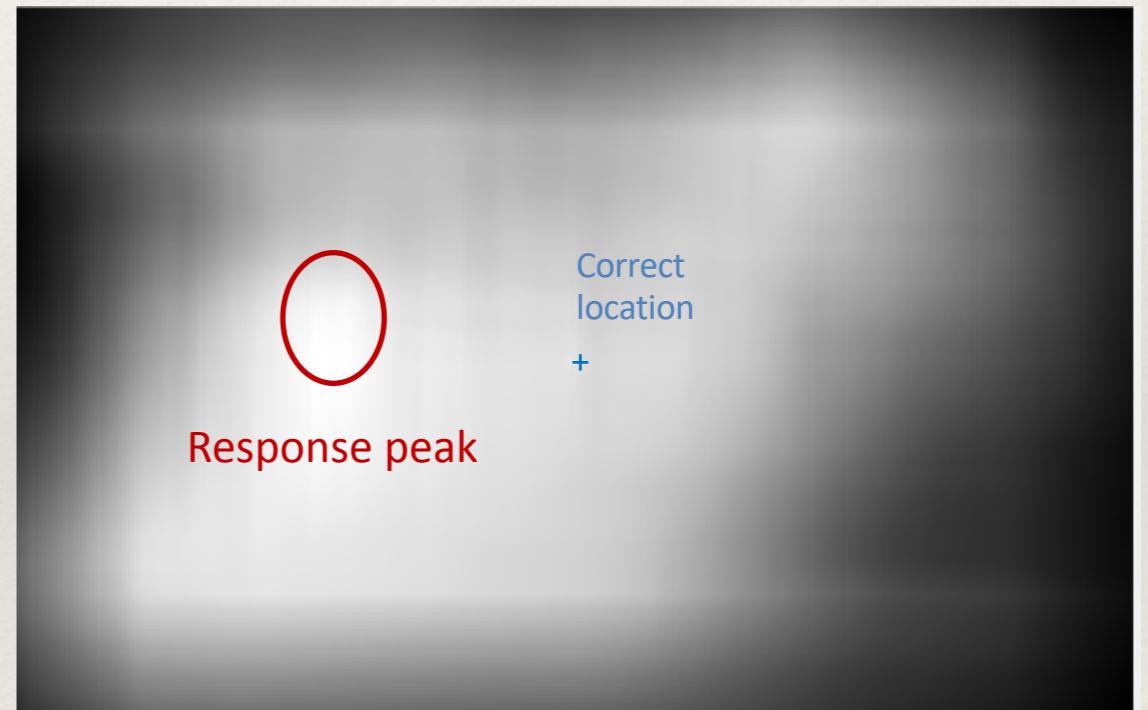
D



f



$I = \text{correlate2d}(D, f, \text{'same'})$



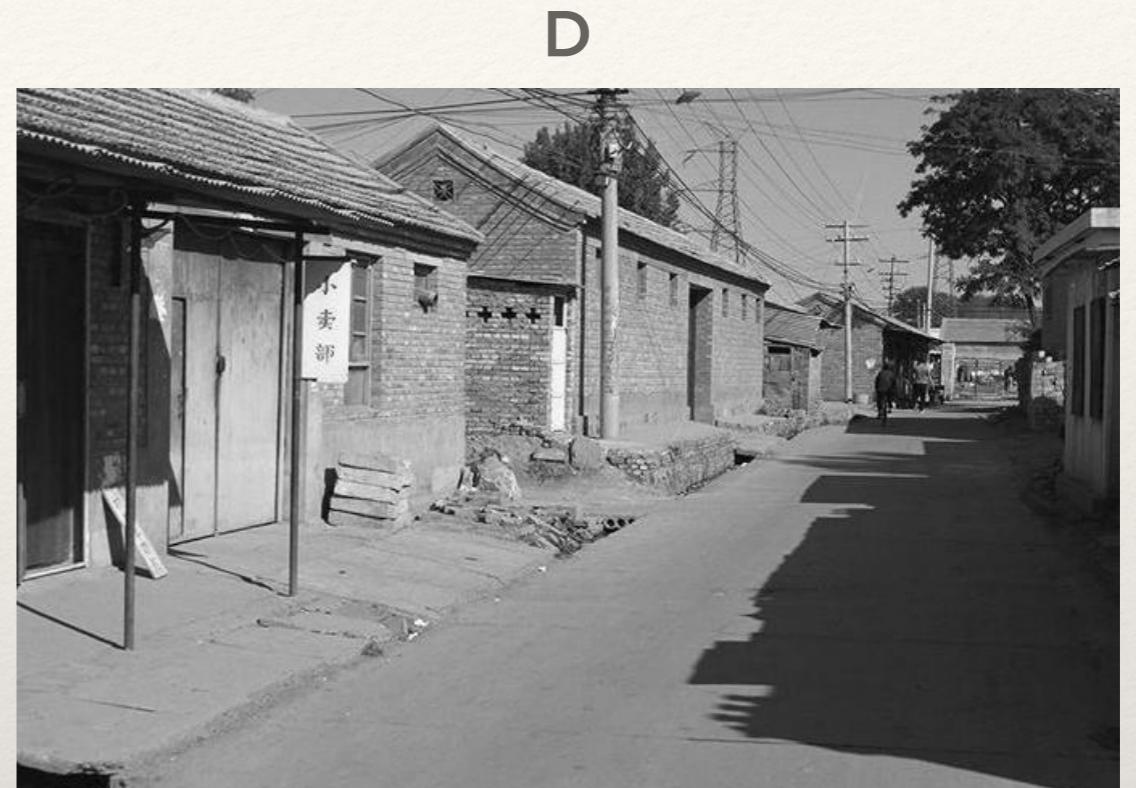
It does't work. Why?

Correlation for template matching

- ❖ It does't work. Why?

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

- ❖ As brightness in I increases, the response in h will increase, as long as f is positive.
- ❖ Overall brighter regions will give higher correlation response.

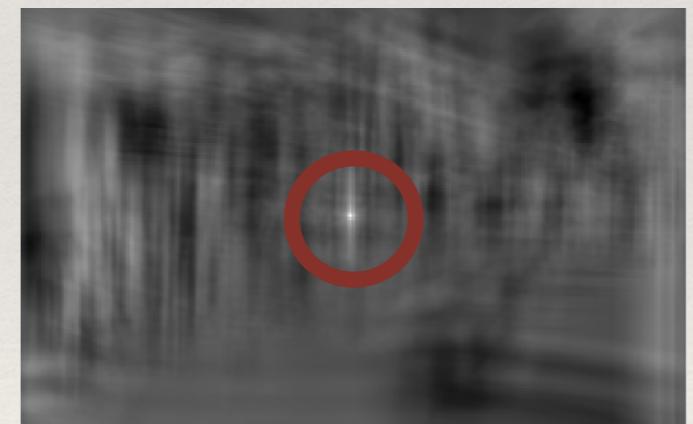
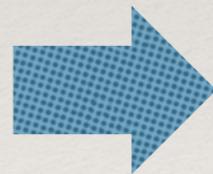
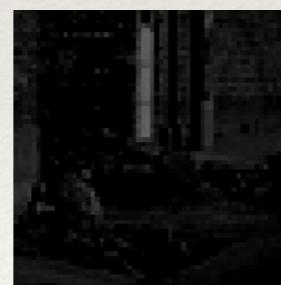


f



Correlation for template matching

- ❖ **Normalization correlation.** Subtract the mean so zero centered. Score is higher only when dark parts match and when light parts match.
- ❖ This is a great way to build simple, fast pattern detectors that require little computation.



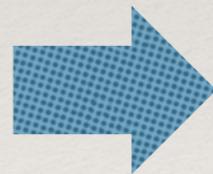
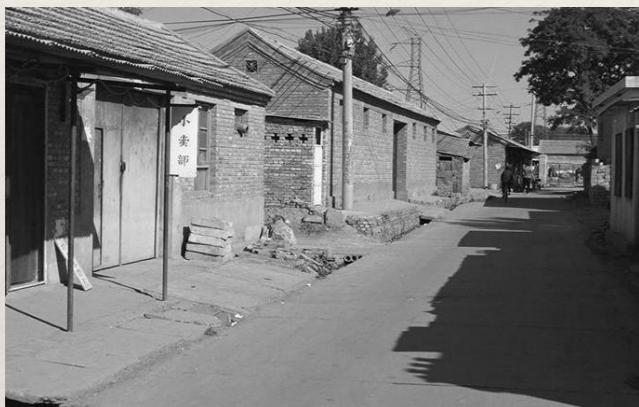
$I2 = \text{correlate2d}(D2, f2, 'same')$

$D2 = D - \text{np.mean}(D)$

$f2 = f - \text{np.mean}(f)$

Correlation for template matching

- ❖ What happens with convolution?



$I2 = \text{convolve2d}(D2, f2, \text{'same'})$

$D2 = D - \text{np.mean}(D)$

$f2 = f - \text{np.mean}(f)$

Correlation for template matching

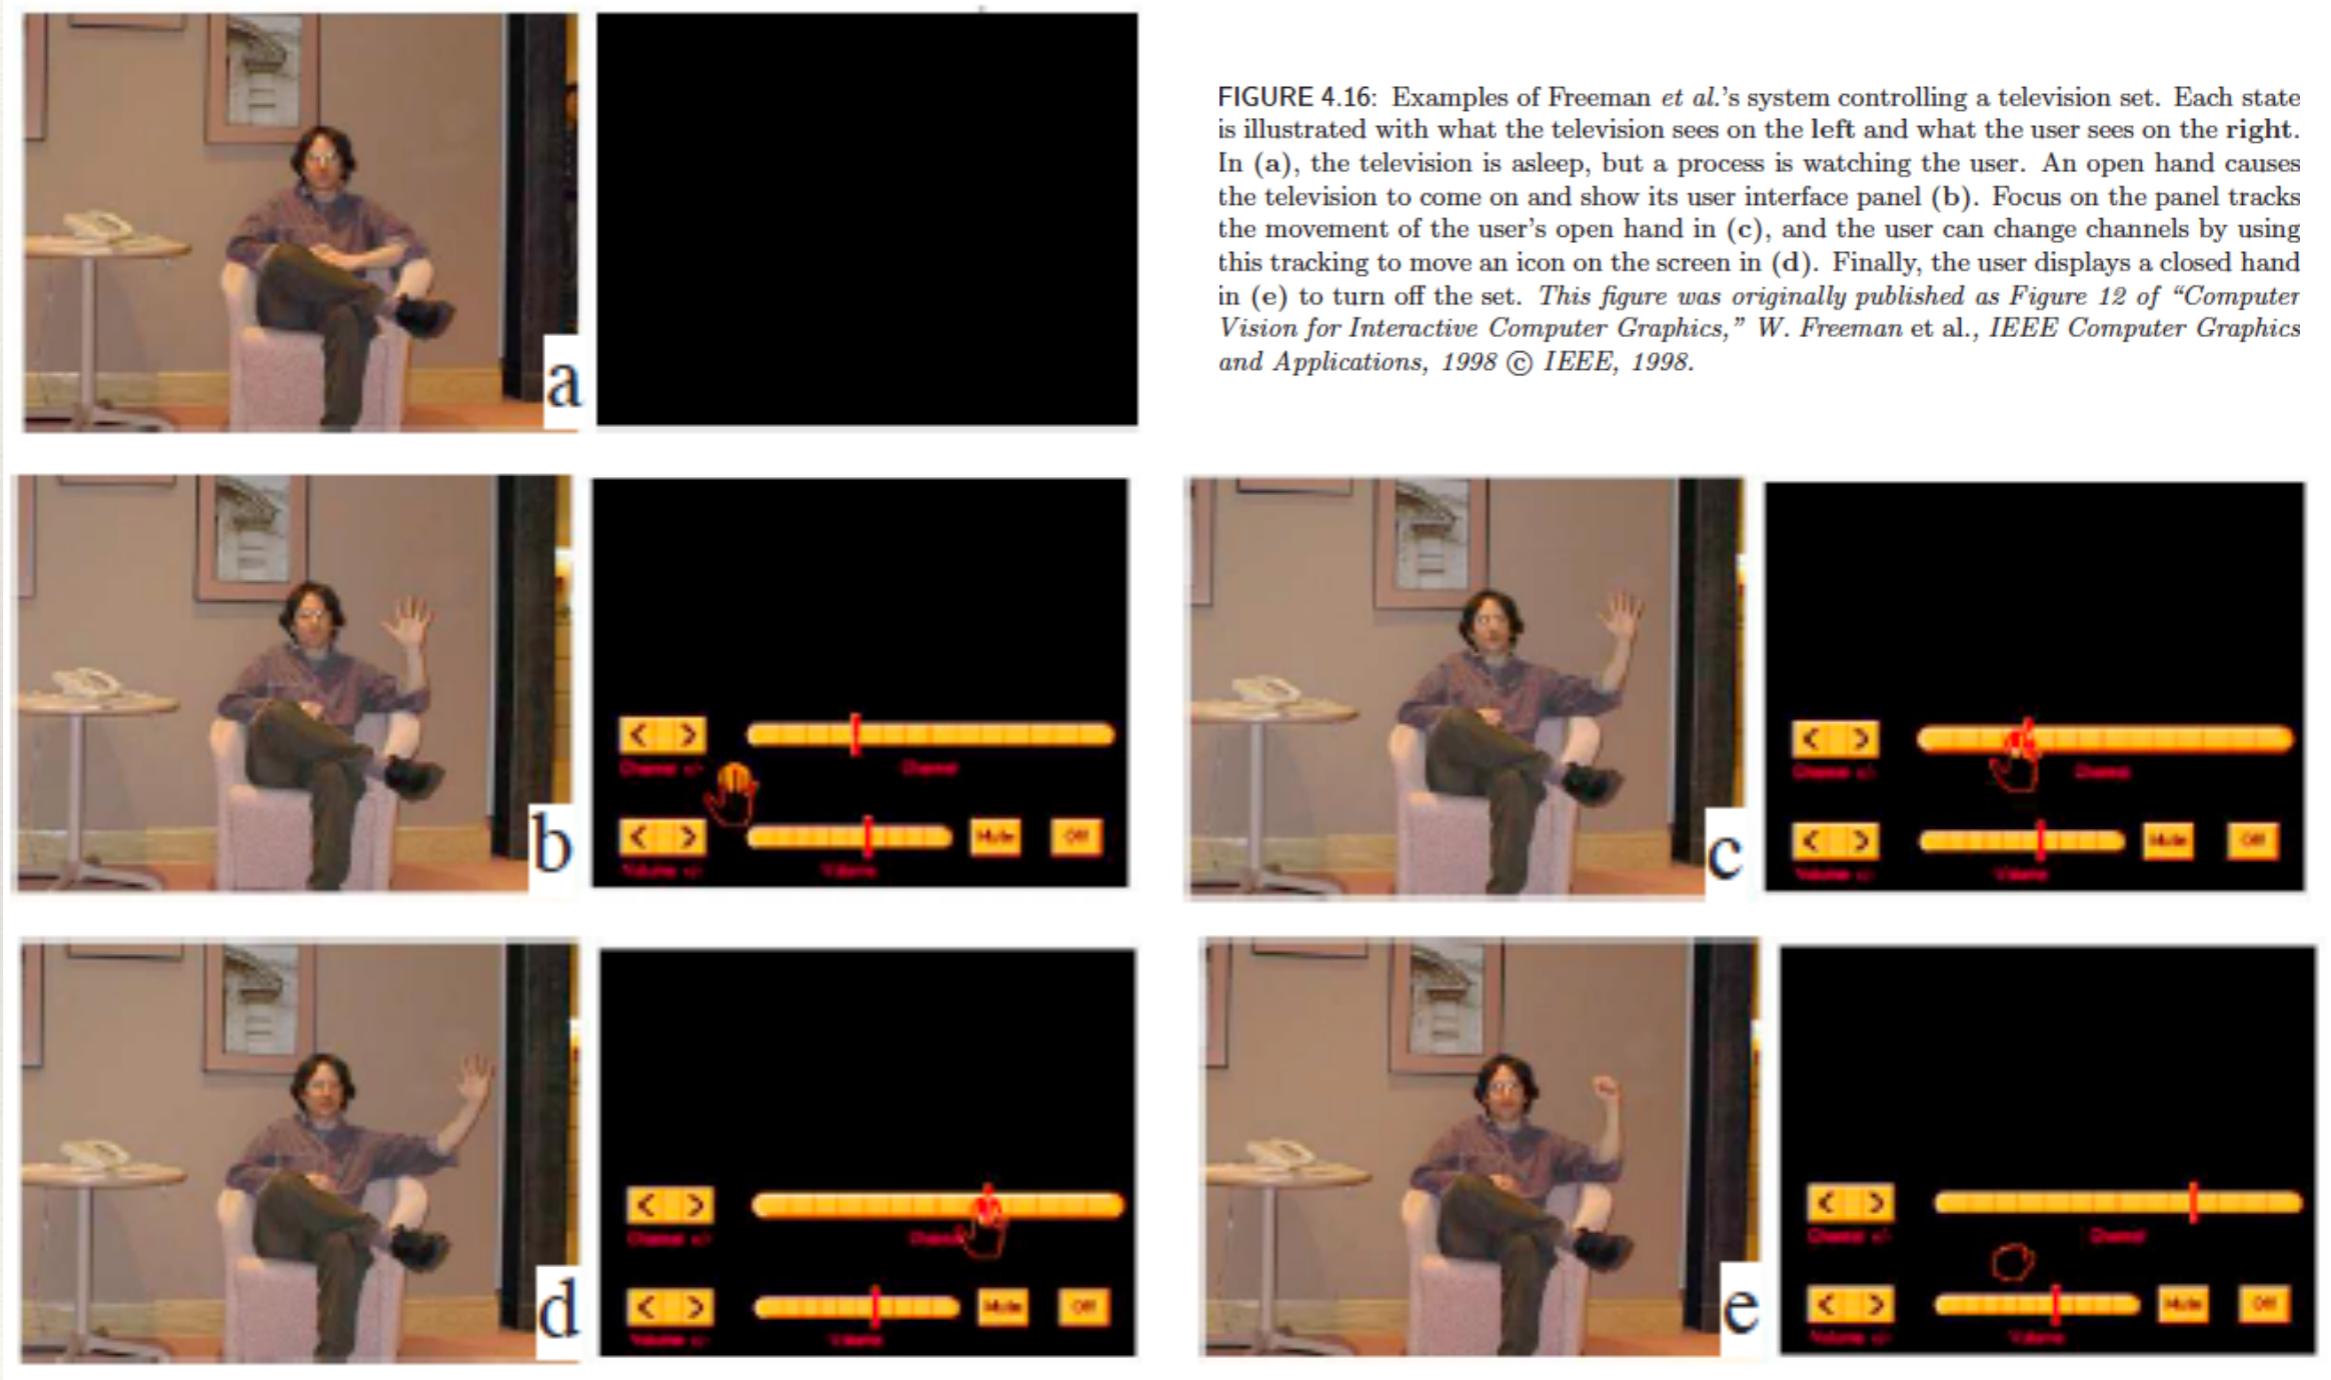
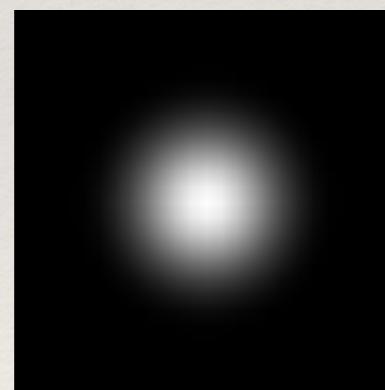
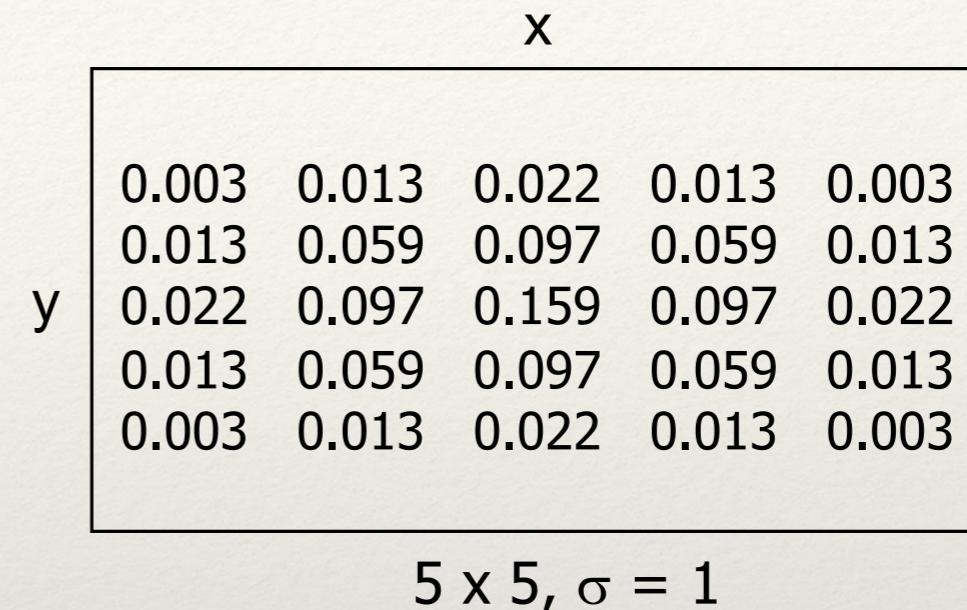
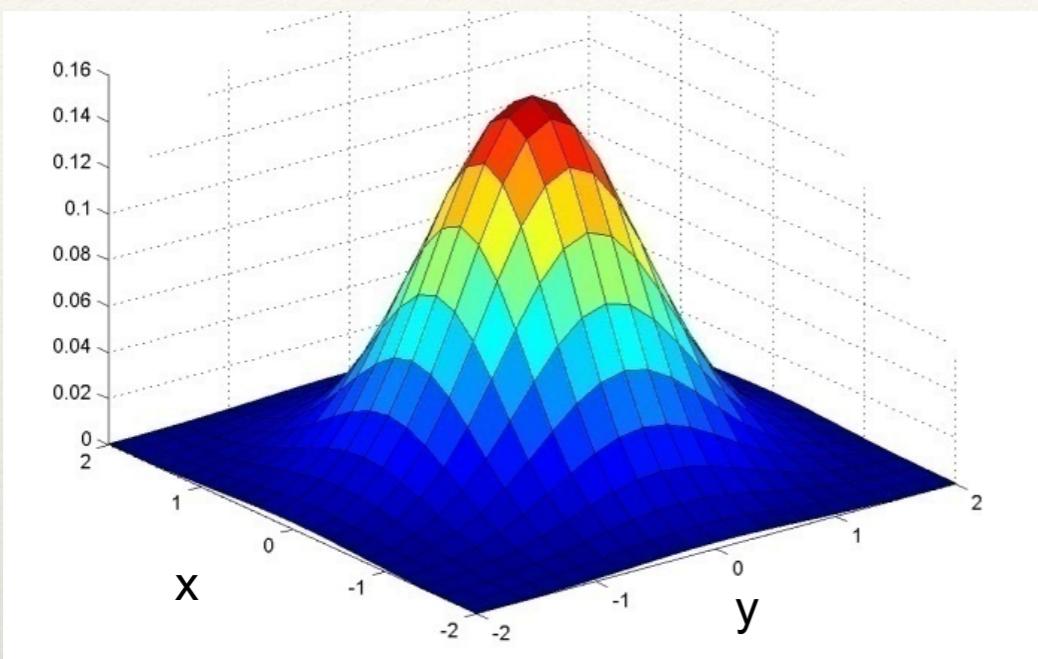


FIGURE 4.16: Examples of Freeman *et al.*'s system controlling a television set. Each state is illustrated with what the television sees on the left and what the user sees on the right. In (a), the television is asleep, but a process is watching the user. An open hand causes the television to come on and show its user interface panel (b). Focus on the panel tracks the movement of the user's open hand in (c), and the user can change channels by using this tracking to move an icon on the screen in (d). Finally, the user displays a closed hand in (e) to turn off the set. *This figure was originally published as Figure 12 of “Computer Vision for Interactive Computer Graphics,” W. Freeman et al., IEEE Computer Graphics and Applications, 1998 © IEEE, 1998.*

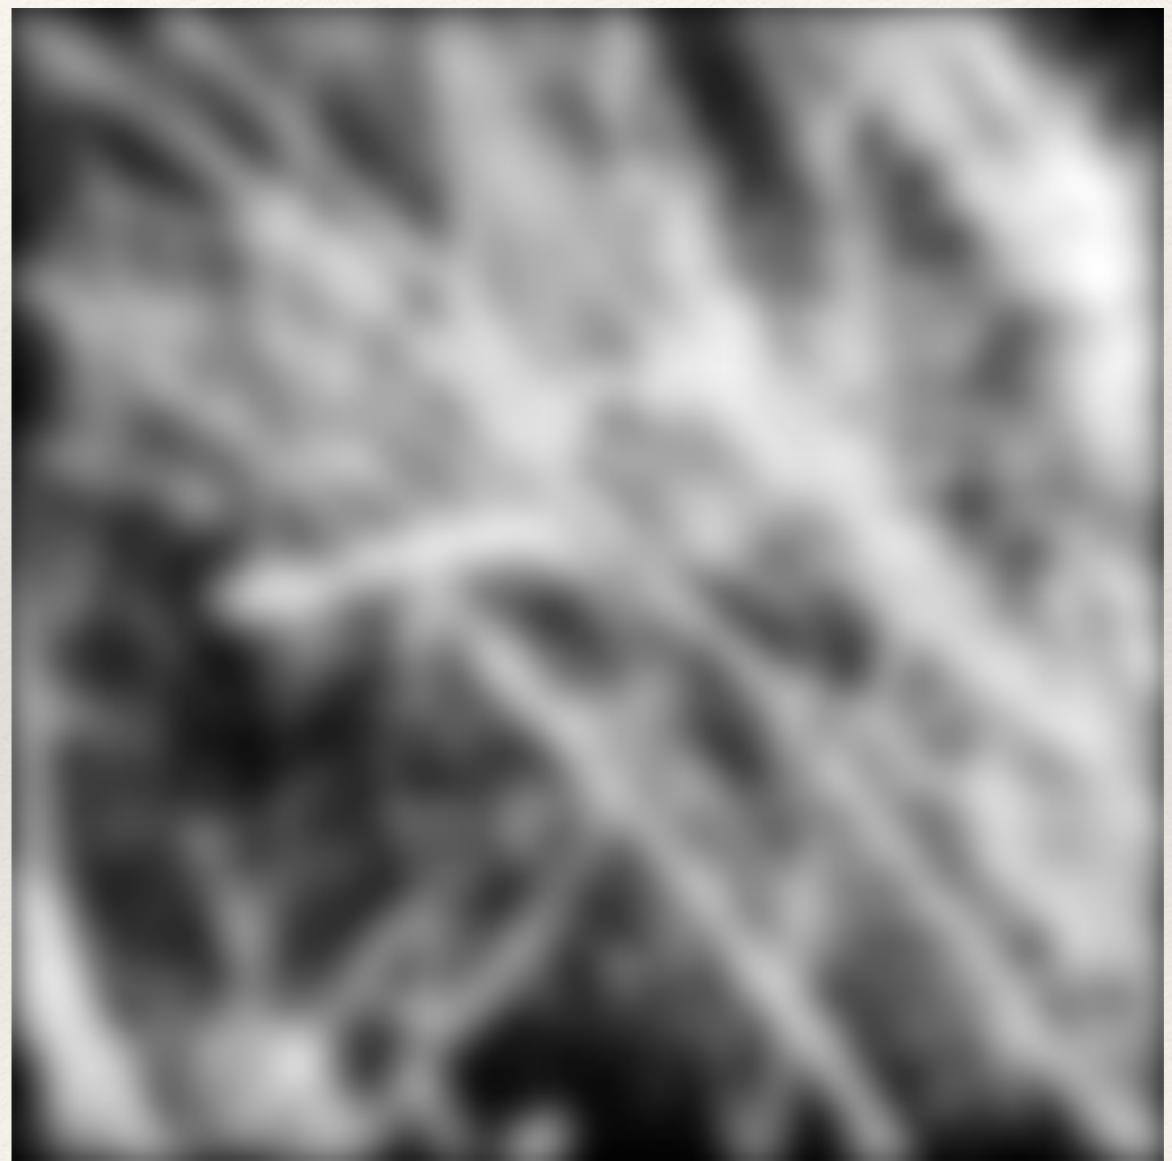
Gaussian filters



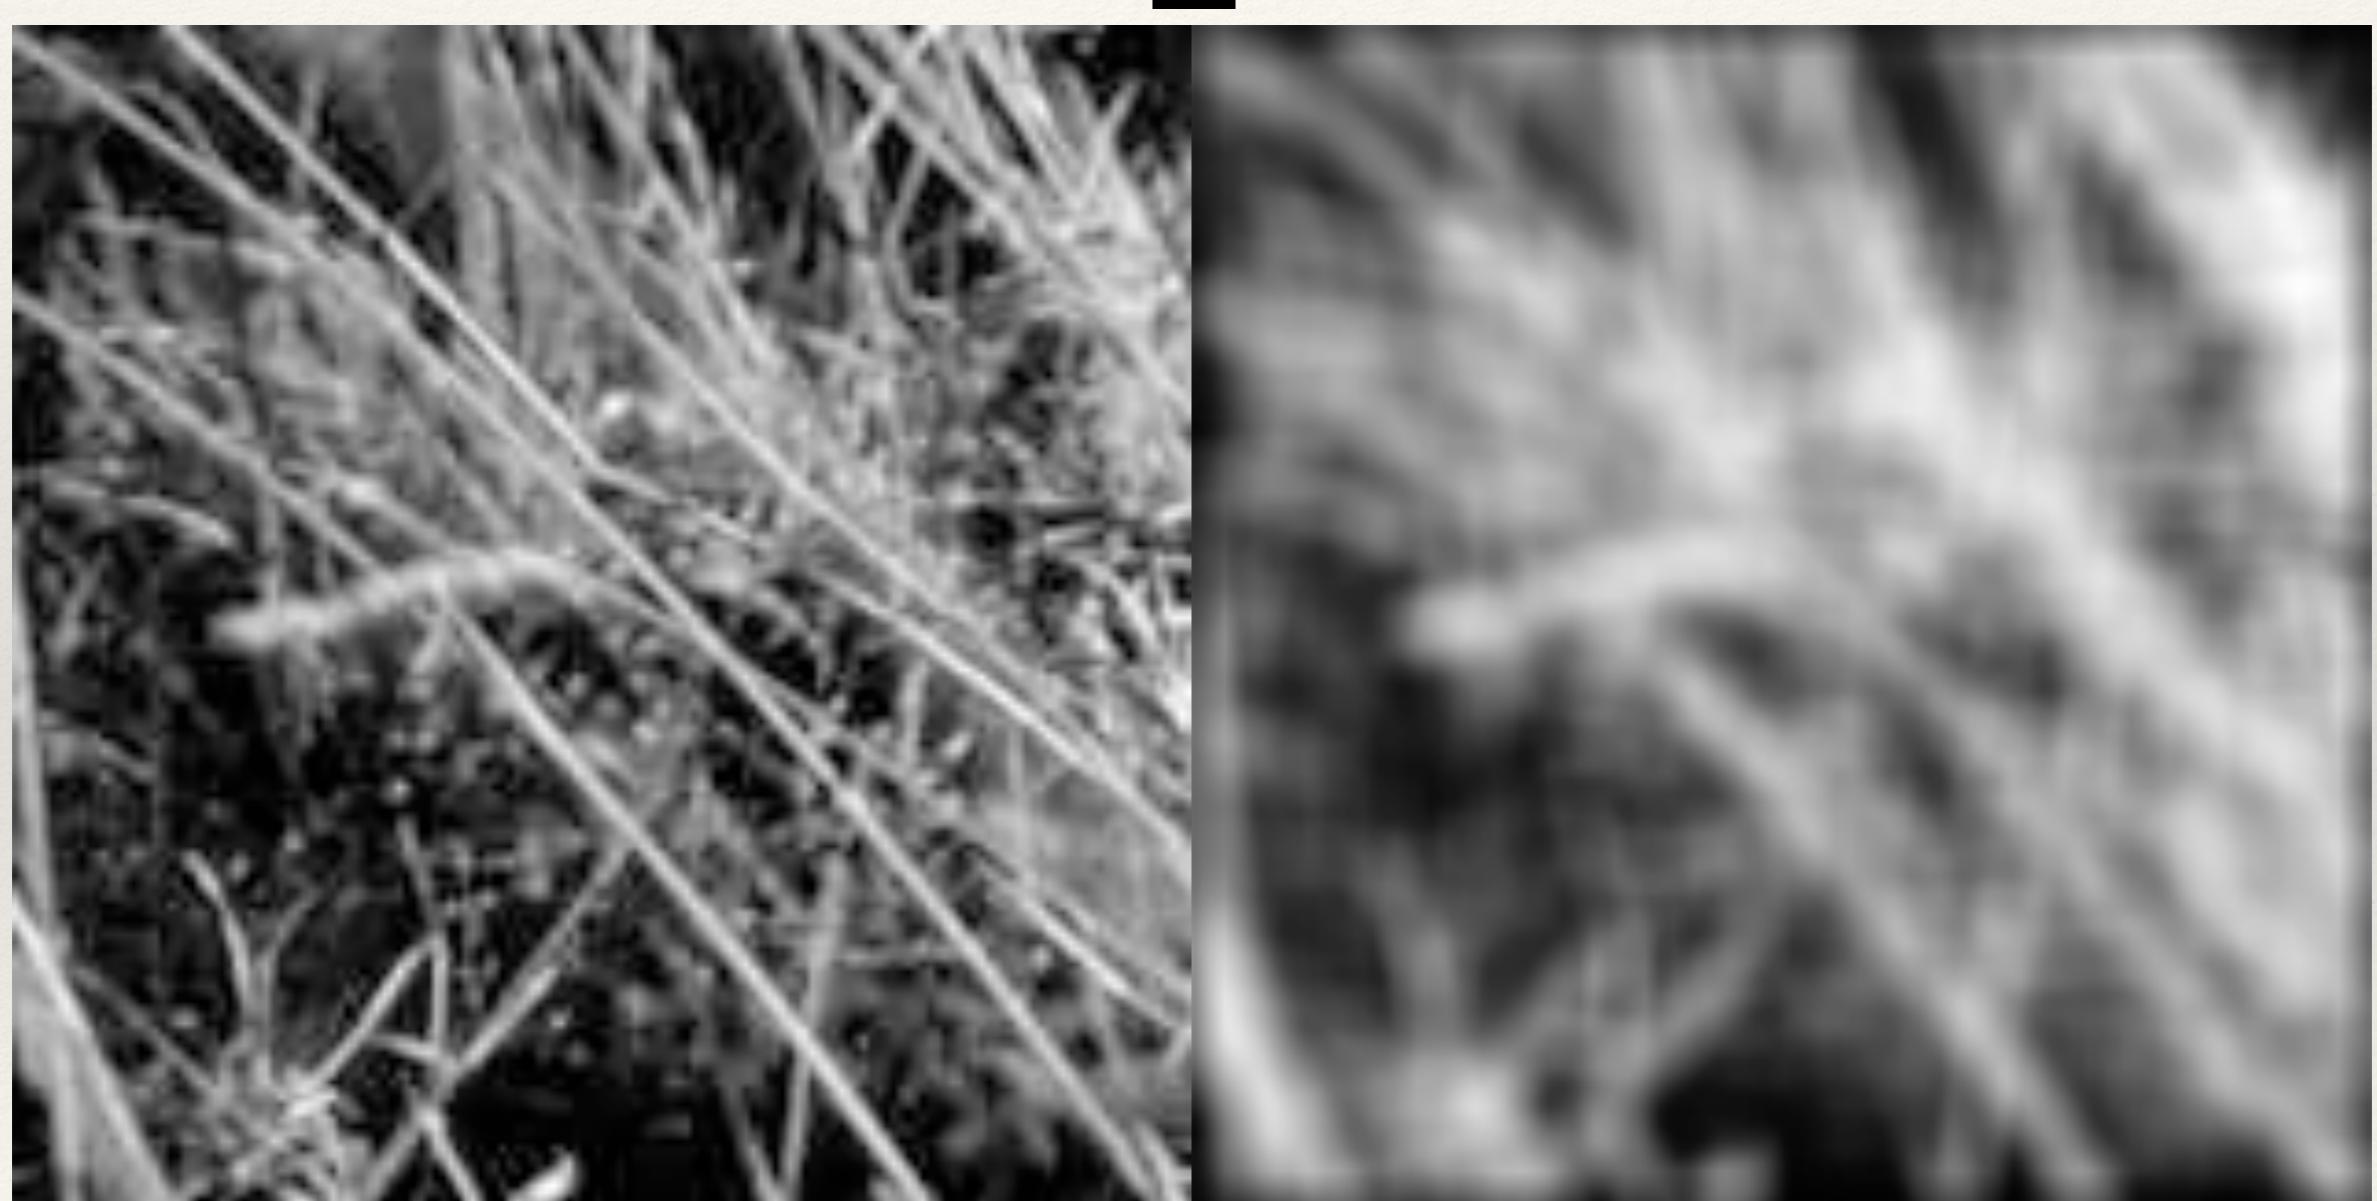
$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

Weight contributions of neighboring pixels by nearness

Smoothing with Gaussian filter



Smoothing with box filter

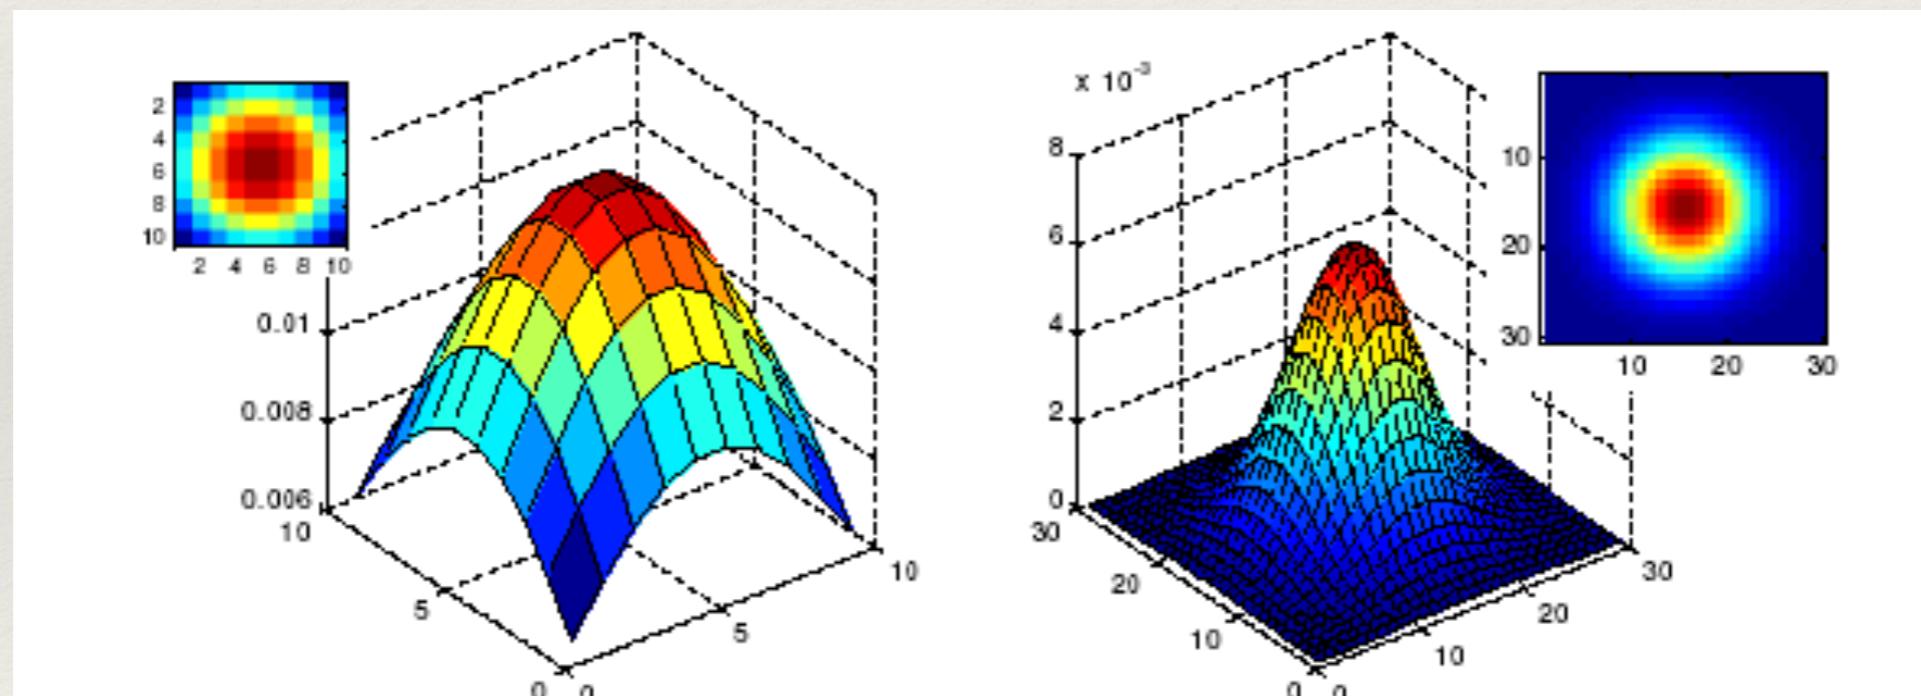


Gaussian filters

- ❖ Remove “high-frequency” components from the image (low-pass filter)
 - ❖ Images become more smooth
- ❖ Gaussian convolved with Gaussian is another Gaussian
 - ❖ So can smooth with small-width kernel, repeat, and get same result as larger-width kernel would have
 - ❖ Convolving twice with Gaussian kernel of width σ is same as convolving once with kernel of width $\sigma\sqrt{2}$
- ❖ Separable kernel
 - ❖ Factors into product of two 1D Gaussians

Gaussian filters

- ❖ What parameters matter here?
 - ❖ **Size** of kernel or mask

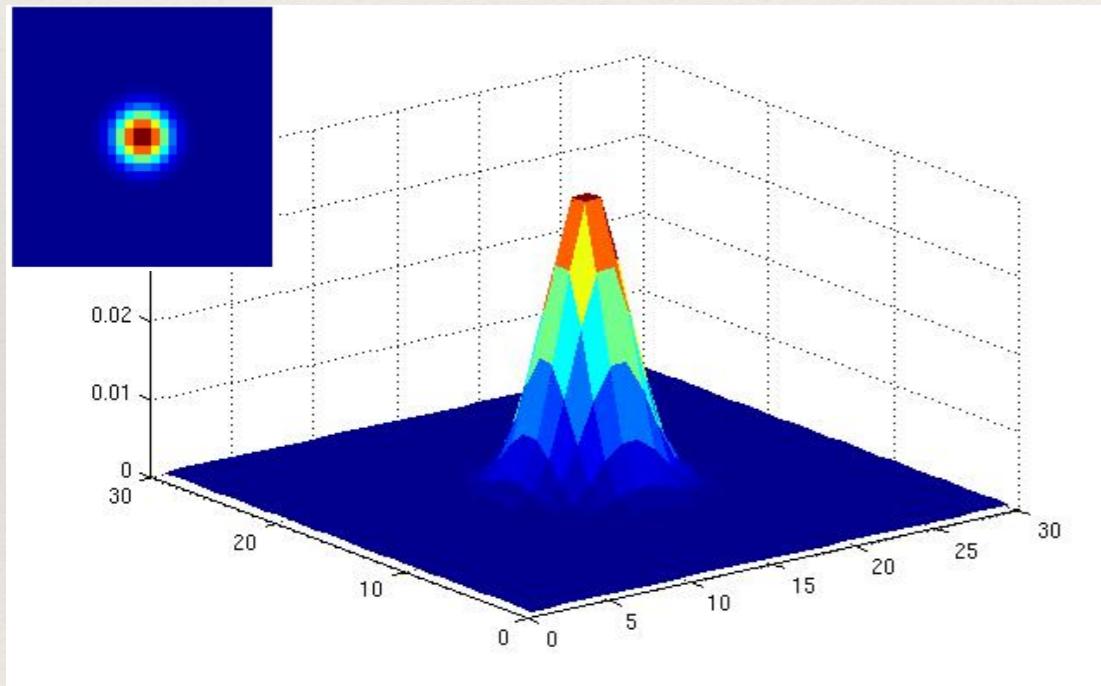


$\sigma = 5$ with 10×10 kernel

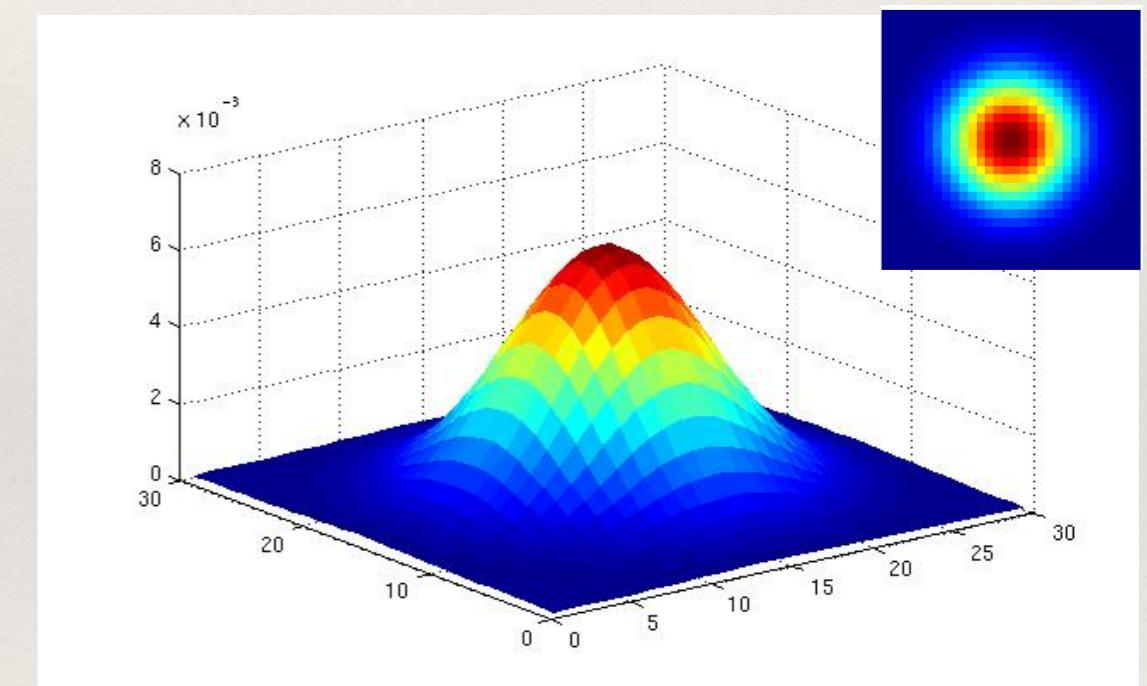
$\sigma = 5$ with 30×30 kernel

Gaussian filters

- ❖ What parameters matter here?
 - ❖ **Variance** of Gaussian: determines extent of smoothing



$\sigma = 5$ with 10×10 kernel



$\sigma = 5$ with 30×30 kernel

Gaussian filters: Separability

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y

In this case, the two functions are the (identical) 1D Gaussian

Separability

2D convolution
(center location only)

The filter factors
into a product of 1D
filters:

Perform convolution
along rows, followed by
convolution
along the remaining
column:

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 2 & 3 & 3 \\ \hline 3 & 5 & 5 \\ \hline 4 & 4 & 6 \\ \hline \end{array} = \begin{array}{l} = 2 + 6 + 3 = 11 \\ = 6 + 20 + 10 = 36 \\ = 4 + 8 + 6 = 18 \\ \hline \end{array} \quad 65$$

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} \times \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 2 & 3 & 3 \\ \hline 3 & 5 & 5 \\ \hline 4 & 4 & 6 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline 11 & & \\ \hline & 18 & \\ \hline & 18 & \\ \hline \end{array}$$

$$\begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 11 & & \\ \hline & 18 & \\ \hline & 18 & \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & 65 & \\ \hline \end{array}$$

Source: K. Grauman

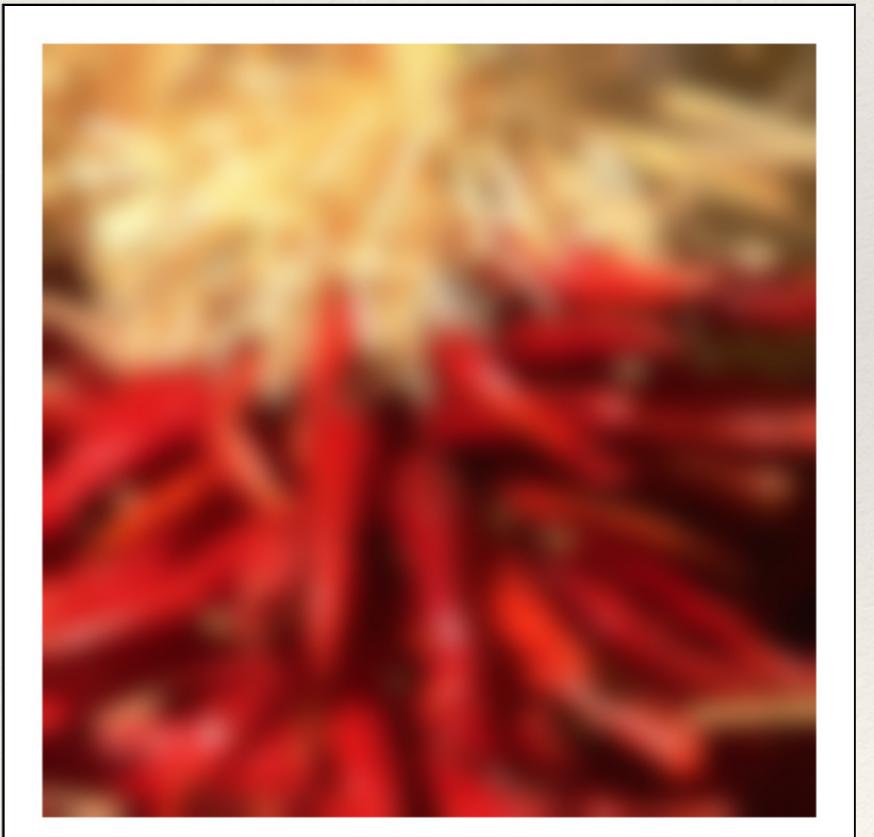
Separability

- ❖ Why is separability useful in practice?
 - ❖ MxN image, PxQ filter
 - ❖ 2D convolution: $\sim MNPQ$ multiply-adds
 - ❖ Separable 2D: $\sim MN(P+Q)$ multiply-adds
 - ❖ Speed up = $PQ/(P+Q)$
 - ❖ 9x9 filter = $\sim 4.5x$ faster

Practical matters

- ❖ What about near the edge?

- ❖ The filter window falls off the edge of the image
- ❖ Need to extrapolate
- ❖ Methods
 - ❖ clip filter (black)
 - ❖ wrap around
 - ❖ copy edge
 - ❖ reflect across edge



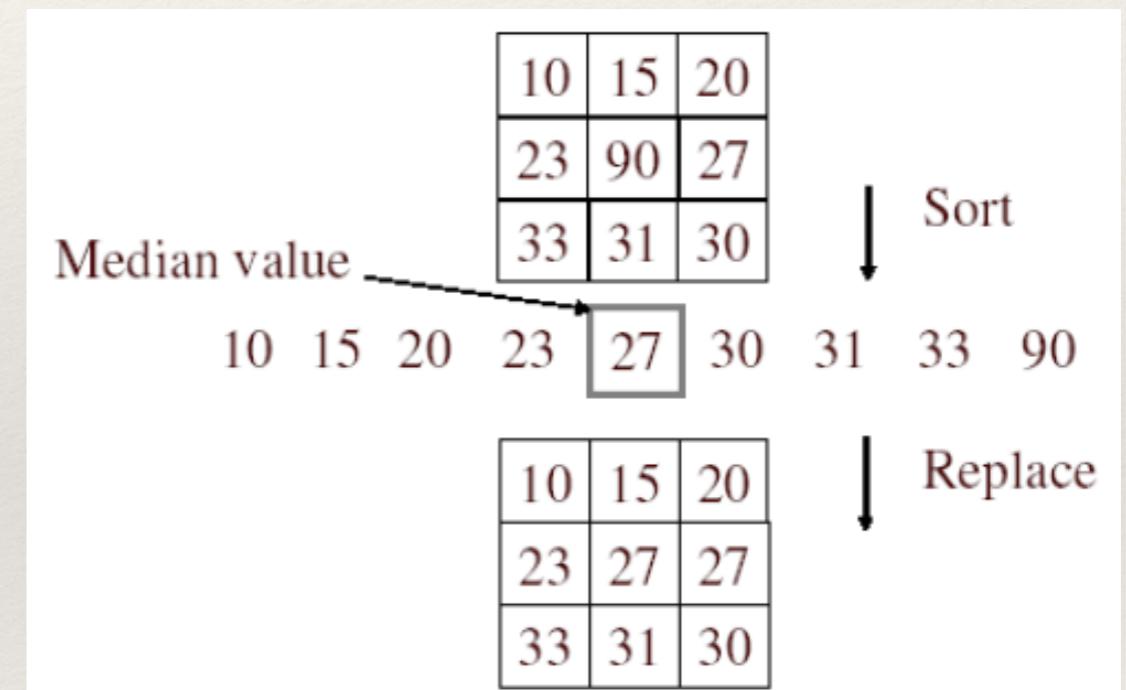
Source: S. Marschner

Convolution in Convolutional Neural Networks

- ❖ Convolution is the basic operation in CNNs
- ❖ Learning convolution kernels allows us to learn which ‘features’ provide useful information in images

Non-linear filters

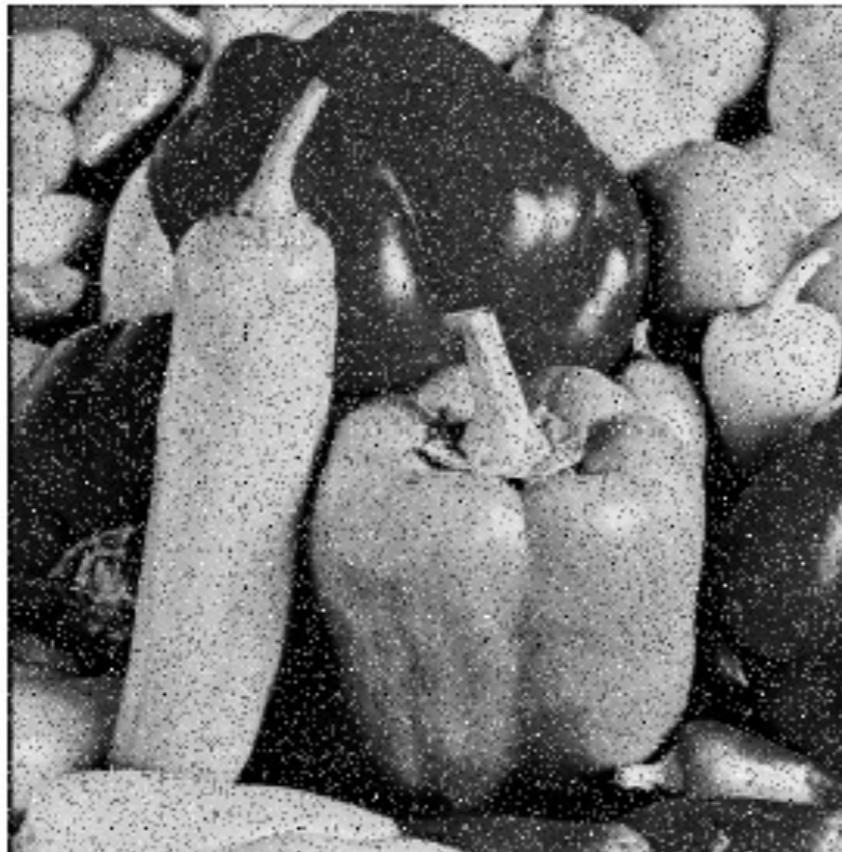
- ❖ Median filter
 - ❖ Operates over a window by selecting the median intensity in the window.
 - ❖ ‘Rank’ filter as based on ordering of gray levels
 - ❖ E.G., min, max, range filters



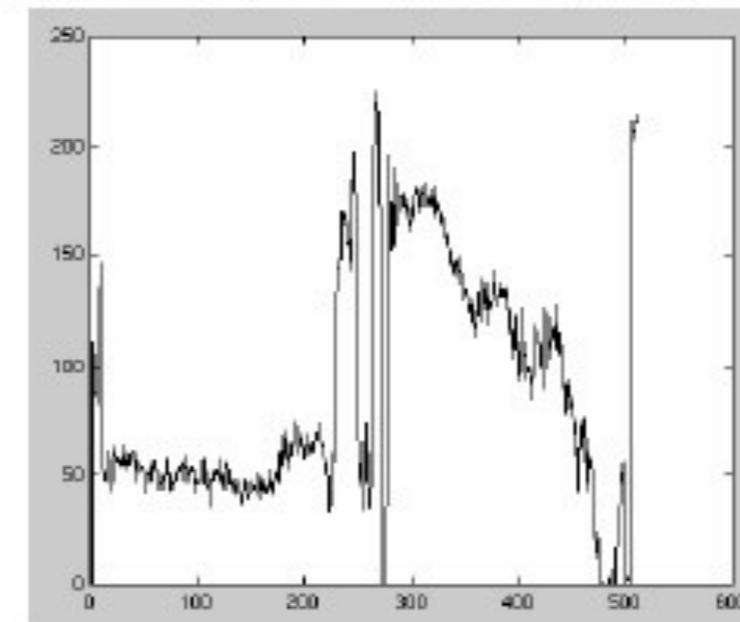
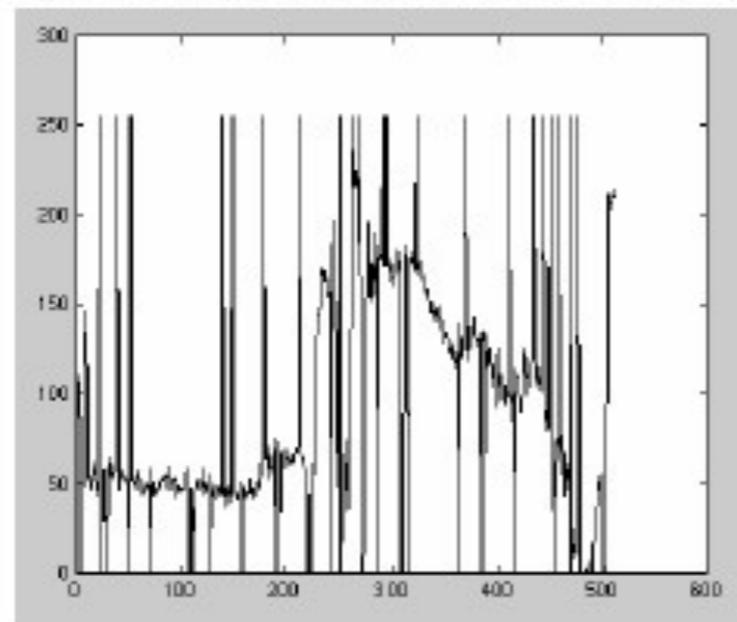
Non-linear filters

- ❖ Median filter
 - ❖ Operates over a window by selecting the median intensity in the window.
 - ❖ ‘Rank’ filter as based on ordering of gray levels
 - ❖ E.G., min, max, range filter
 - ❖ Very efficient to remove salt-and-pepper noise
 - ❖ Variants: α -trimmed mean, weighted median

Salt and pepper noise



Median filtered

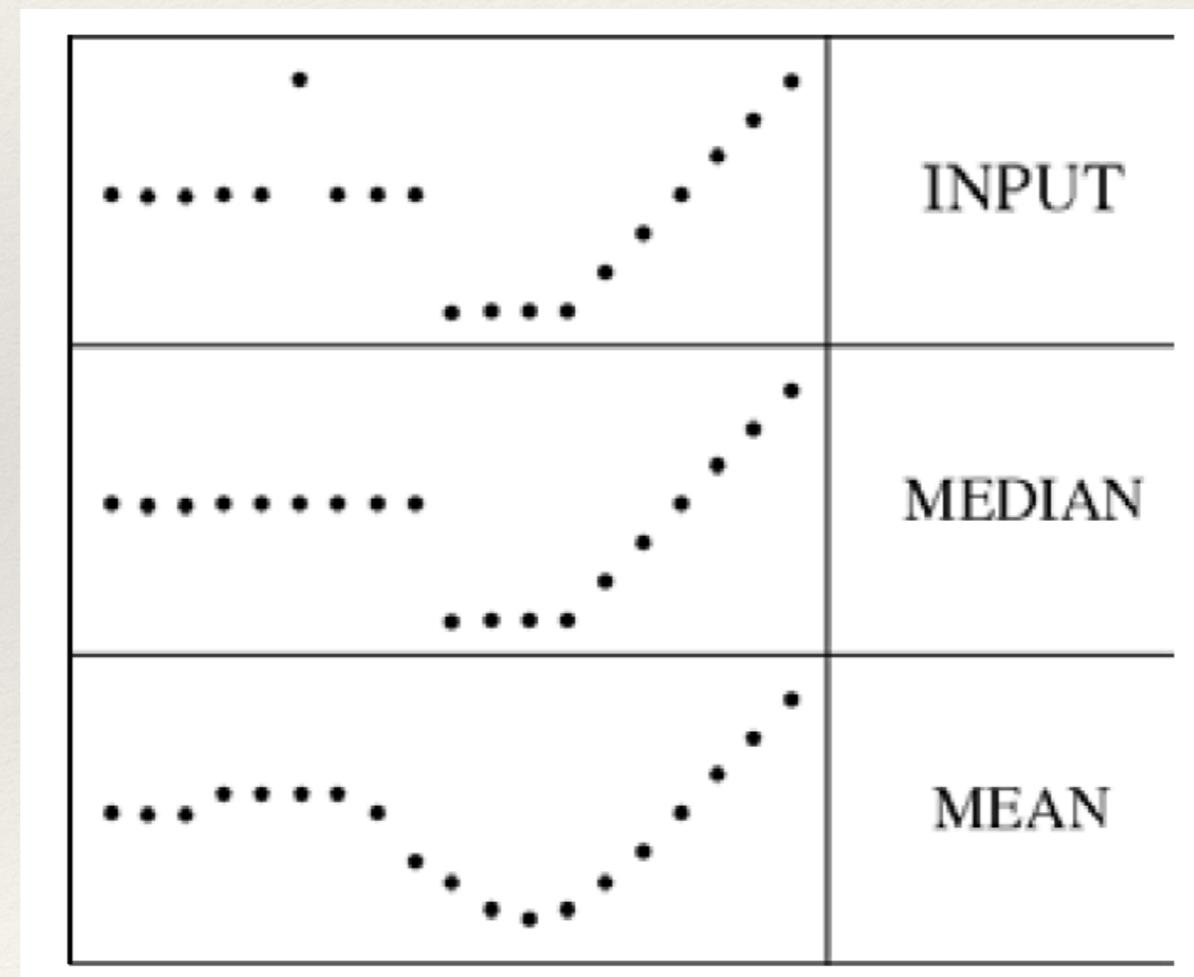


Plots of a row of the image

Source: M. Hebert

Non-linear filters

- ❖ Median filter
 - ❖ Median filter is edge preserving



Source: K. Grauman

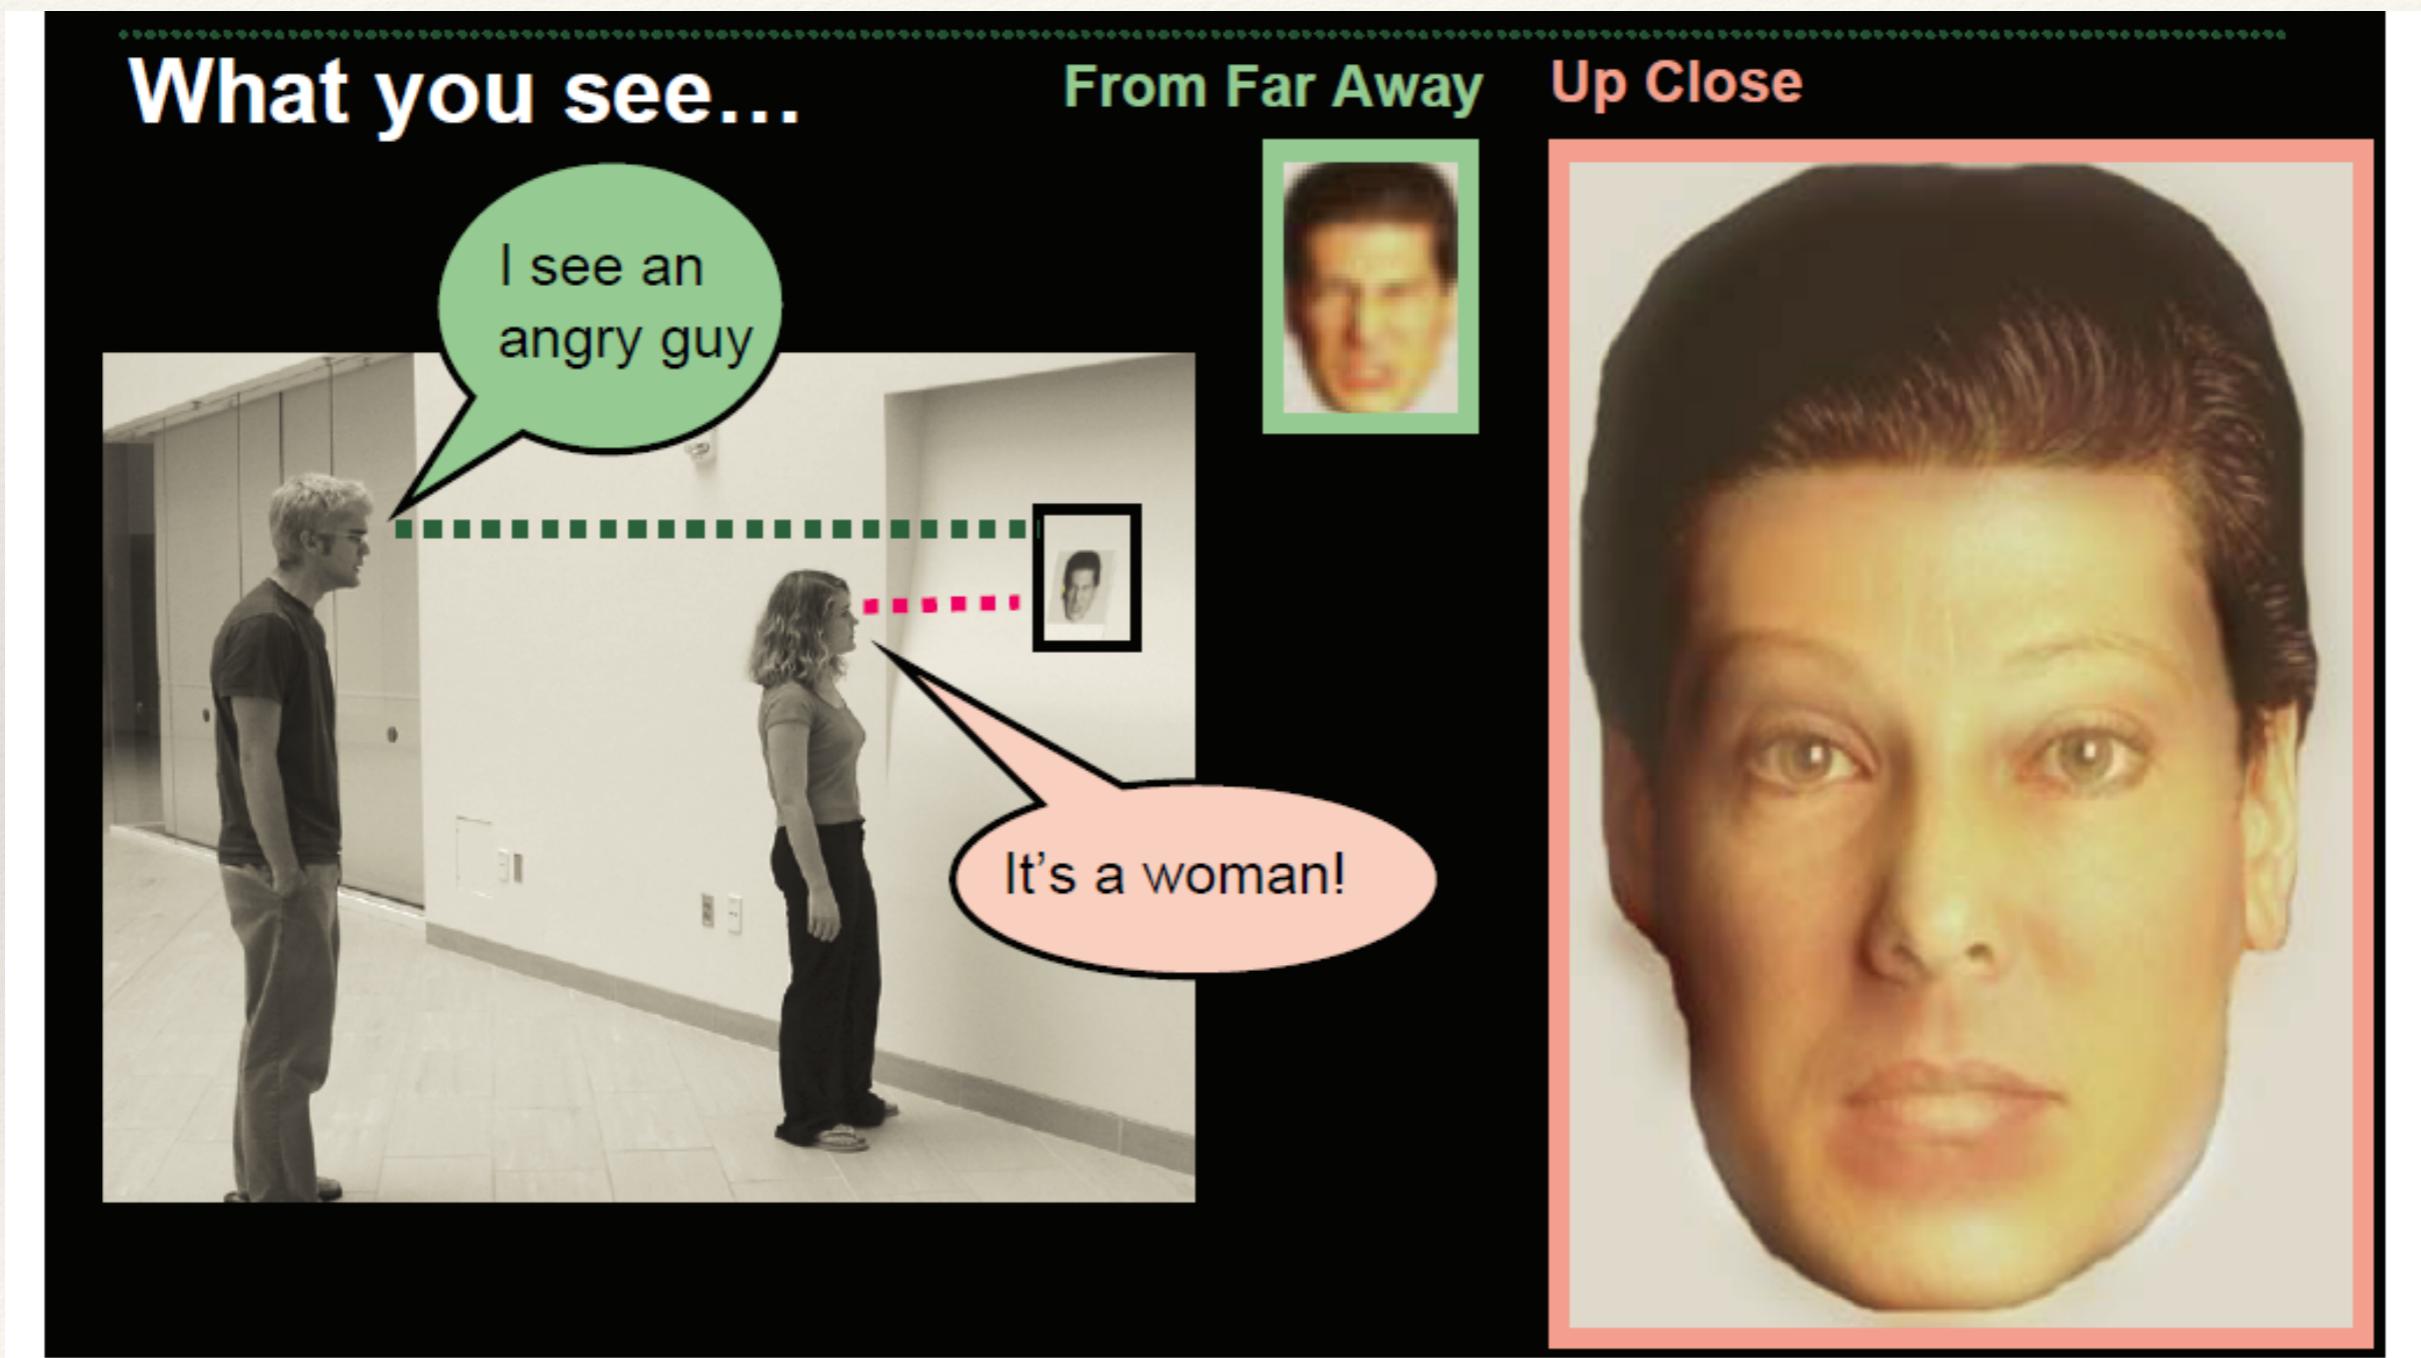
Frequency perspective



What can you see?

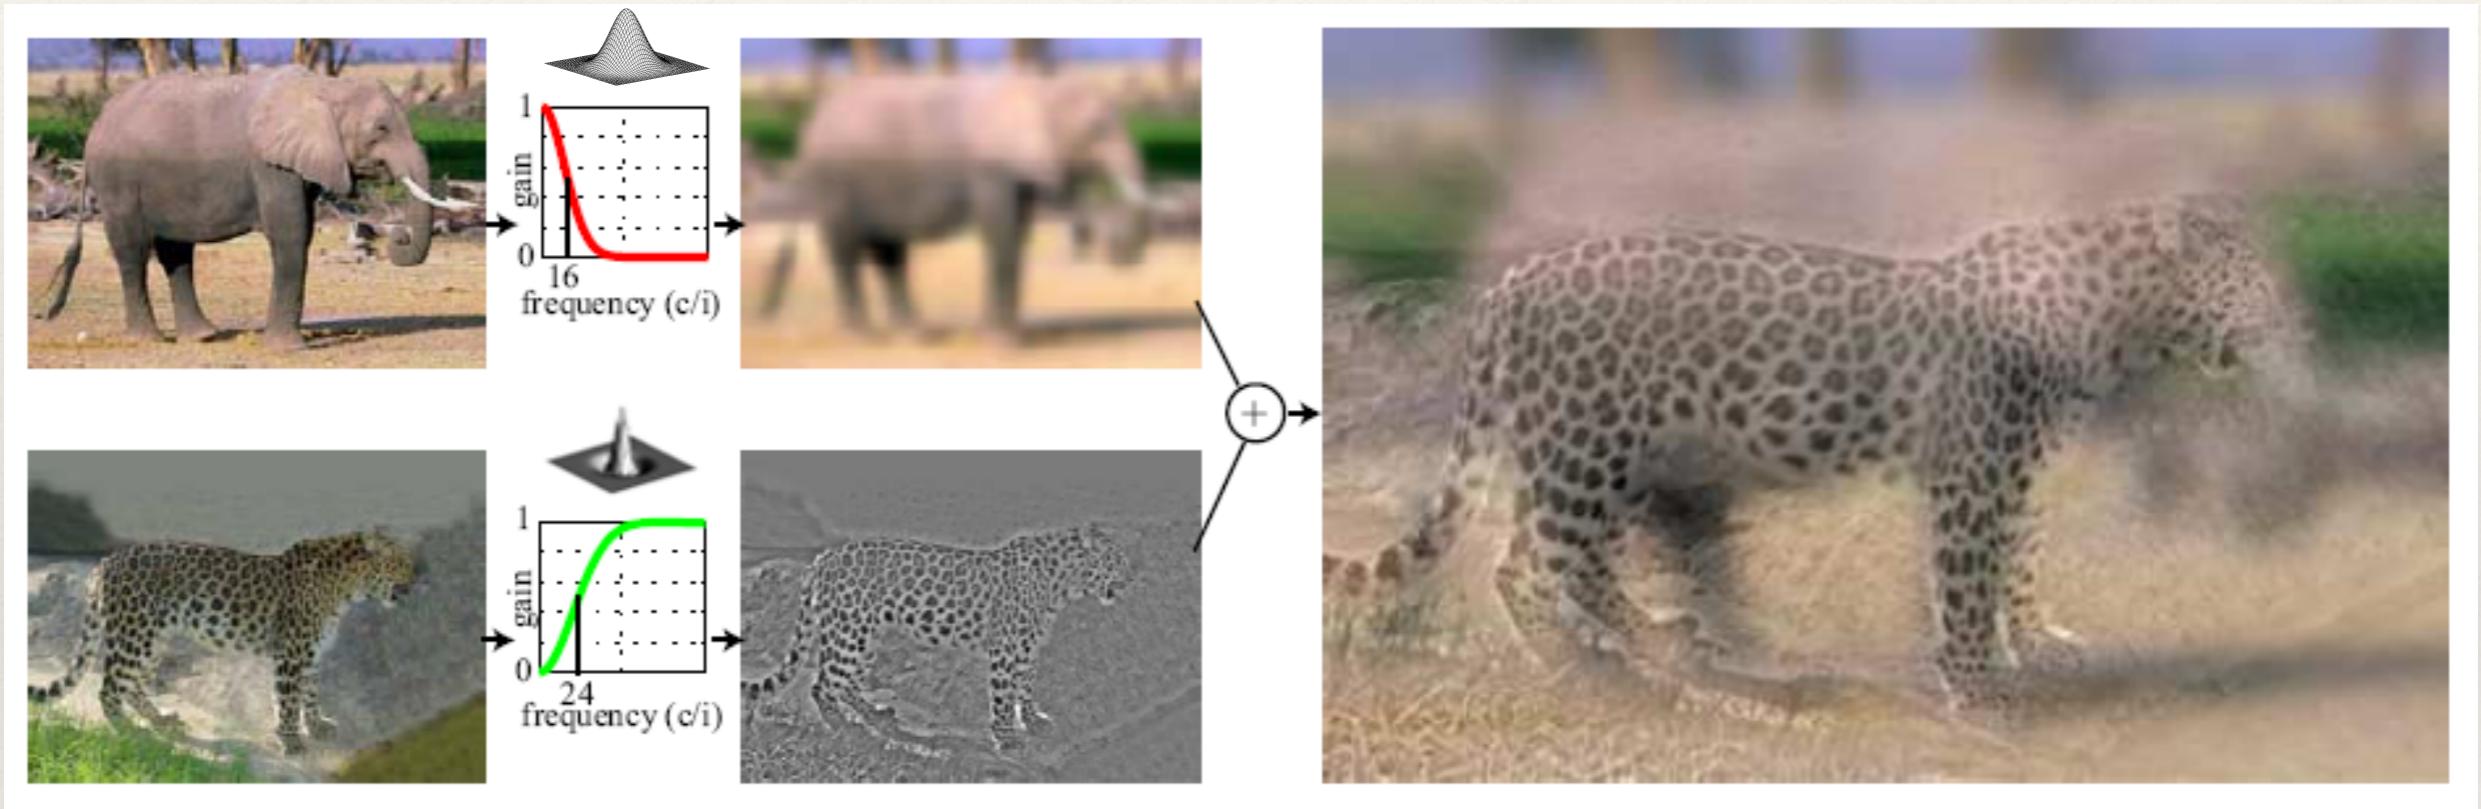
Salvador Dali (萨尔瓦多·达利), 1976

Hybrid Images



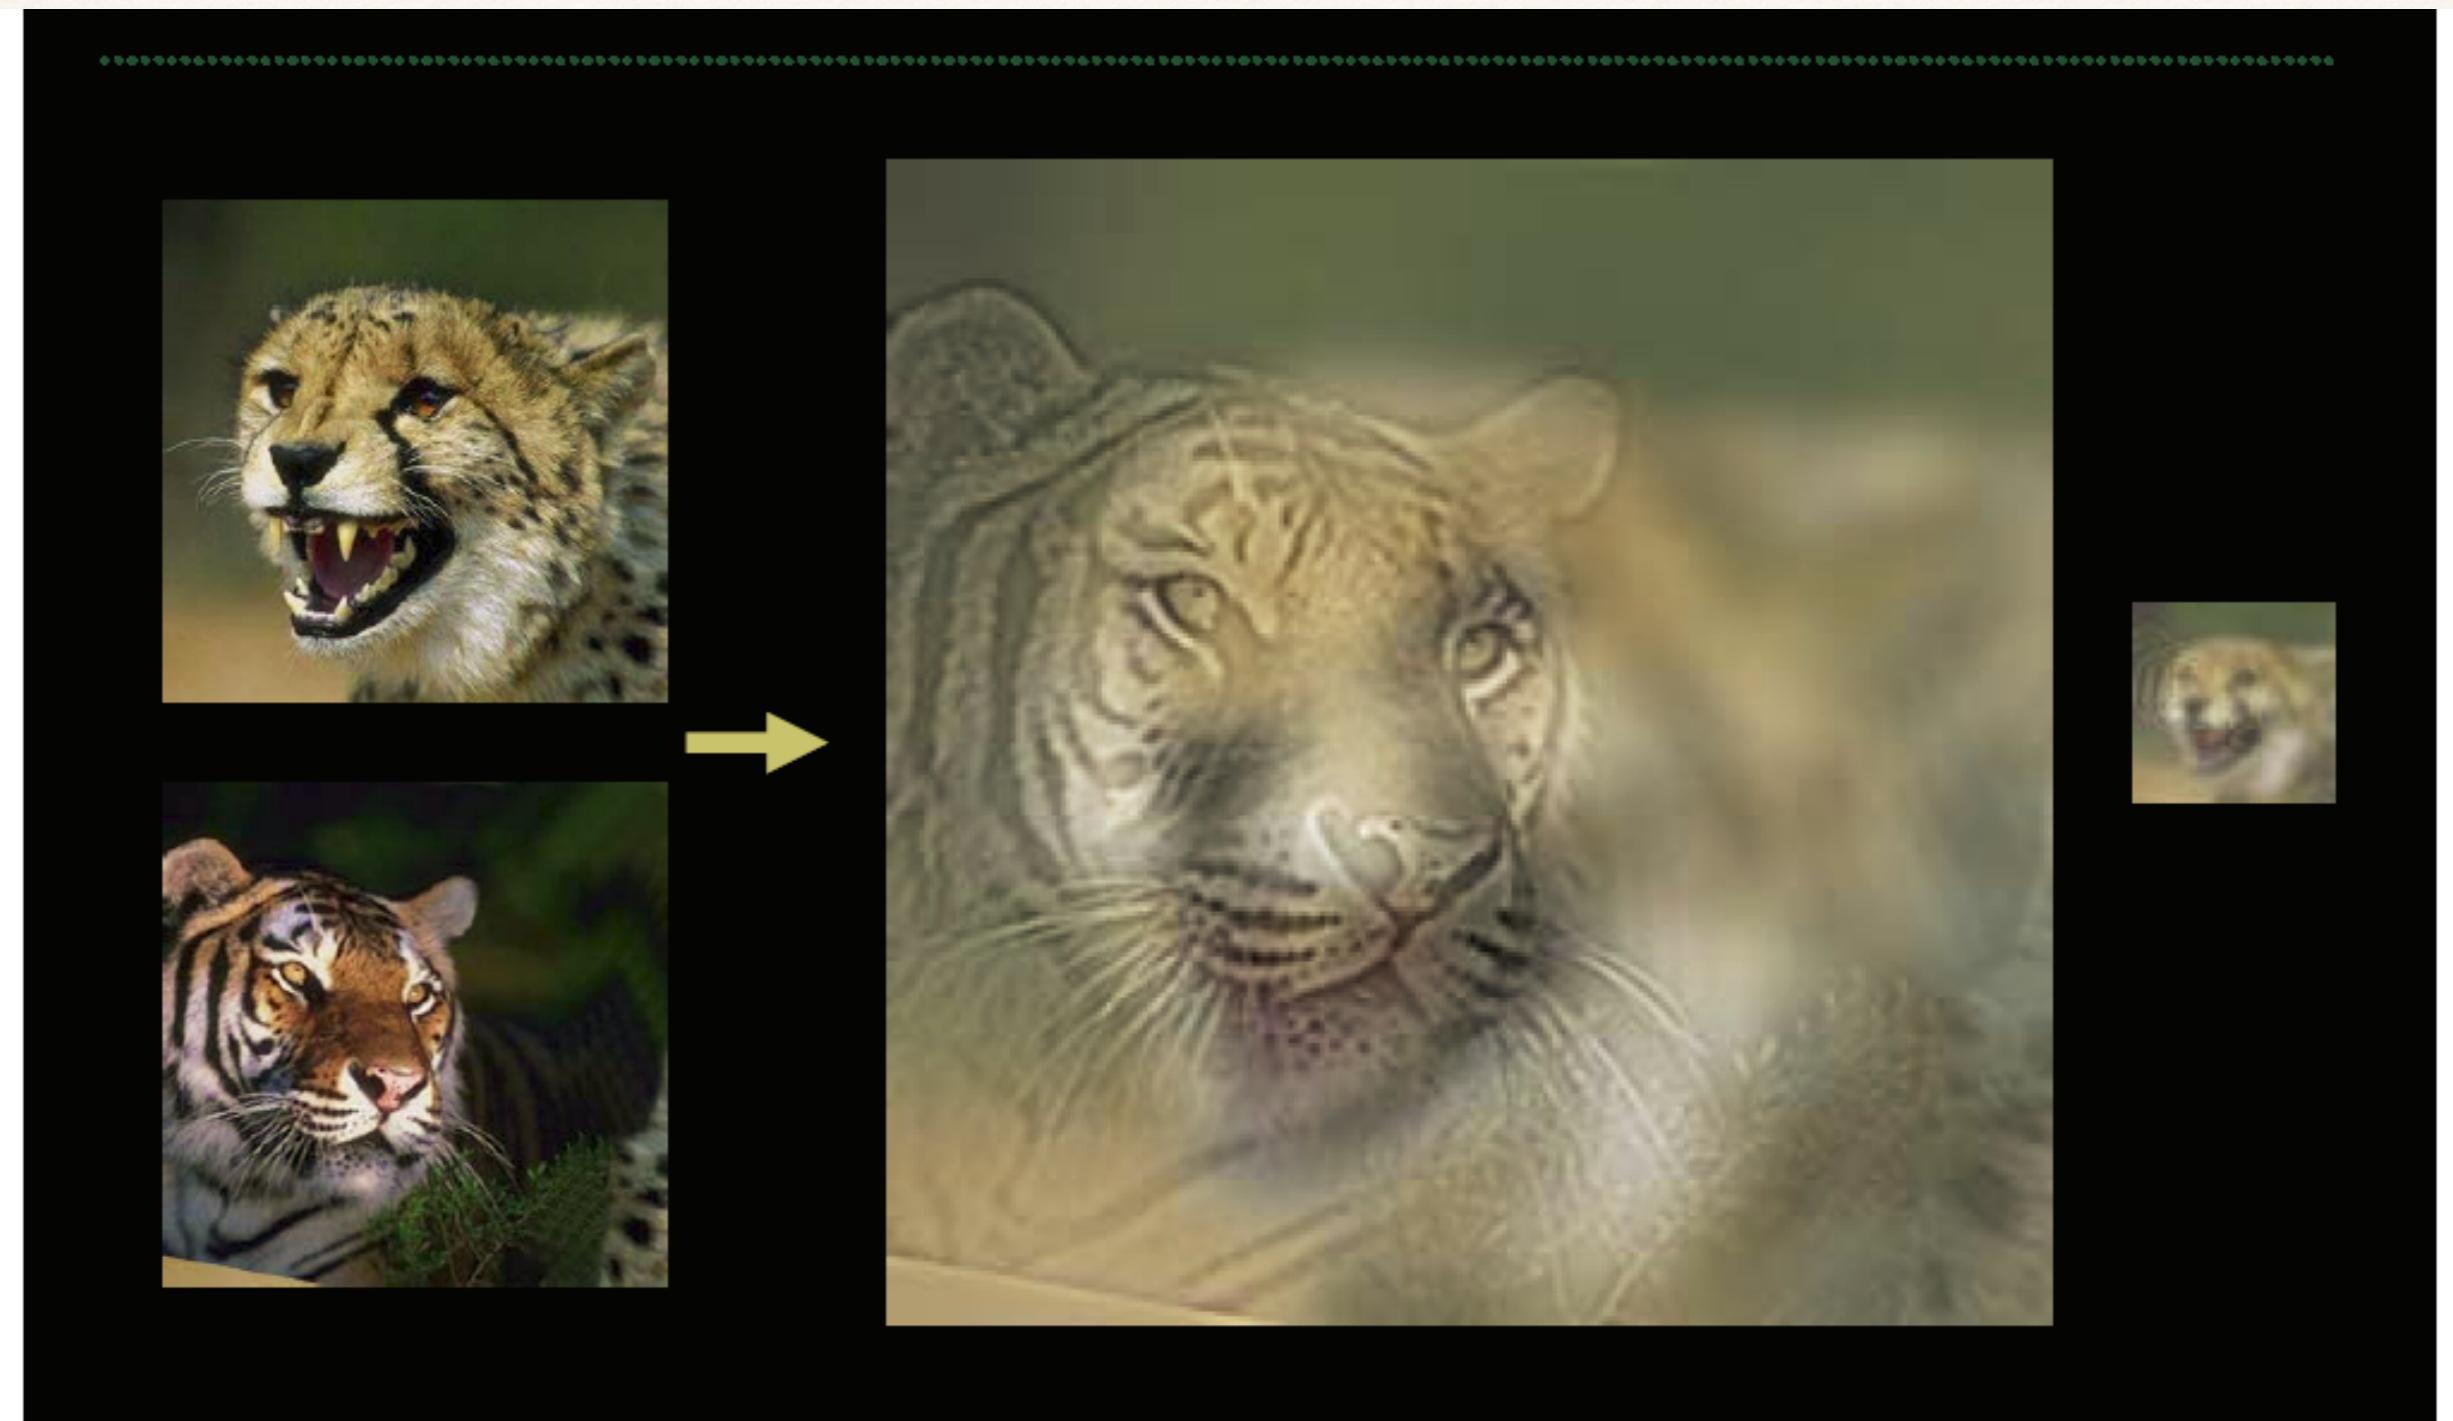
Aude Oliva & Antonio Torralba & Philippe G Schyns, SIGGRAPH 2006

Hybrid Images



Aude Oliva & Antonio Torralba & Philippe G Schyns, SIGGRAPH 2006

Hybrid Images



Aude Oliva & Antonio Torralba & Philippe G Schyns, SIGGRAPH 2006

Changing expression



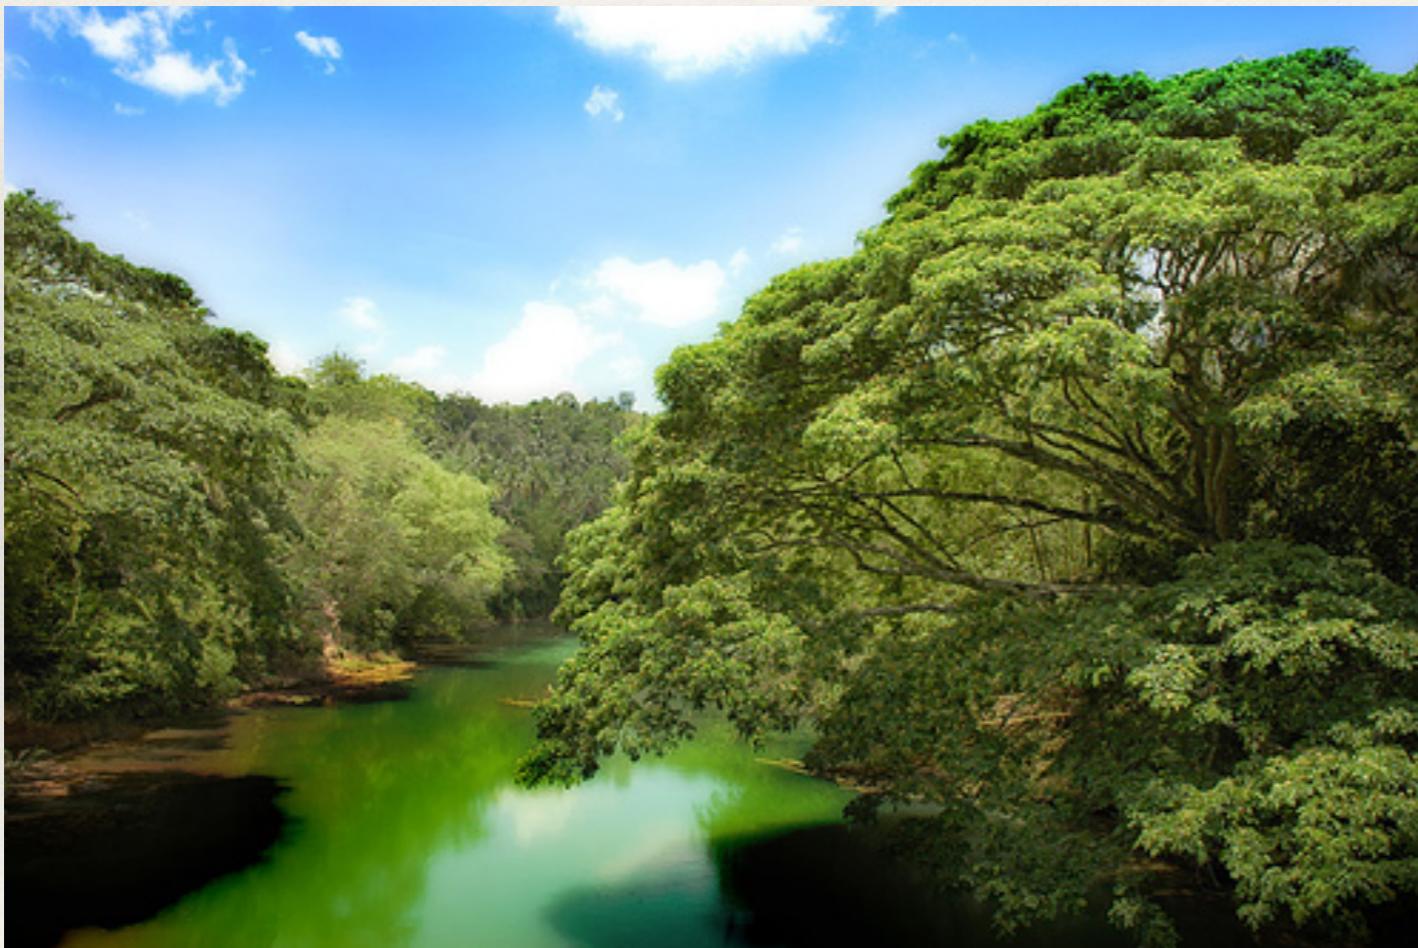
Sad

Surprised



Sampling

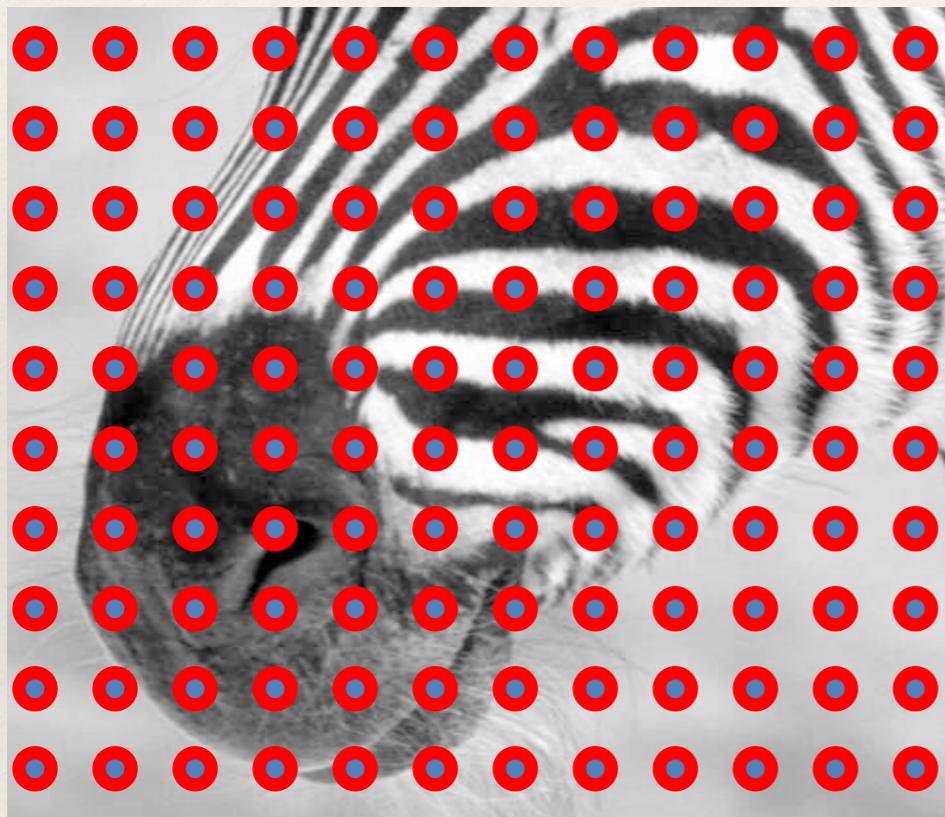
- ❖ Why does a lower resolution image still make sense to us?
What do we lose?



Source: James Tompkin

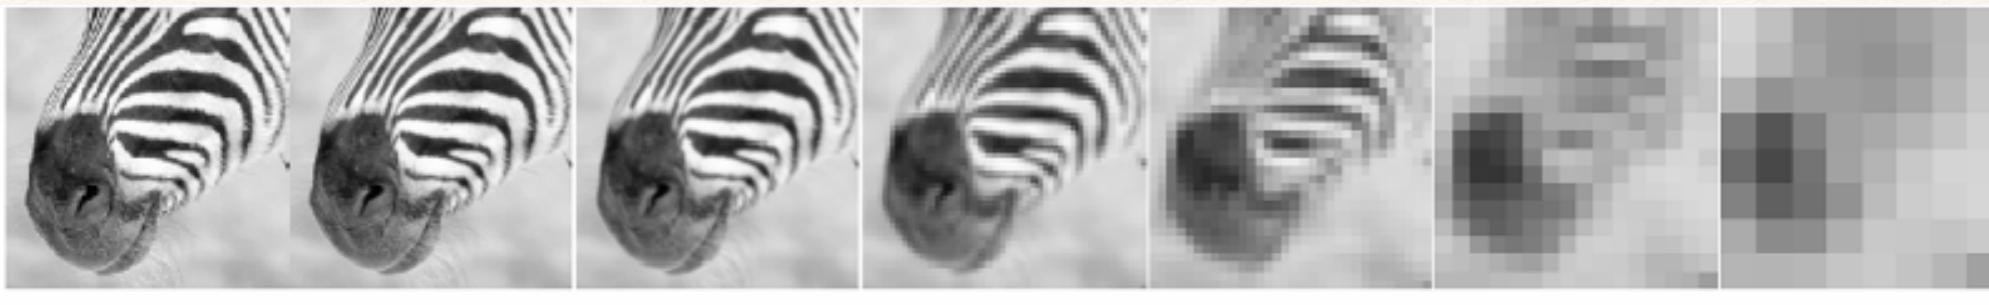
Subsampling

- ❖ Subsampling by a factor of 2



Throw away every other row
and column to create a 1/2 size
image

Source: James Tompkin



512

256

128

64

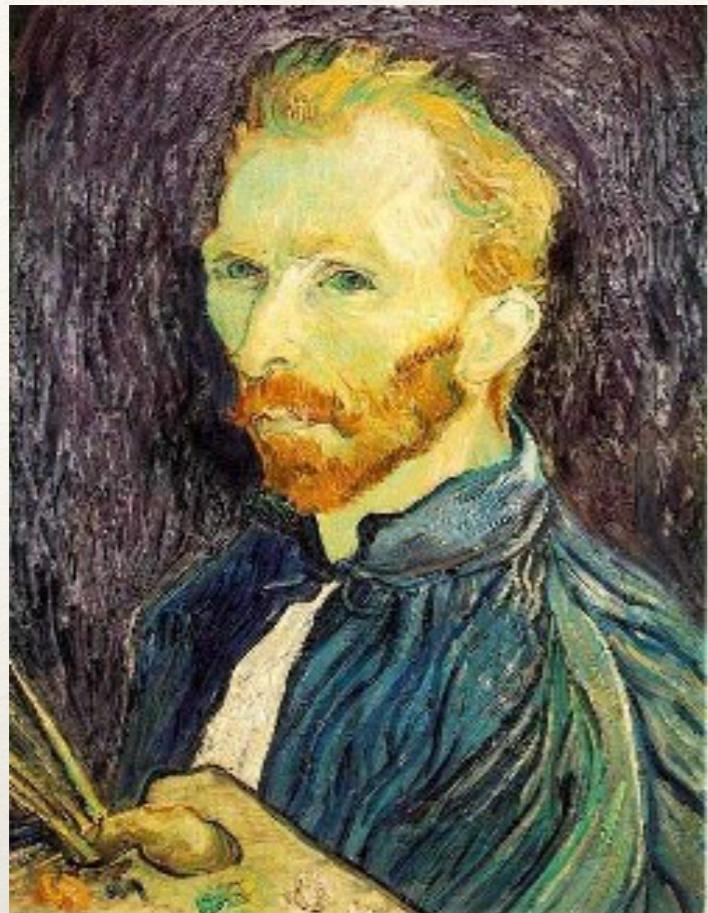
32

16

8



Figure from David Forsyth



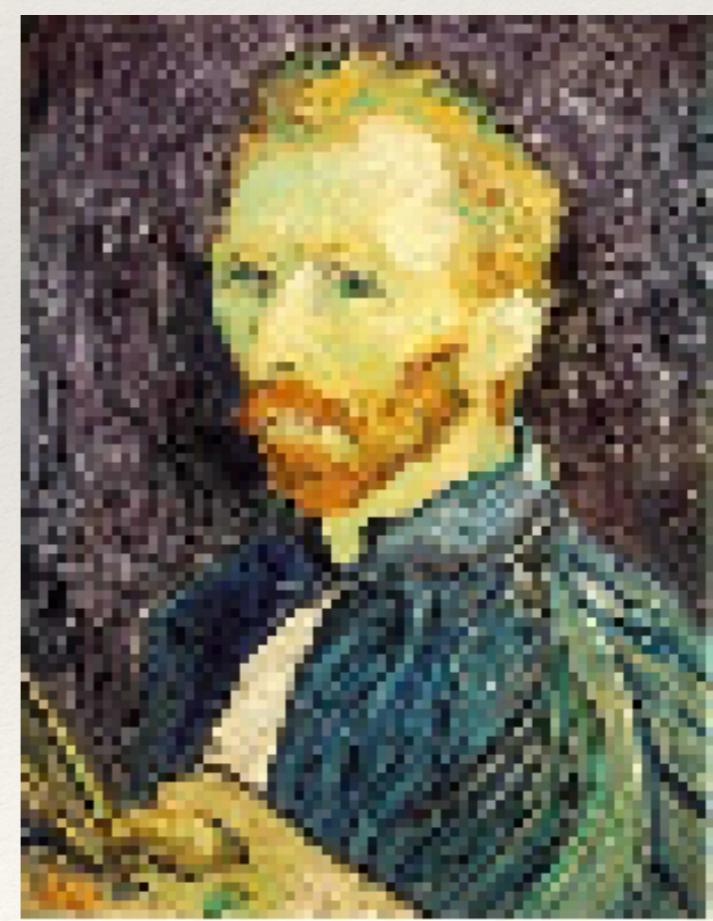
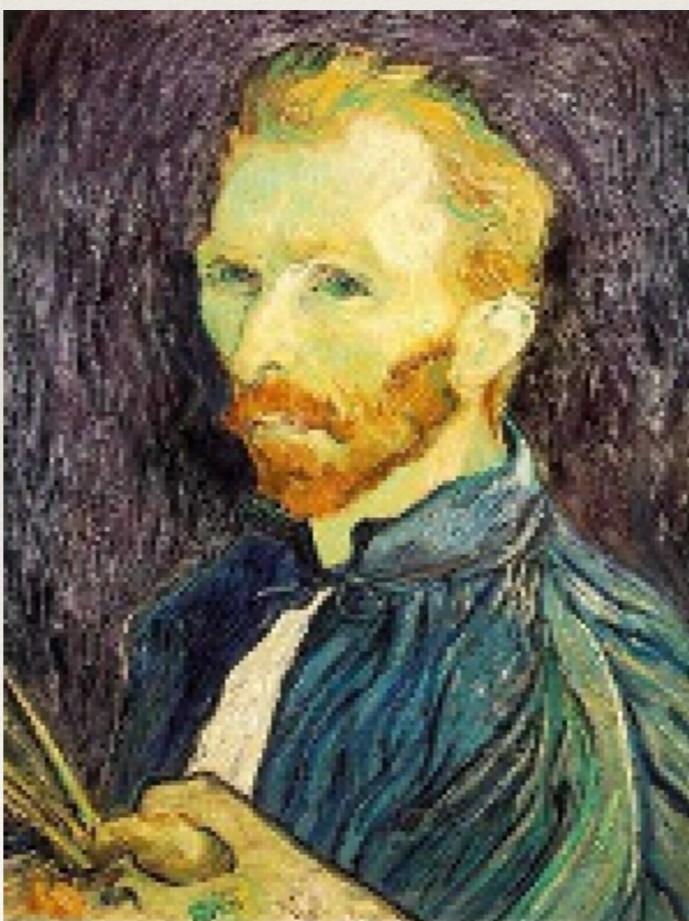
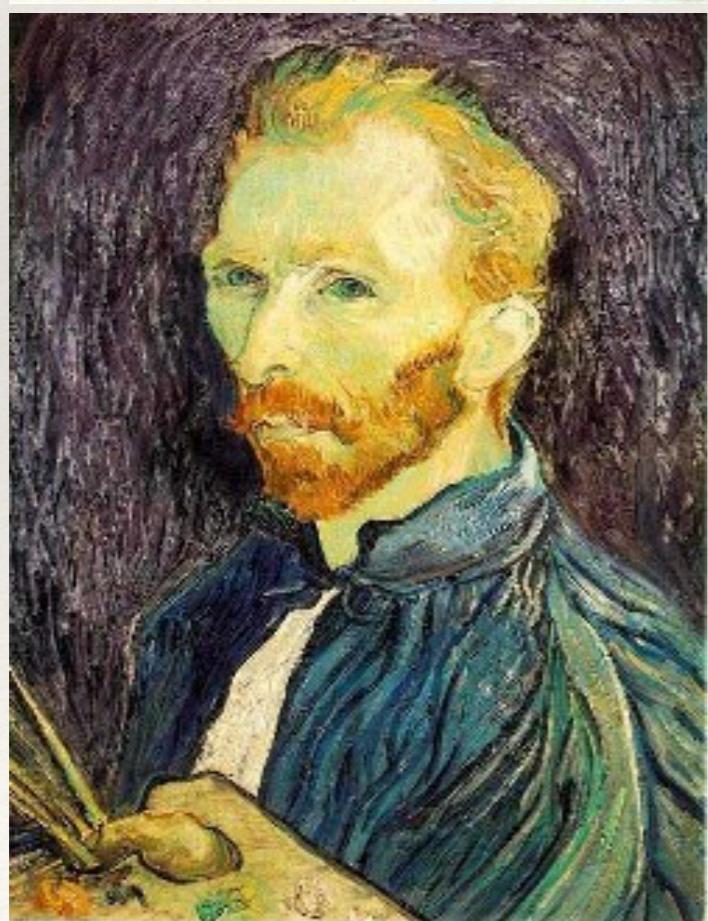
1/2



1/4 2x subsample



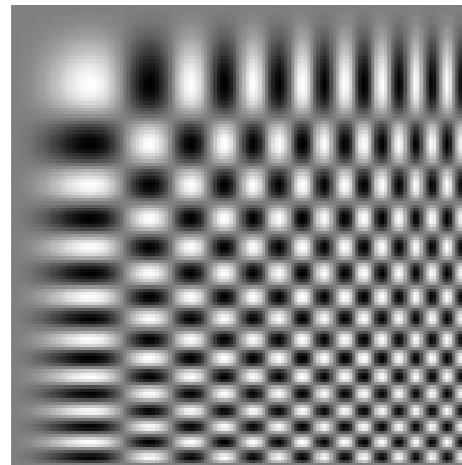
1/8 4x subsample



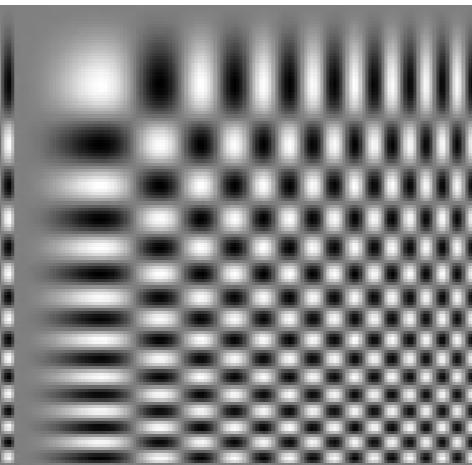
Steve Seitz

Sampling and aliasing

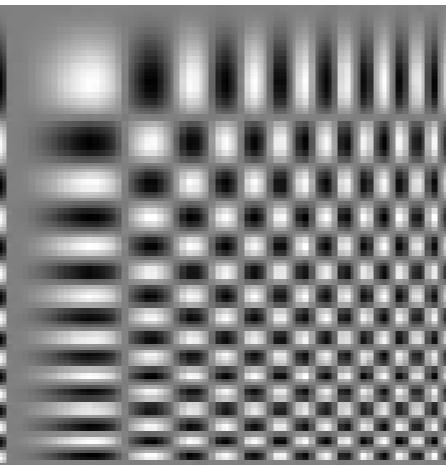
256x256



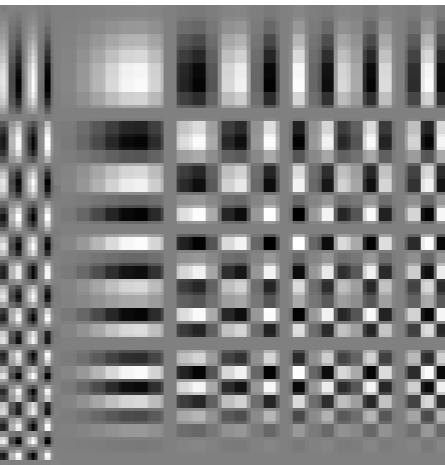
128x128



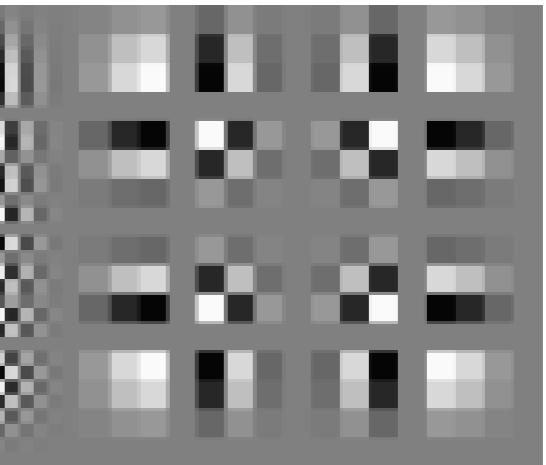
64x64



32x32

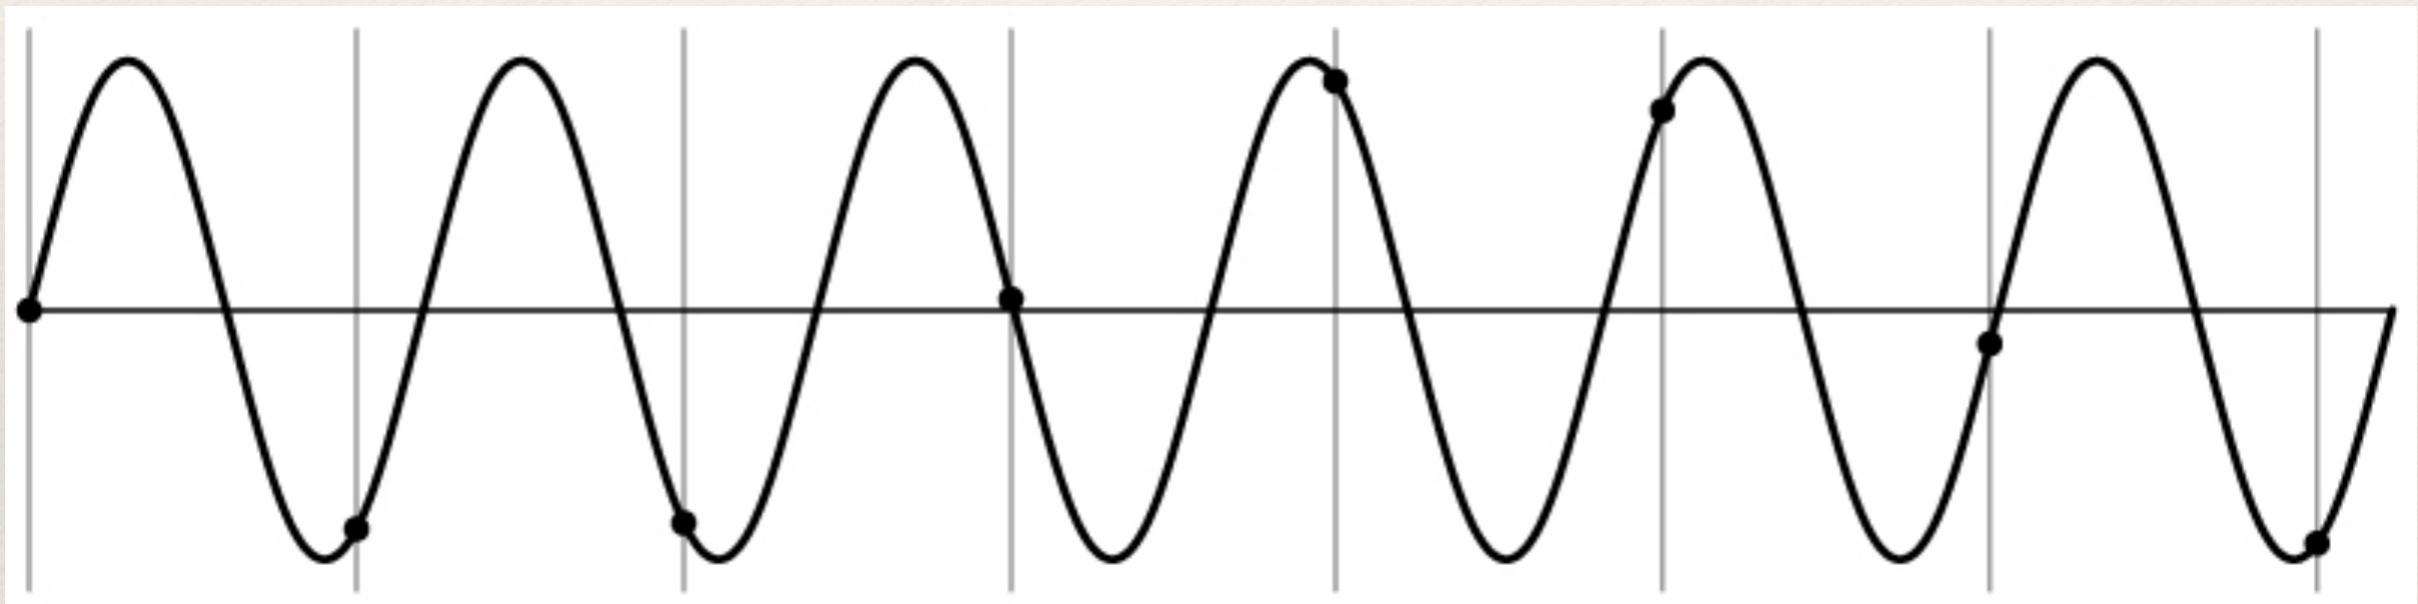


16x16



Sampling and aliasing

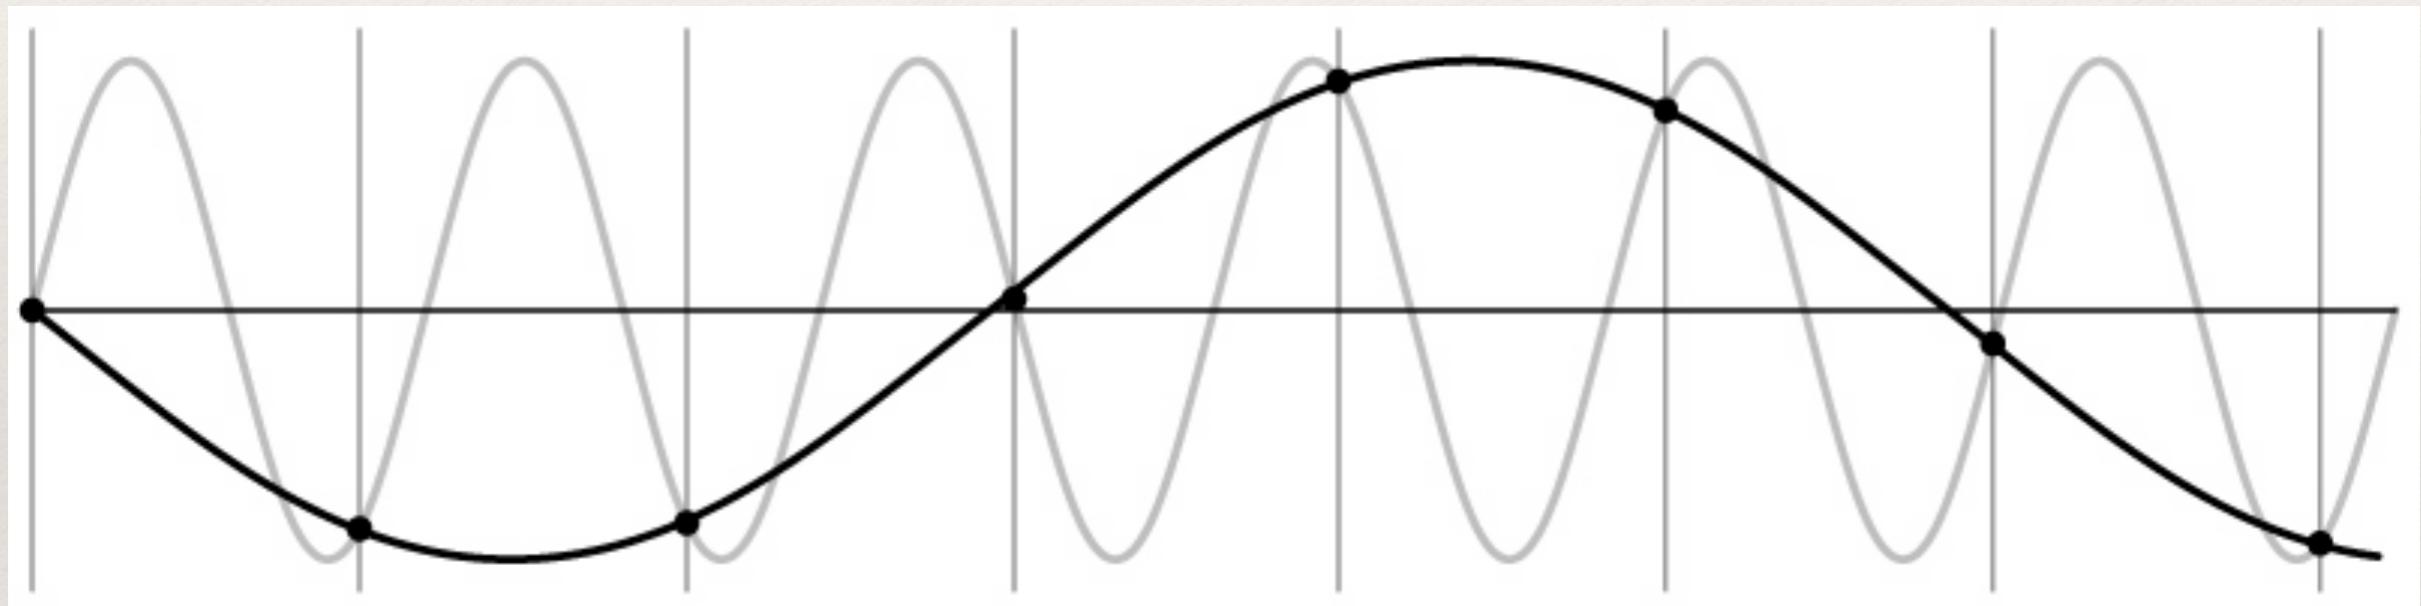
1D example (sinewave):



Source: S. Marschner

Sampling and aliasing

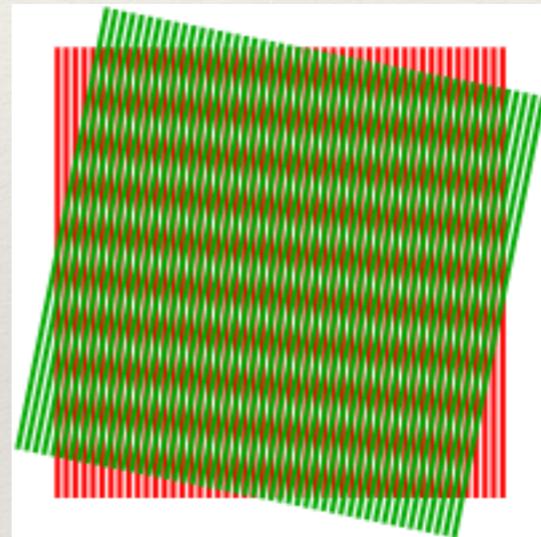
1D example (sinewave):



Source: S. Marschner

Aliasing problem

- ❖ Sub-sampling may be dangerous....
- ❖ Characteristic errors may appear:
 - ❖ “car wheels rolling the wrong way in movies”
 - ❖ “checkerboards disintegrate in ray tracing”
 - ❖ “striped shirts look funny on color television”
 - ❖ Moiré patterns



Source: James Tompkin

Aliasing and Moiré patterns

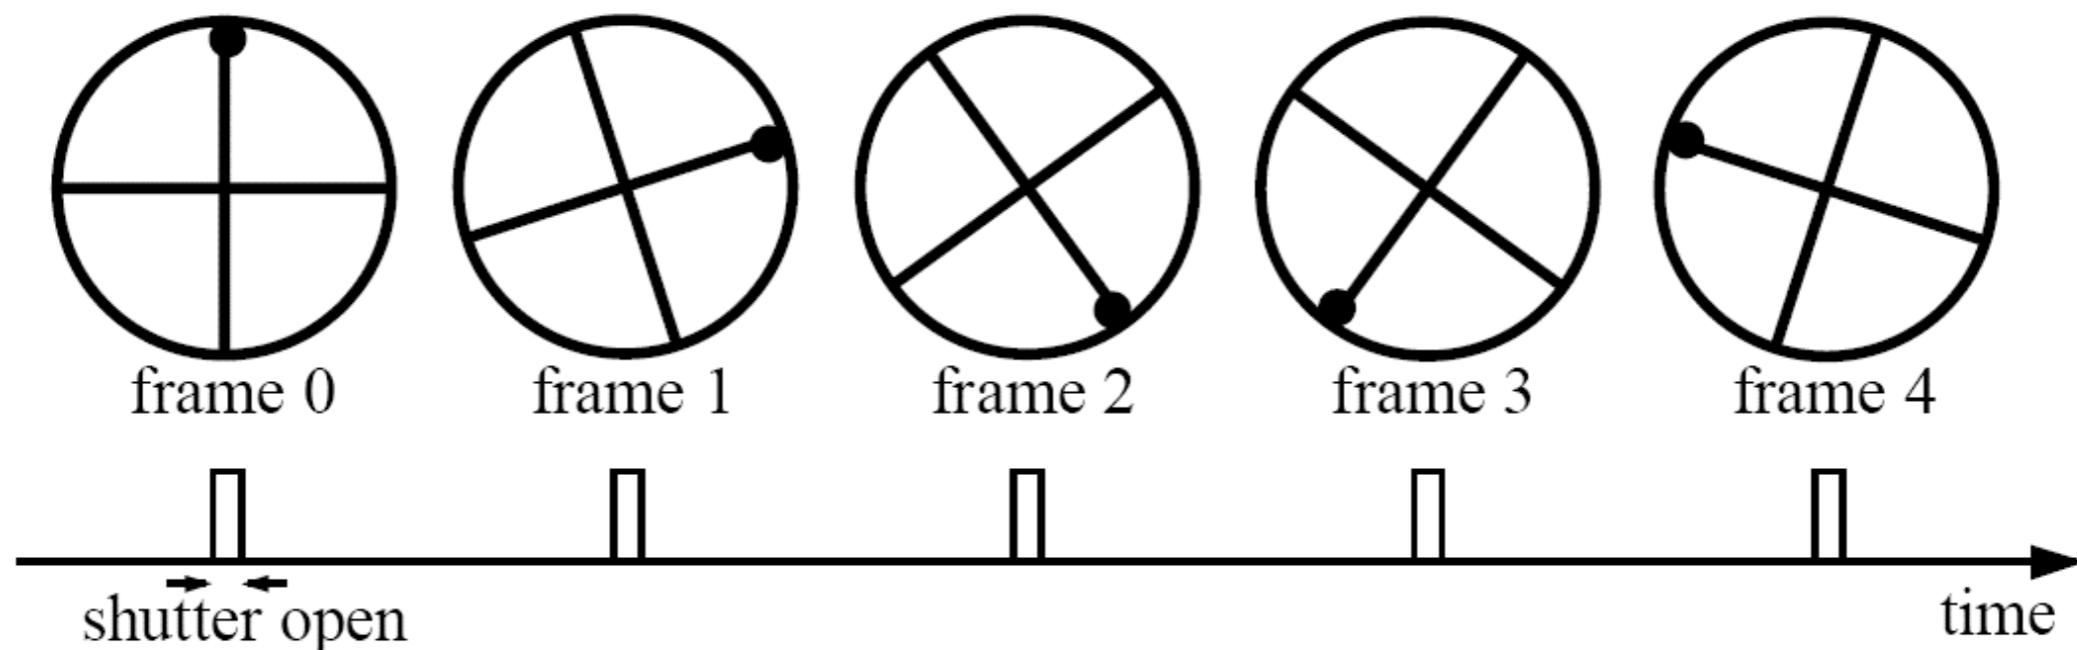
- ❖ Moire pattern, known as the moire effect
- ❖ Useful in metrology, strain analysis or even document authentication and anti-counterfeiting. In other situations moire patterns may be unwanted.
- ❖ This visual perception occurs when a fine pattern on your subject meshes with the pattern on the imaging chip of your camera. Difference in frequency.



Aliasing in video

Imagine a spoked wheel moving to the right (rotating clockwise).
Mark wheel with dot so we can see what's happening.

If camera shutter is only open for a fraction of a frame time (frame time = 1/30 sec. for video, 1/24 sec. for film):

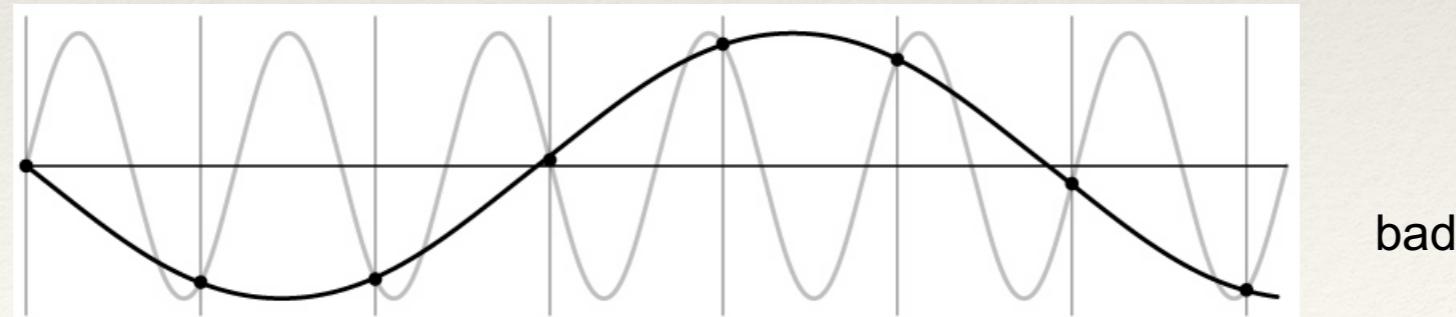
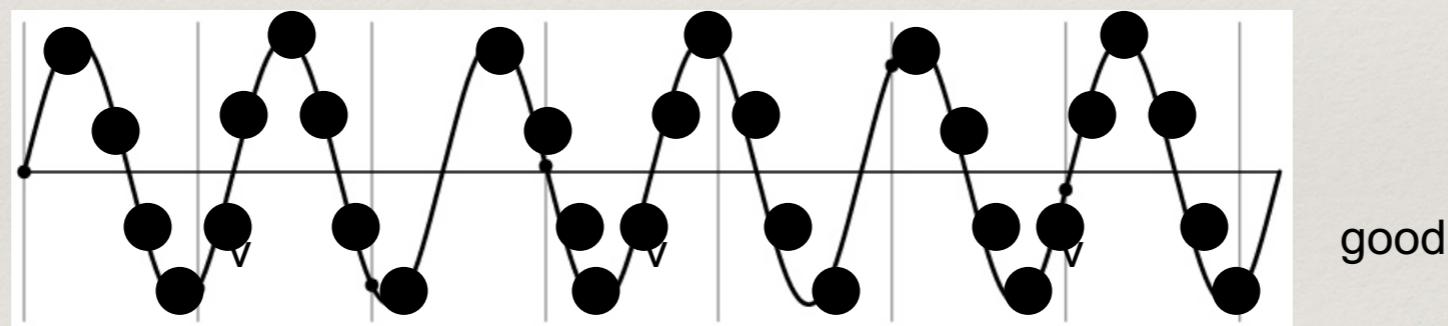


Without dot, wheel appears to be rotating slowly backwards!
(counterclockwise)



Nyquist-Shannon Sampling Theorem

- ❖ When sampling a signal at discrete intervals, the sampling frequency must be $\geq 2 \times f_{max}$
- ❖ f_{max} = max frequency of the input signal
- ❖ This will allow to reconstruct the original perfectly from the sampled version

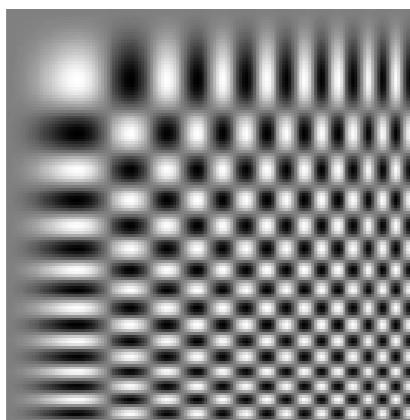


How to fix aliasing?

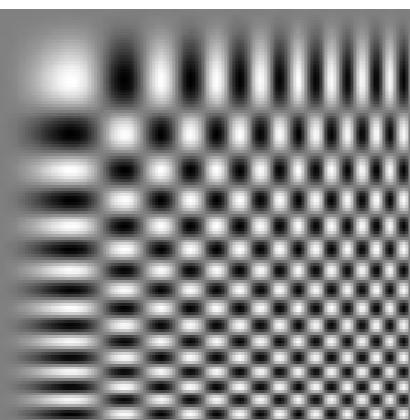
- ❖ Sample more often (better sensor)
- ❖ Get rid of all frequencies that are greater than half the new sampling frequency
 - ❖ Will lose information
 - ❖ But it's better than aliasing
 - ❖ Apply a smoothing (*low pass*) filter

Anti-aliasing

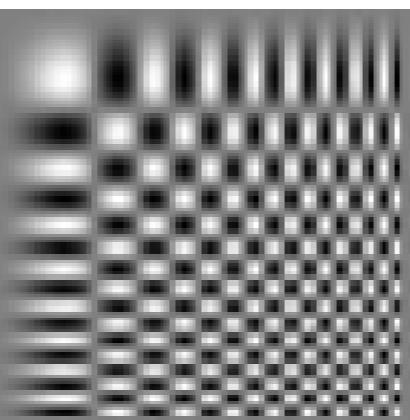
256x256



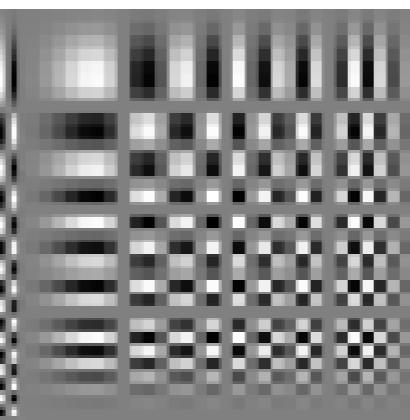
128x128



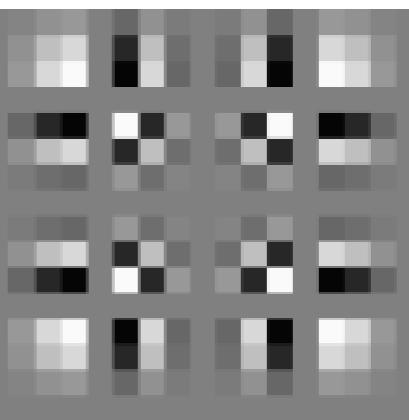
64x64



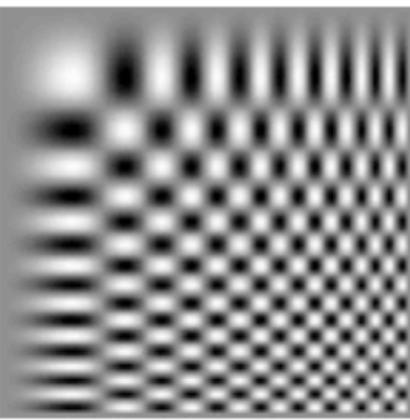
32x32



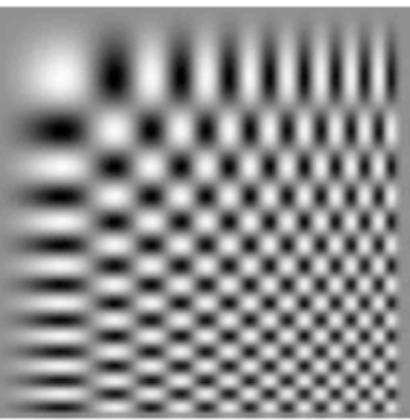
16x16



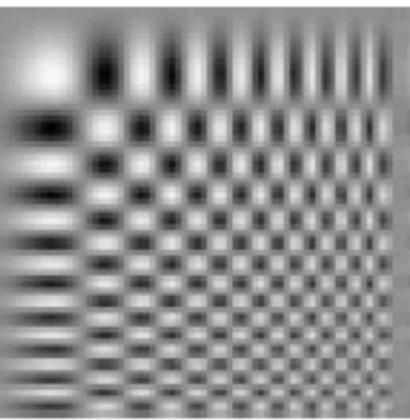
256x256



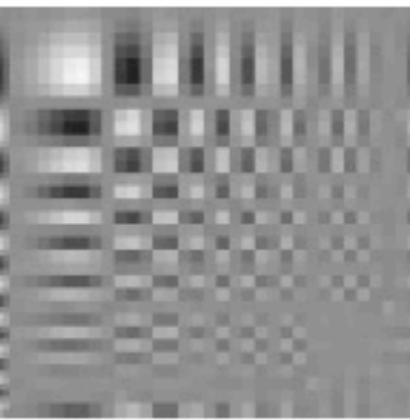
128x128



64x64



32x32



16x16

Algorithm for downsampling by factor of 2

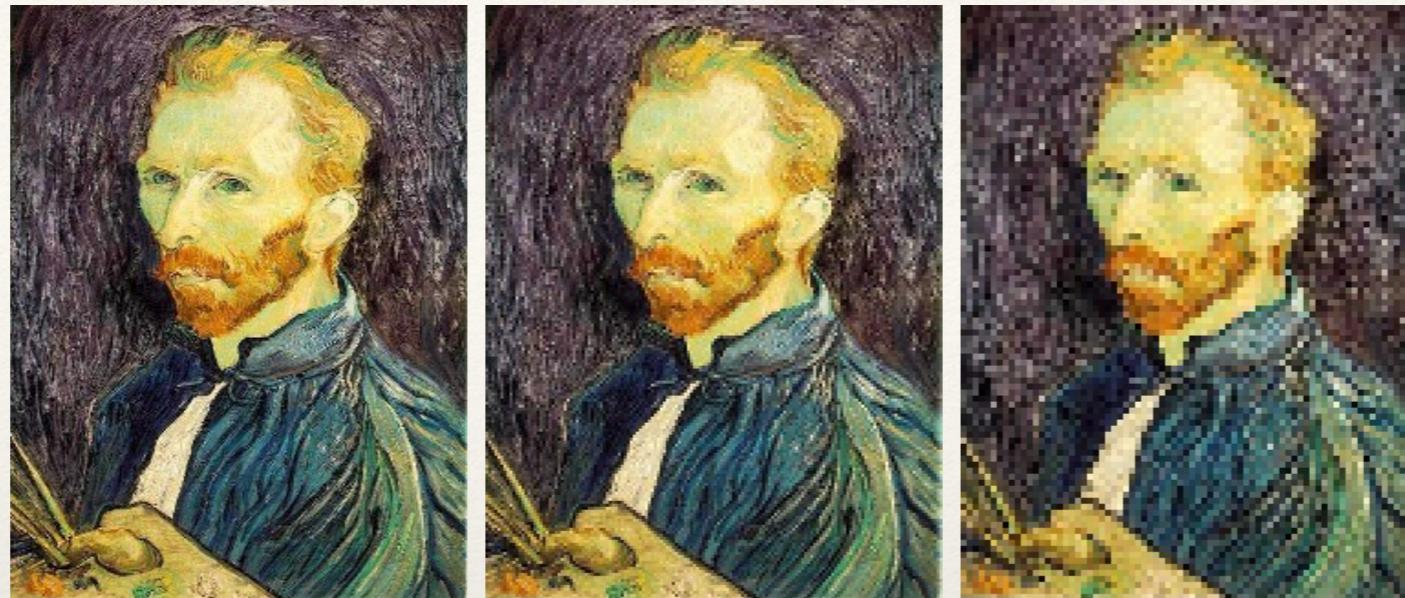
- ❖ Start with $\text{image}(h, w)$

- ❖ Apply low-pass filter:

$$\text{im_blur} = \text{imfilter}(\text{image}, \text{fspecial('gaussian', 7, 1)})$$

- ❖ Sample every other pixel

$$\text{im_small} = \text{im_blur}(1:2:\text{end}, 1:2:\text{end})$$



1/2

1/4 (2x subsample)

1/8 (4x subsample)



Gaussian 1/2

G 1/4

G 1/8

Subsampling with Gaussian pre-filtering

Human spatial frequency sensitivity

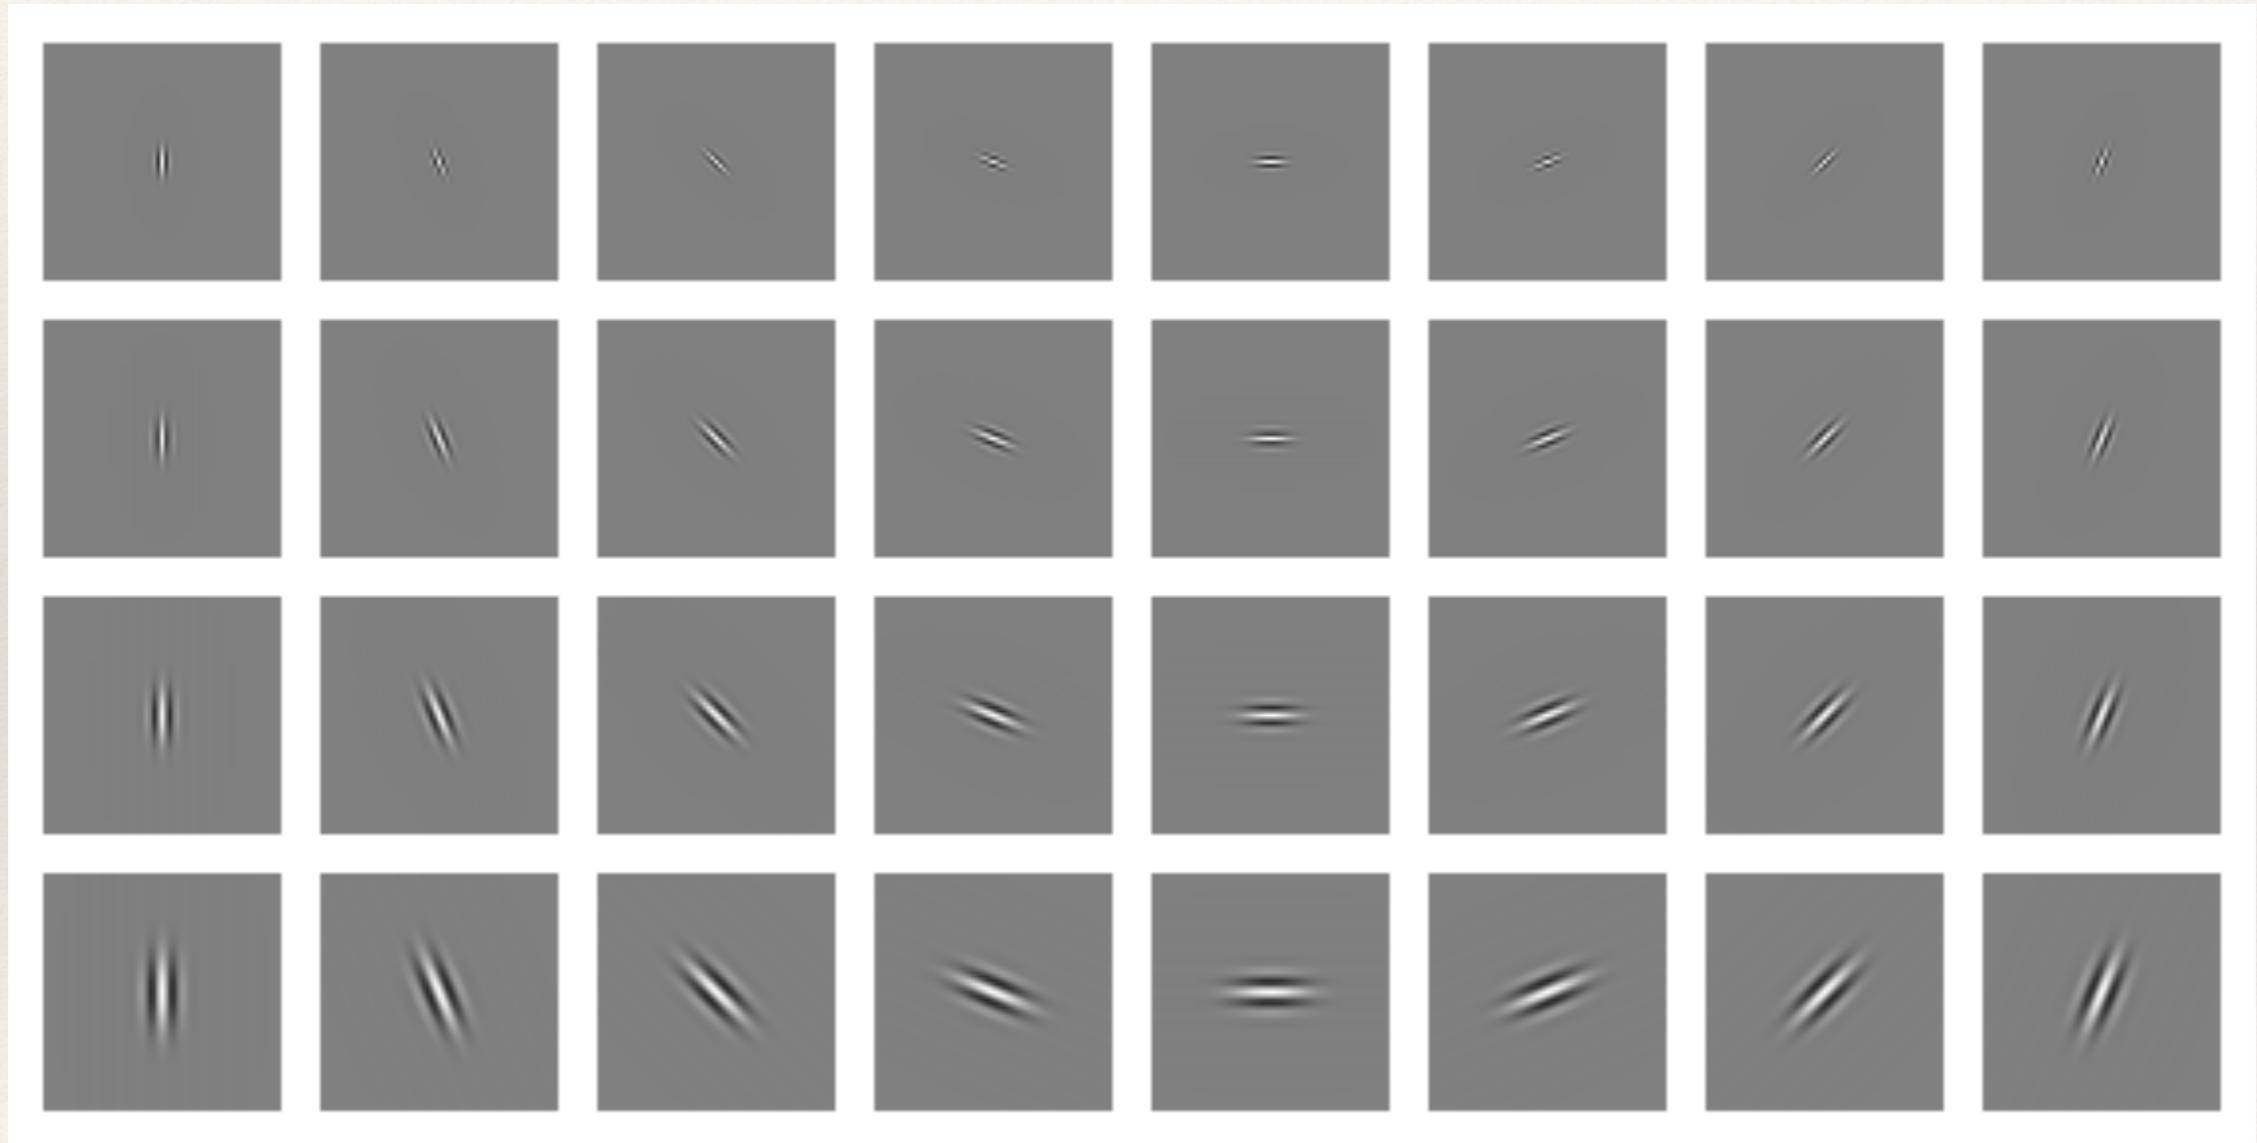
Why do we get different, distance-dependent interpretations of hybrid images?



Hays

Clues from Human Perception

Early processing in humans filters for orientations and scales of frequency.

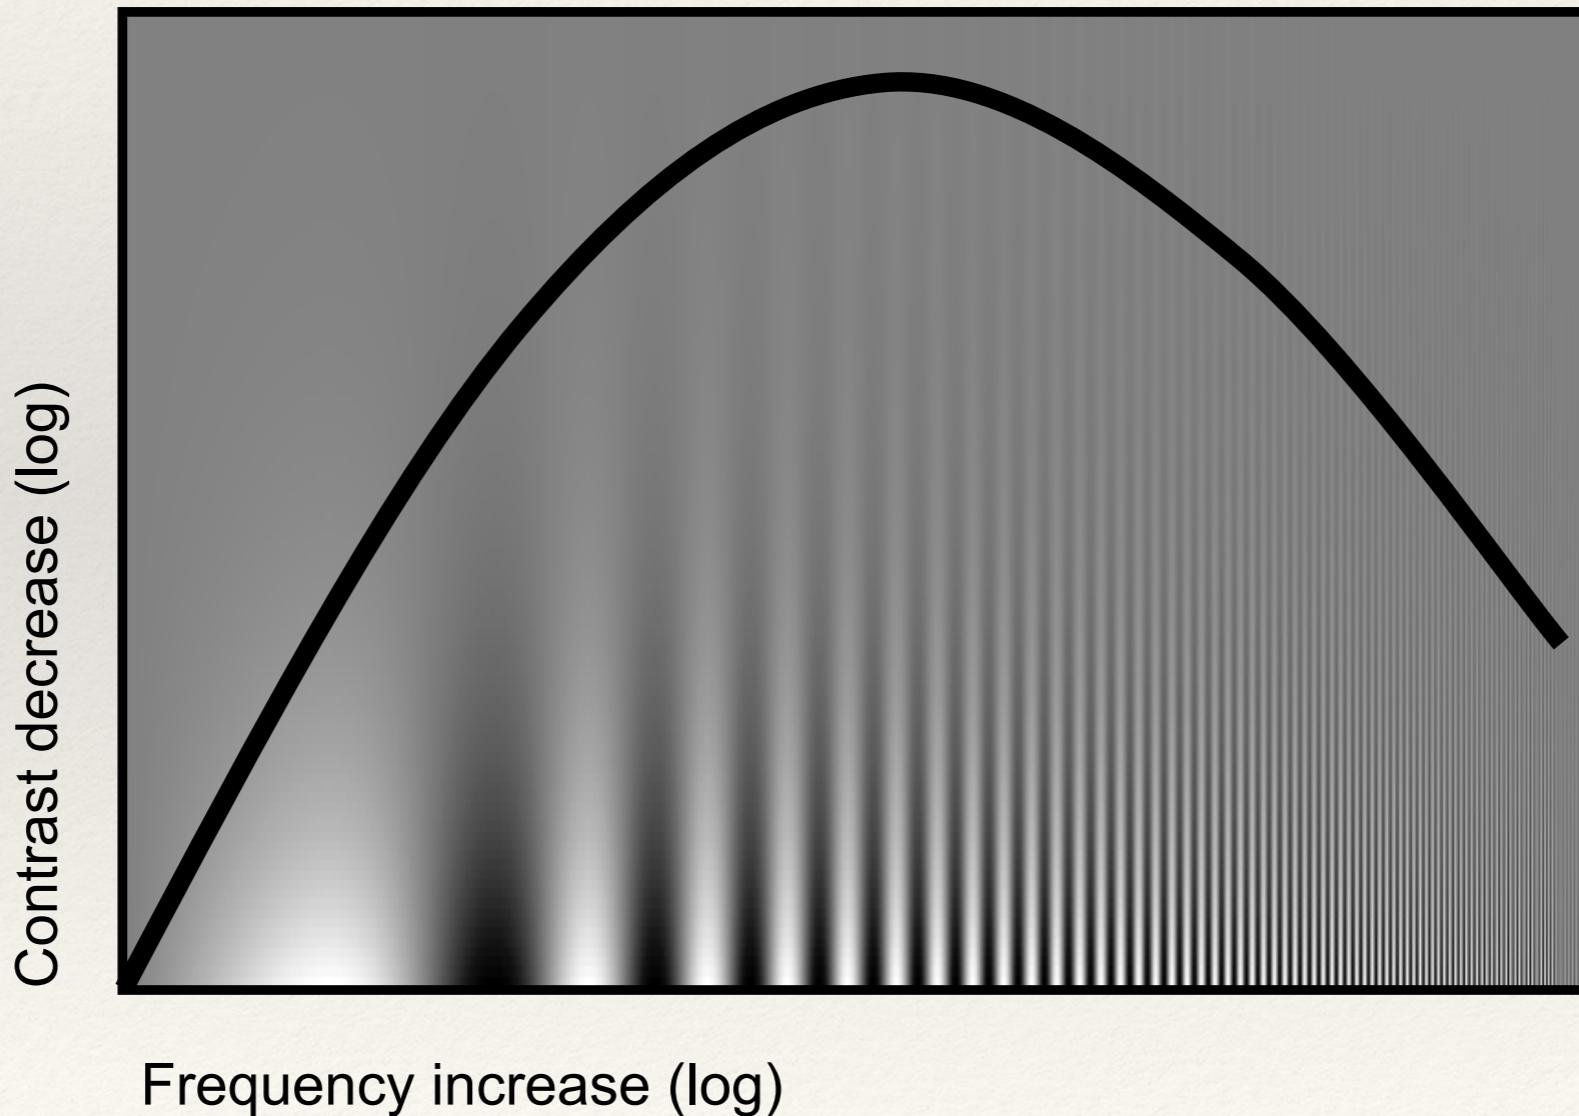


Early Visual Processing: Multi-scale edge and blob filters

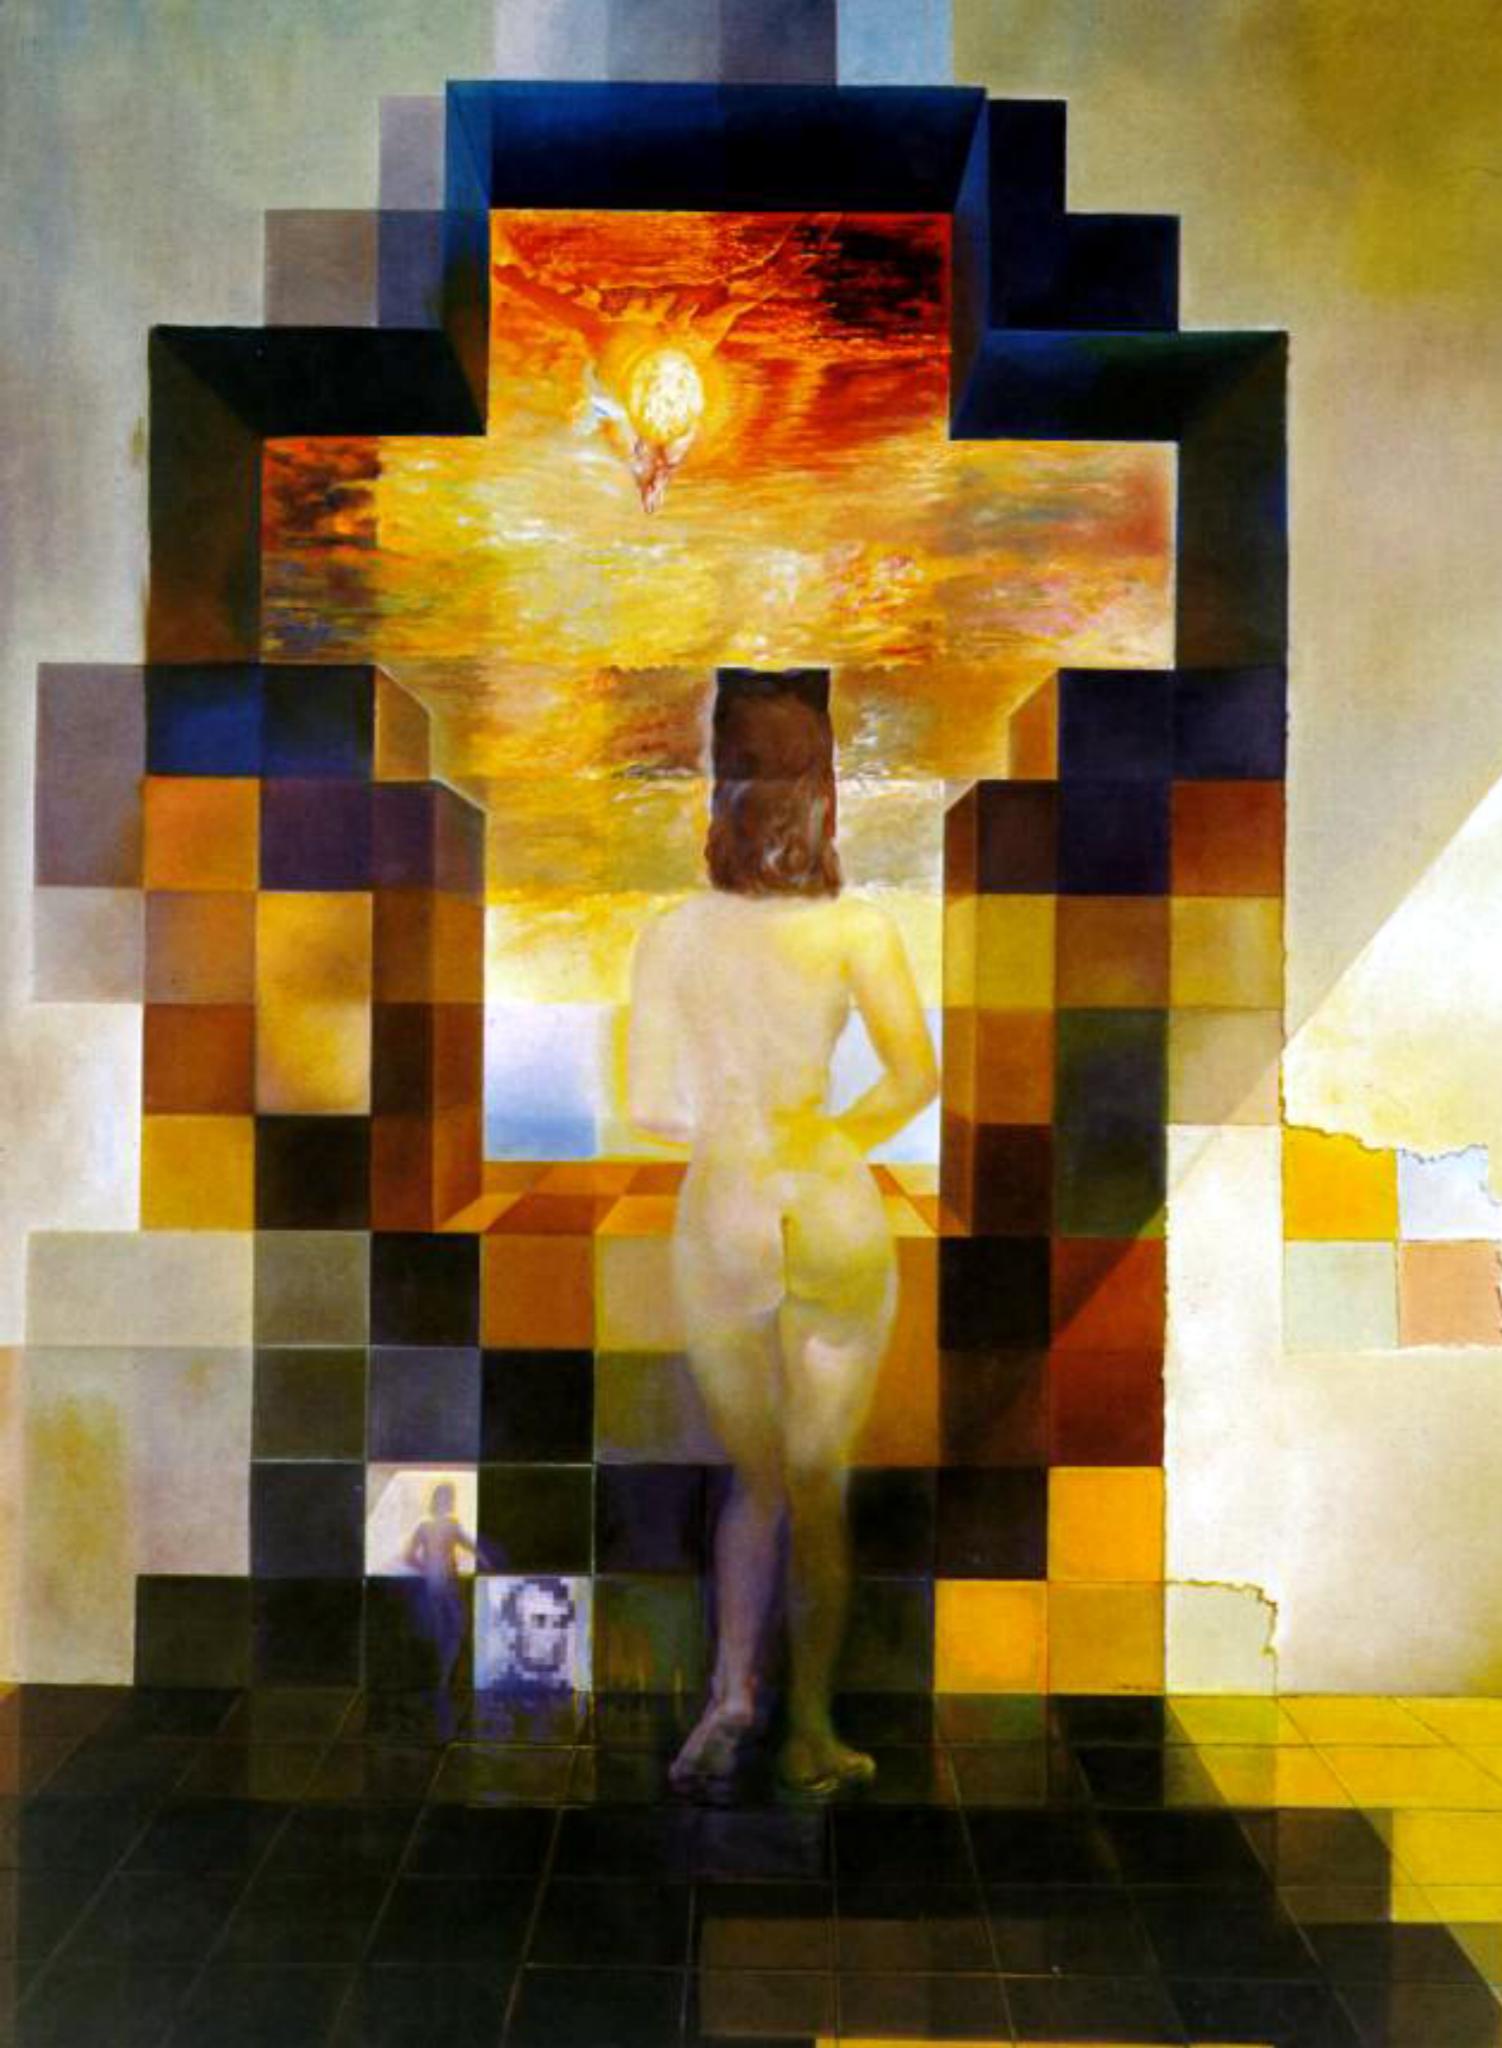
Source: James Tompkin

Campbell-Robson contrast sensitivity curve.

Perceptual cues in the mid-high frequencies dominate perception.

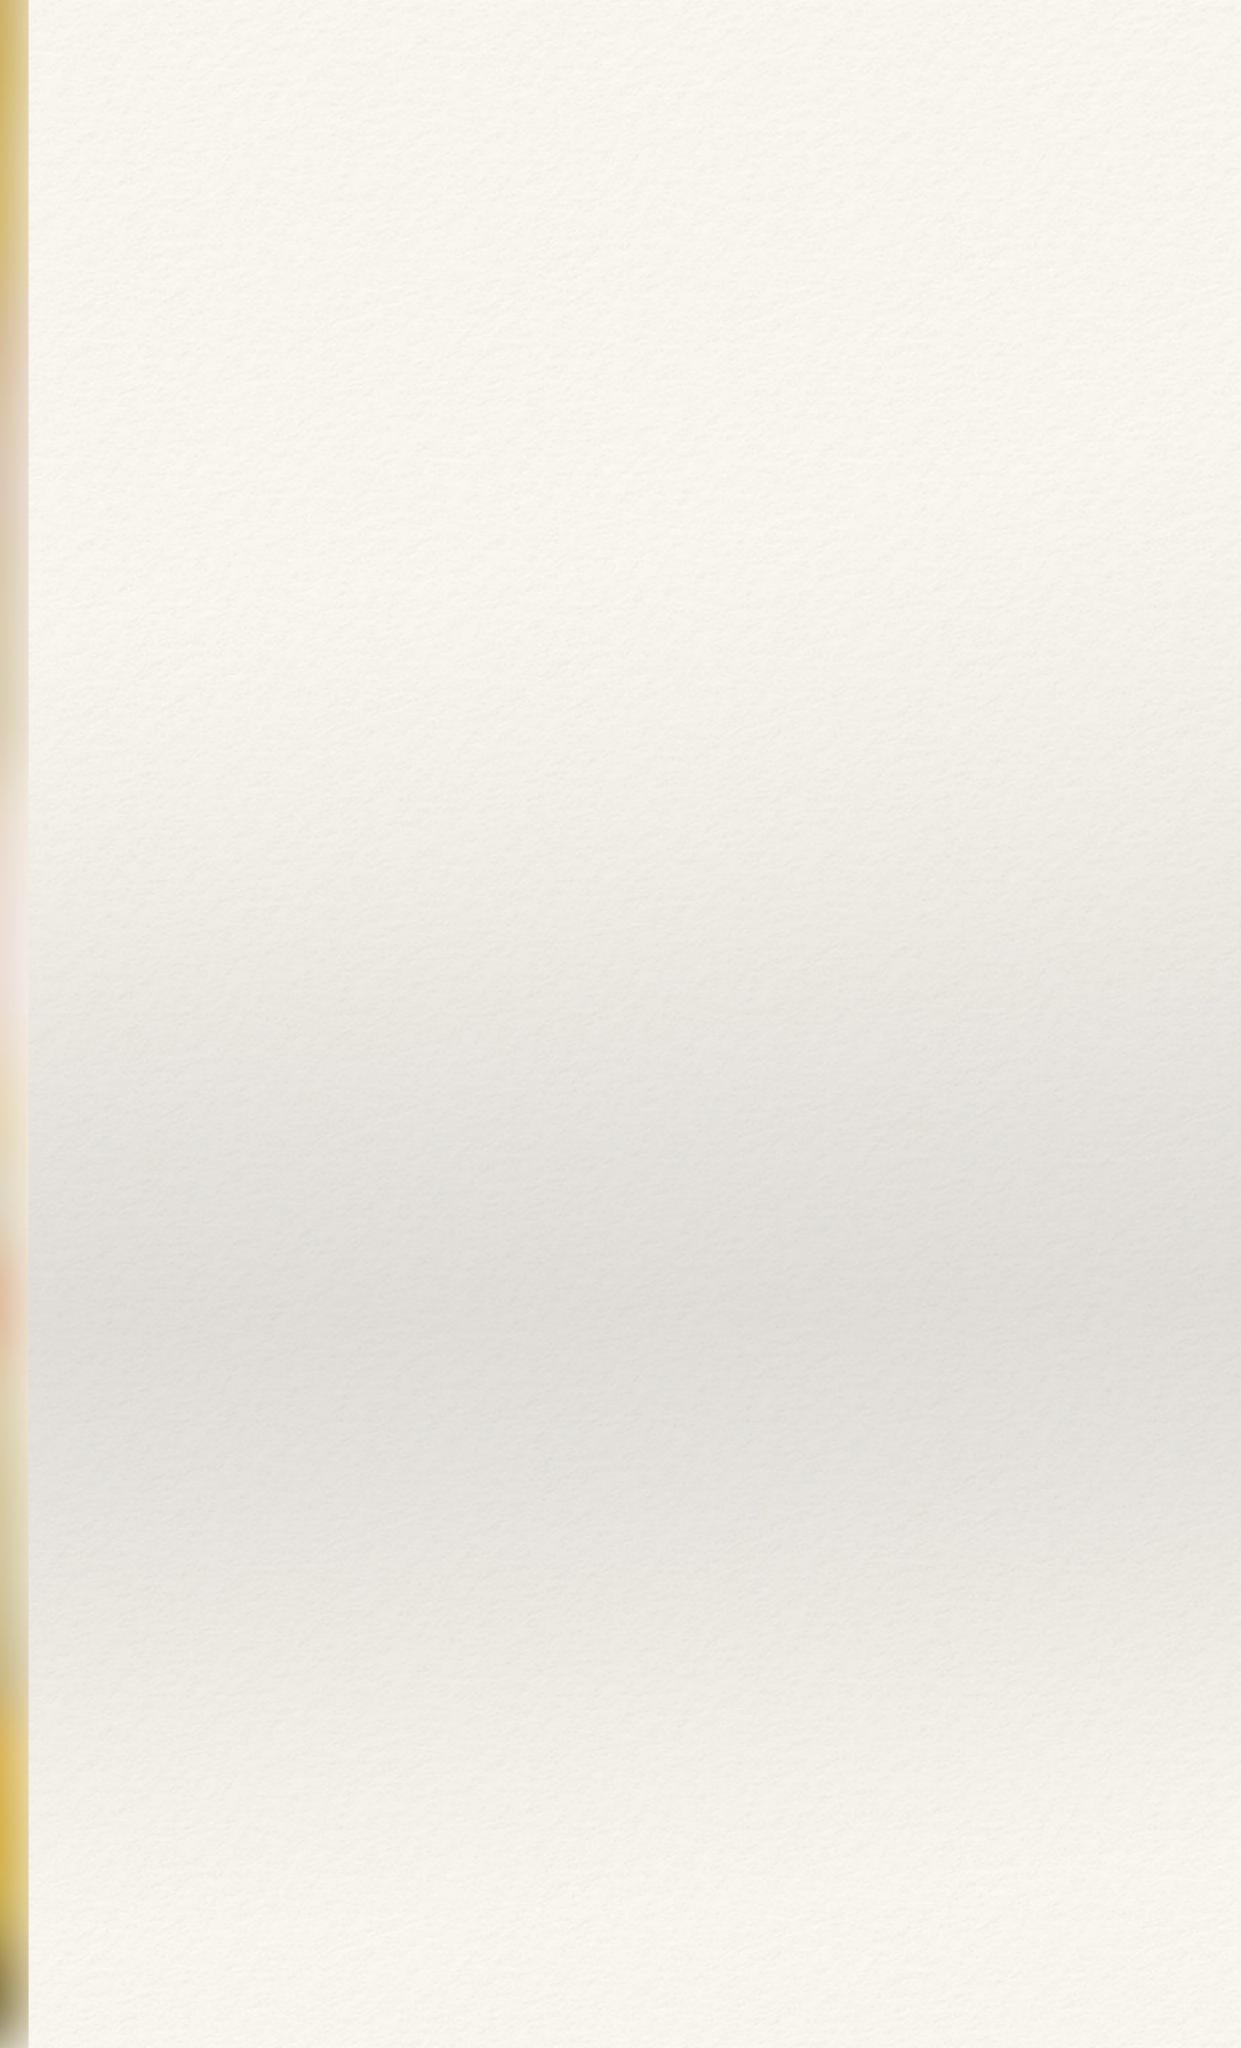
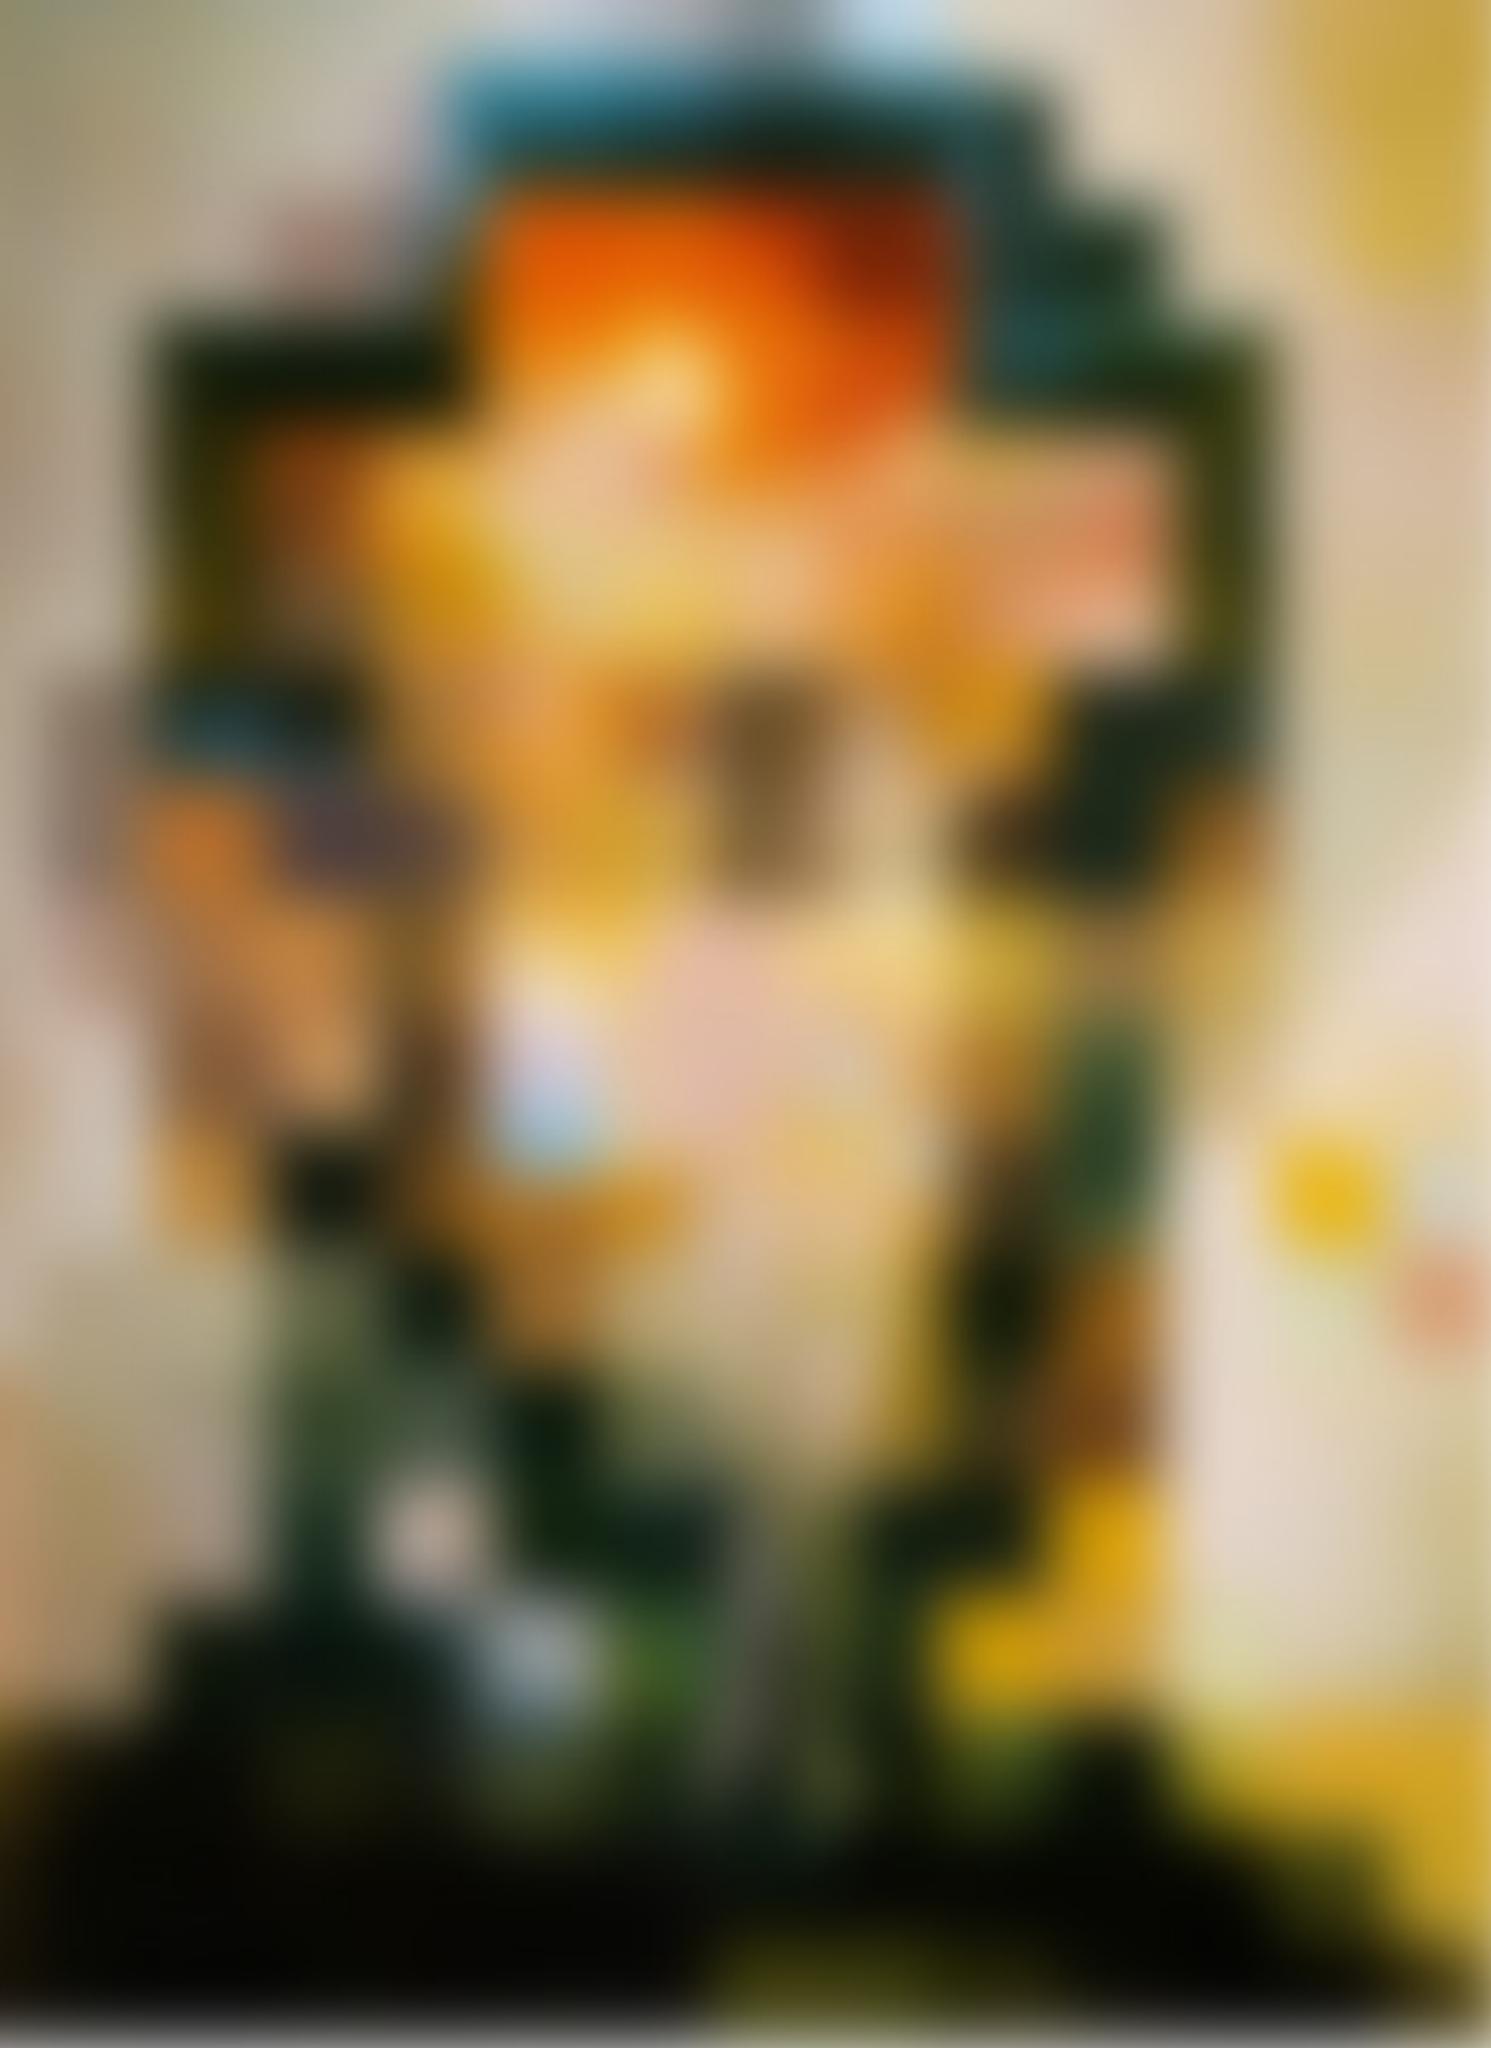


Source: James Tompkin



Salvador Dali

"Gala Contemplating the Mediterranean Sea, which at 30 meters becomes the portrait of Abraham Lincoln", 1976



Fourier Transforms

- ❖ The tool manipulating frequency. Sampling, De-noise, filtering...
- ❖ Linear, space invariant operations are just diagonal operations in the frequency domain
- ❖ Linear convolution is multiplication in the frequency domain

Fourier series

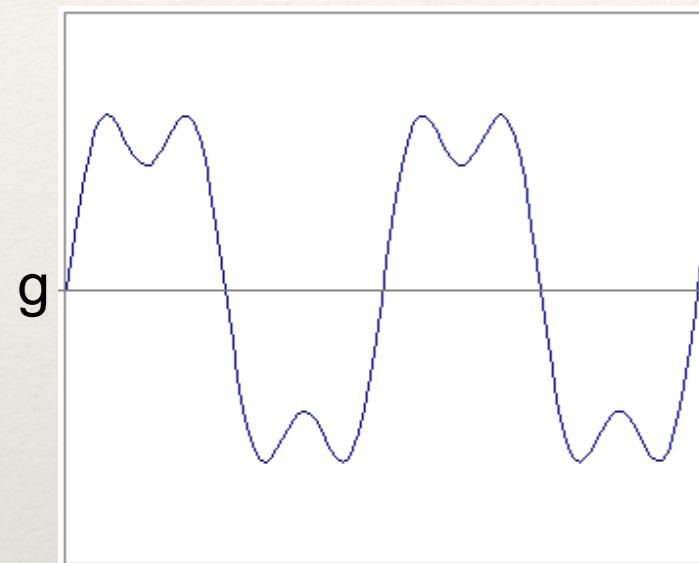
- ❖ A bold idea (1807):
 - ❖ **Any** univariate function can be rewritten as a weighted sum of sines and cosines of different frequencies.
$$\sum_{n=1}^{\infty} (a_n \cos(nt) + b_n \sin(nt))$$
 - ❖ Add enough of them to get any signal $g(t)$ you want!



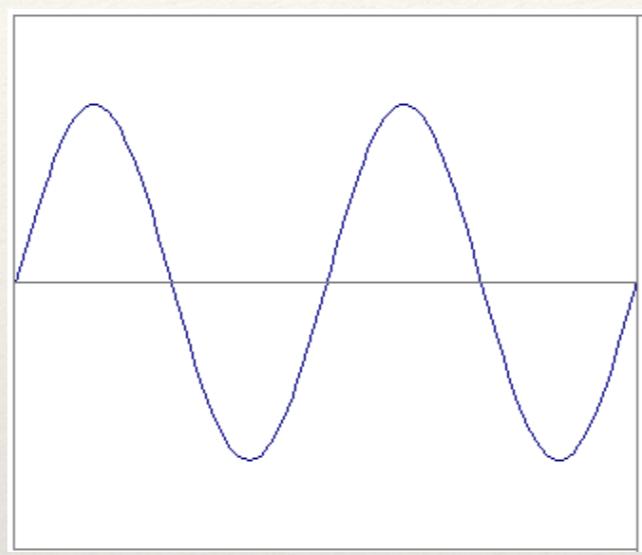
Jean Baptiste Joseph Fourier (1768-1830)

$$t = [0,2], f = 1$$

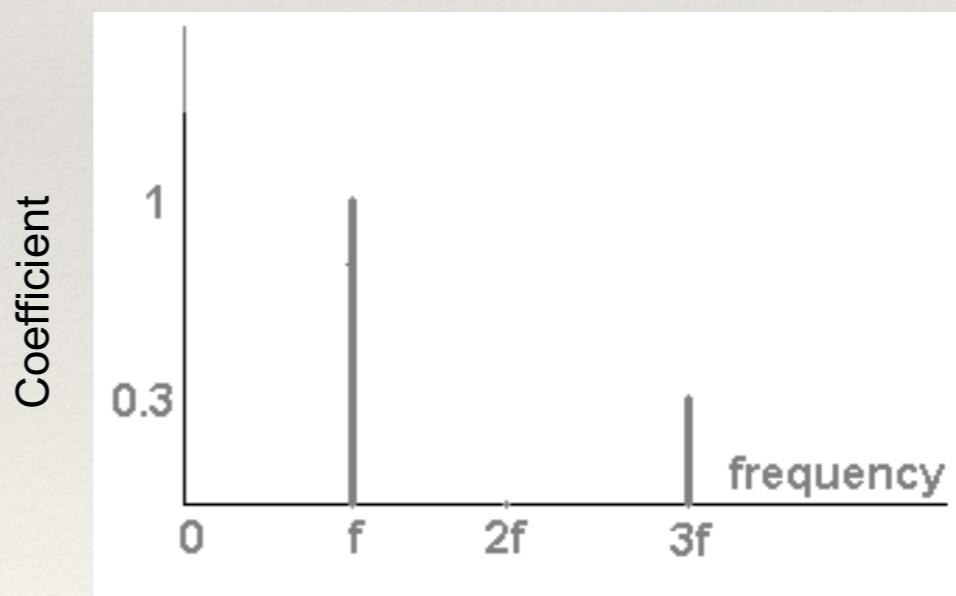
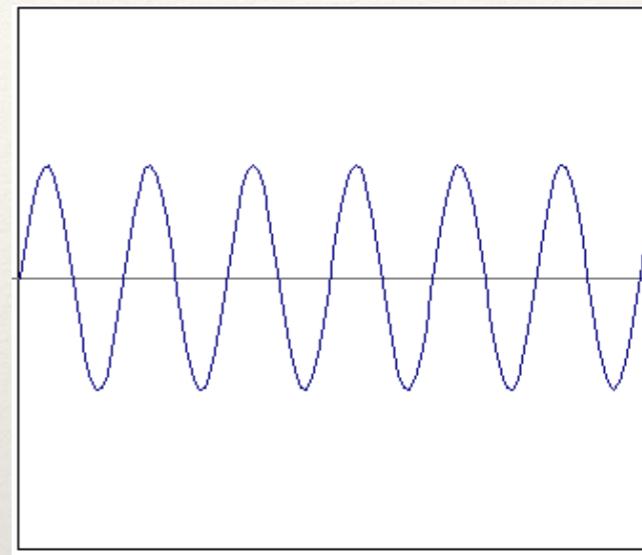
$$g(t) = (1)\sin(2\pi f t) + (1/3)\sin(2\pi(3f)t)$$



=



+

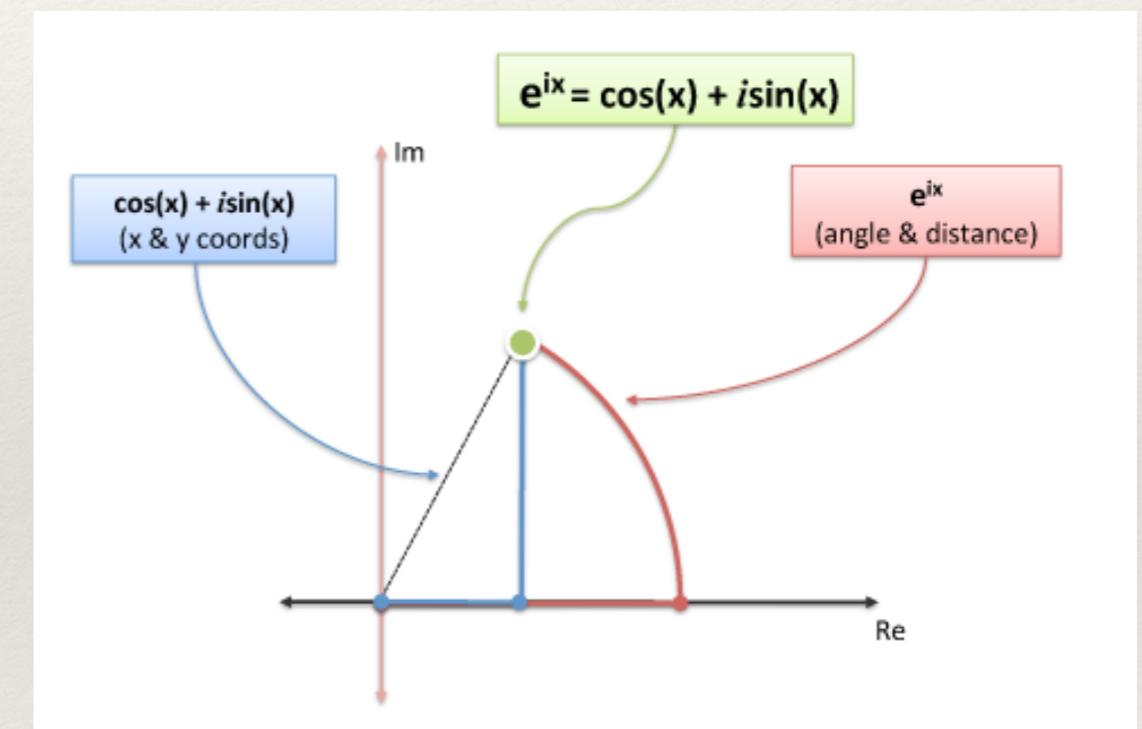
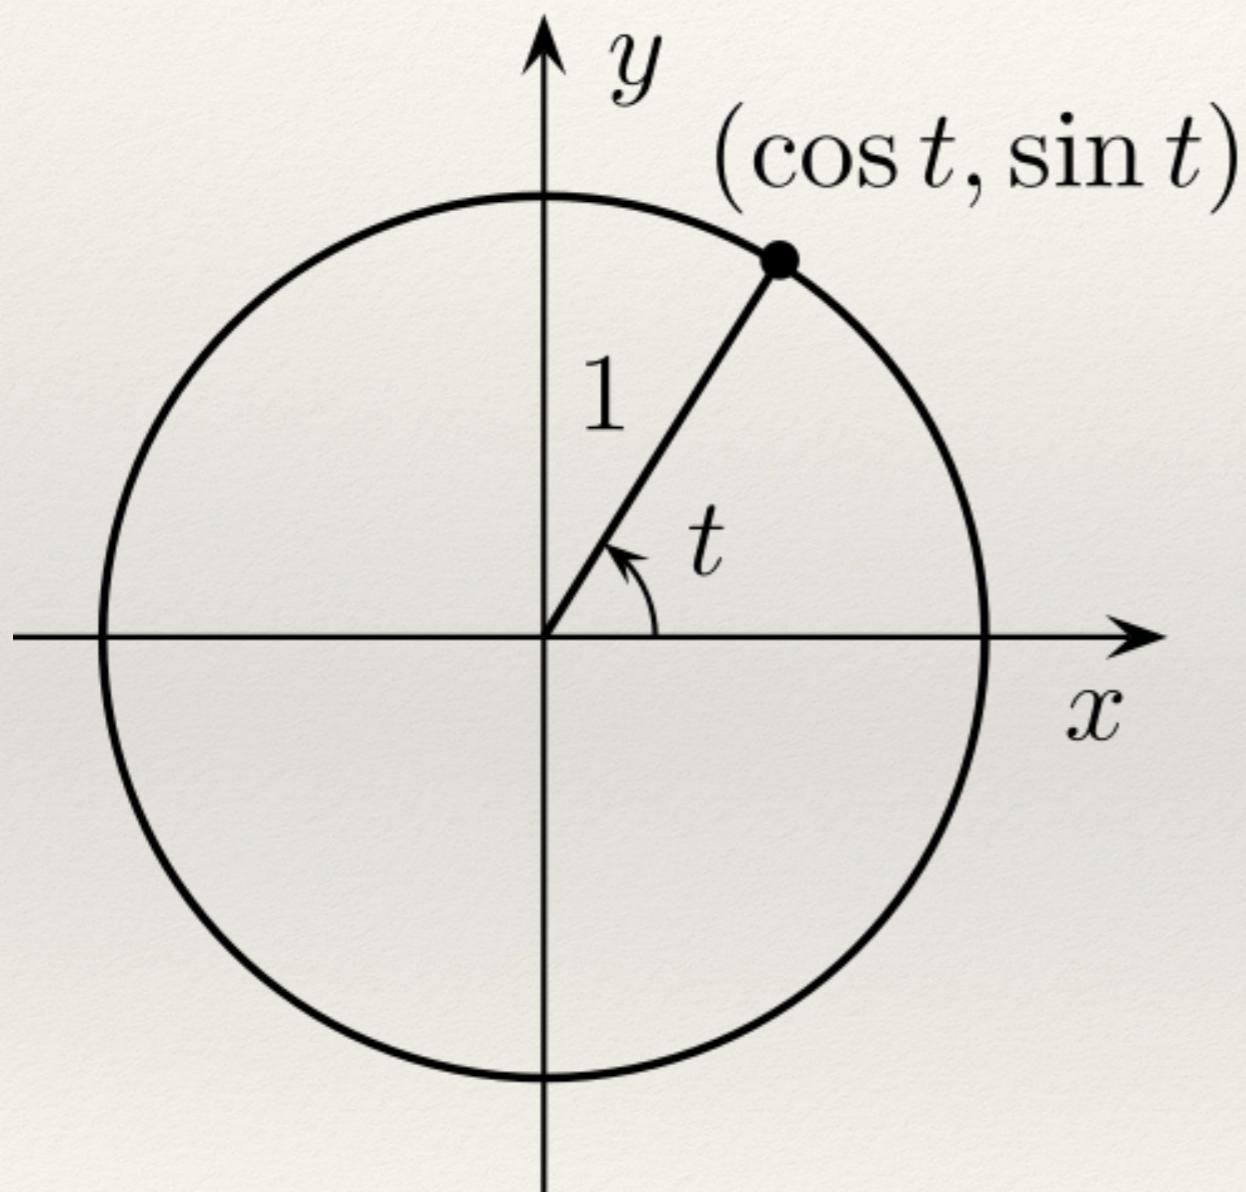


Slides: Efros

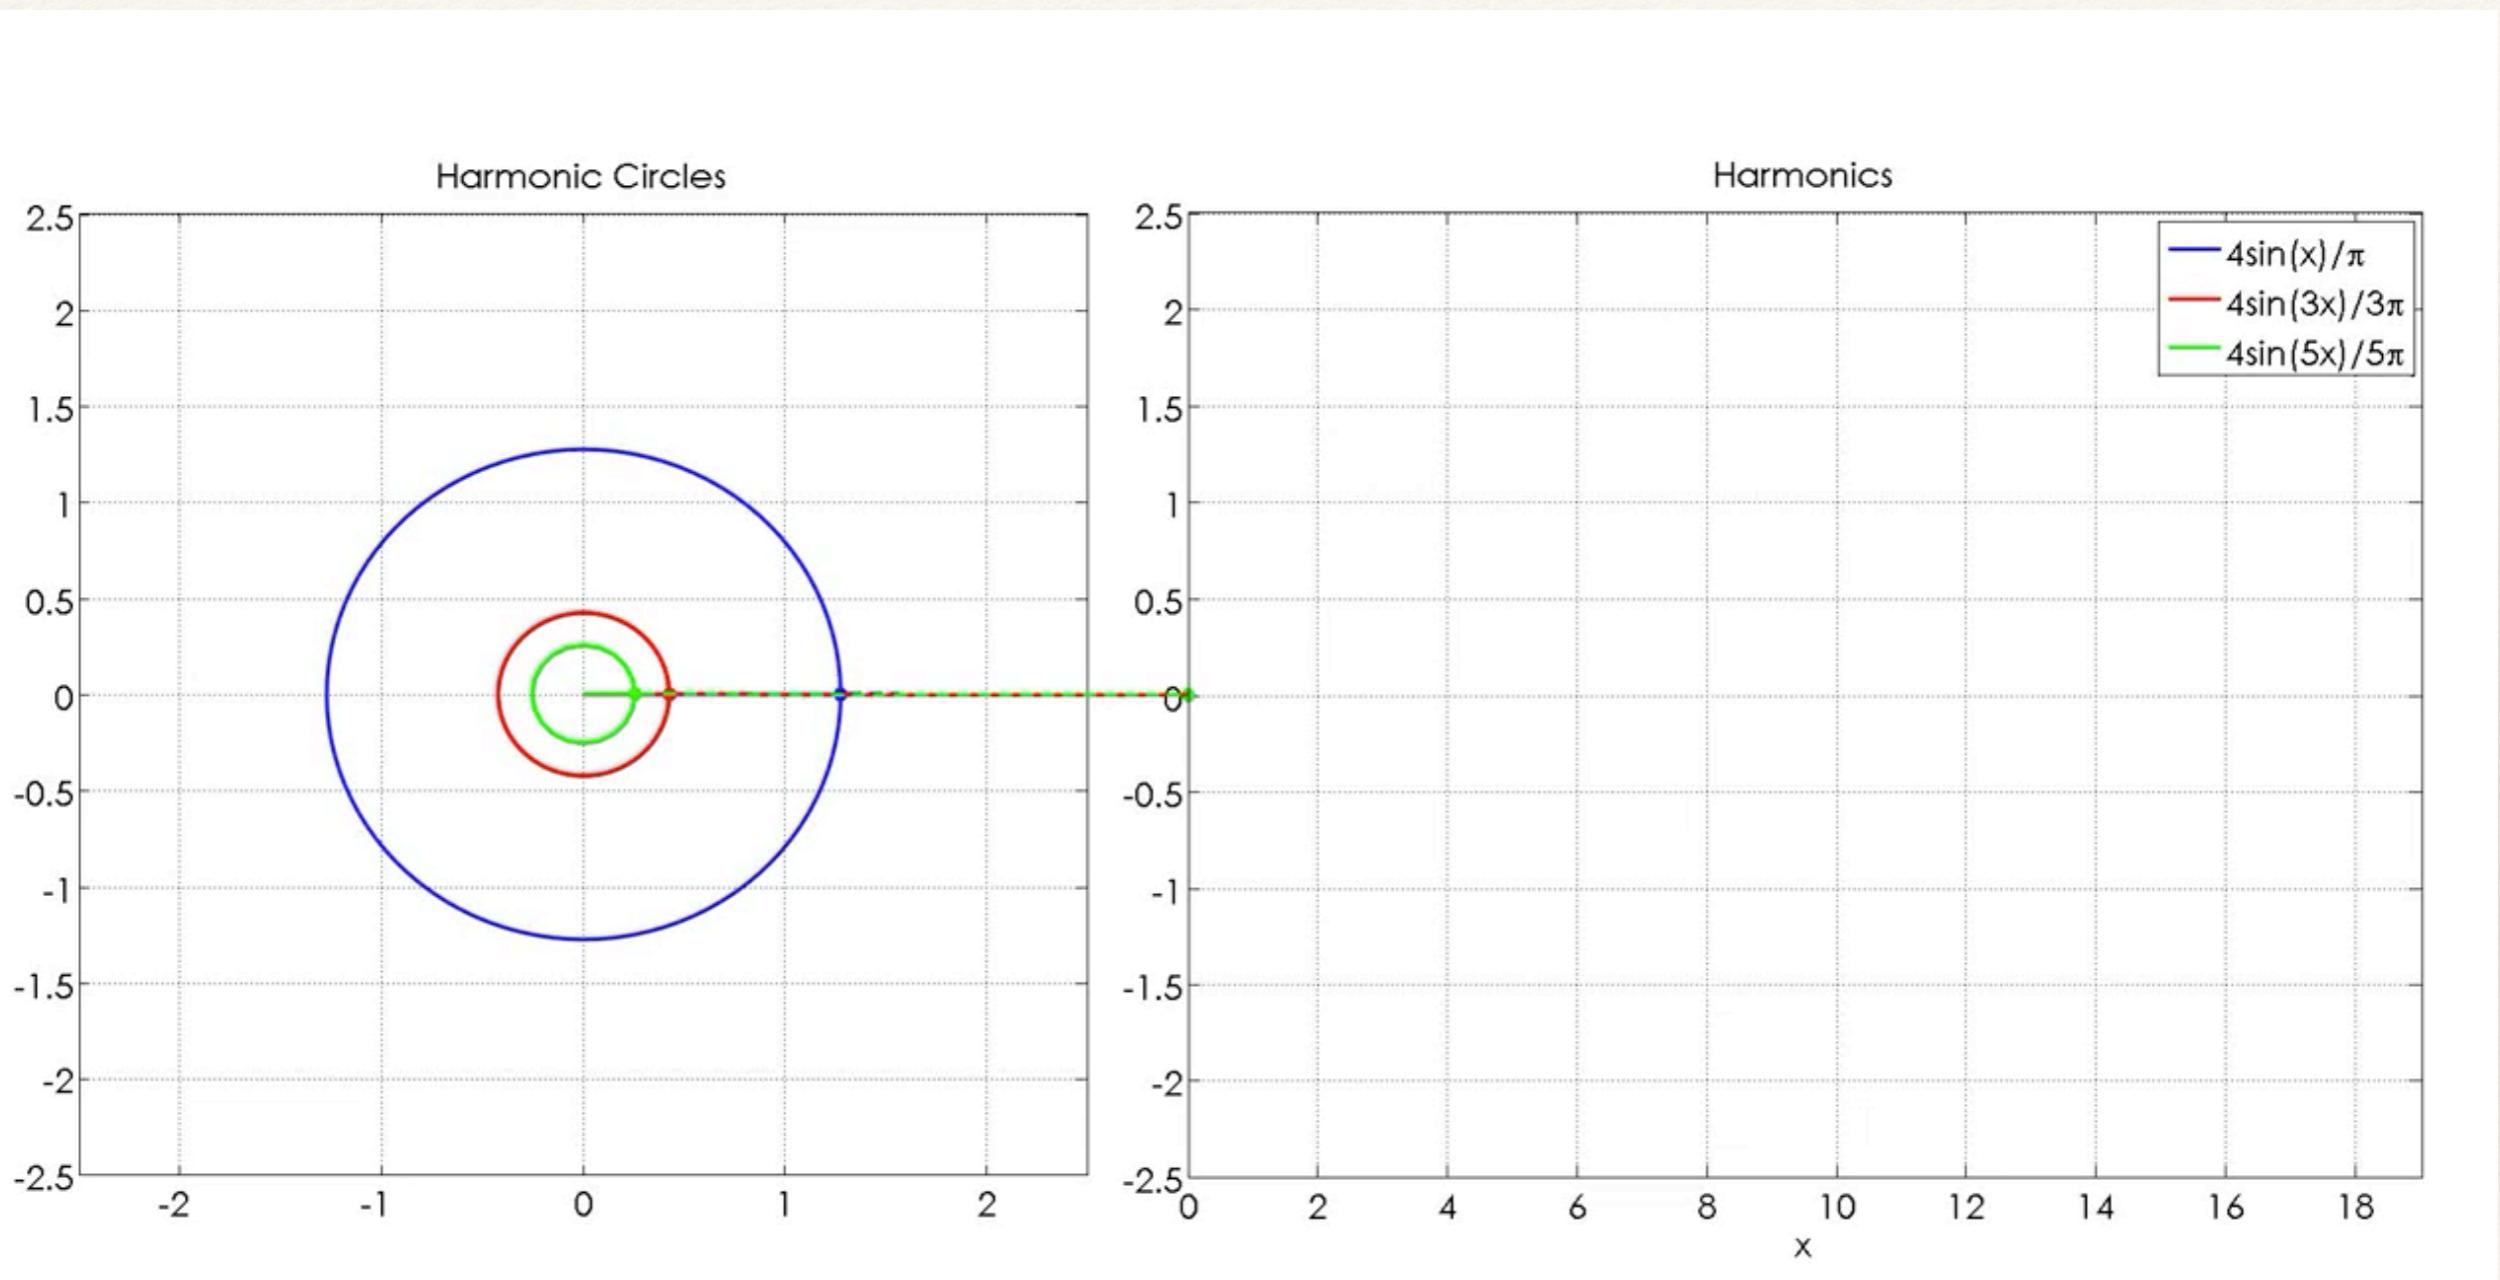


Wikipedia-Fourier transform

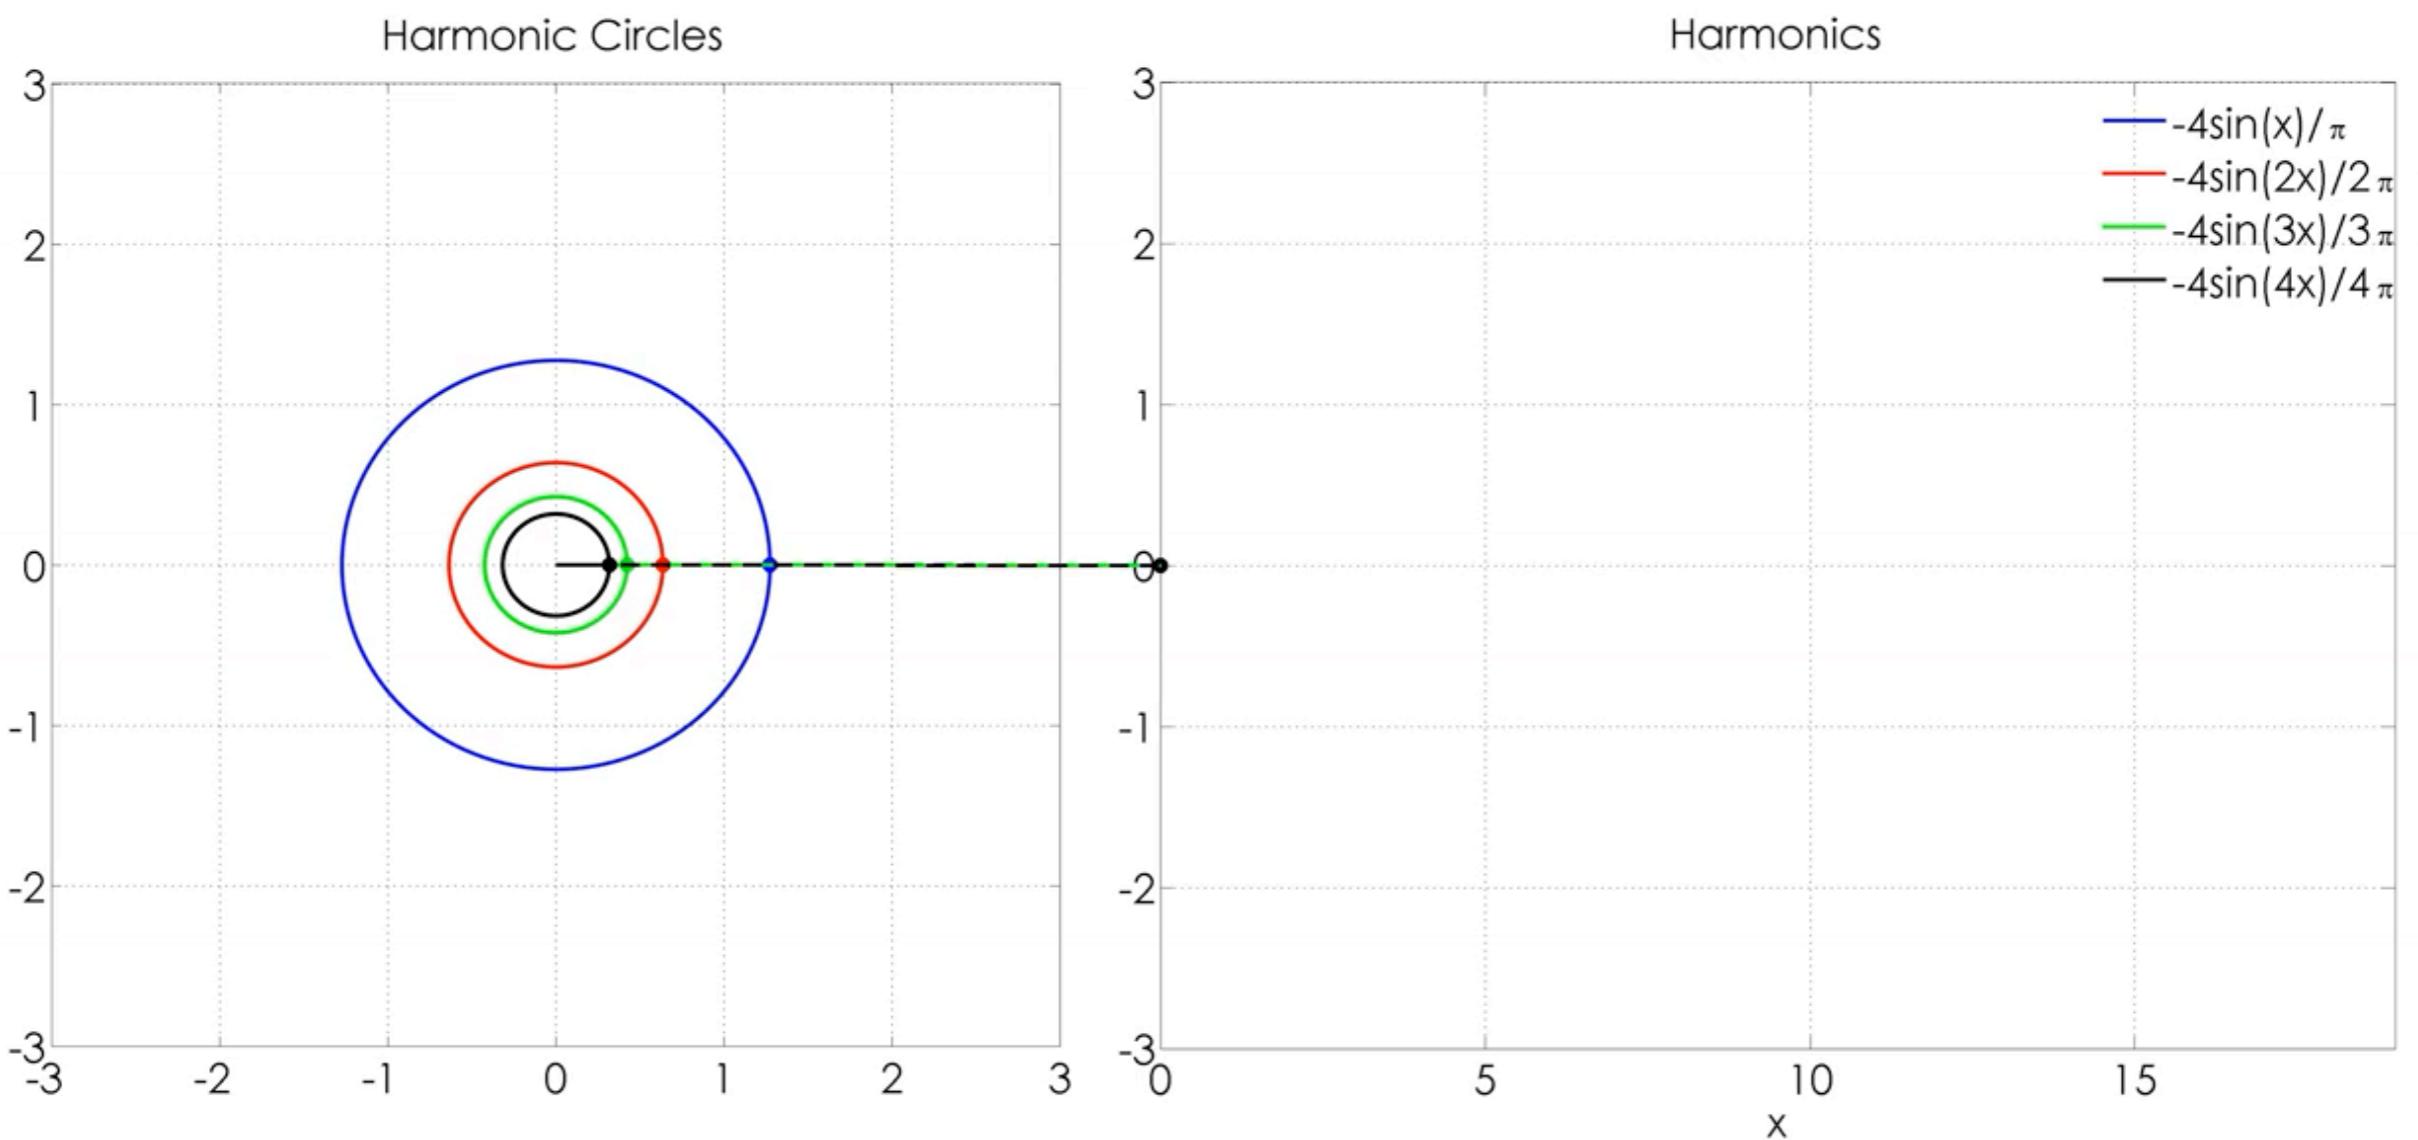
Sine/cosine and circle



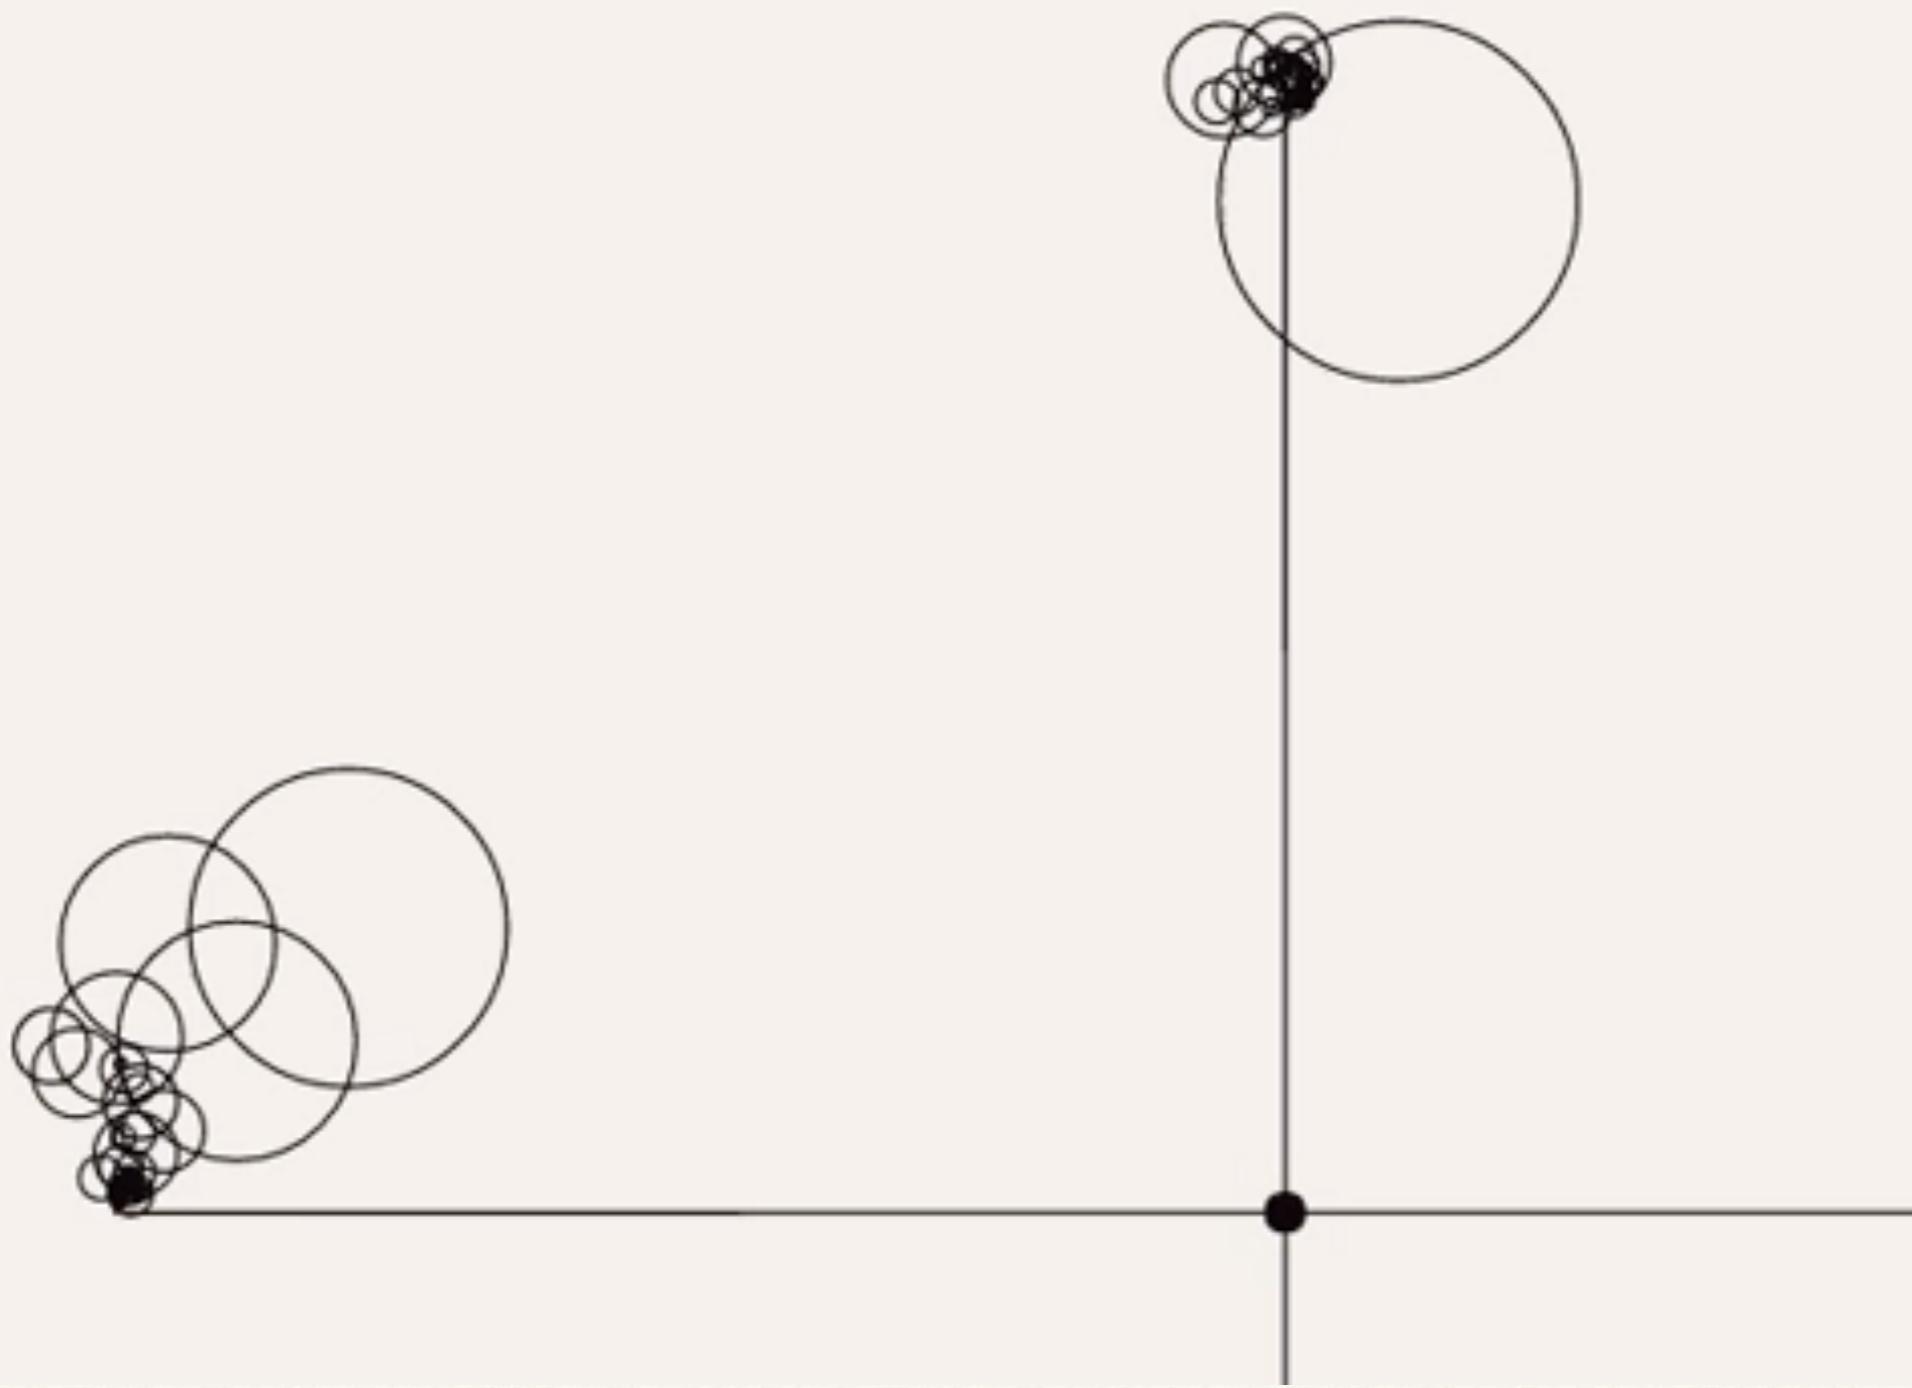
Square wave (approx.)



Square wave (approx.)



One series in each of x and y

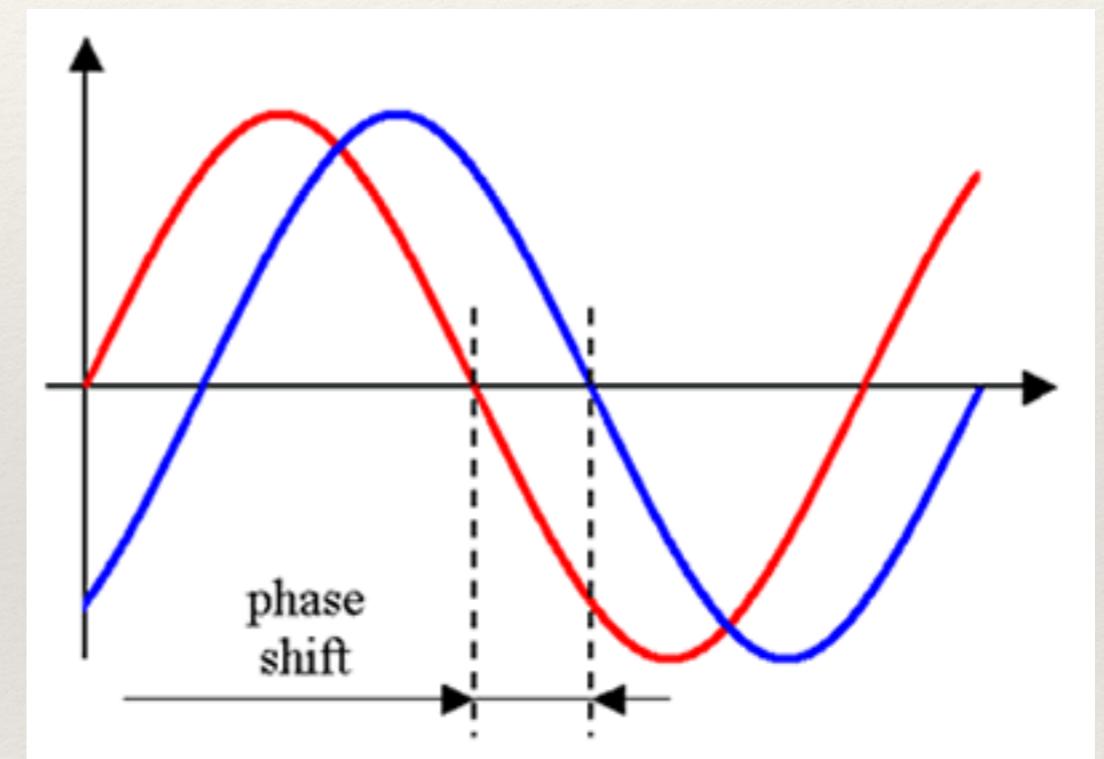


Amplitude-phase form

- ❖ Add phase term to shift cos values into sin values

$$\sum_{n=1}^N (a_n \sin(nx + \phi_n))$$

Phase

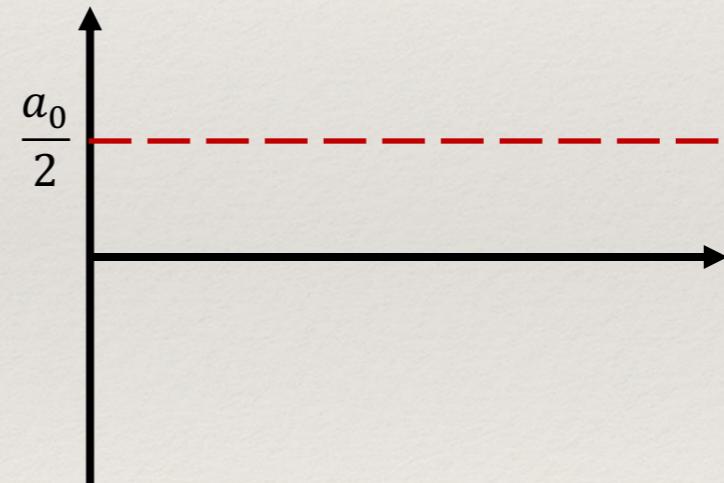


Amplitude-phase form

- ❖ Add component of infinite frequency
 - ▶ = mean of signal over period
 - ▶ = value around which signal fluctuates

$$\frac{a_0}{2} + \sum_{n=1}^N (a_n \sin(nx + \phi_n))$$

Average of signal
over period

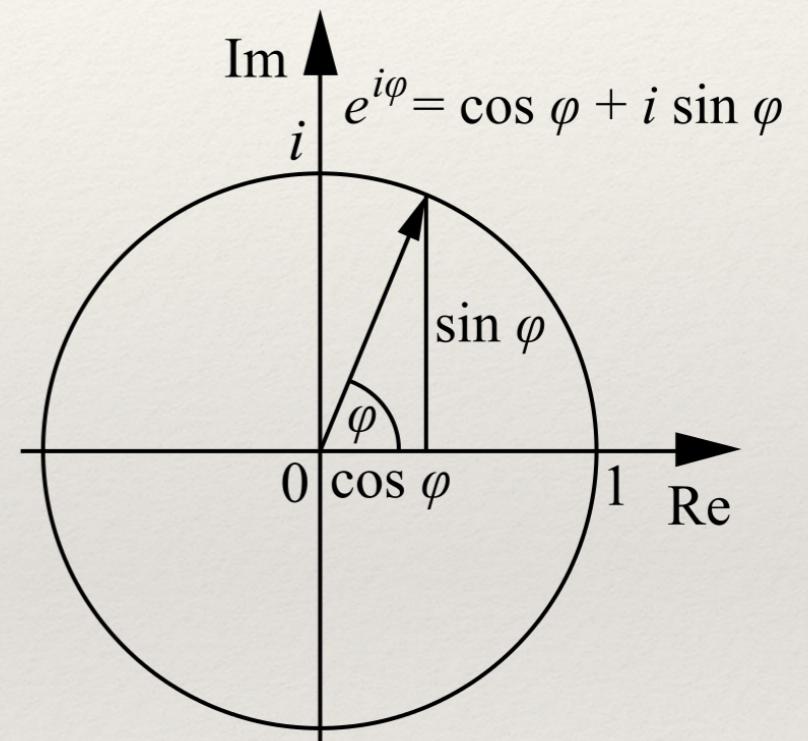


Amplitude-phase form

$$s_N(x) \triangleq \frac{A_0}{2} + \sum_{n=1}^N A_n \cdot \sin\left(\frac{2\pi n x}{P} + \phi_n\right), \quad \text{for integer } N \geq 1. \quad (\text{Eq.1})$$

$$s_N(x) = \overbrace{a_0}^{A_0}/2 + \sum_{n=1}^N \left(\overbrace{a_n}^{A_n \sin(\phi_n)} \cos\left(\frac{2\pi n x}{P}\right) + \overbrace{b_n}^{A_n \cos(\phi_n)} \sin\left(\frac{2\pi n x}{P}\right) \right), \quad (\text{Eq.2})$$

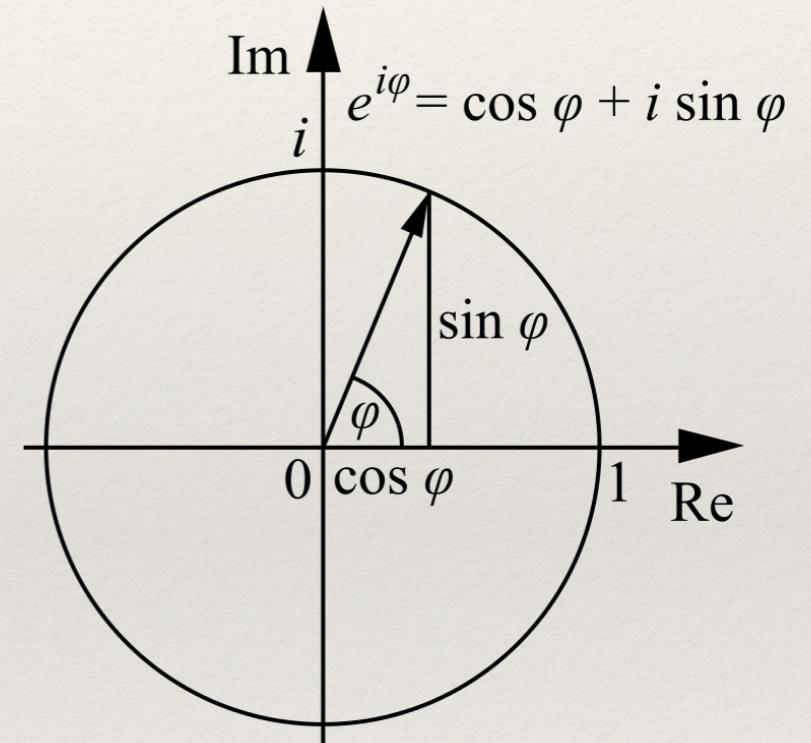
$$s_N(x) = \sum_{n=-N}^N c_n \cdot e^{i \frac{2\pi n x}{P}}, \quad (\text{Eq.3})$$



Fourier Transform

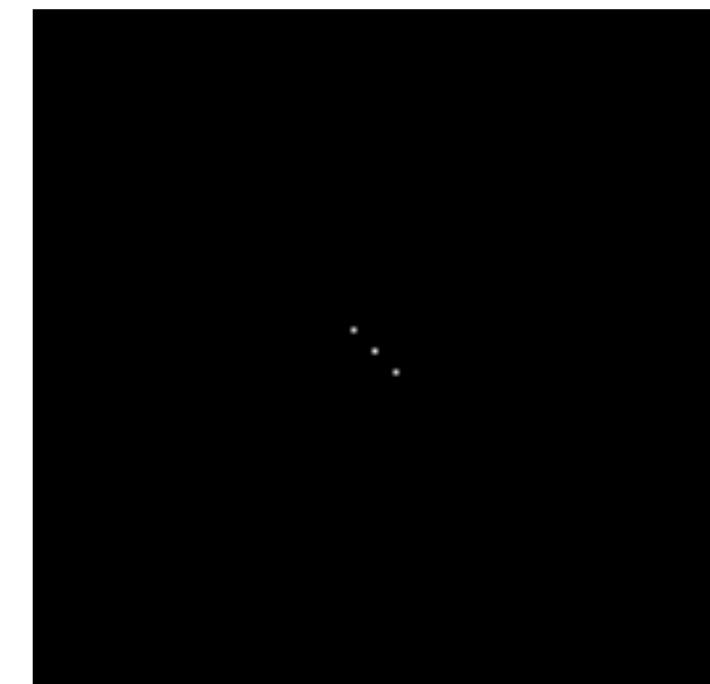
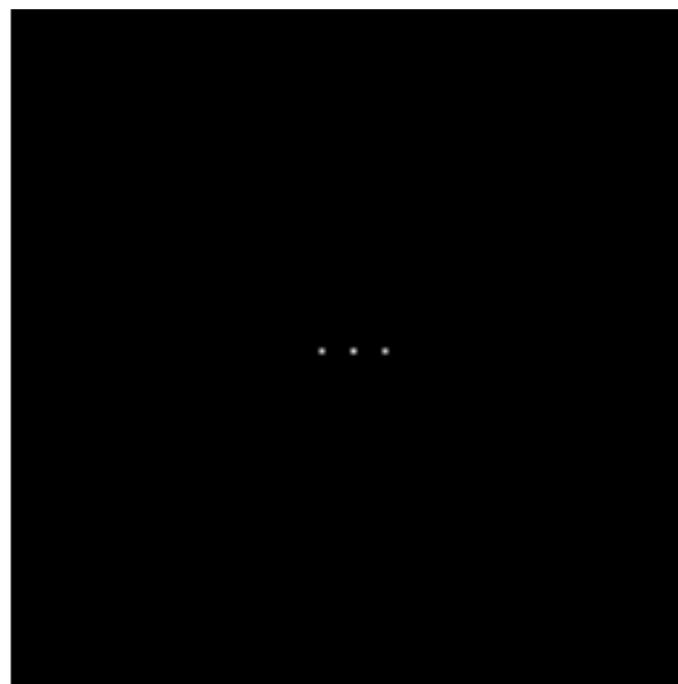
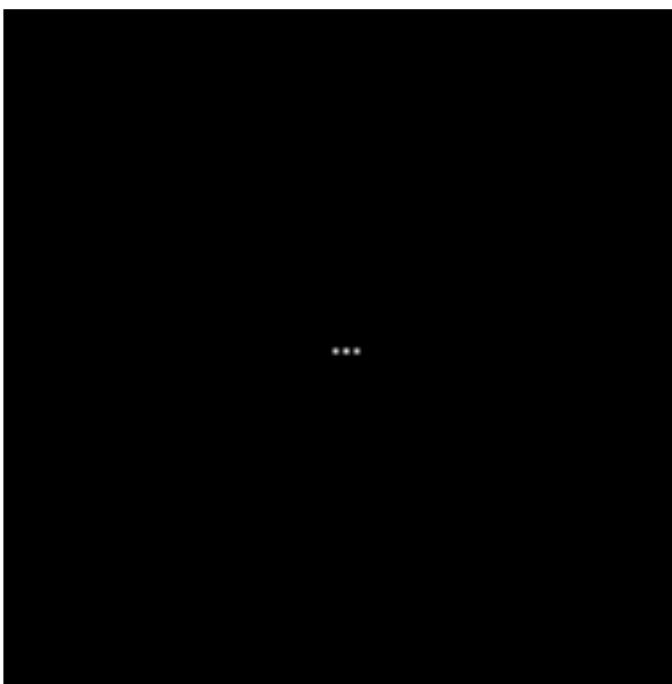
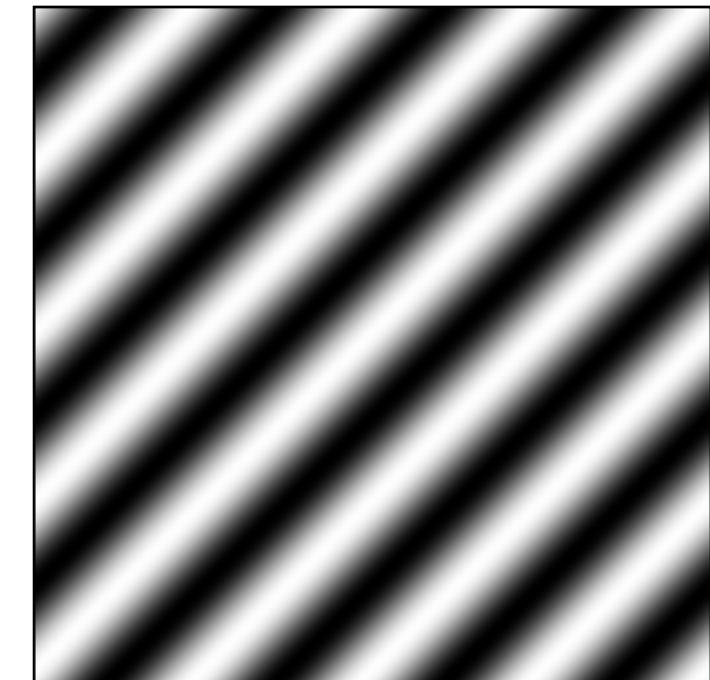
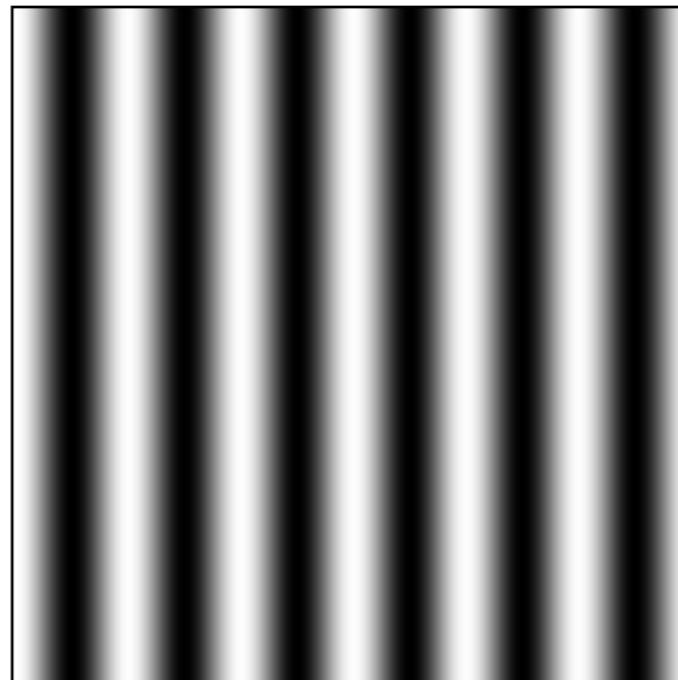
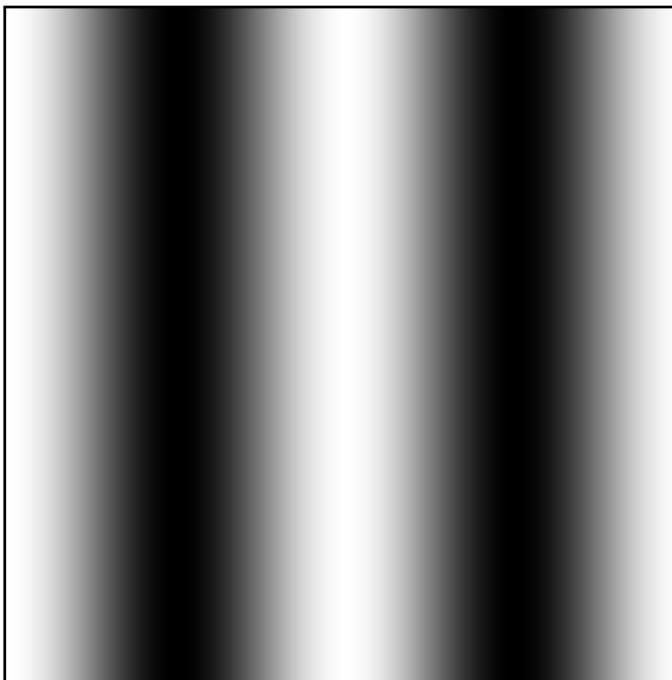
- ❖ Stores the amplitude and phase at each frequency:
 - ❖ For mathematical convenience, this is often notated in terms of real and complex numbers
 - ❖ Related by Euler's formula
- ❖ Amplitude encodes how much signal there is at a particular frequency
$$A = \pm \sqrt{\text{Re}(\varphi)^2 + \text{Im}(\varphi)^2}$$
- ❖ Phase encodes spatial information (indirectly)

$$\phi = \tan^{-1} \frac{\text{Im}(\varphi)}{\text{Re}(\varphi)}$$



Fourier analysis in images

Spatial domain images

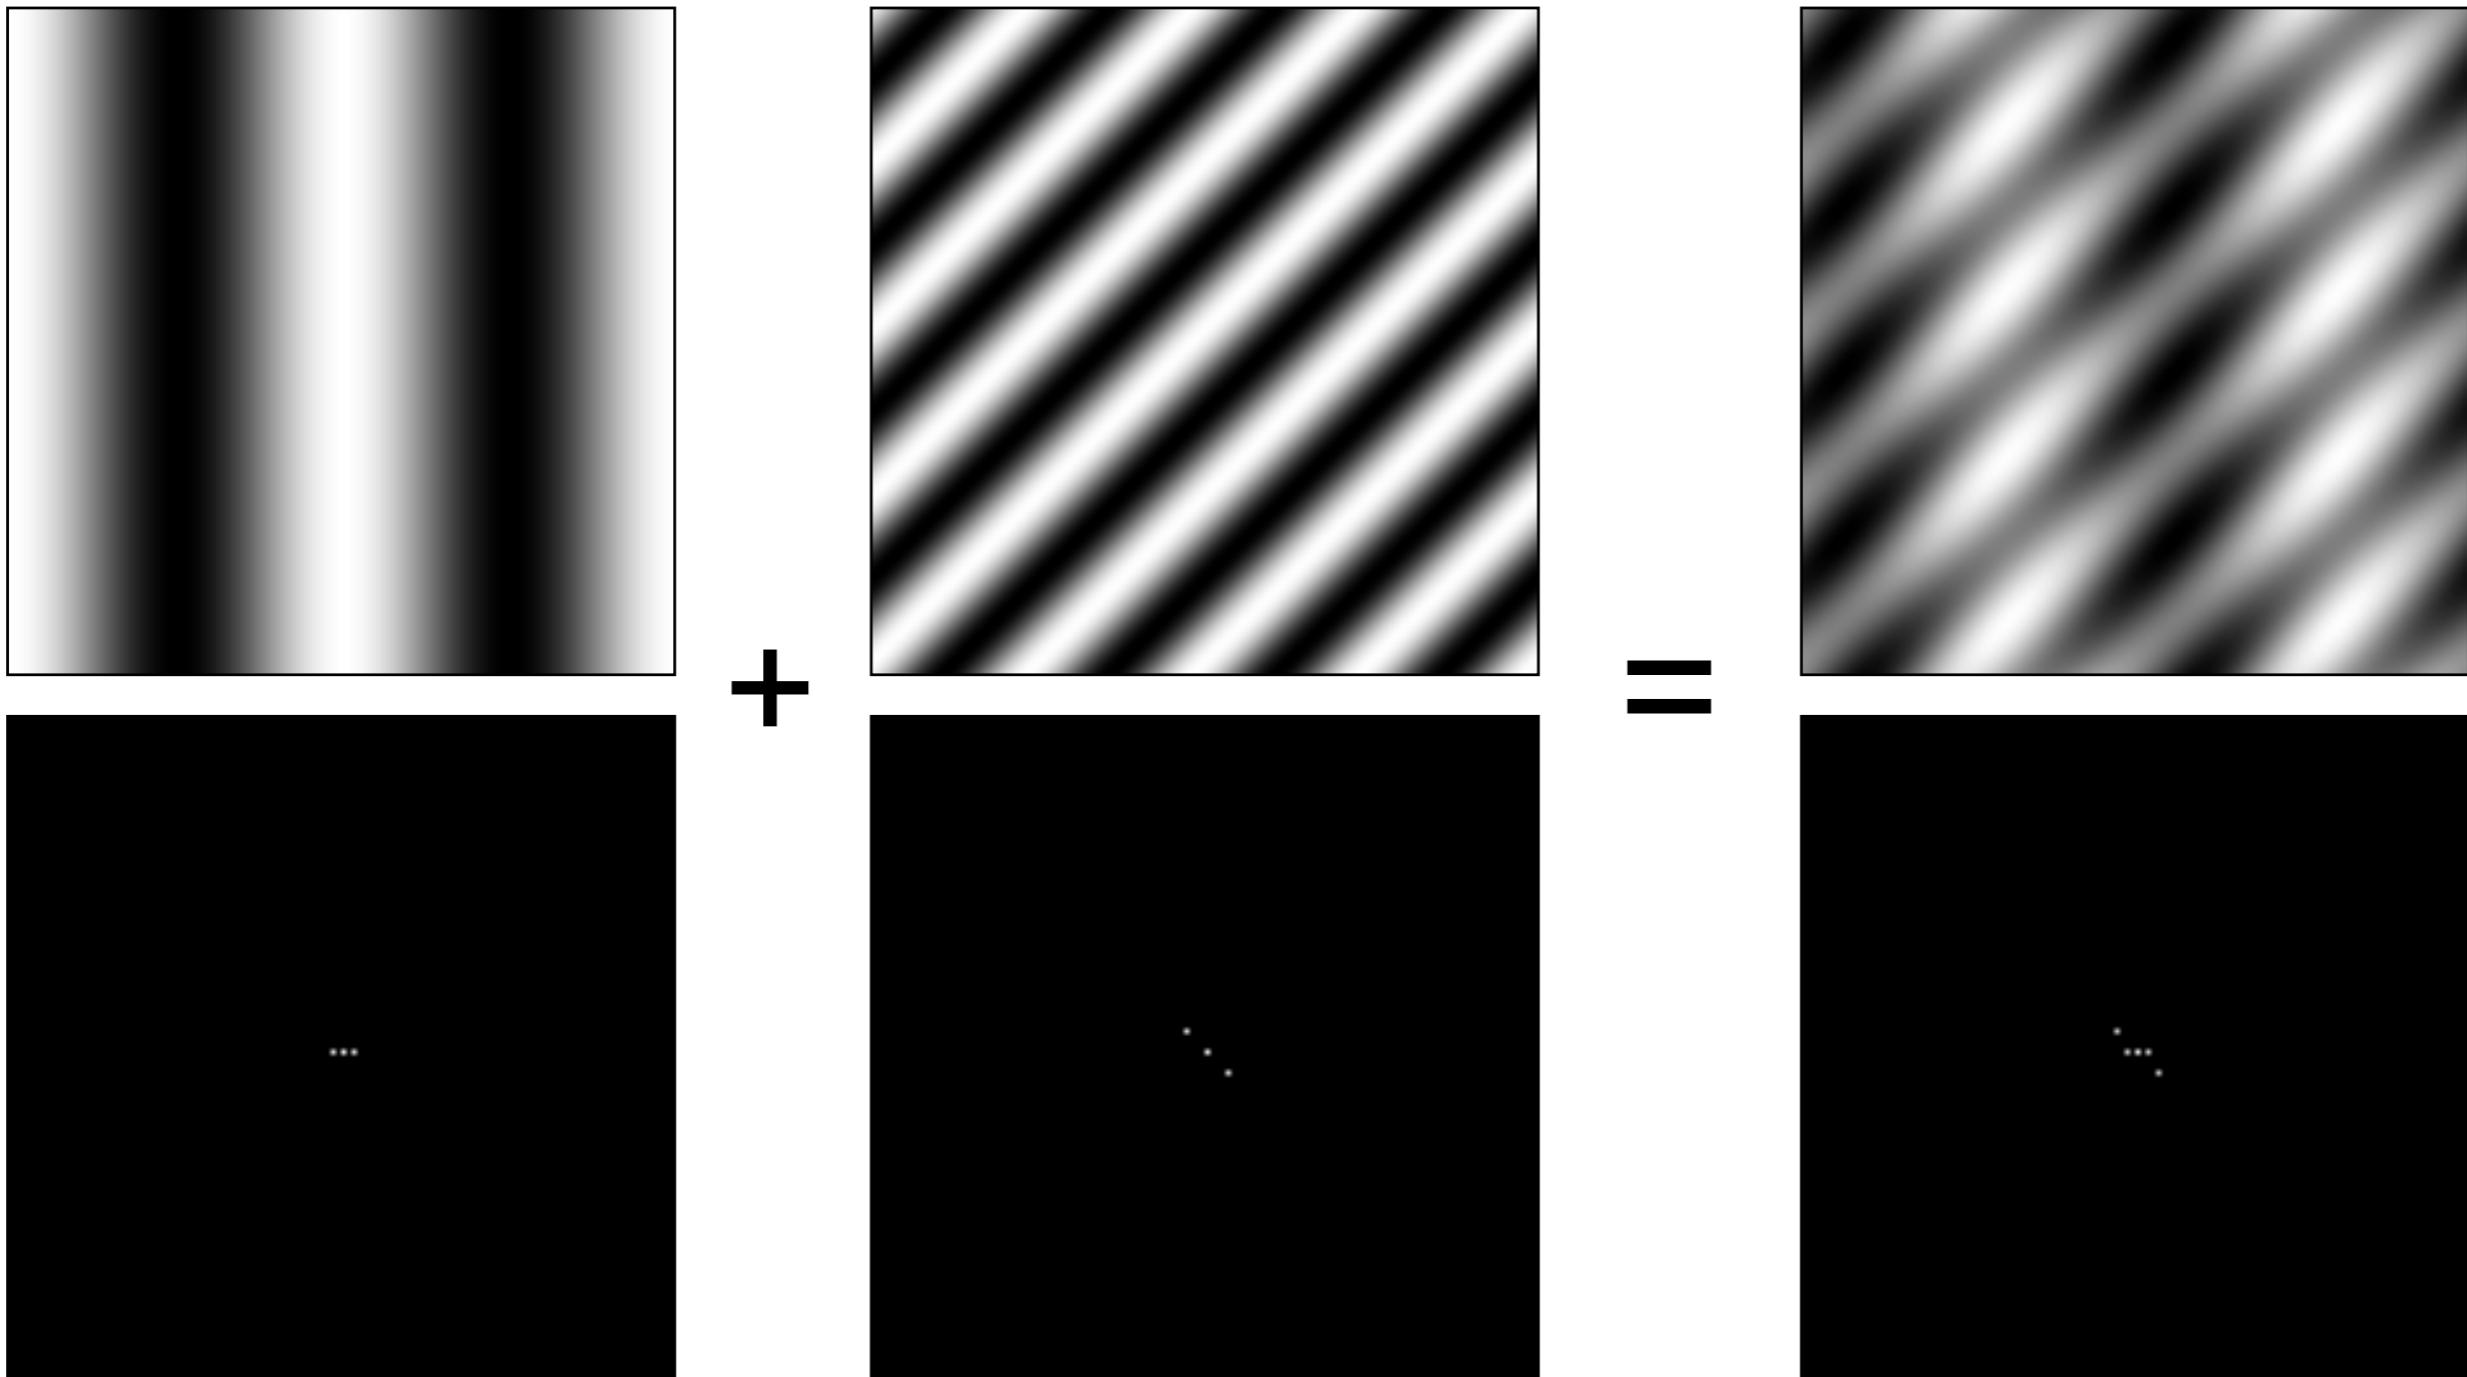


Fourier decomposition amplitude images

<http://sharp.bu.edu/~slehar/fourier/fourier.html#filtering> More: <http://www.cs.unm.edu/~brayer/vision/fourier.html>

Signals can be composed

Spatial domain images



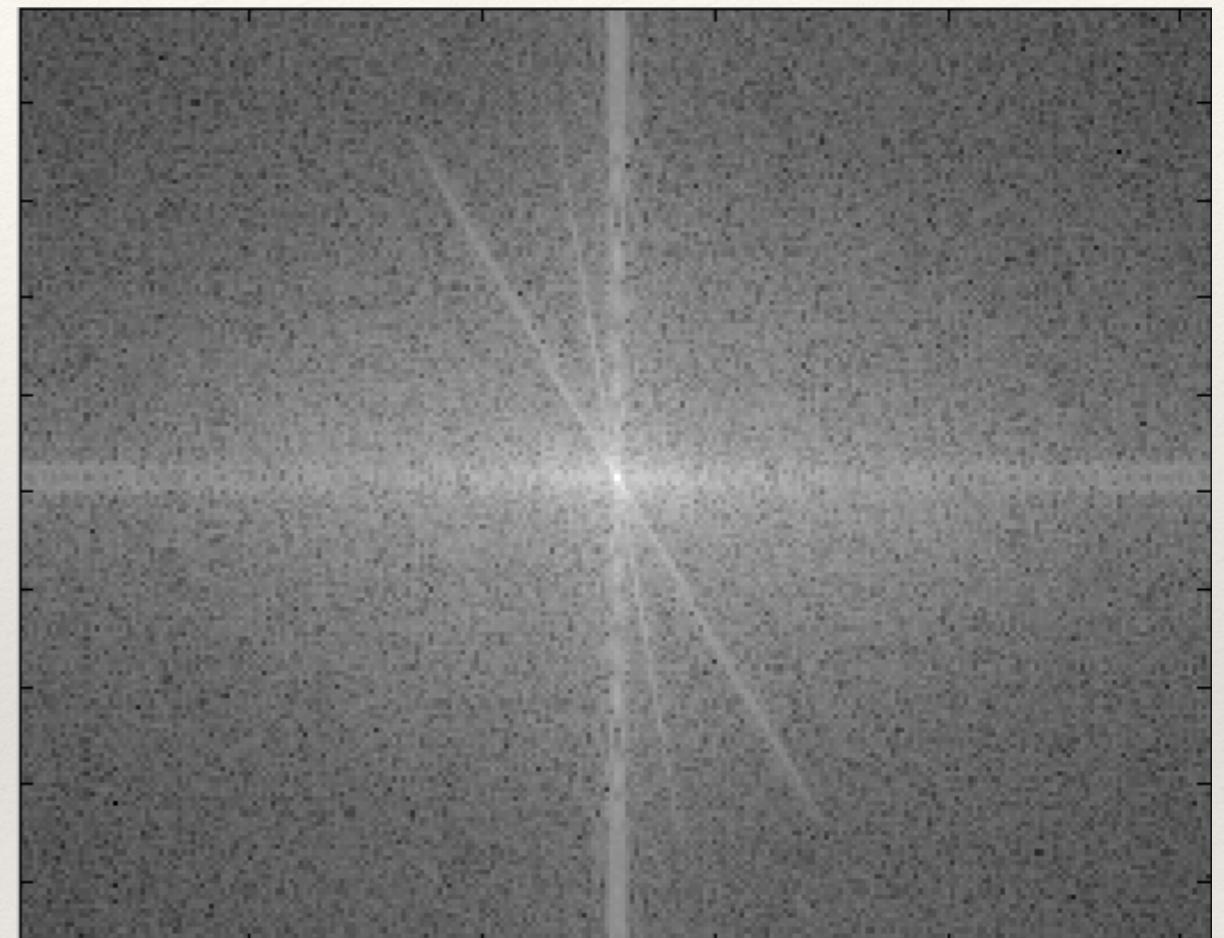
Fourier decomposition amplitude images

Natural images

Natural image



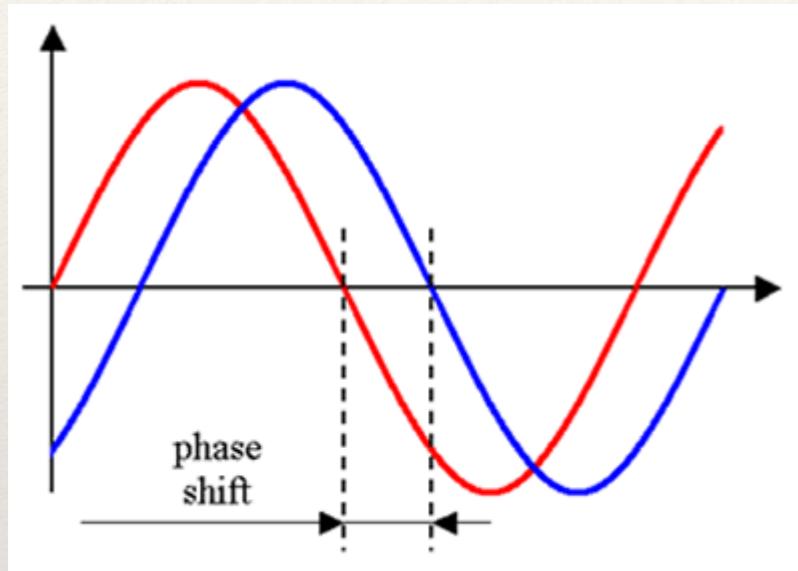
Fourier decomposition amplitude image



- ① What does it mean to be at pixel x,y ?
- ② What does it mean to be more or less bright in the Fourier decomposition image?

Source: James Tompkin

Amplitude / Phase



- ❖ Amplitude tells you “how much”
- ❖ Phase tells you “where”
- ❖ Translate the image?
 - ❖ Amplitude unchanged
 - ❖ Adds a constant to the phase.

❖

[[Peppergrower](#); Wikipedia]

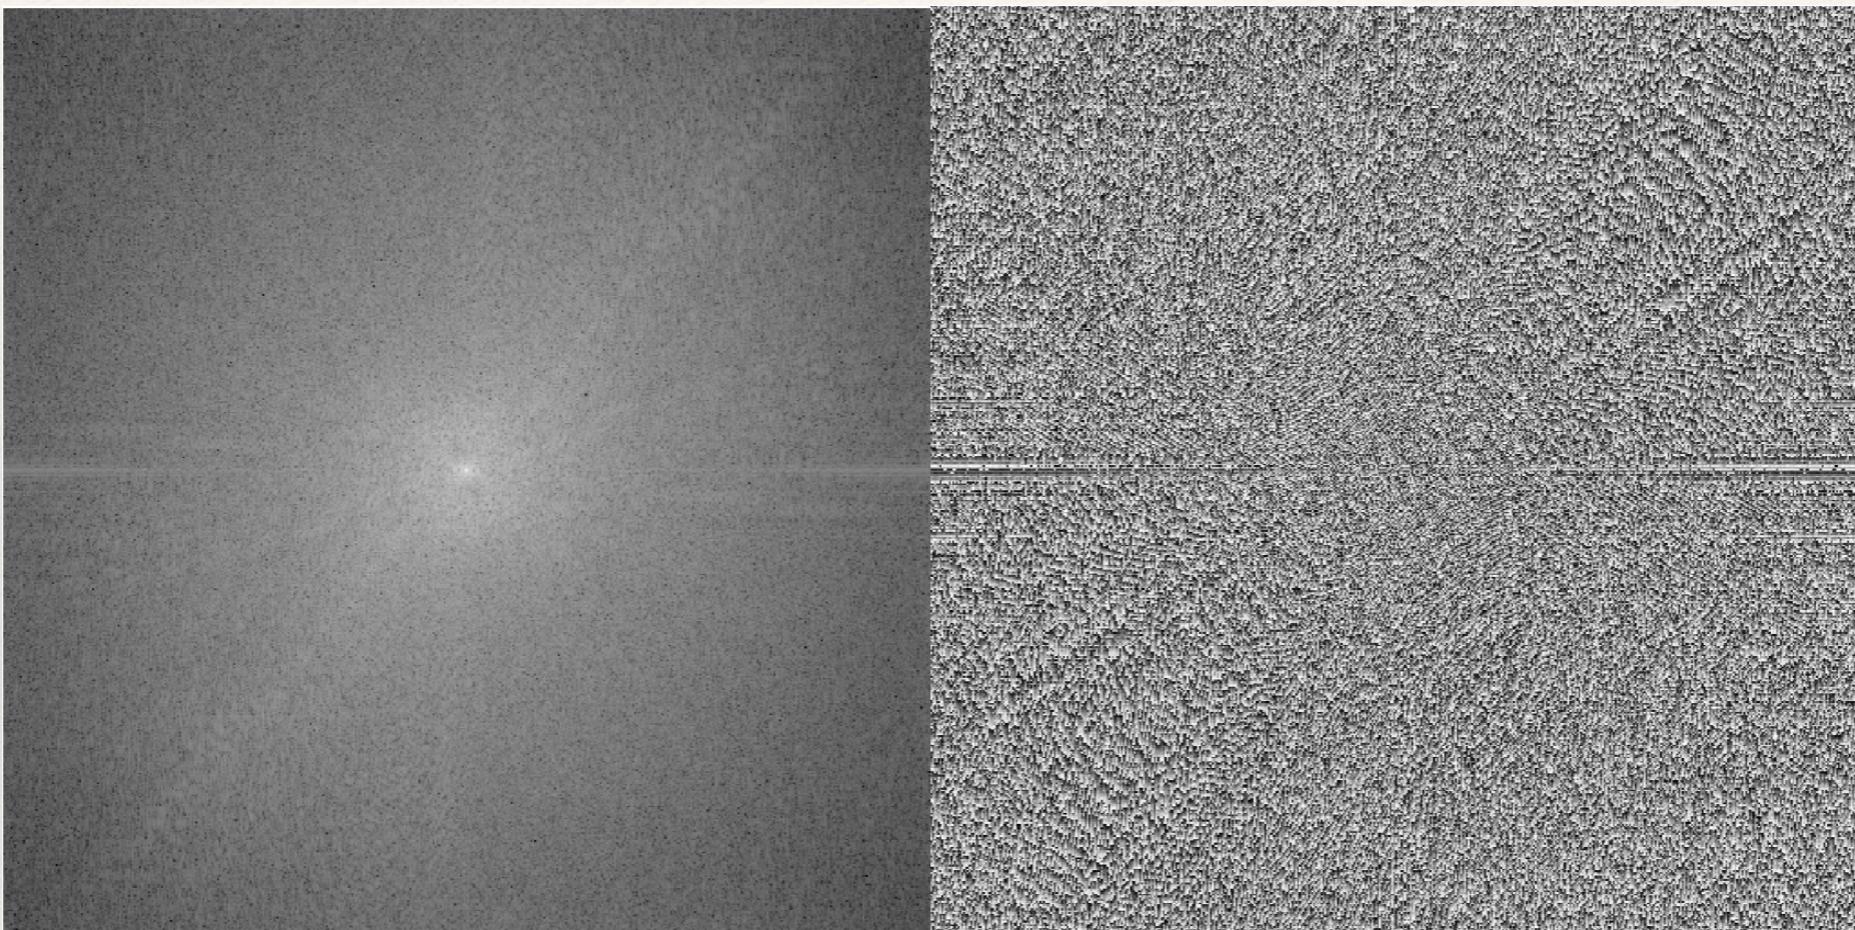
Amplitude / Phase



Efros

Amplitude / Phase

Amplitude



Phase

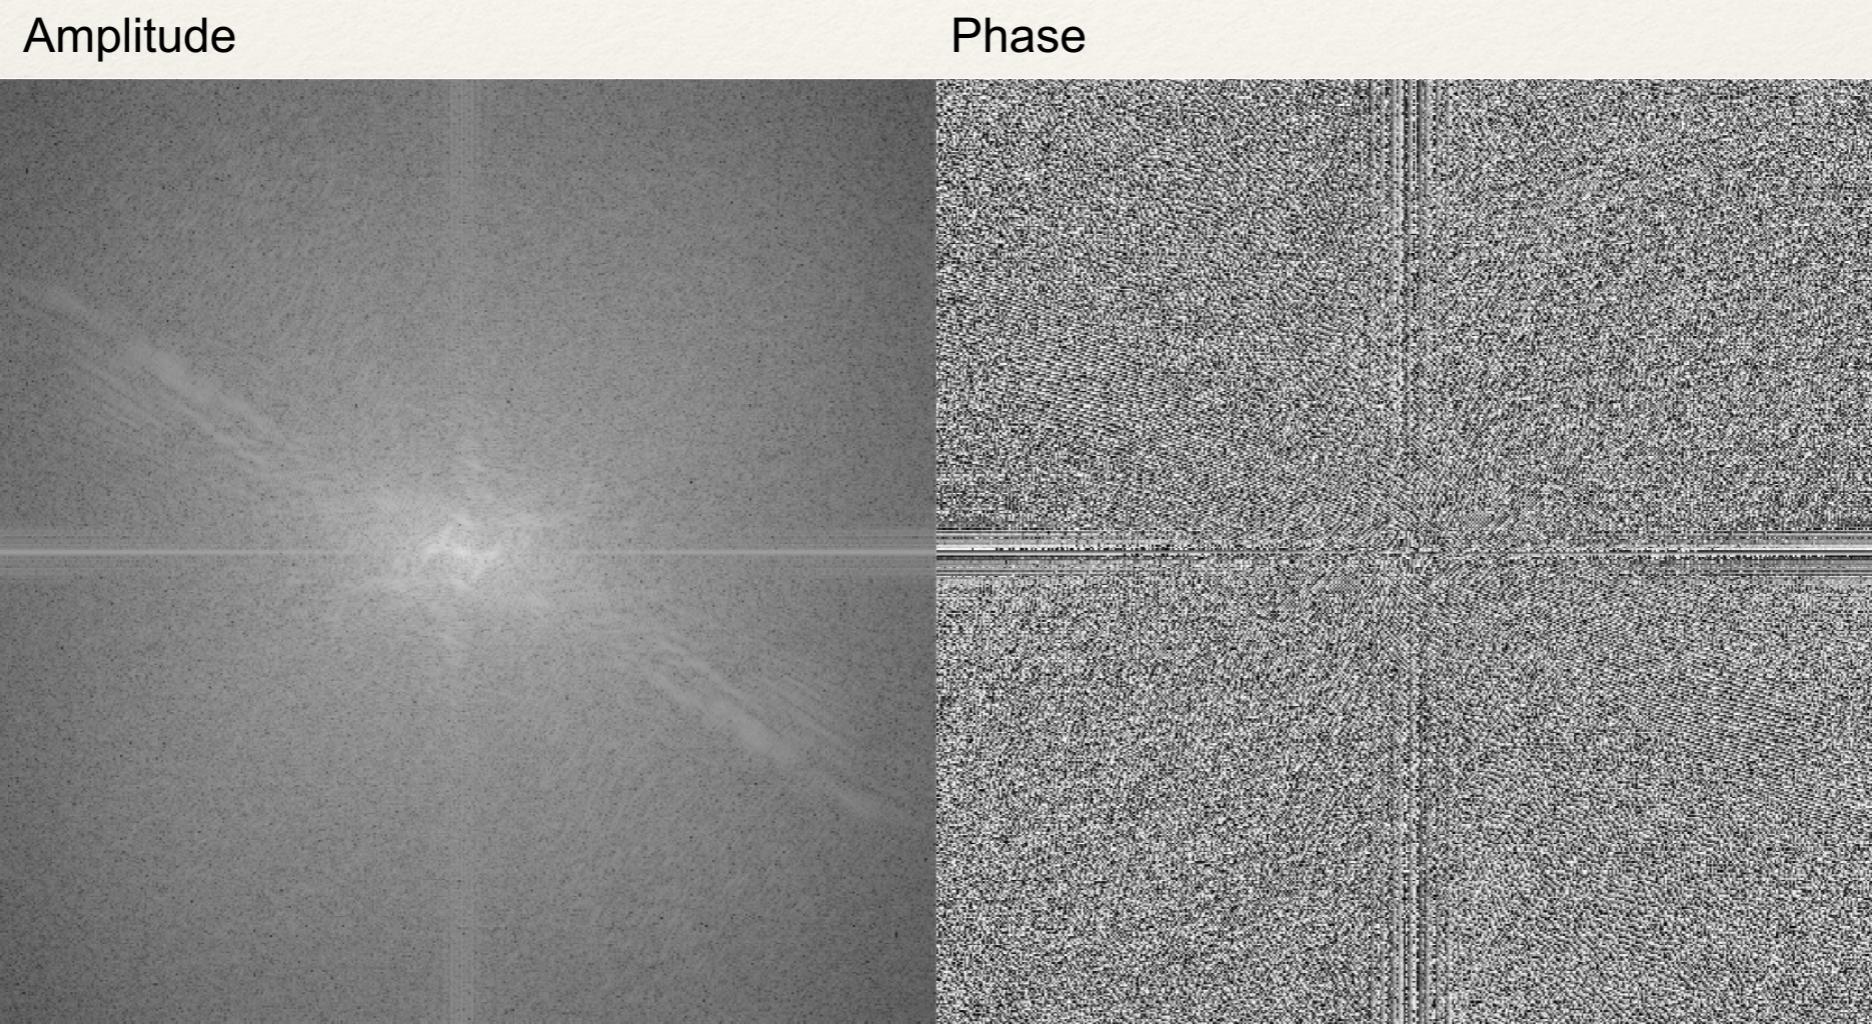
Efros

Amplitude / Phase



Efros

Amplitude / Phase

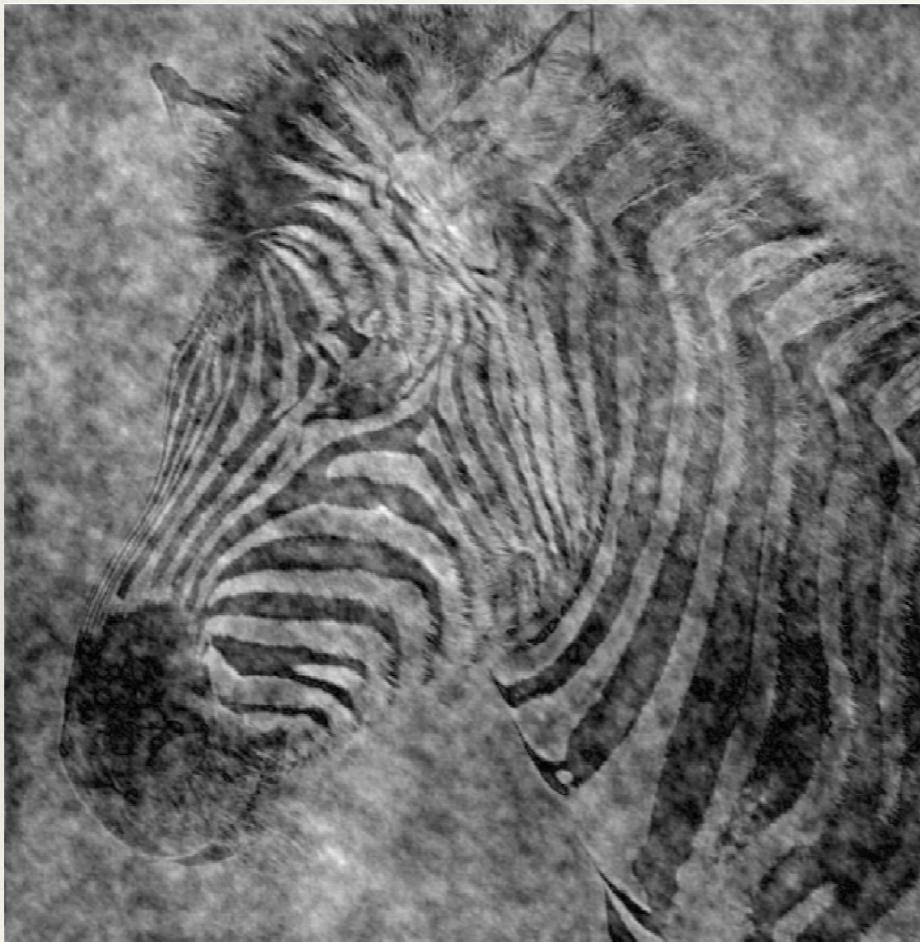


Efros

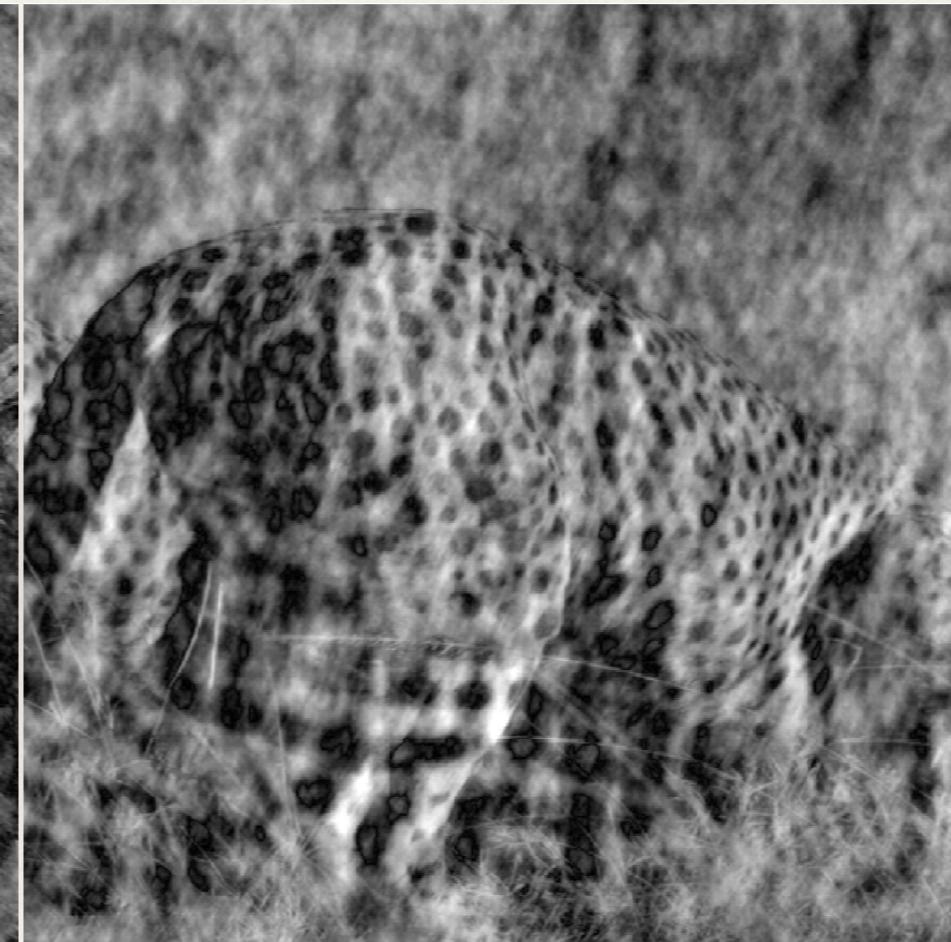
Amplitude / Phase

Cheebra

Zebra phase, cheetah amplitude



Cheetah phase, zebra amplitude



Efros

Amplitude / Phase

- ❖ The frequency amplitude of natural images are quite similar
 - ❖ Heavy in low frequencies, falling off in high frequencies
 - ❖ Will *any* image be like that, or is it a property of the world we live in?
- ❖ Most information in the image is carried in the phase, not the amplitude
 - ❖ Not quite clear why

Properties of Fourier Transforms

- ❖ Linearity

$$F[ax(t) + by(t)] = a F[x(t)] + b F[y(t)]$$

- ❖ Fourier transform of a real signal is symmetric about the origin
- ❖ The energy of the signal is the same as the energy of its Fourier transform

The Convolution Theorem

- ❖ The Fourier transform of the convolution of two functions is the product of their Fourier transforms

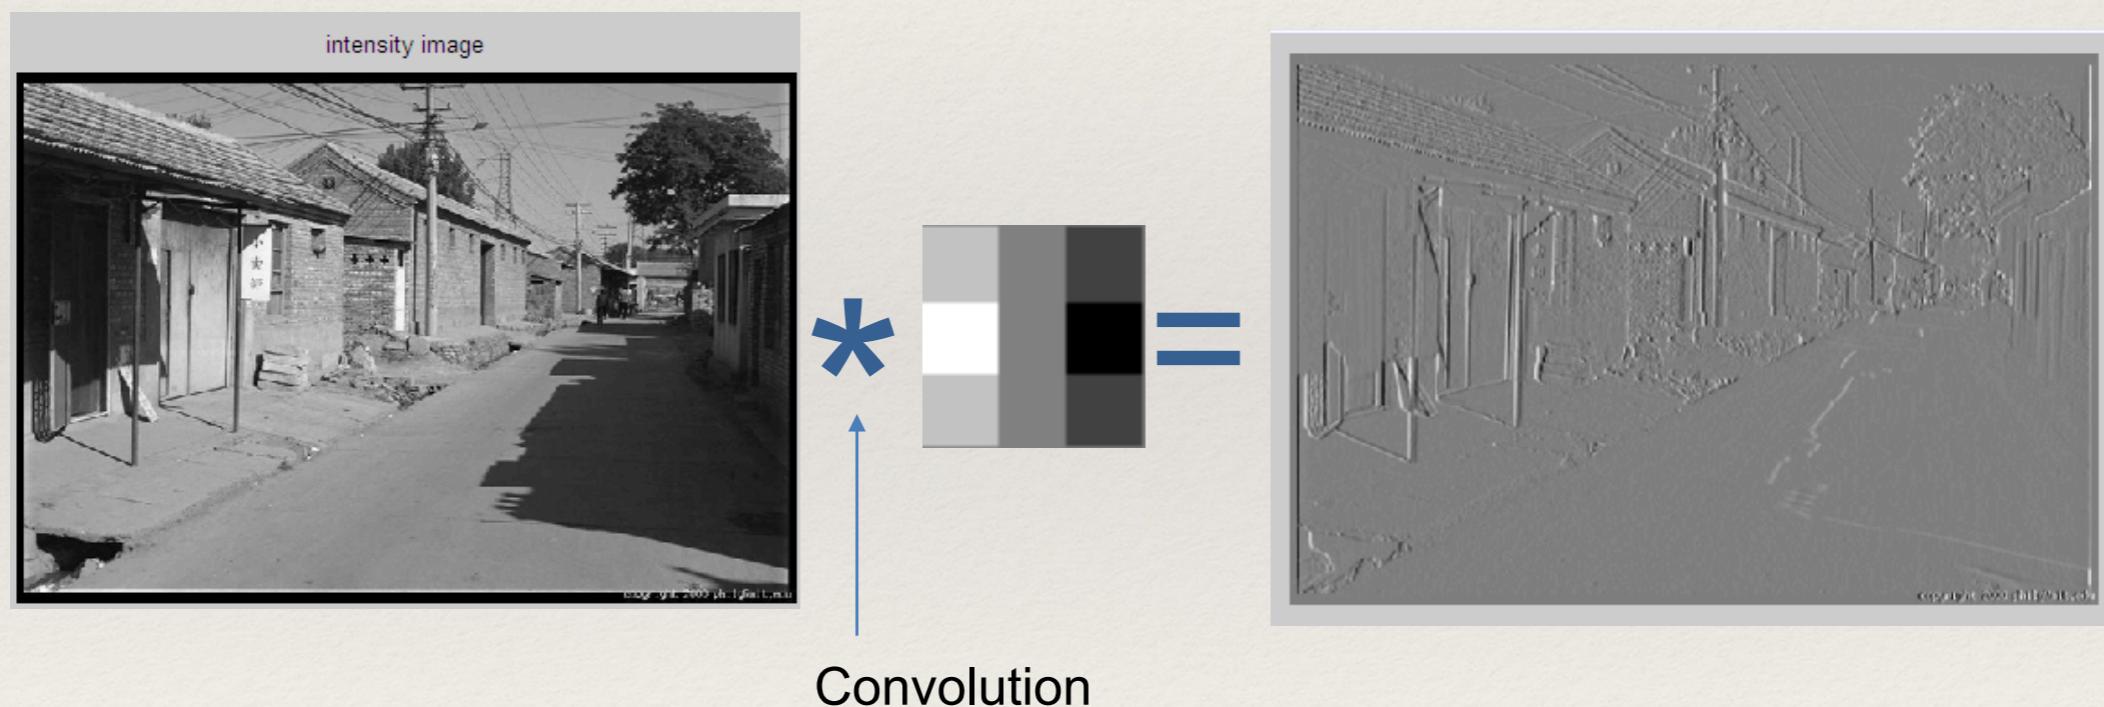
$$F[g * h] = F[g]F[h]$$

- ❖ Convolution in spatial domain is equivalent to multiplication in frequency domain

$$g * h = F^{-1}[F[g]F[h]]$$

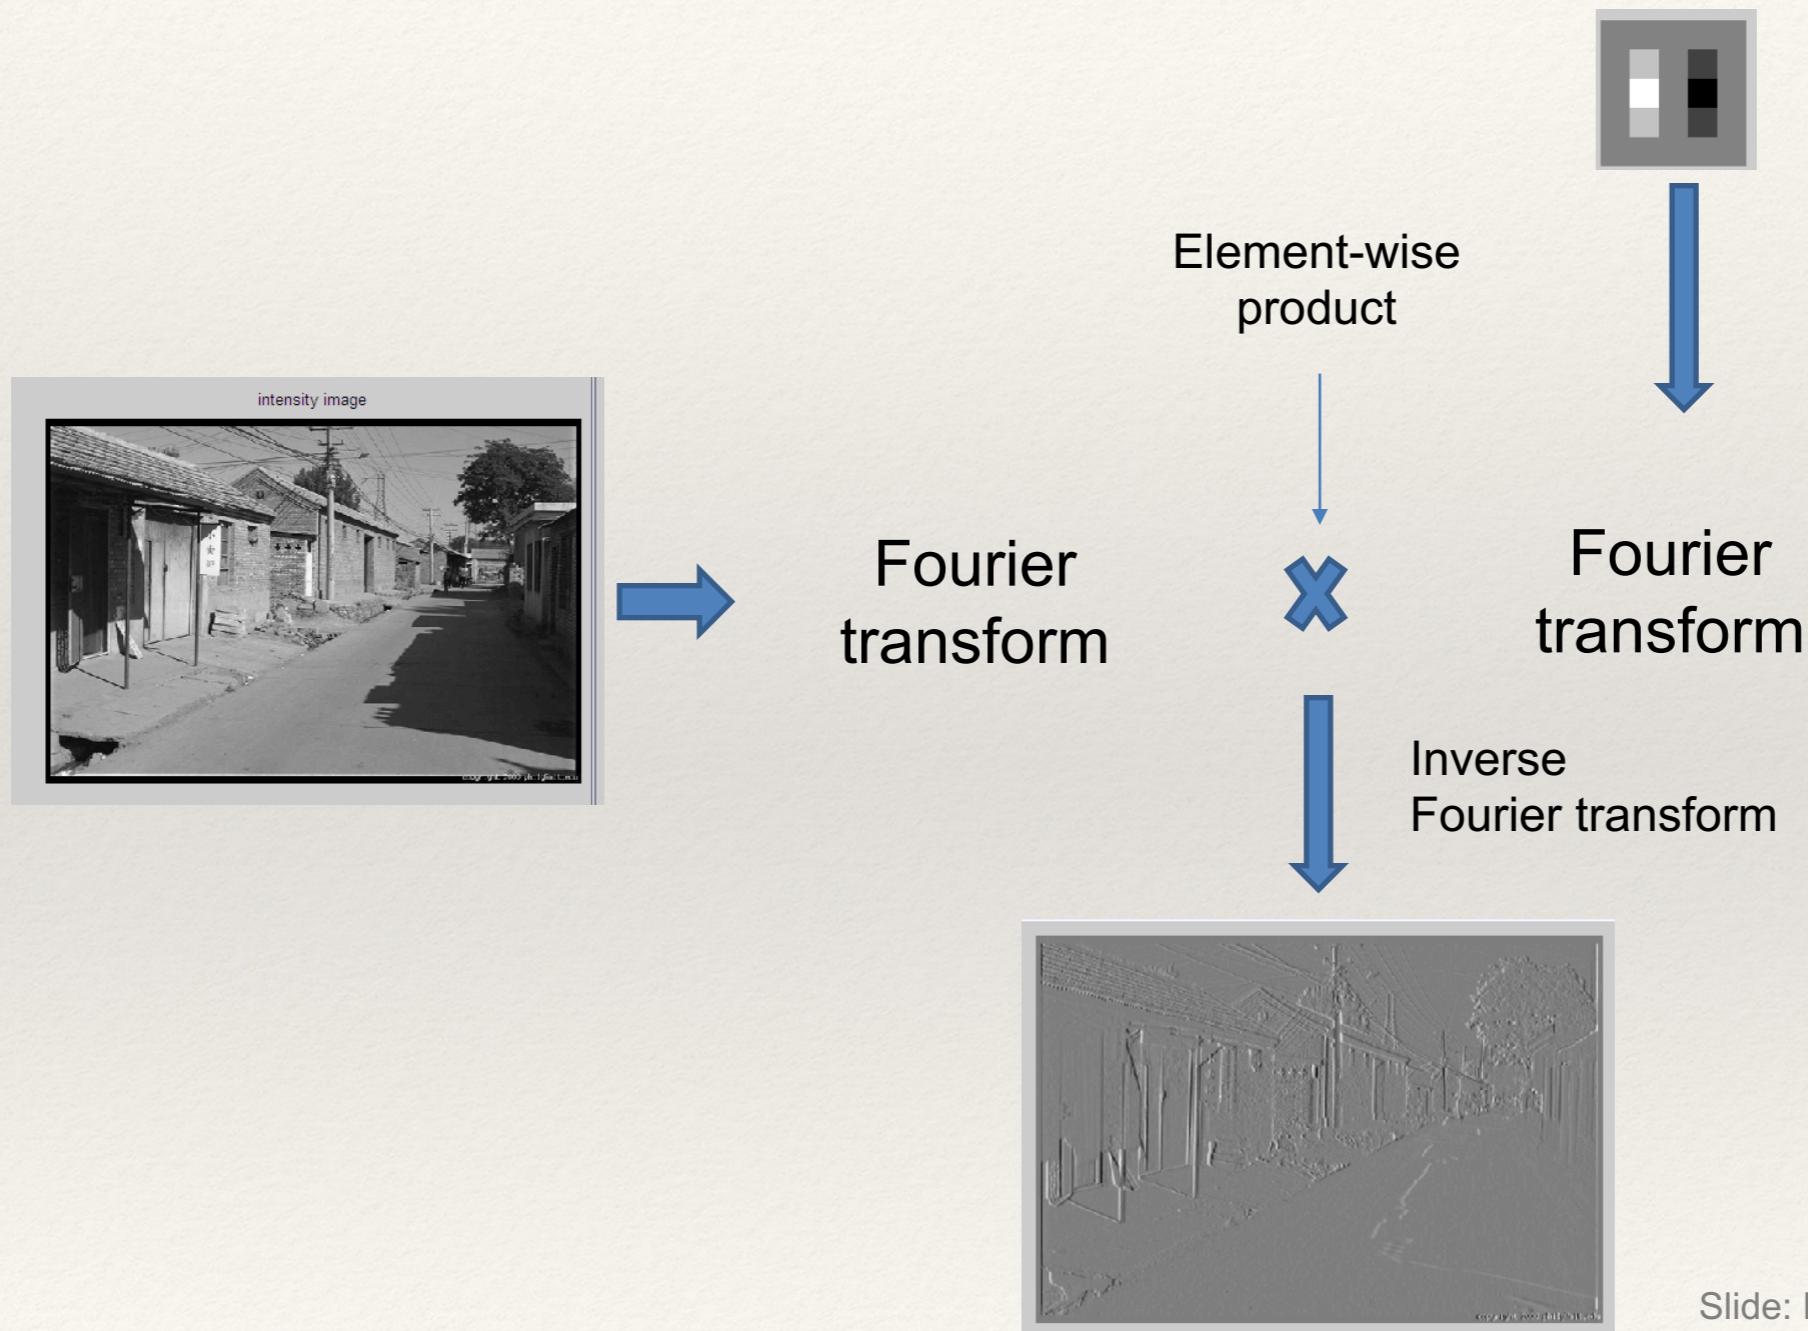
Filtering in spacial domain

1	0	-1
2	0	-2
1	0	-1



Hays

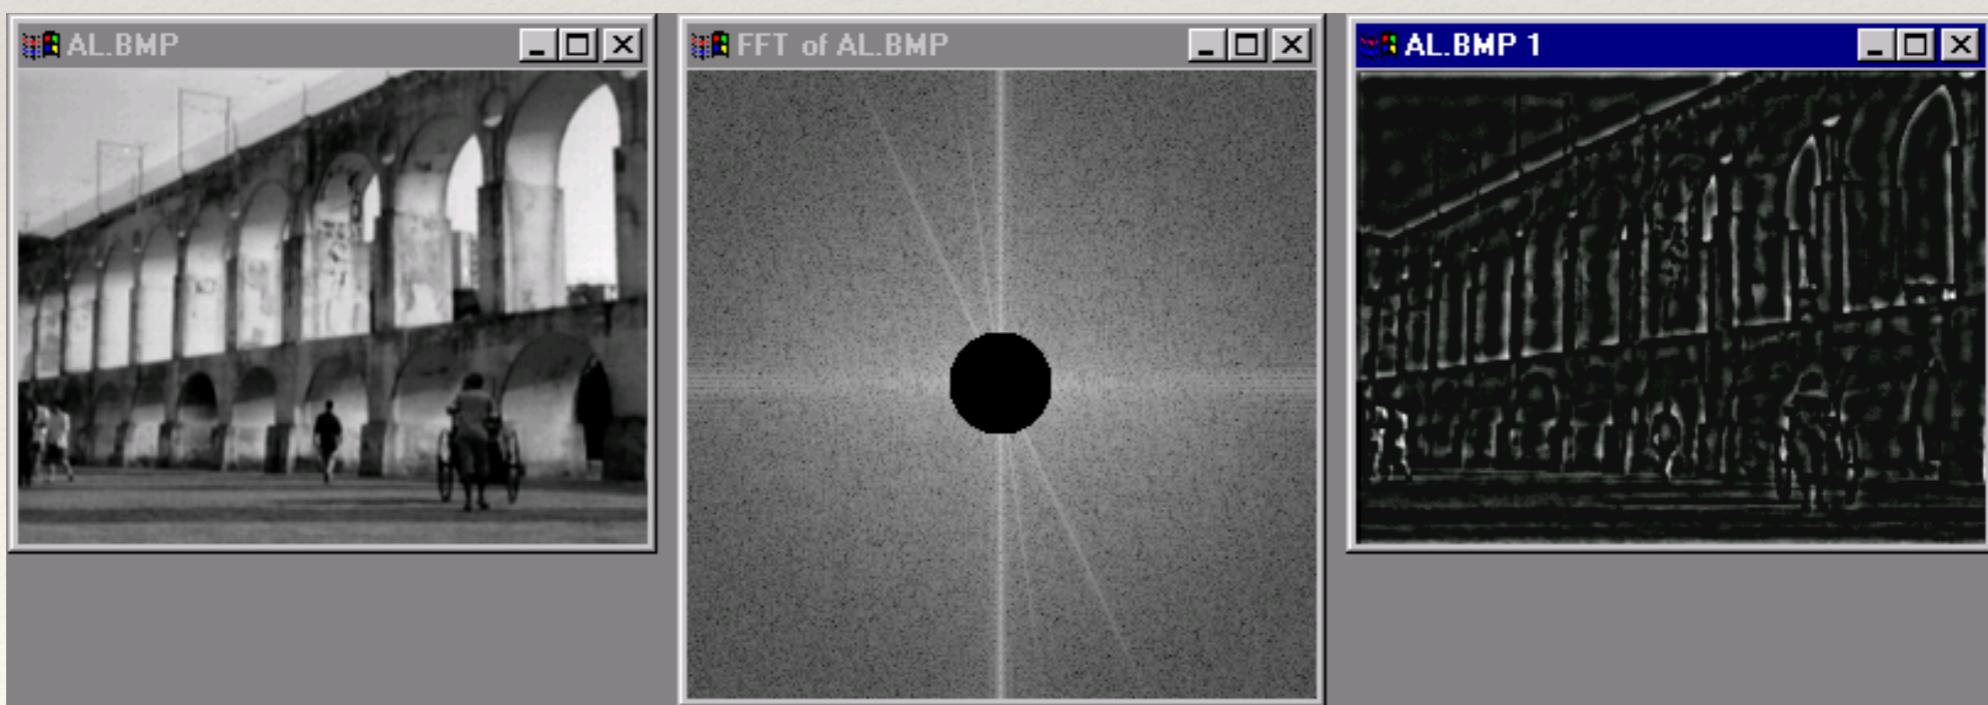
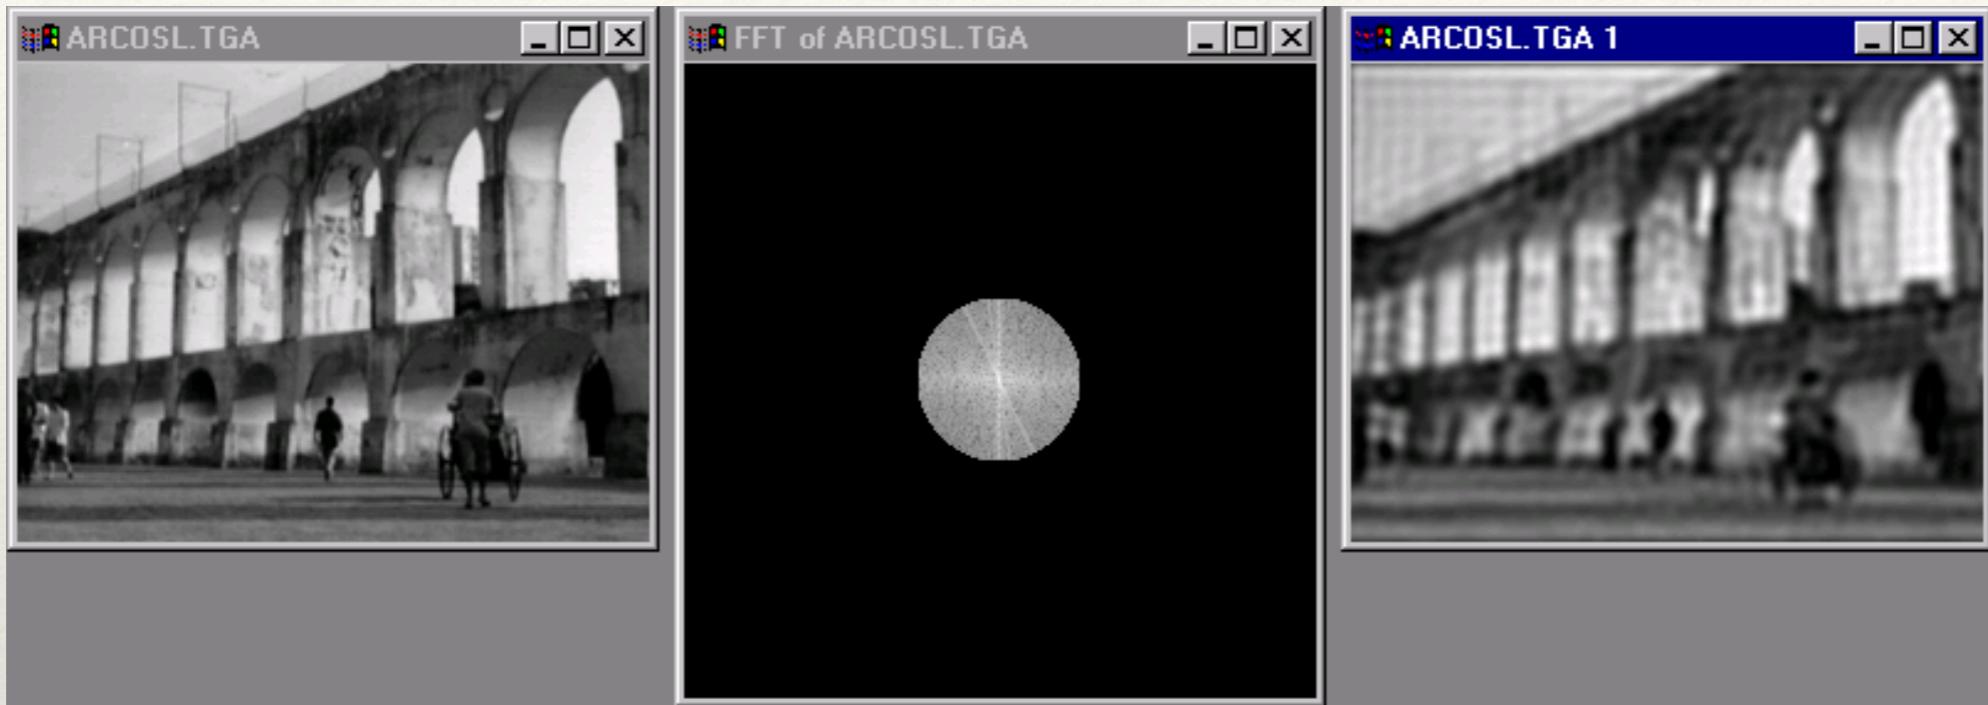
Filtering in frequency domain



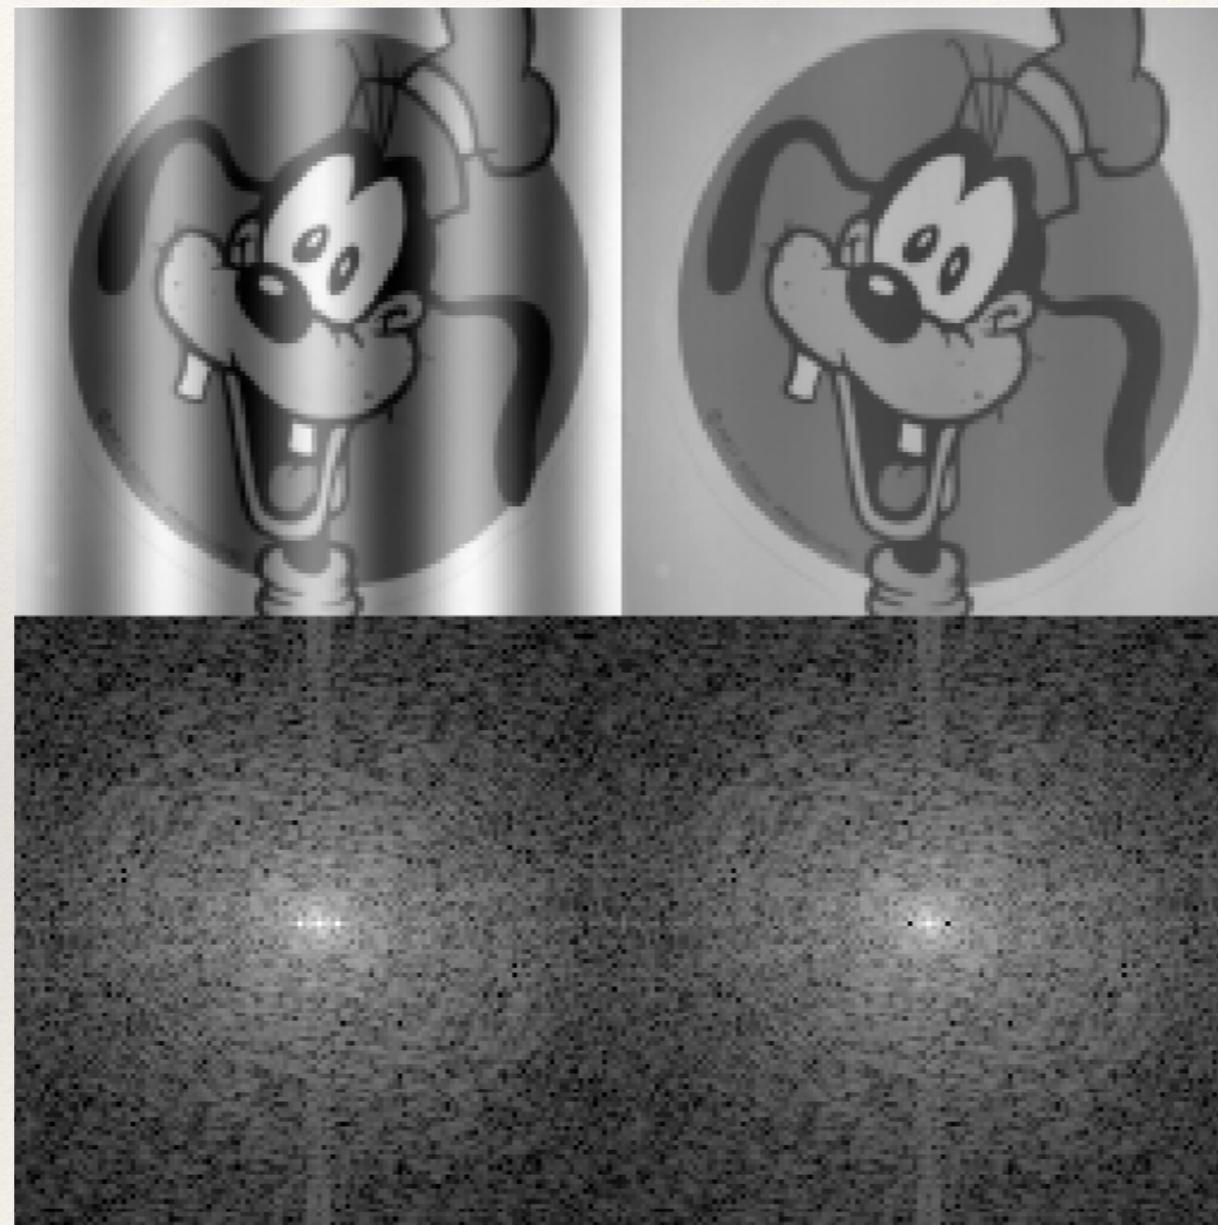
Frequency manipulation



Low and high pass filtering



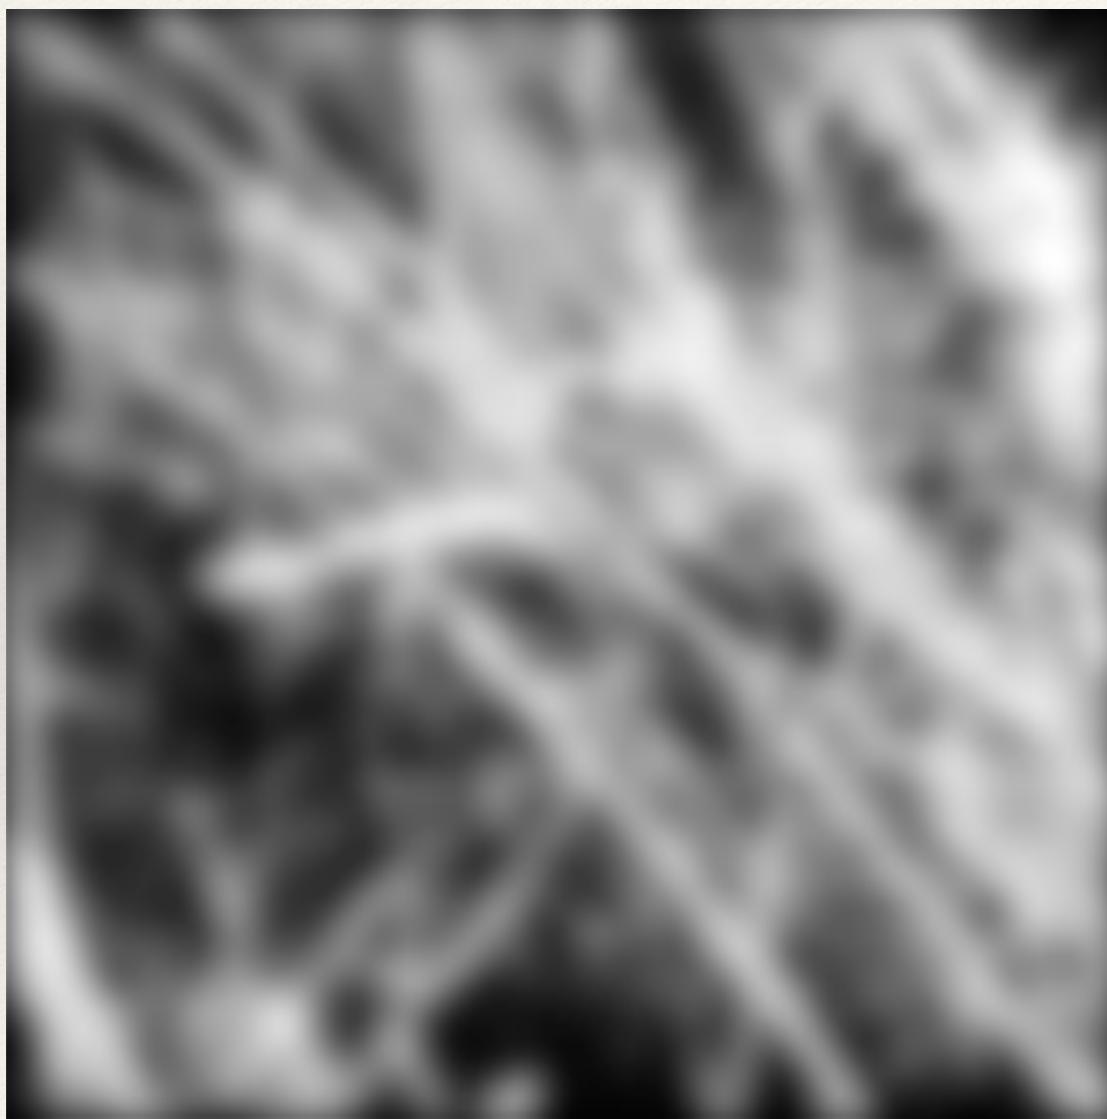
Removing frequency bands



Brayer

Why does the Gaussian filter give a nice smooth image, but
the square filter give edgy artifacts?

Gaussian

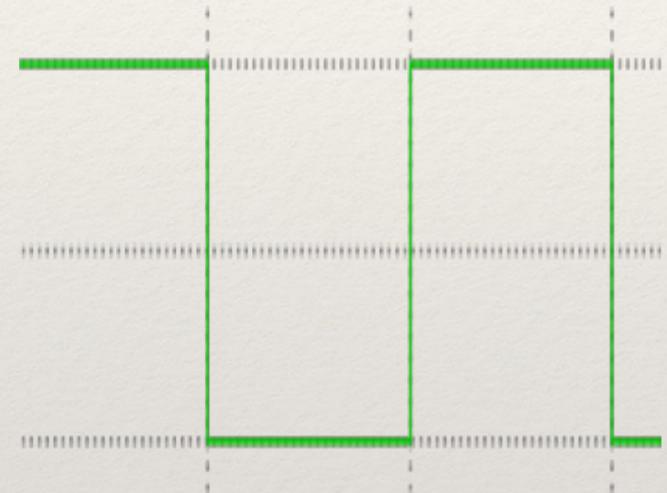


Box filter

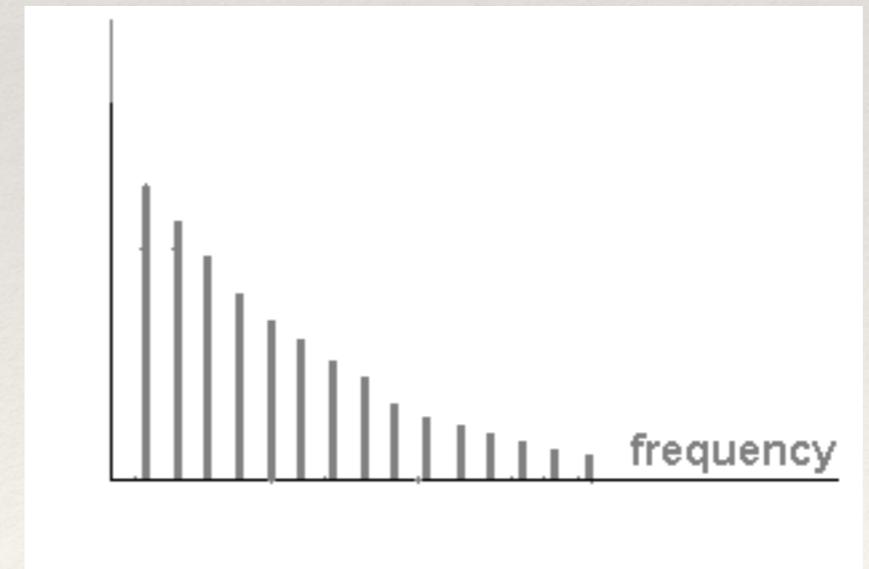


Why do we have those lines in the image?

- ❖ Sharp edges in the image need all frequencies to represent them.



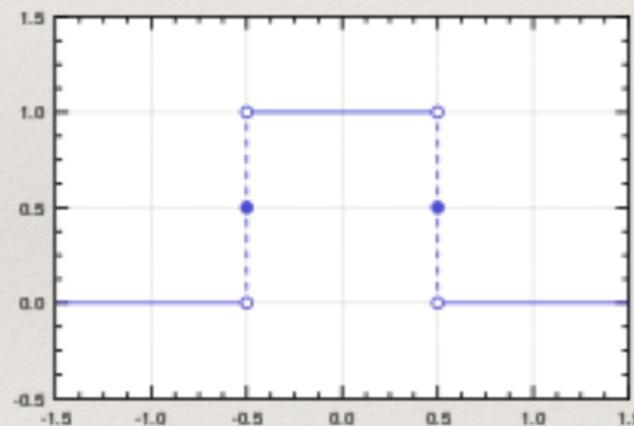
$$= A \sum_{k=1}^{\infty} \frac{1}{k} \sin(2\pi kt)$$



Box filter / sinc filter duality

- ❖ What is the spatial representation of the hard cutoff (box) in the frequency domain?
- ❖ <http://madebyevan.com/dft/>

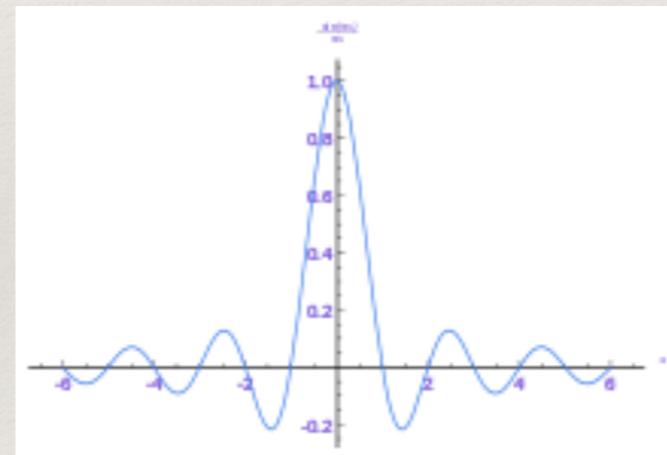
Box filter



Spatial Domain

Frequency Domain

Sinc filter



$$\text{sinc}(x) = \sin(x) / x$$

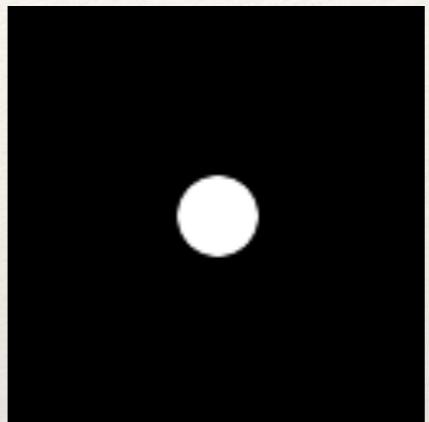
Frequency Domain

Spatial Domain

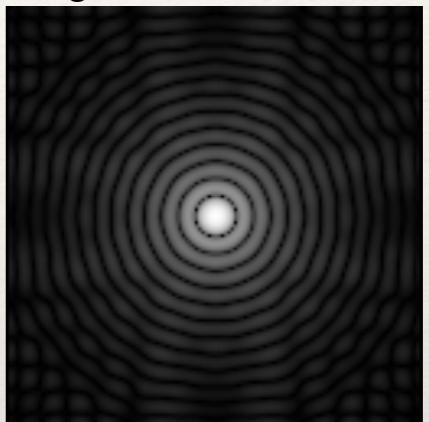
Gaussian filter duality

- ❖ Fourier transform of one Gaussian is another Gaussian (with inverse variance).
- ❖ Why is this useful?
 - ❖ Smooth degradation in frequency components
 - ❖ No sharp cut-off
 - ❖ No negative values
 - ❖ Never zero (infinite extent)

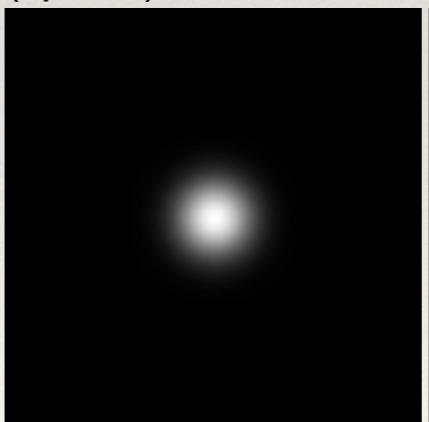
Box filter (spatial)



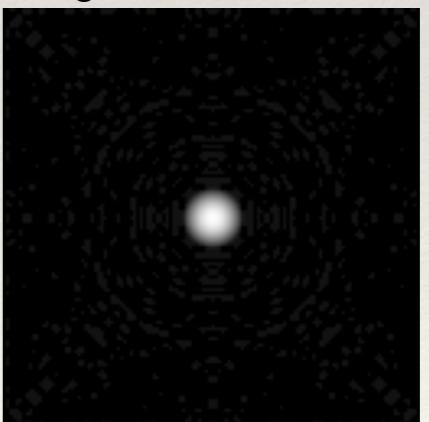
Frequency domain magnitude



Gaussian filter (spatial)



Frequency domain magnitude

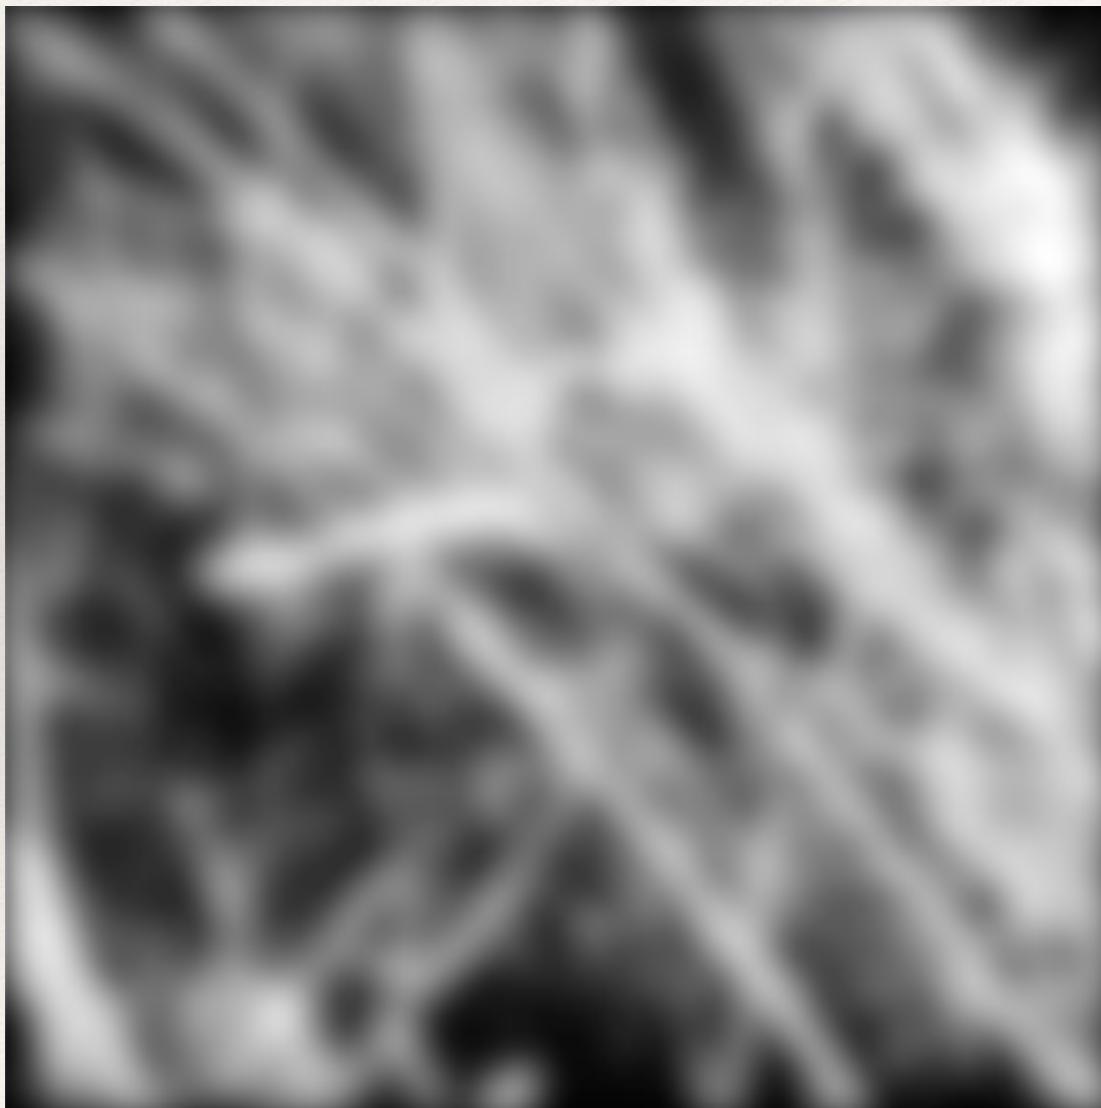


Source: James Tompkin

Ringing artifacts -> 'Gibbs effect'

Where *infinite* series can never be reached

Gaussian



Box filter

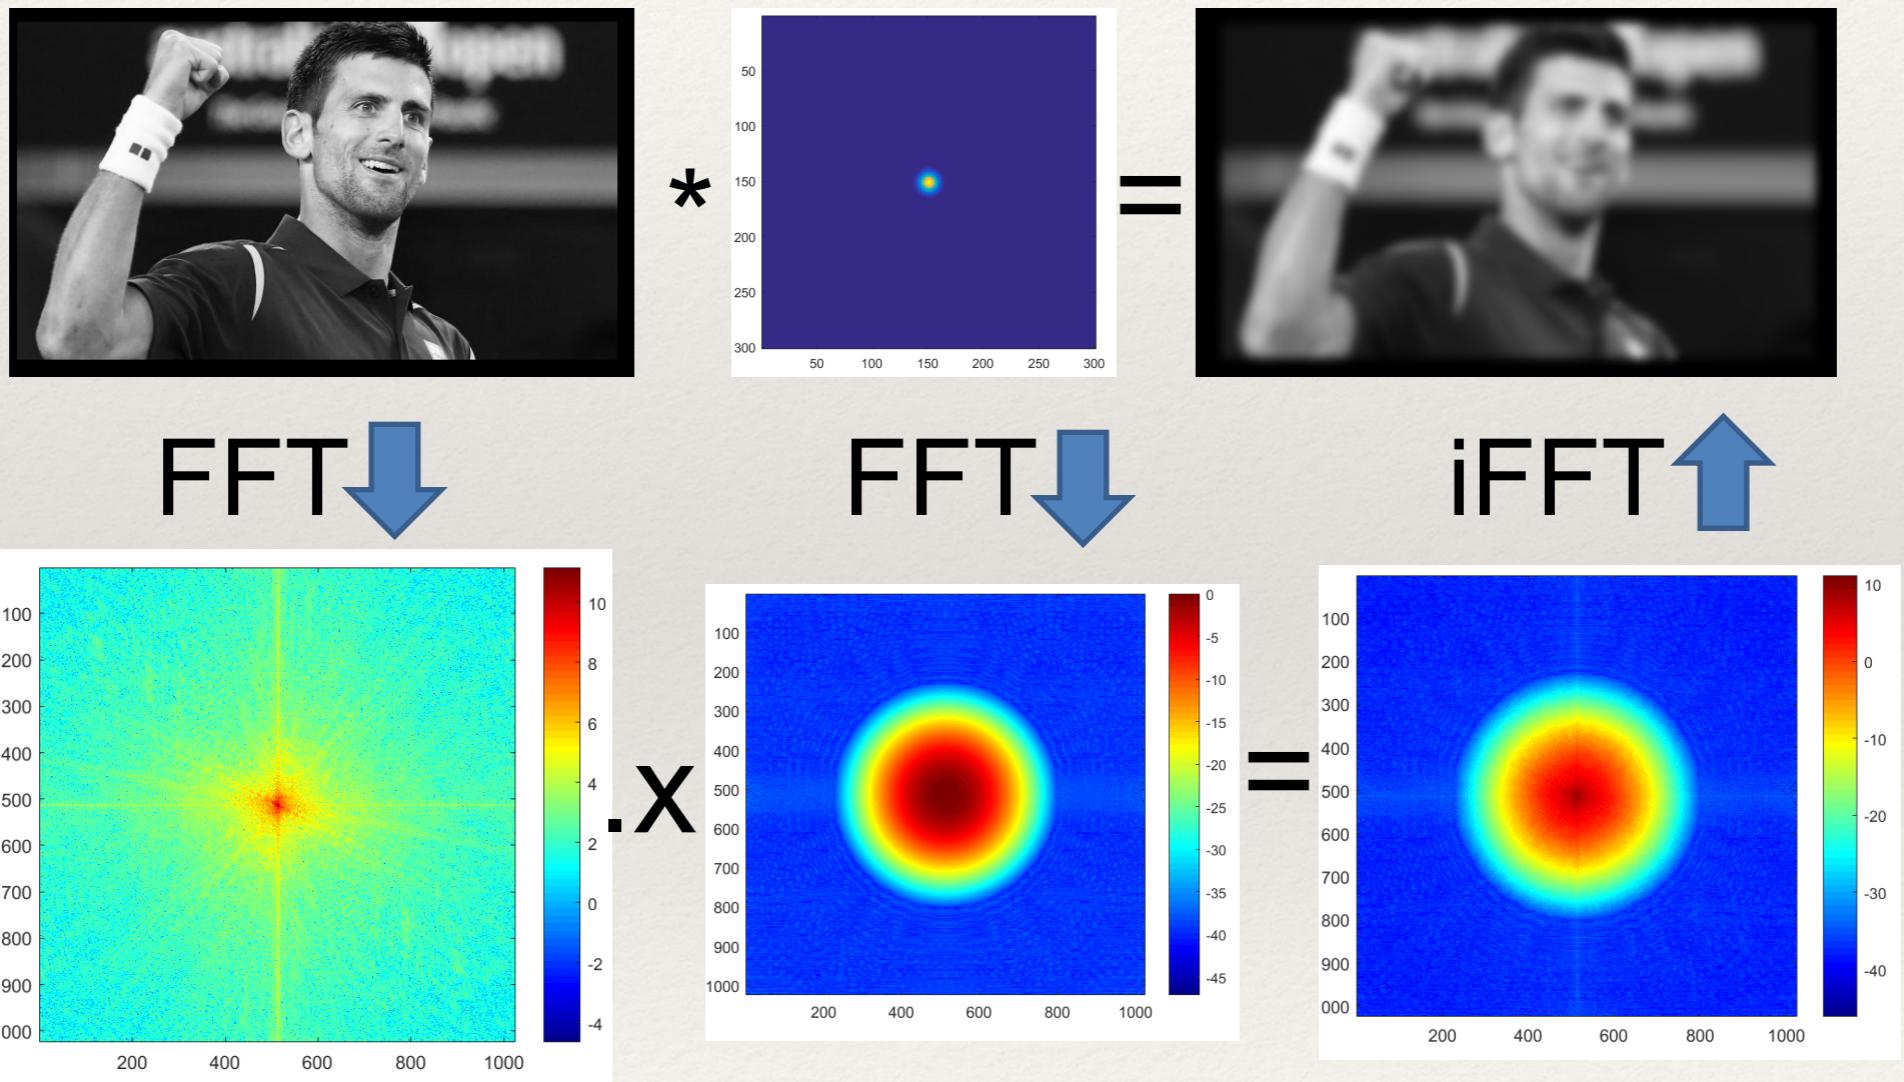


Source: James Tompkin

Is convolution invertible?

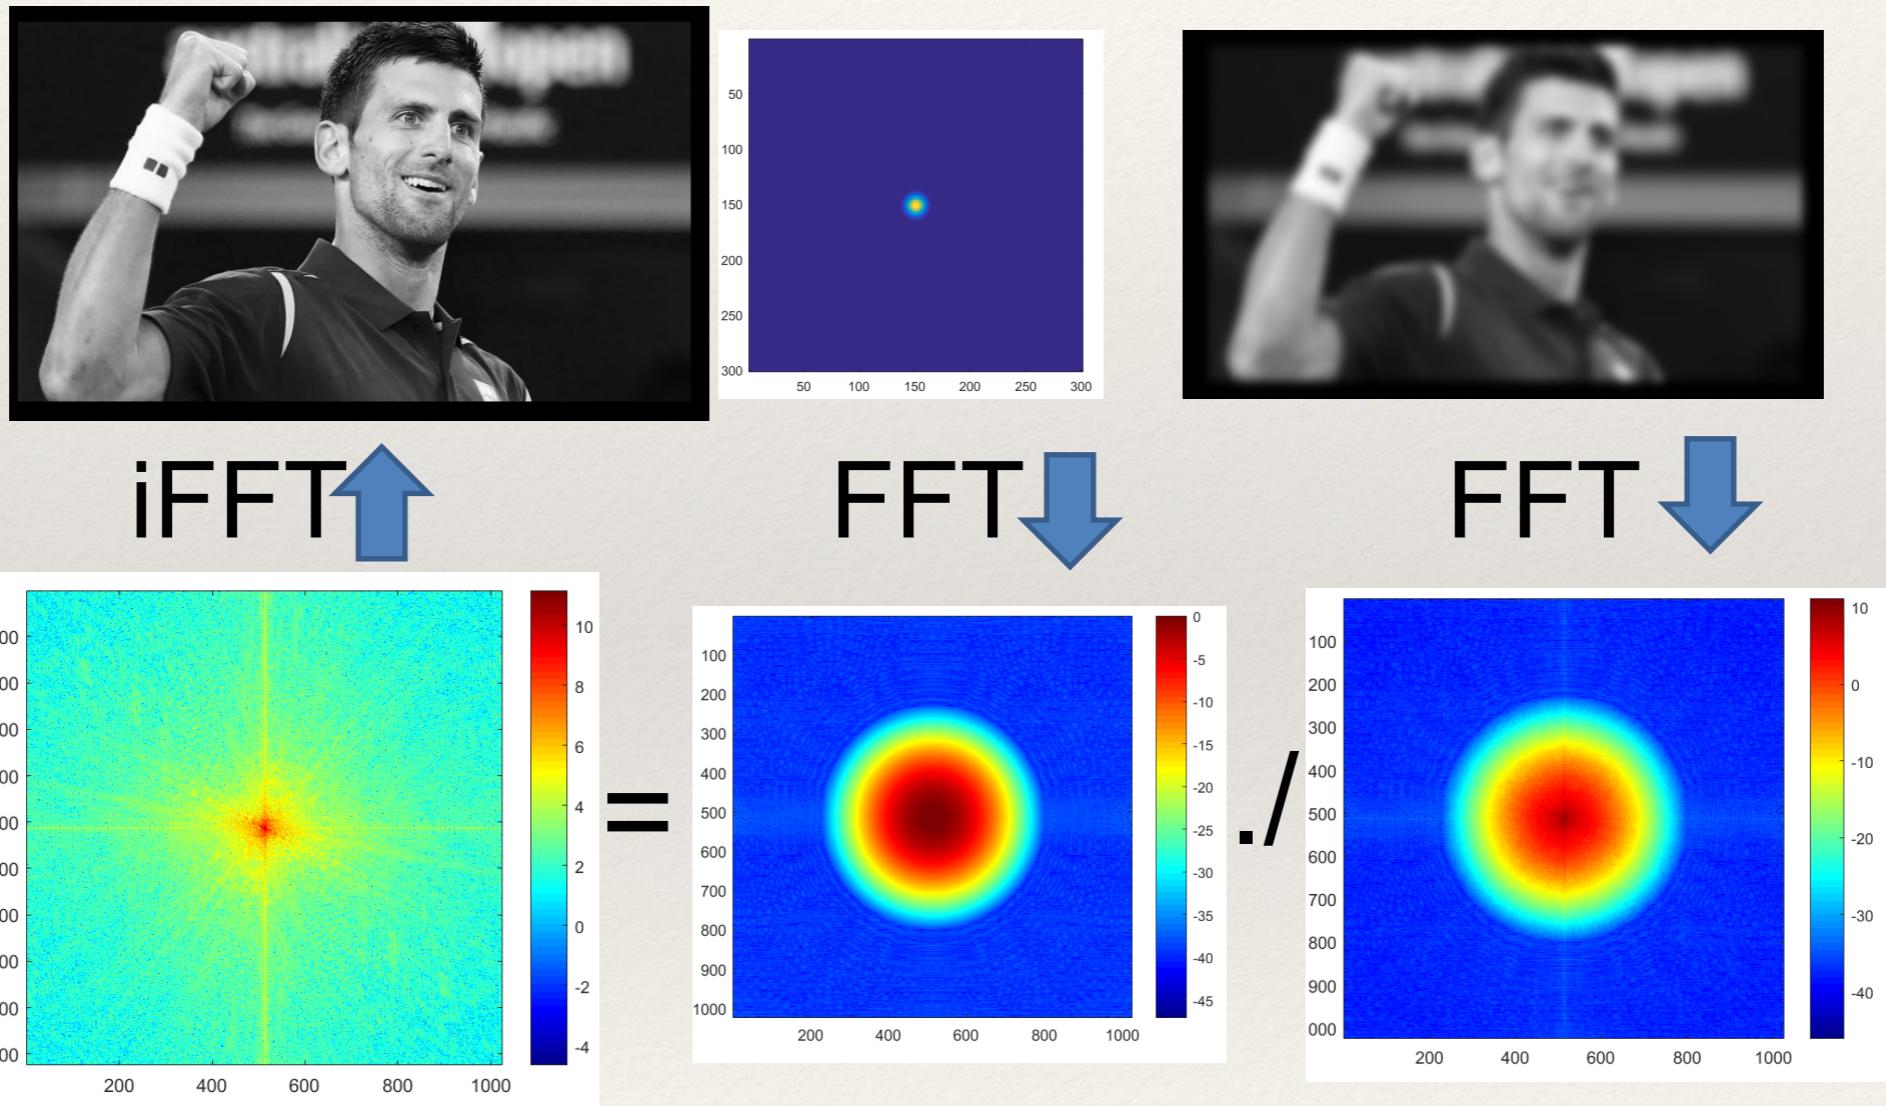
- ❖ If convolution is just multiplication in the Fourier domain, isn't deconvolution just division?
- ❖ Sometimes, it clearly is invertible (e.g. a convolution with an identity filter)
- ❖ In one case, it clearly isn't invertible (e.g. convolution with an all zero filter)
- ❖ What about for common filters like a Gaussian?

Convolution

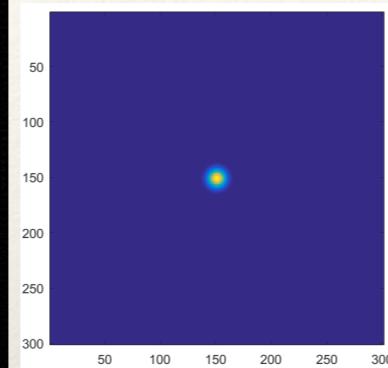


Hays

Deconvolution



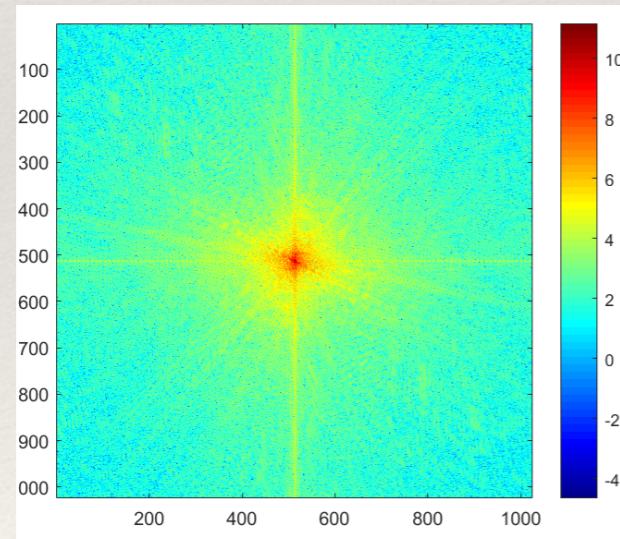
But under more realistic conditions



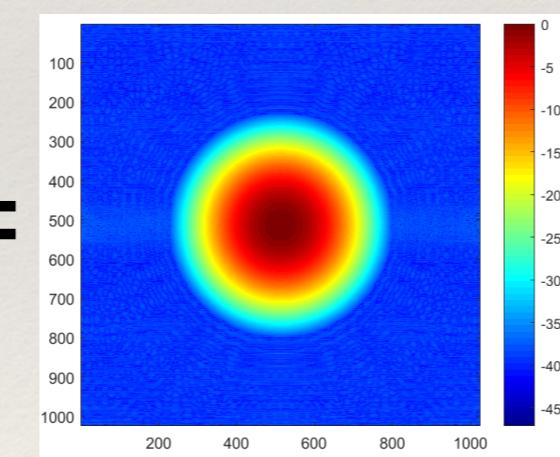
Random noise, .000001 magnitude



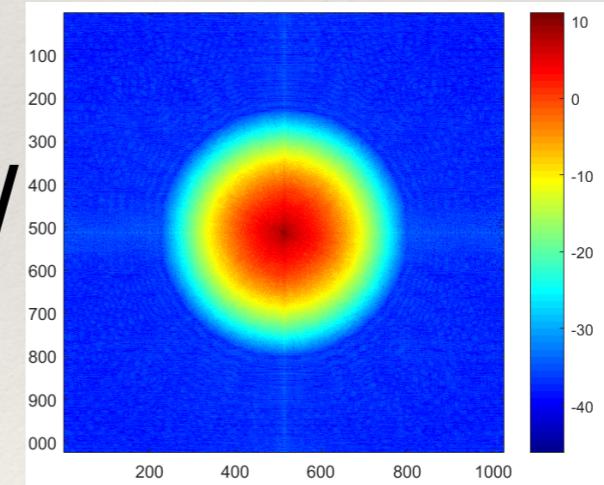
iFFT



FFT

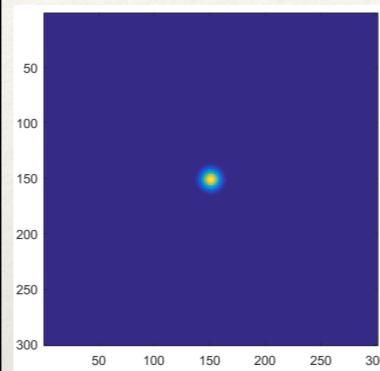


FFT



Hays

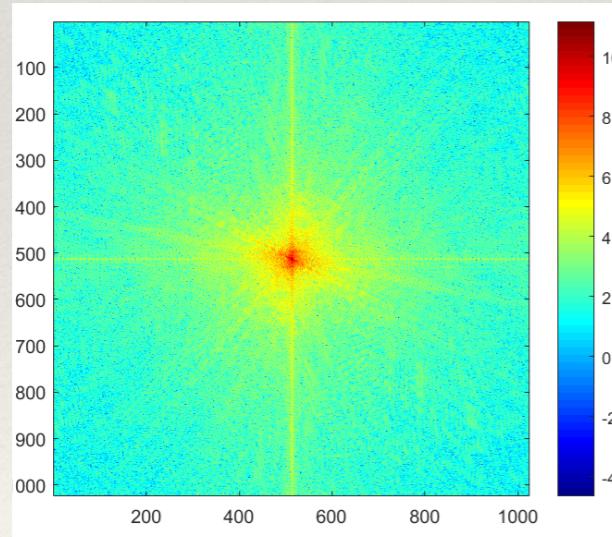
But under more realistic conditions



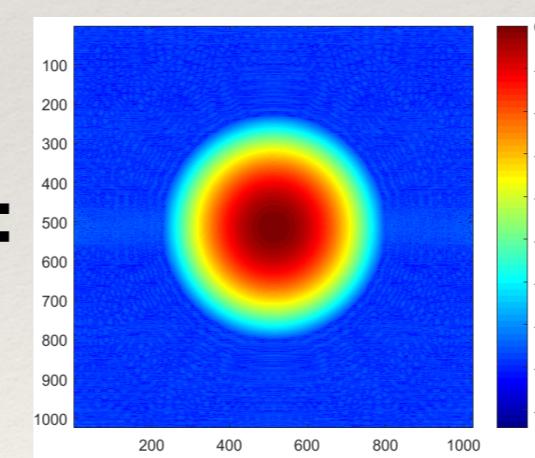
Random noise, .0001 magnitude



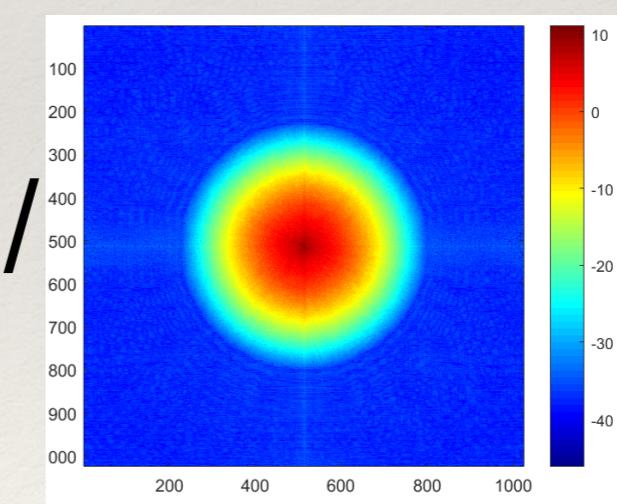
iFFT



FFT

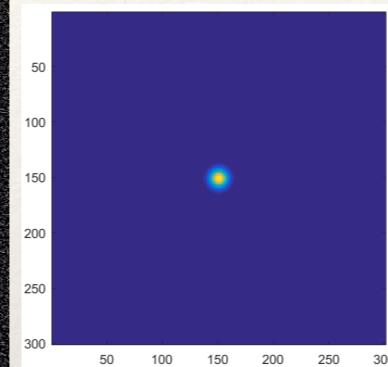


FFT



Hays

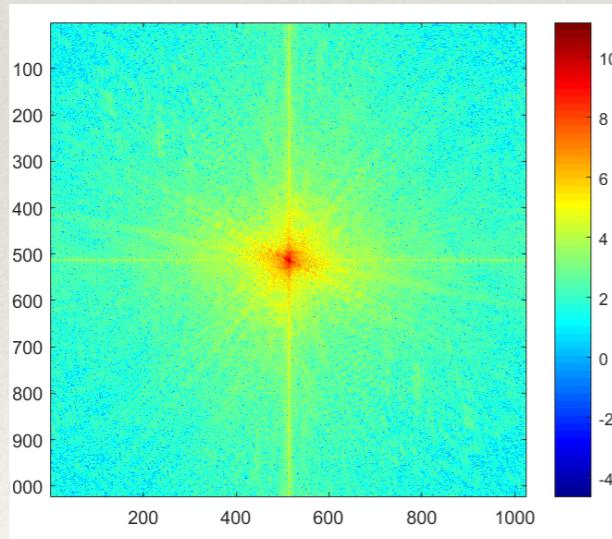
But under more realistic conditions



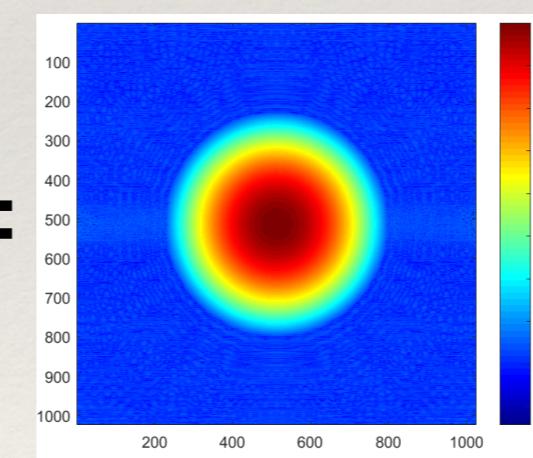
Random noise, .001 magnitude



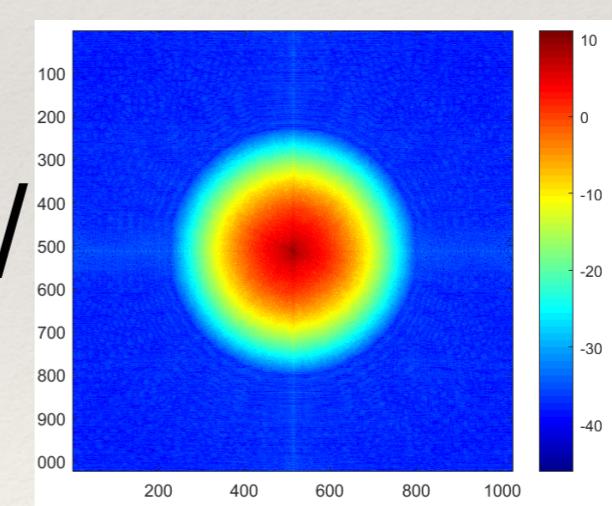
iFFT



FFT



FFT



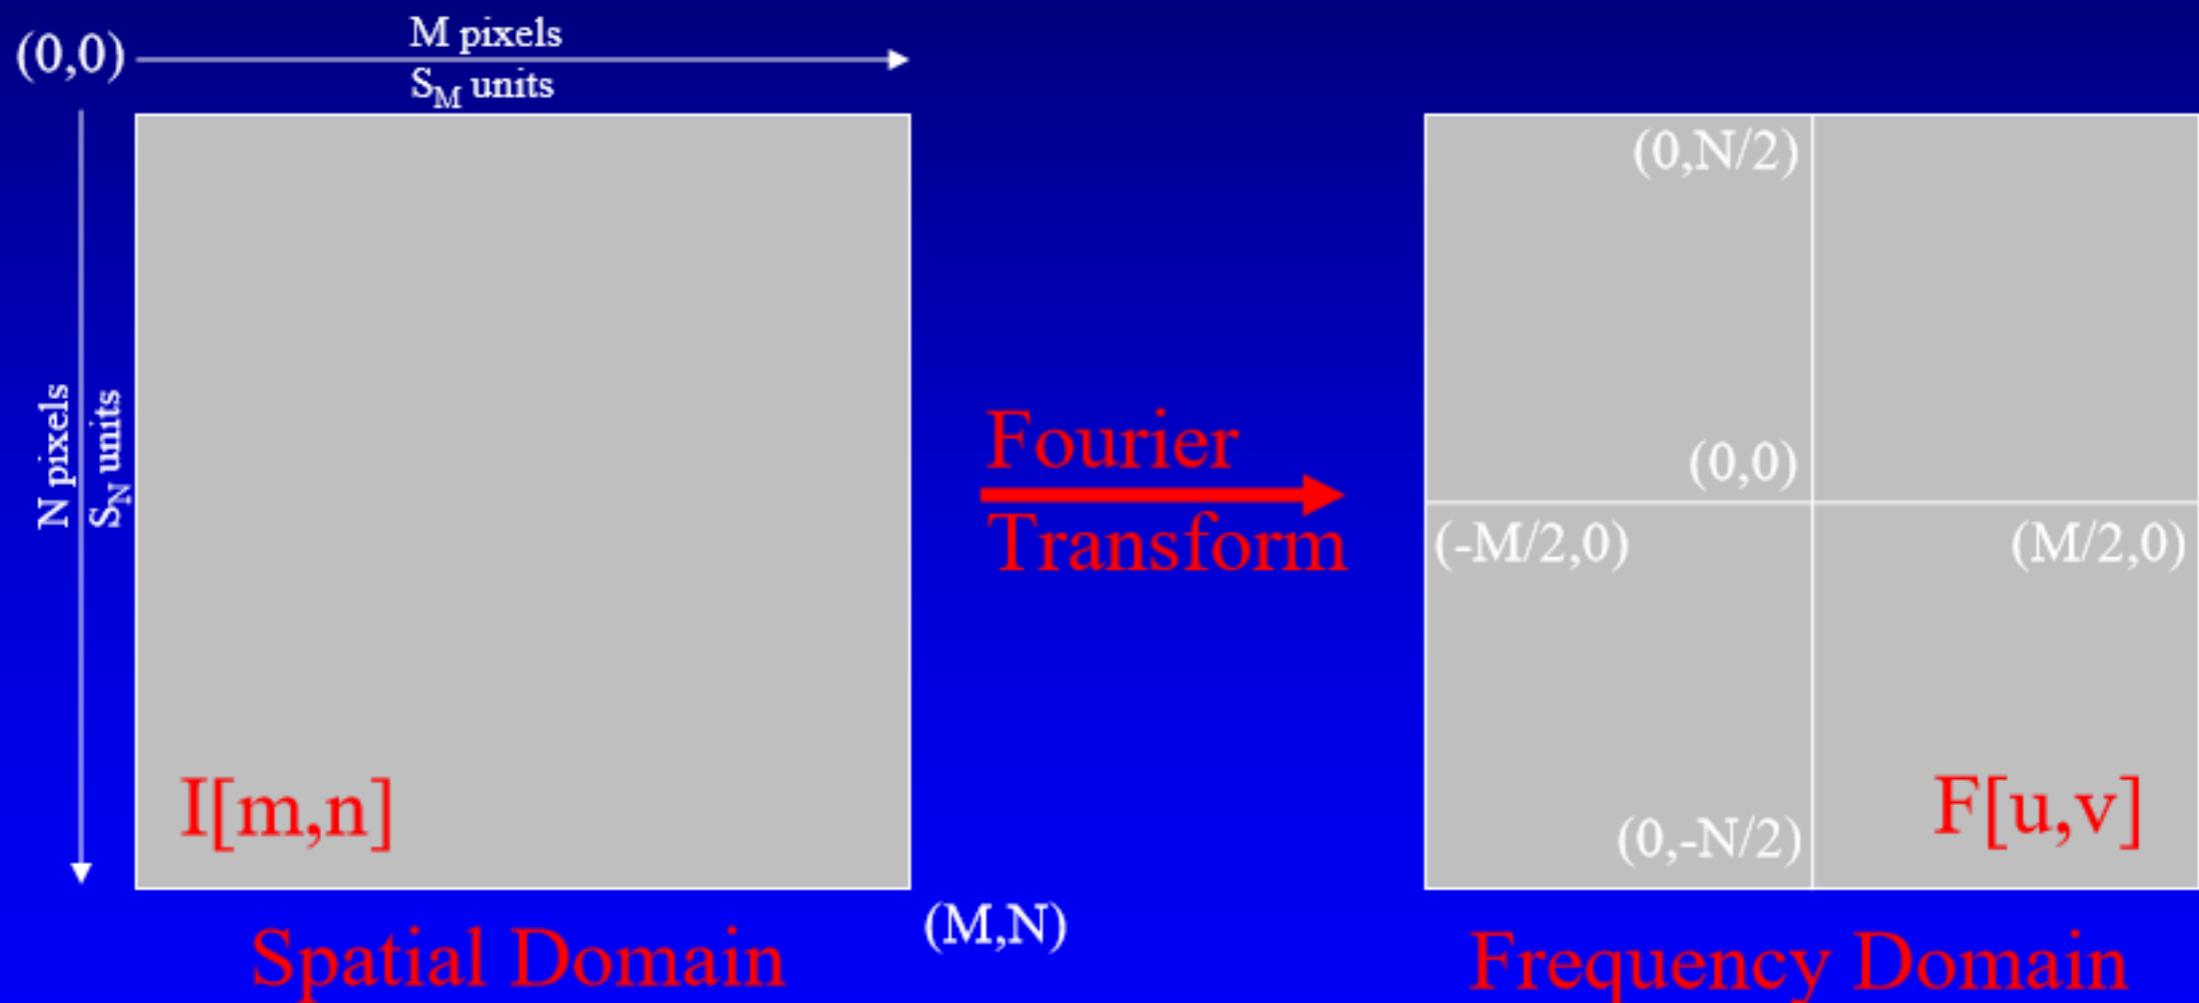
Hays

Deconvolution is hard.

- ❖ Active research area.
- ❖ Even if you know the filter (non-blind deconvolution), it is still hard and requires strong *regularization* to counteract noise.
- ❖ If you don't know the filter (blind deconvolution), then it is harder still.

2D Discrete Fourier Transform

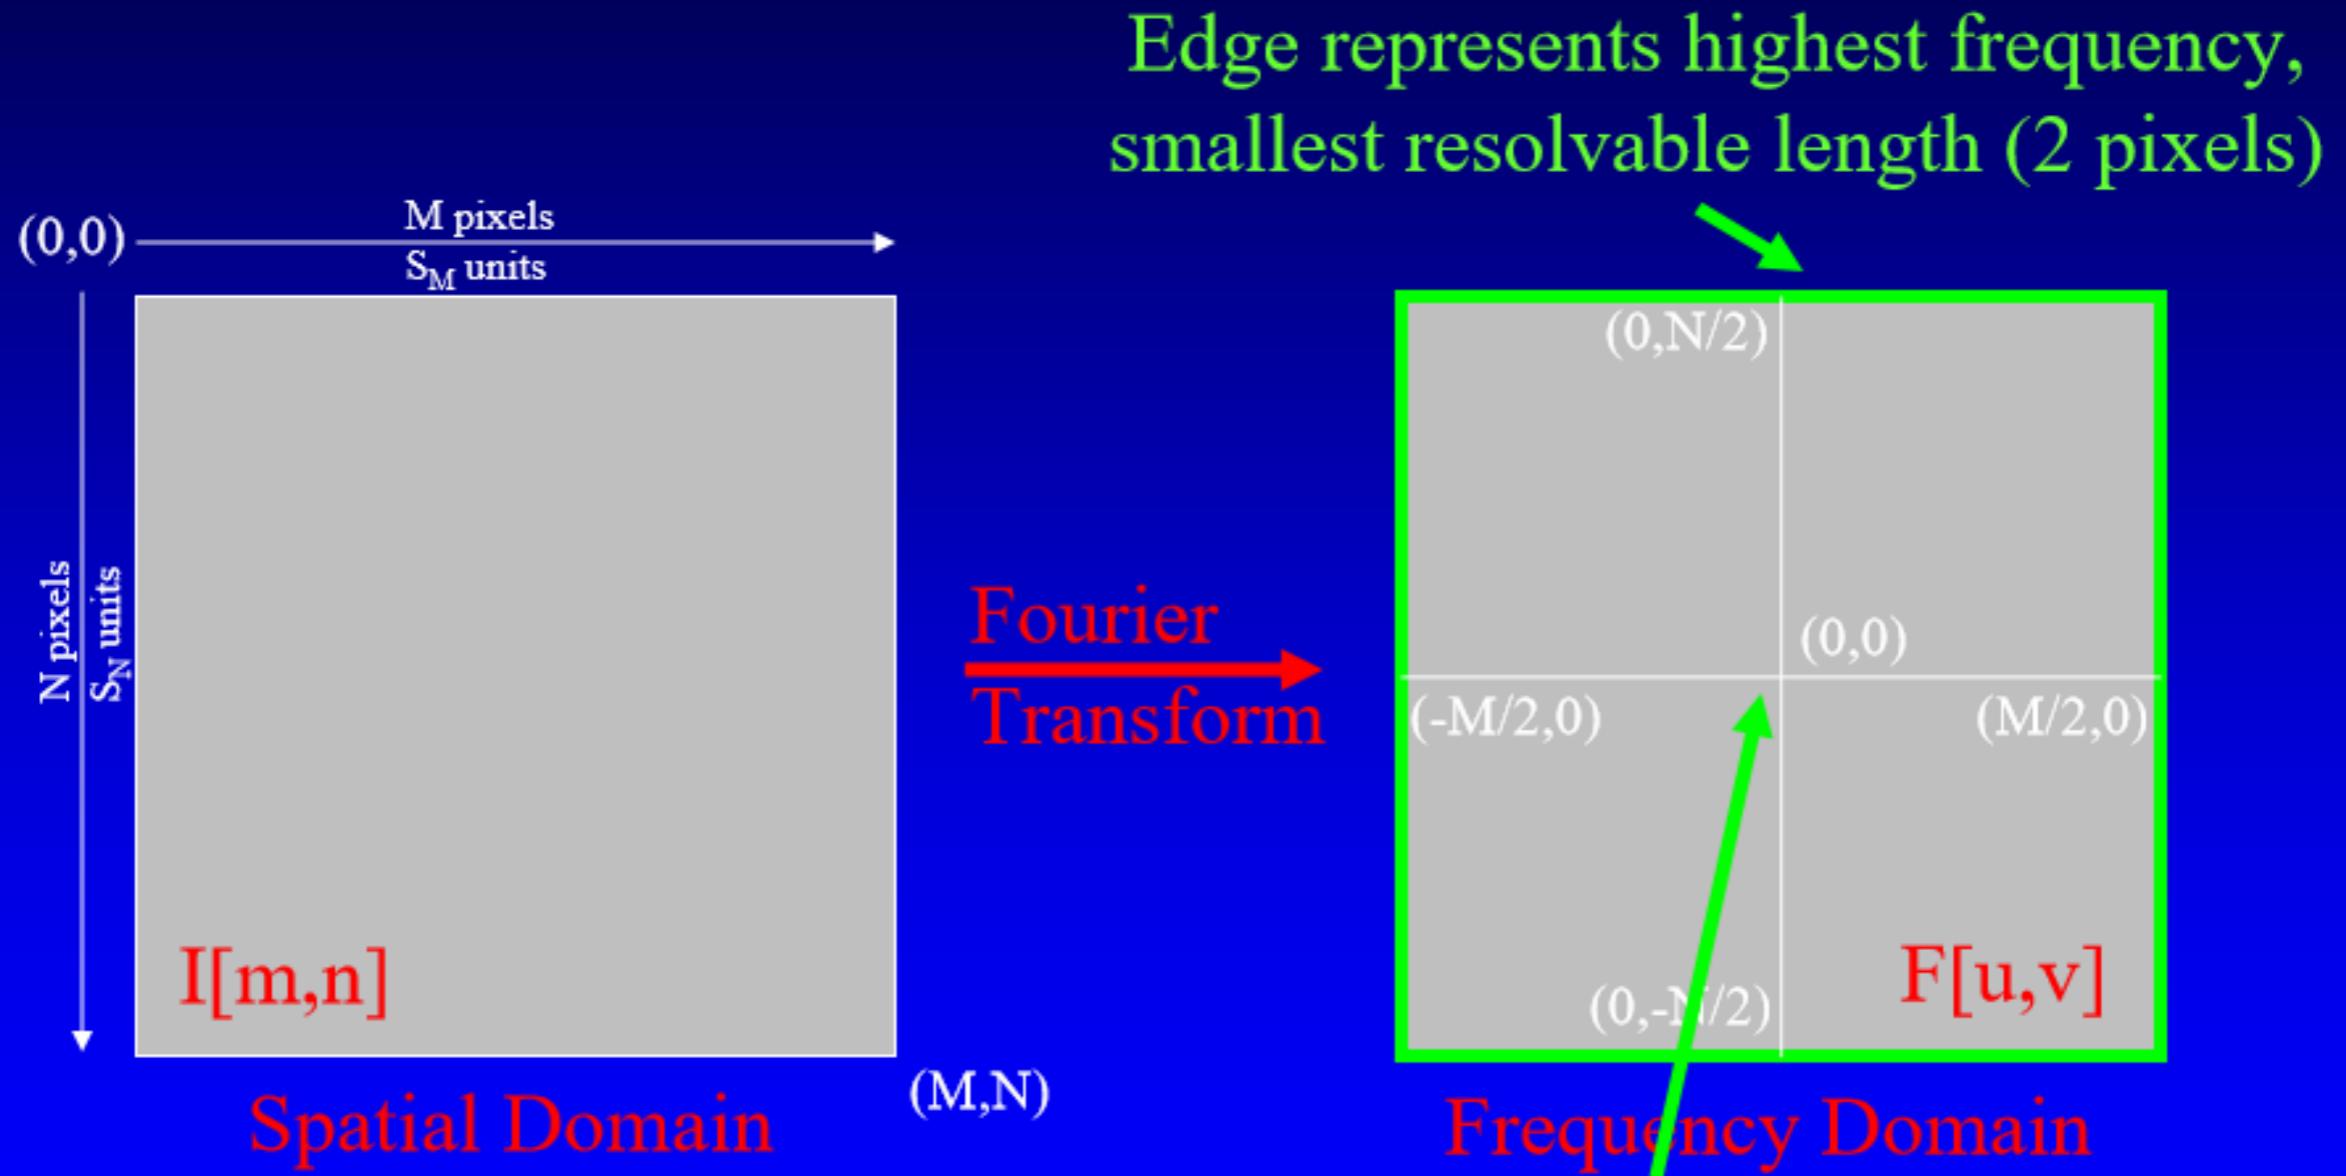
$$F[u,v] = \frac{1}{MN} \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} I[m,n] \cdot e^{-i2\pi \left(\frac{um}{M} + \frac{vn}{N} \right)}$$



Source: Seul et al, *Practical Algorithms for Image Analysis*, 2000, p. 249, 262.

2D FFT can be computed as two discrete Fourier transforms in 1 dimension

2D Discrete Fourier Transform



Edge represents highest frequency,
smallest resolvable length (2 pixels)

Center represents lowest frequency,
which represents average pixel value