

Shanghai Jiao Tong University

Computer Vision

Instructor: Xu Zhao
Class No.: C032703 F032528

Spring 2020



Xu Zhao @ Shanghai Jiao Tong university

Lecture 9: Segmentation

Contents

- ❖ **Segmentation by clustering**
- ❖ **Deformable contours for segmentation**

Grouping in vision

- ❖ Goals:
 - ❖ Gather features that belong together
 - ❖ Obtain an intermediate representation that compactly describes key image (video) parts
- ❖ Top down vs. bottom up segmentation
 - ❖ Top down: pixels belong together because they are from the same object
 - ❖ Bottom up: pixels belong together because they look similar
- ❖ Hard to measure success
 - ❖ What is interesting depends on the app



[Figure by J. Shi]

Determine image regions



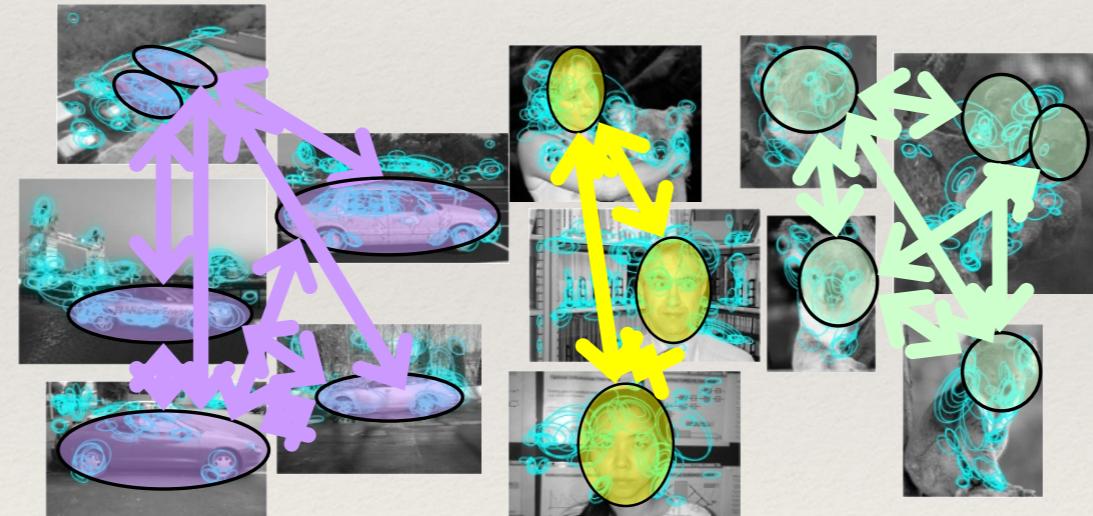
[http://poseidon.csd.auth.gr/LAB_RESEARCH/Latest/imgs/SpeakDepVidIndex_img2.jpg]

Group video frames into shots



[Figure by Wang & Suter]

Figure-ground

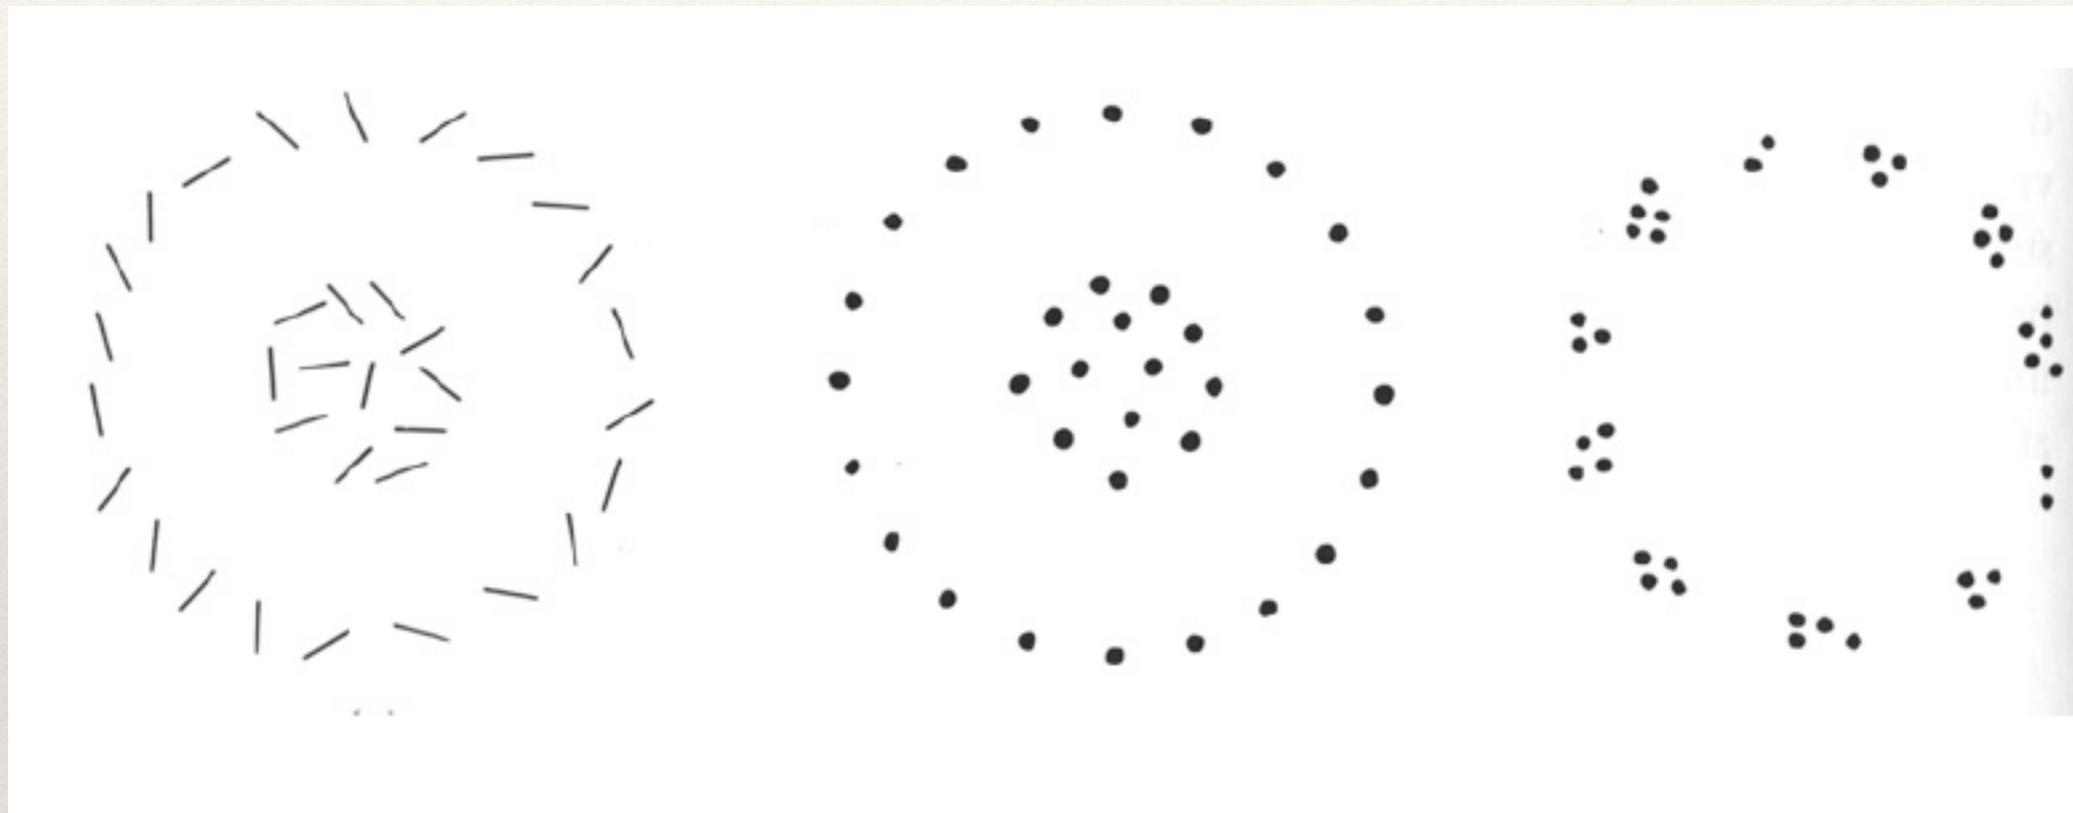


[Figure by Grauman & Darrell]

Object-level grouping

Slide credit: Kristen Grauman

What are meta-cues for grouping?



Slide credit: Kristen Grauman

Gestalt

- ❖ Gestalt: whole or group
 - ❖ Whole is greater than sum of its parts
 - ❖ Relationships among parts can yield new properties/features
 - Humans interpret image information collectively
- ❖ Psychologists identified series of factors that predispose set of elements to be grouped (by human visual system)
 - ❖ Humans interpret image information collectively



Not grouped



Proximity



Similarity



Similarity



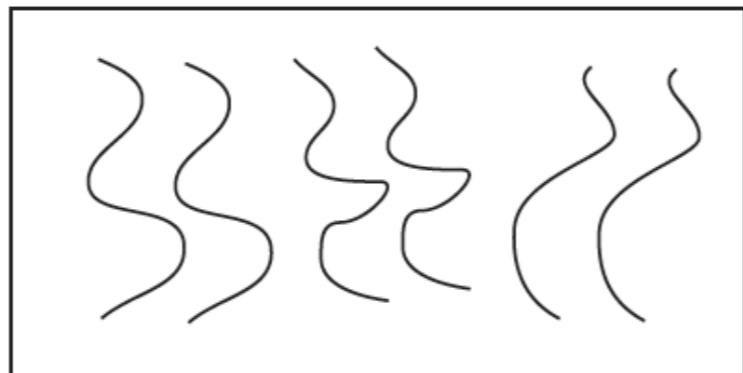
Common Fate



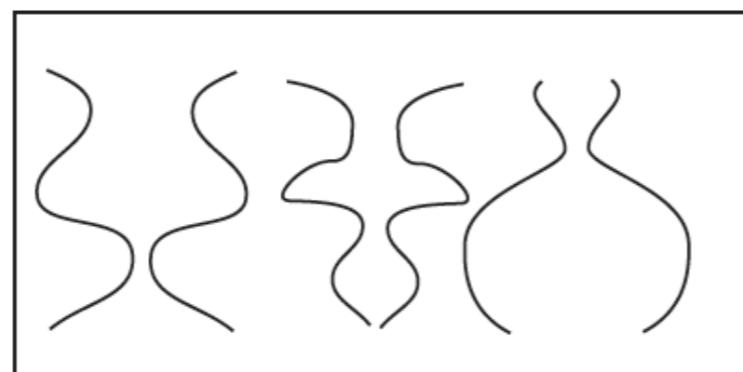
Common Region



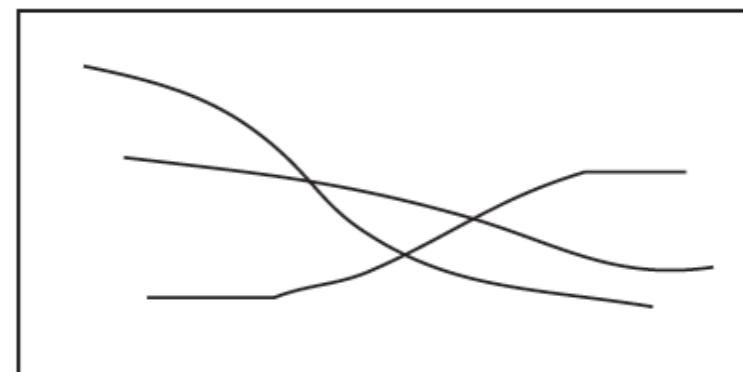
FIGURE 9.4: Examples of Gestalt factors that lead to grouping (which are described in greater detail in the text).



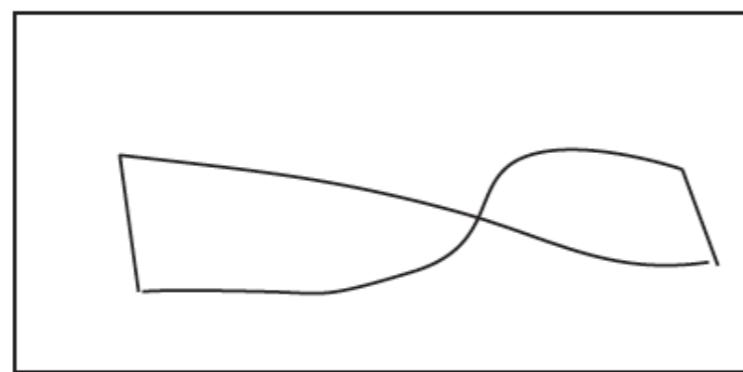
Parallelism



Symmetry



Continuity



Closure

FIGURE 9.5: Examples of Gestalt factors that lead to grouping (which are described in greater detail in the text).

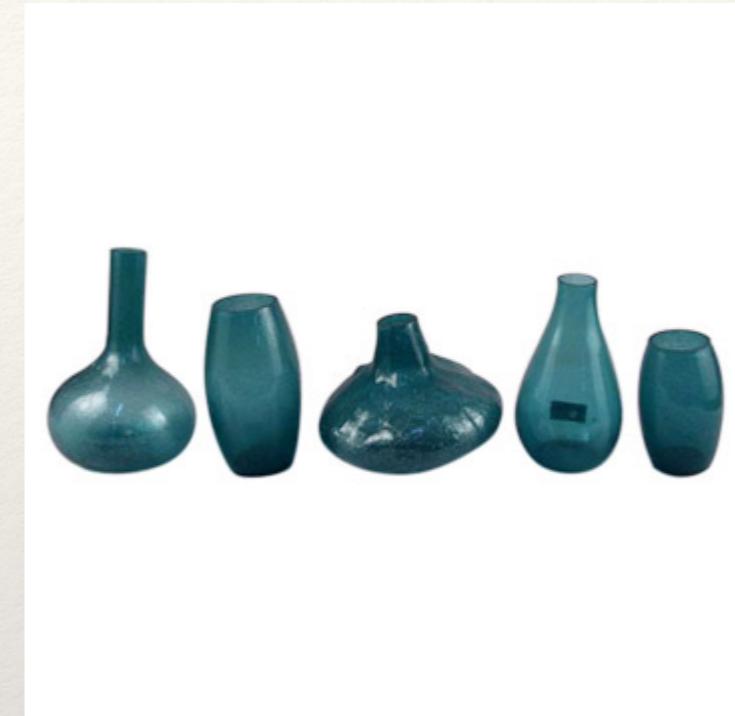
Similarity



Slide credit: Kristen Grauman

http://chicagoist.com/attachments/chicagoist_alicia/GEESE.jpg, http://www.delivery.superstock.com/WI/223/1532/PreviewComp/SuperStock_1532R-0831.jpg

Symmetry



Slide credit: Kristen Grauman

Common fate

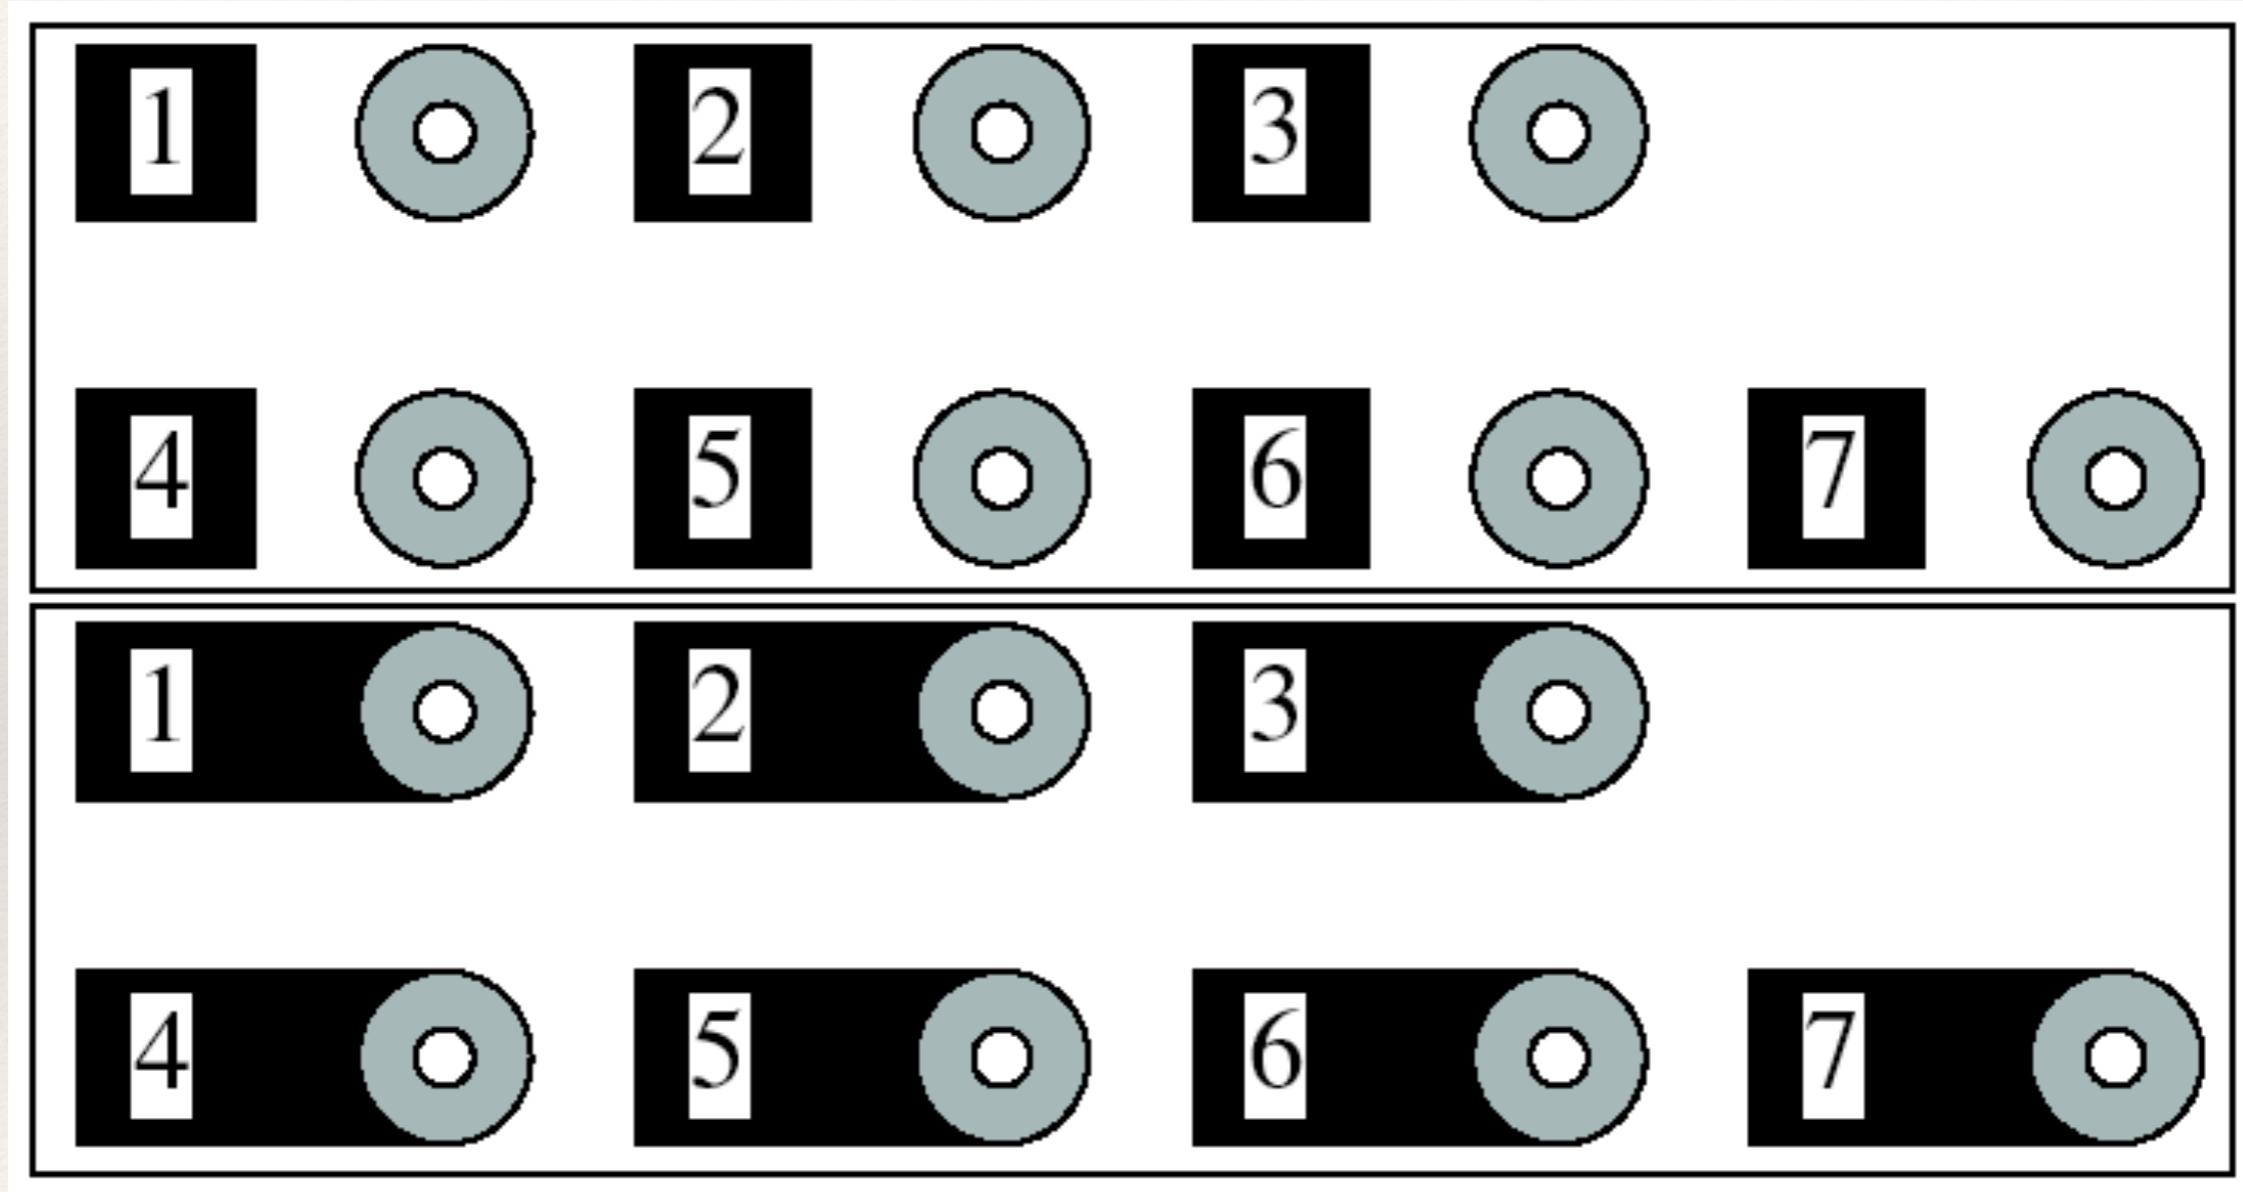


Image credit: Arthus-Bertrand (via F. Durand)



(c) 2005 Heiko Burkhardt, iliano.com

Grouping phenomena in real life



Forsyth & Ponce, Figure 14.7

Gestalt

- ❖ Gestalt: whole or group
 - ❖ Whole is greater than sum of its parts
 - ❖ Relationships among parts can yield new properties/features
 - Humans interpret image information collectively
- ❖ Psychologists identified series of factors that predispose set of elements to be grouped (by human visual system)
 - ❖ Humans interpret image information collectively
- ❖ Inspiring observations/explanations; challenge remains how to best map to algorithms.

Bottom-up segmentation via clustering

- ❖ Algorithms:
 - ❖ Mode finding and mean shift: k-means, EM, mean-shift
 - ❖ Graph-based: normalized cuts
- ❖ Features: color, texture, ...
 - ❖ Quantization for texture summaries

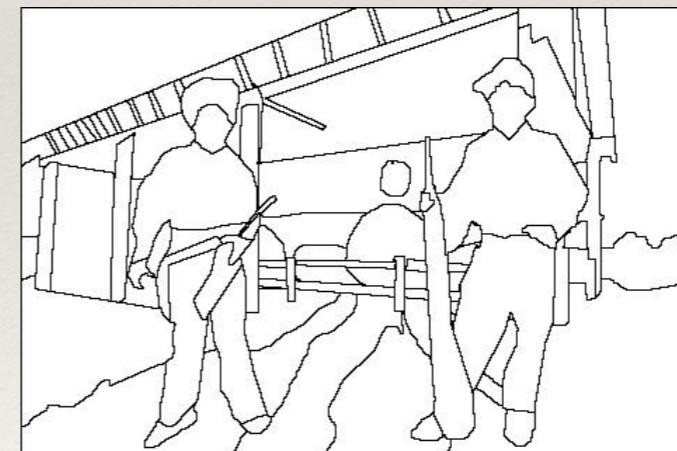
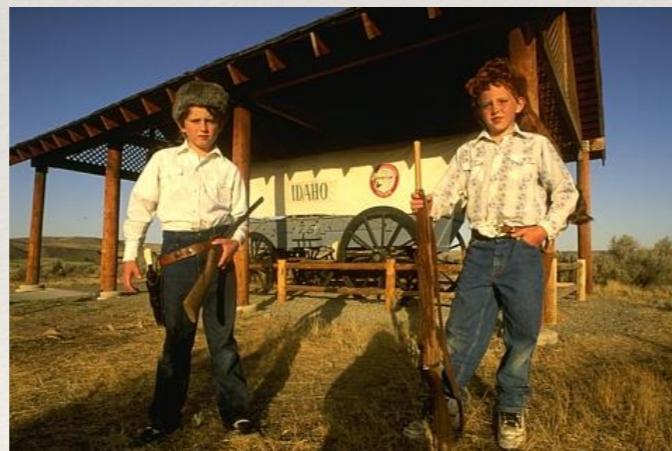
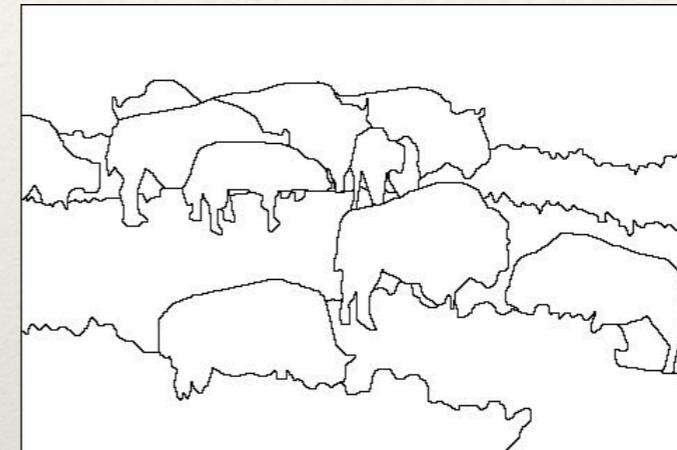
The goals of segmentation

Separate image into coherent “objects”

image



human segmentation

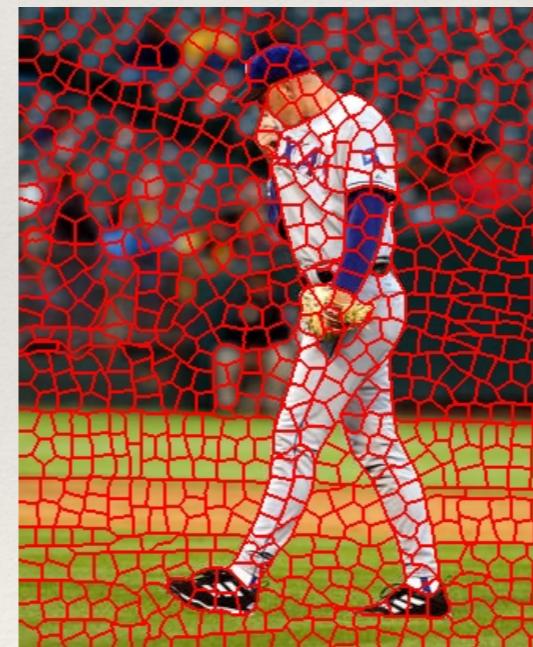
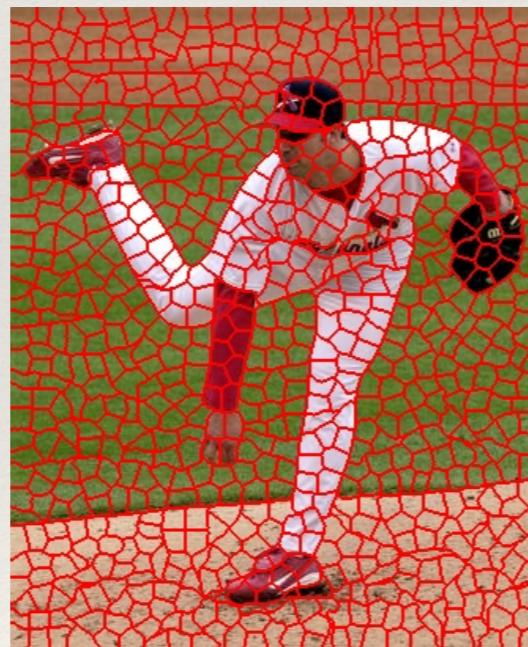


The goals of segmentation

Separate image into coherent “objects”

Group together similar-looking pixels for efficiency of further processing

“superpixels”



X. Ren and J. Malik. [Learning a classification model for segmentation](#). ICCV 2003.

Source: Lana Lazebnik

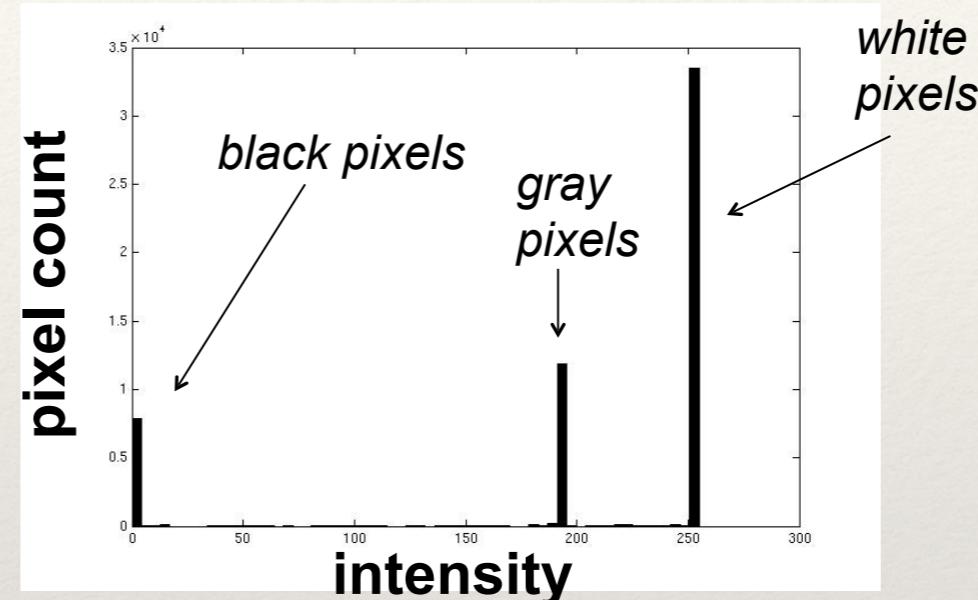
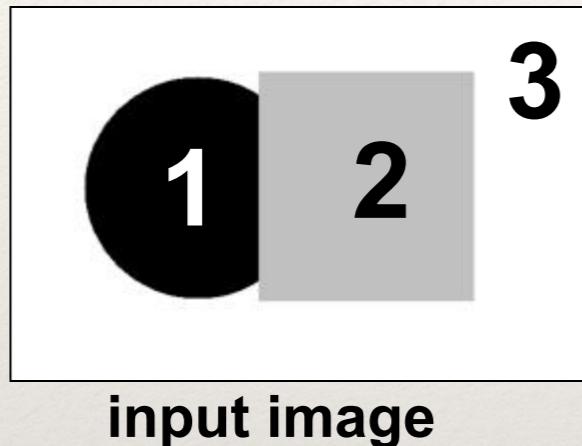
Clustering

- ❖ Clustering algorithms:
 - ❖ Unsupervised learning
 - ❖ Detect patterns in unlabeled data
 - ❖ E.g. group emails or search results
 - ❖ E.g. find categories of customers
 - ❖ E.g. group pixels into regions
 - ❖ Useful when don't know what you're looking for
 - ❖ Requires data, but no labels

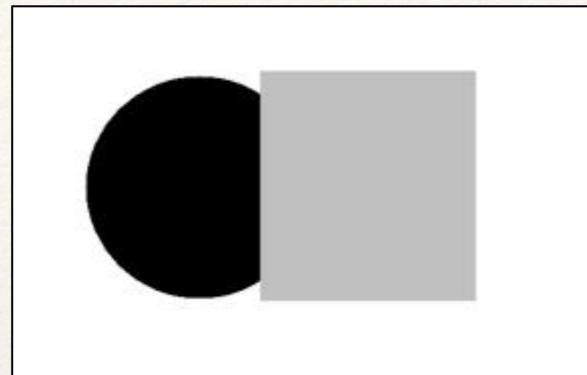


Slide credit: Dan Klein

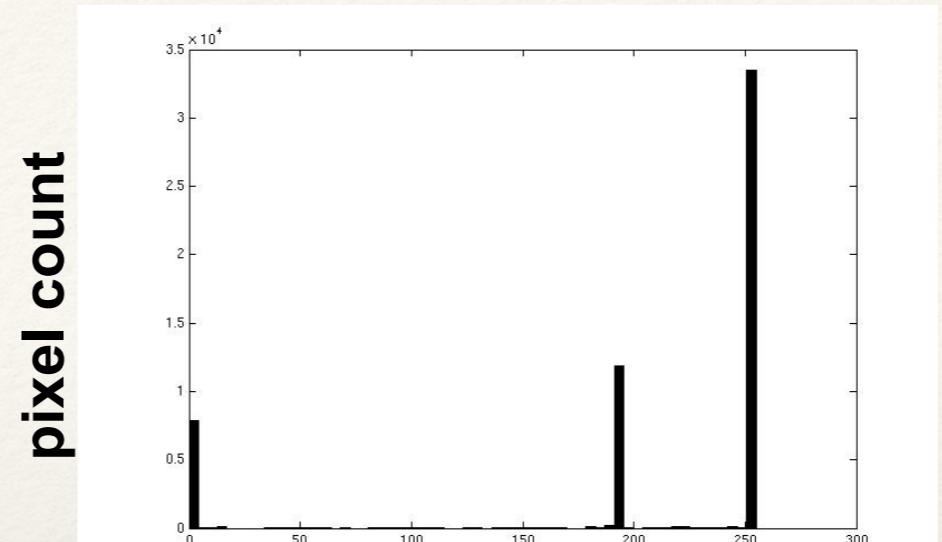
Image segmentation: toy example



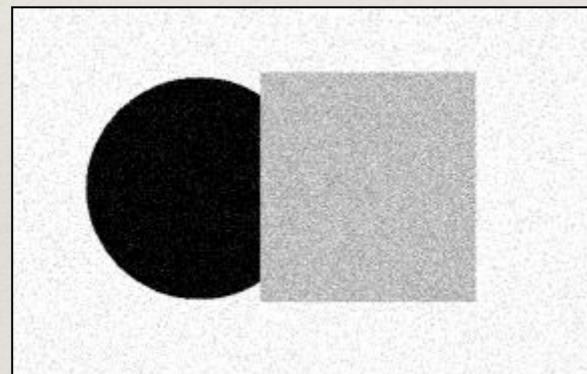
- ❖ These intensities define the three groups.
- ❖ We could label every pixel in the image according to which of these primary intensities it is.
 - ❖ i.e., segment the image based on the intensity feature.
- ❖ What if the image isn't quite so simple?



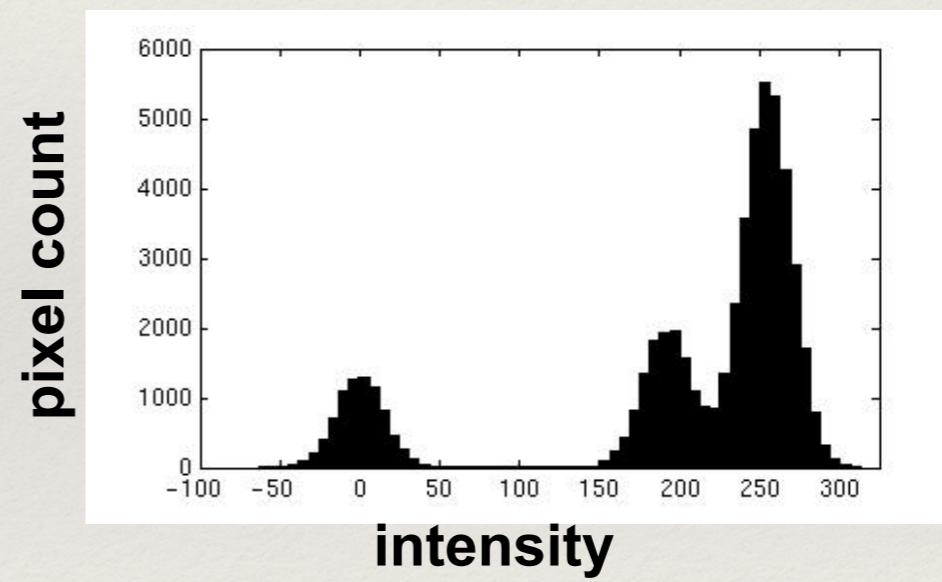
input image



intensity

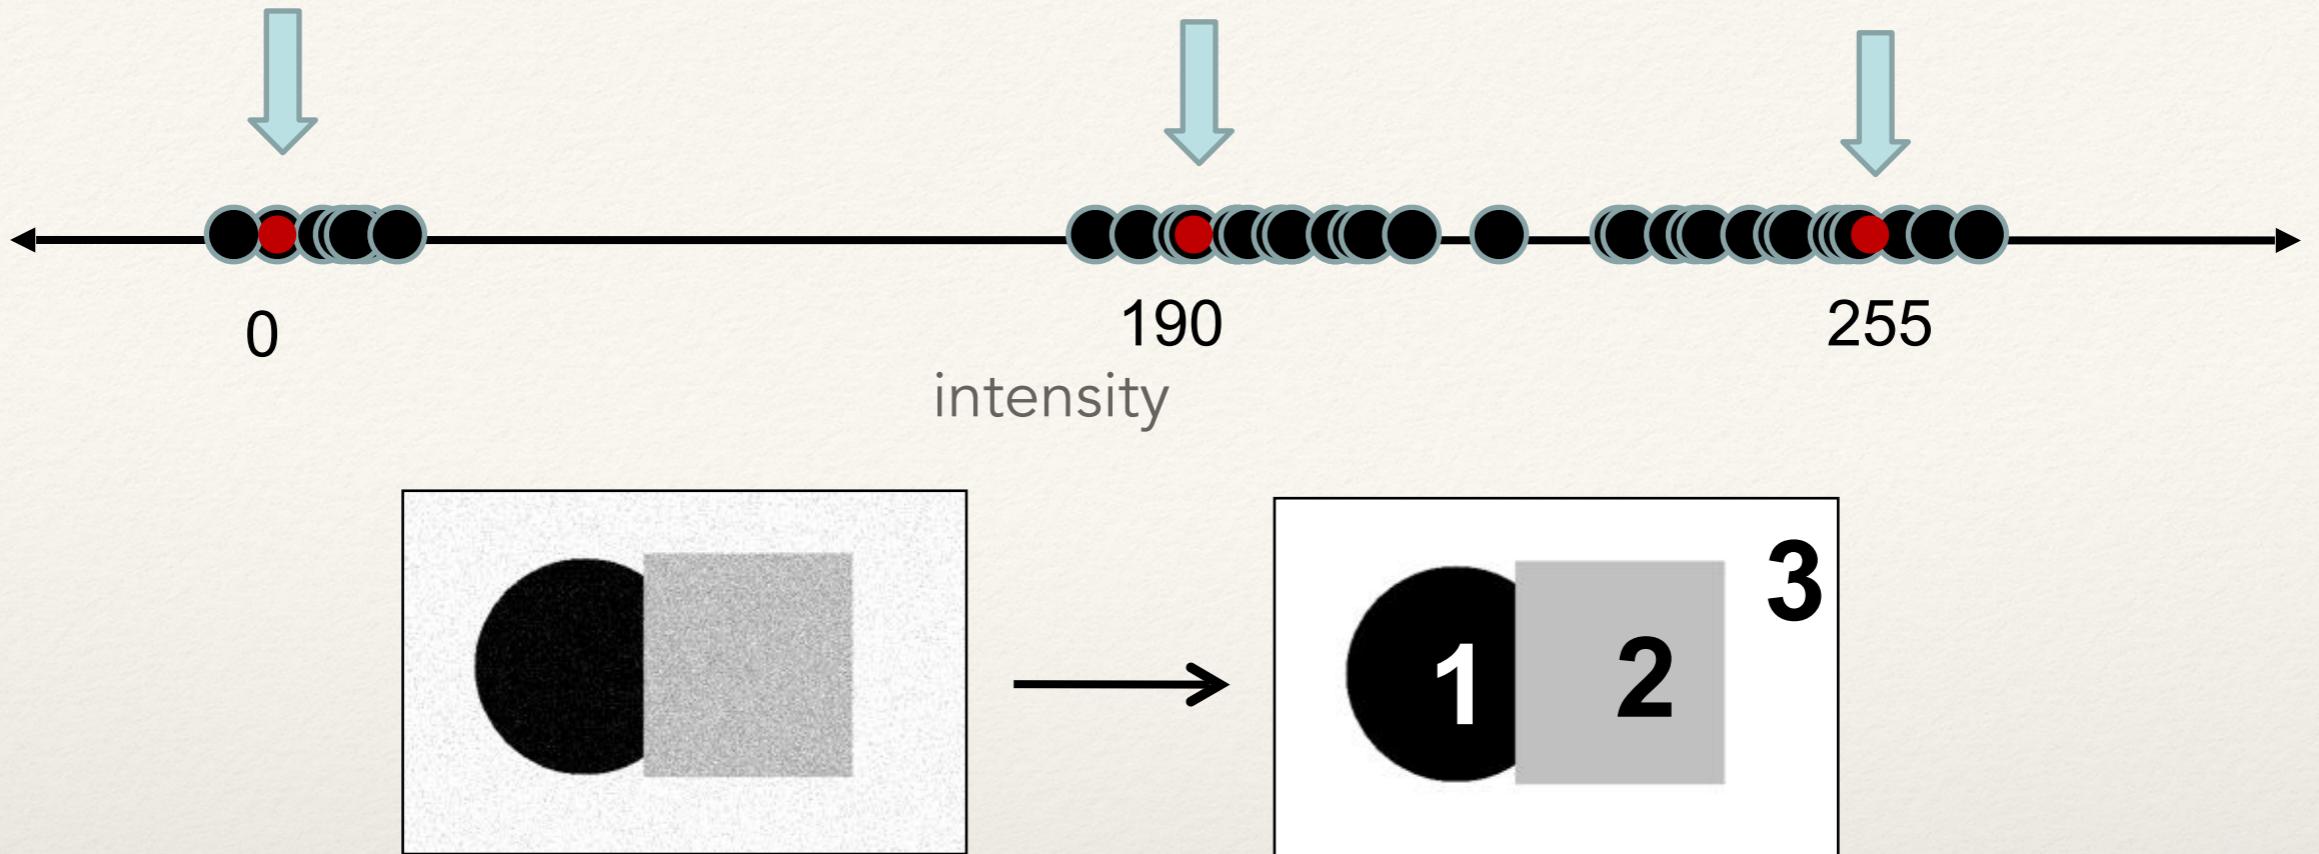


input image



intensity

Now how to determine the three main intensities that define our groups?
We need to cluster.

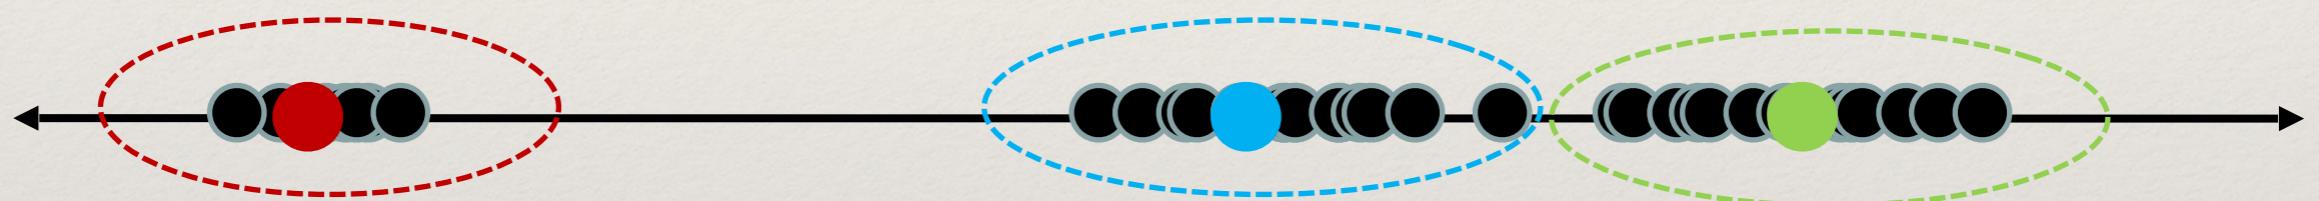


- ❖ Goal: choose three “centers” as the representative intensities, and label every pixel according to which of these centers it is nearest to.
- ❖ Best cluster centers are those that minimize SSD between all points and their nearest cluster center c_i :

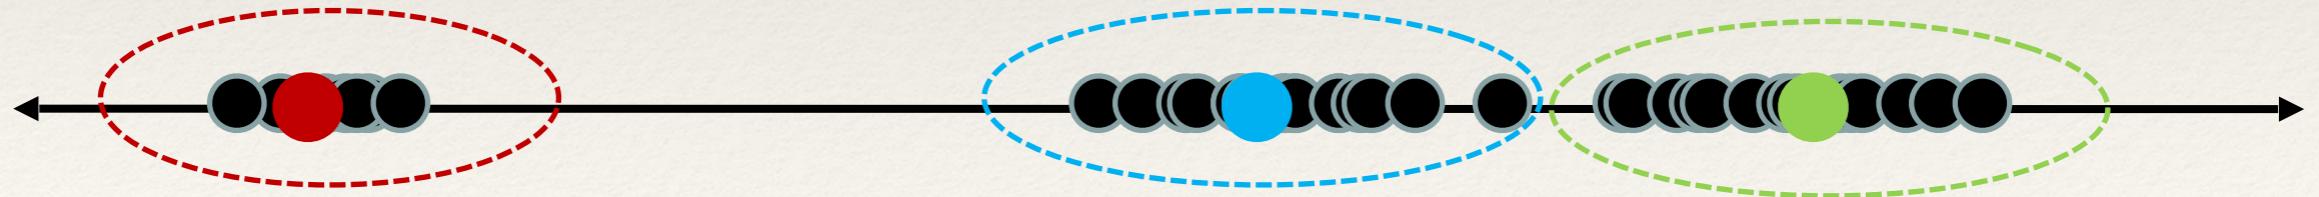
$$\sum_{\text{clusters } i} \sum_{\text{points } p \text{ in cluster } i} \|p - c_i\|^2$$

Clustering

- ❖ With this objective, it is a “chicken and egg” problem:
 - ❖ If we knew the **cluster centers**, we could allocate points to groups by assigning each to its closest center.



- ❖ If we knew the **group memberships**, we could get the centers by computing the mean per group.



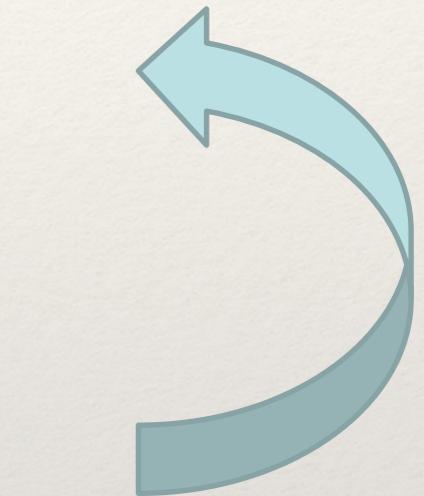
K-means clustering

- Basic idea: randomly initialize the k cluster centers, and iterate between the two steps we just saw.

1. Randomly initialize the cluster centers, c_1, \dots, c_K
2. Given cluster centers, determine points in each cluster
 - For each point p , find the closest c_i . Put p into cluster i
3. Given points in each cluster, solve for c_i
 - Set c_i to be the mean of points in cluster i
4. If c_i have changed, repeat Step 2

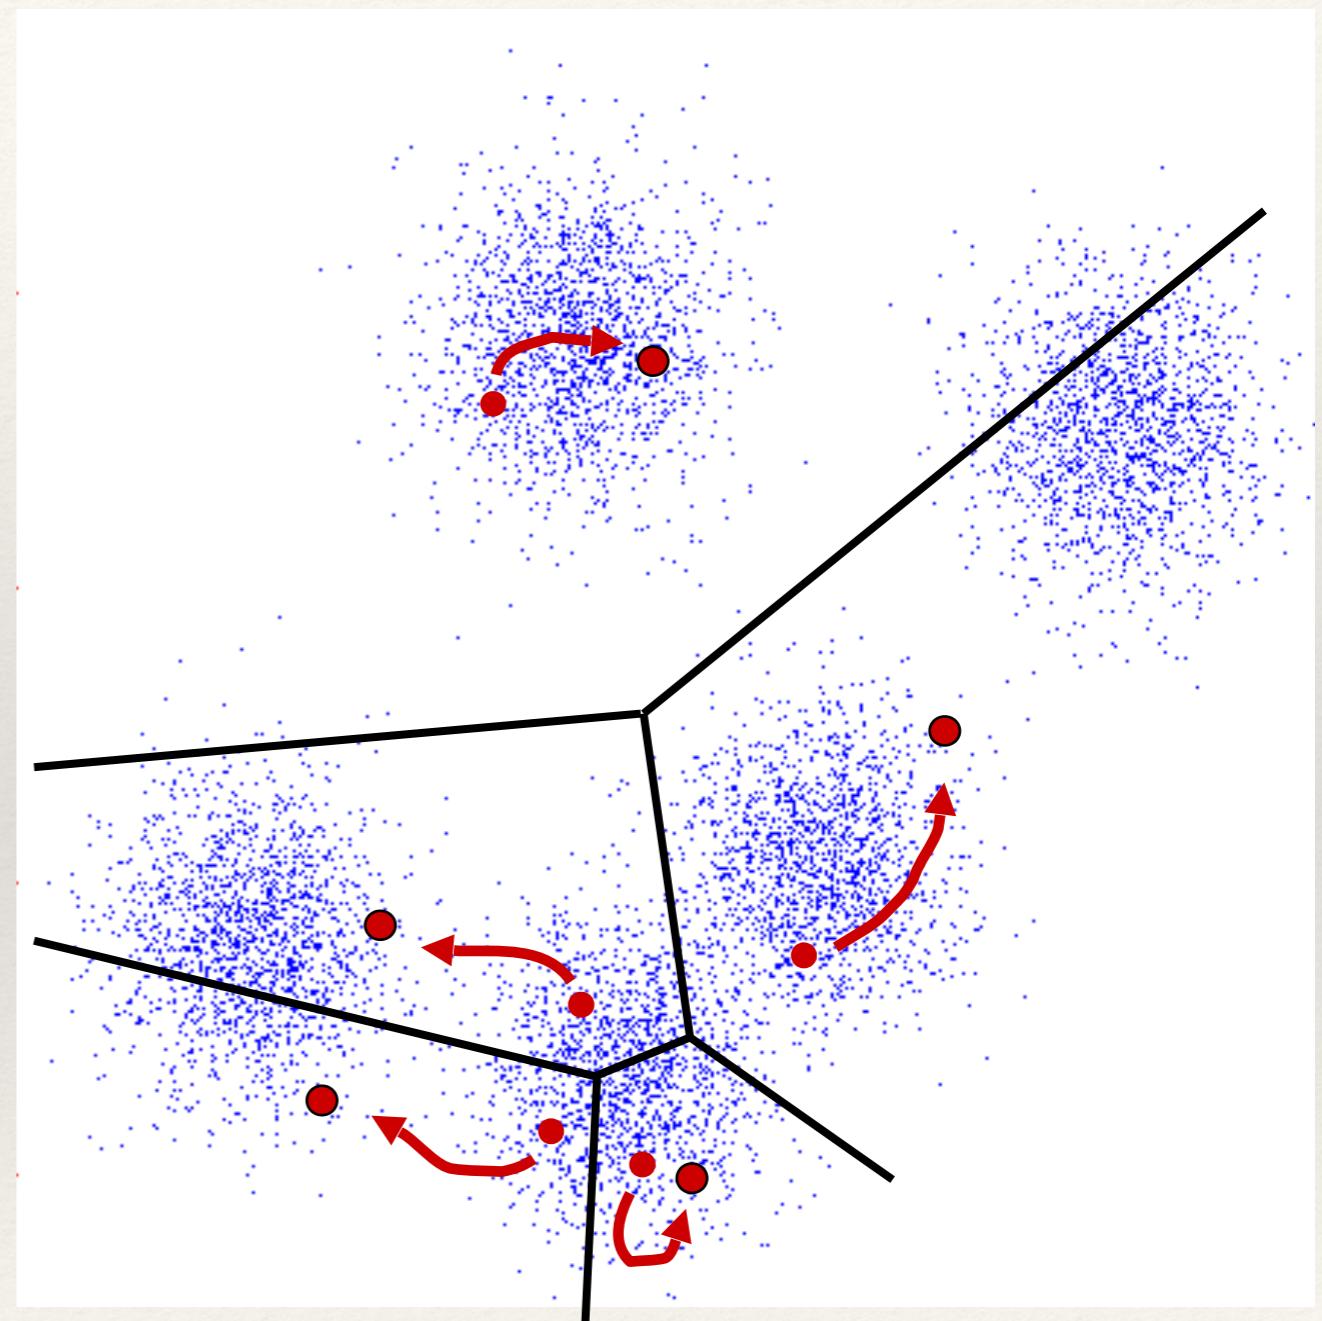
- Properties
 - Will always converge to some solution
 - Can be a “local minimum”
 - does not always find the global minimum of objective function:

$$\sum_{\text{clusters } i} \sum_{\text{points } p \text{ in cluster } i} \|p - c_i\|^2$$



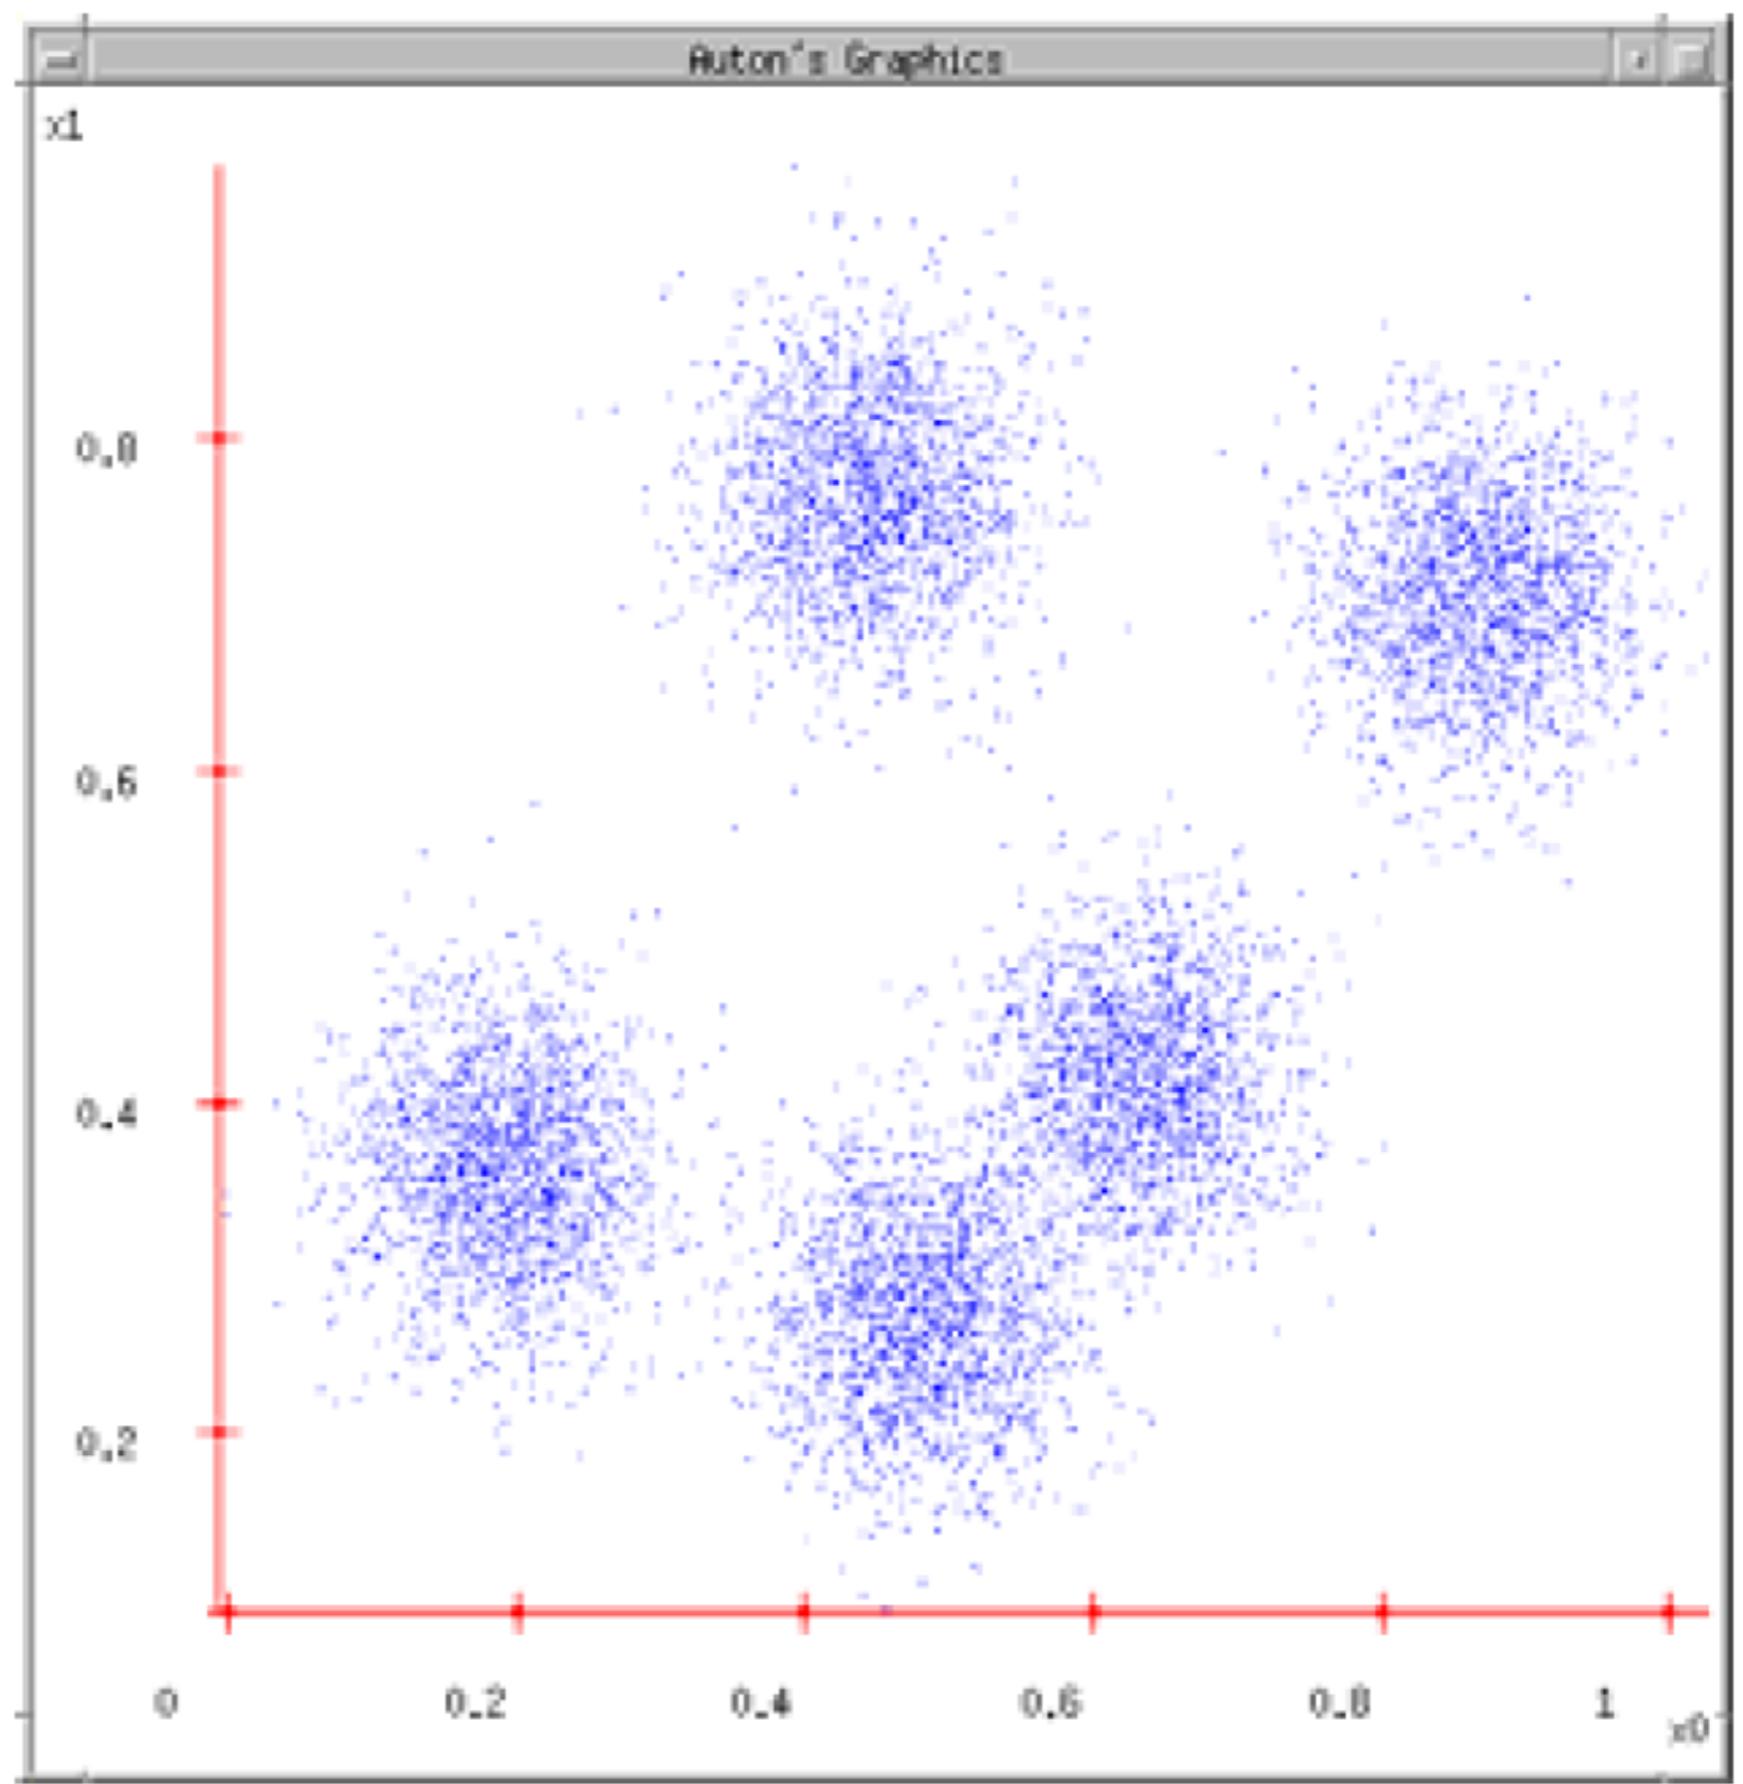
K-Means

- ❖ An iterative clustering algorithm
 - ❖ Pick K random points as cluster centers (means)
 - ❖ Alternate:
 - ❖ Assign data instances to closest mean
 - ❖ Assign each mean to the average of its assigned points
 - ❖ Stop when no points' assignments change



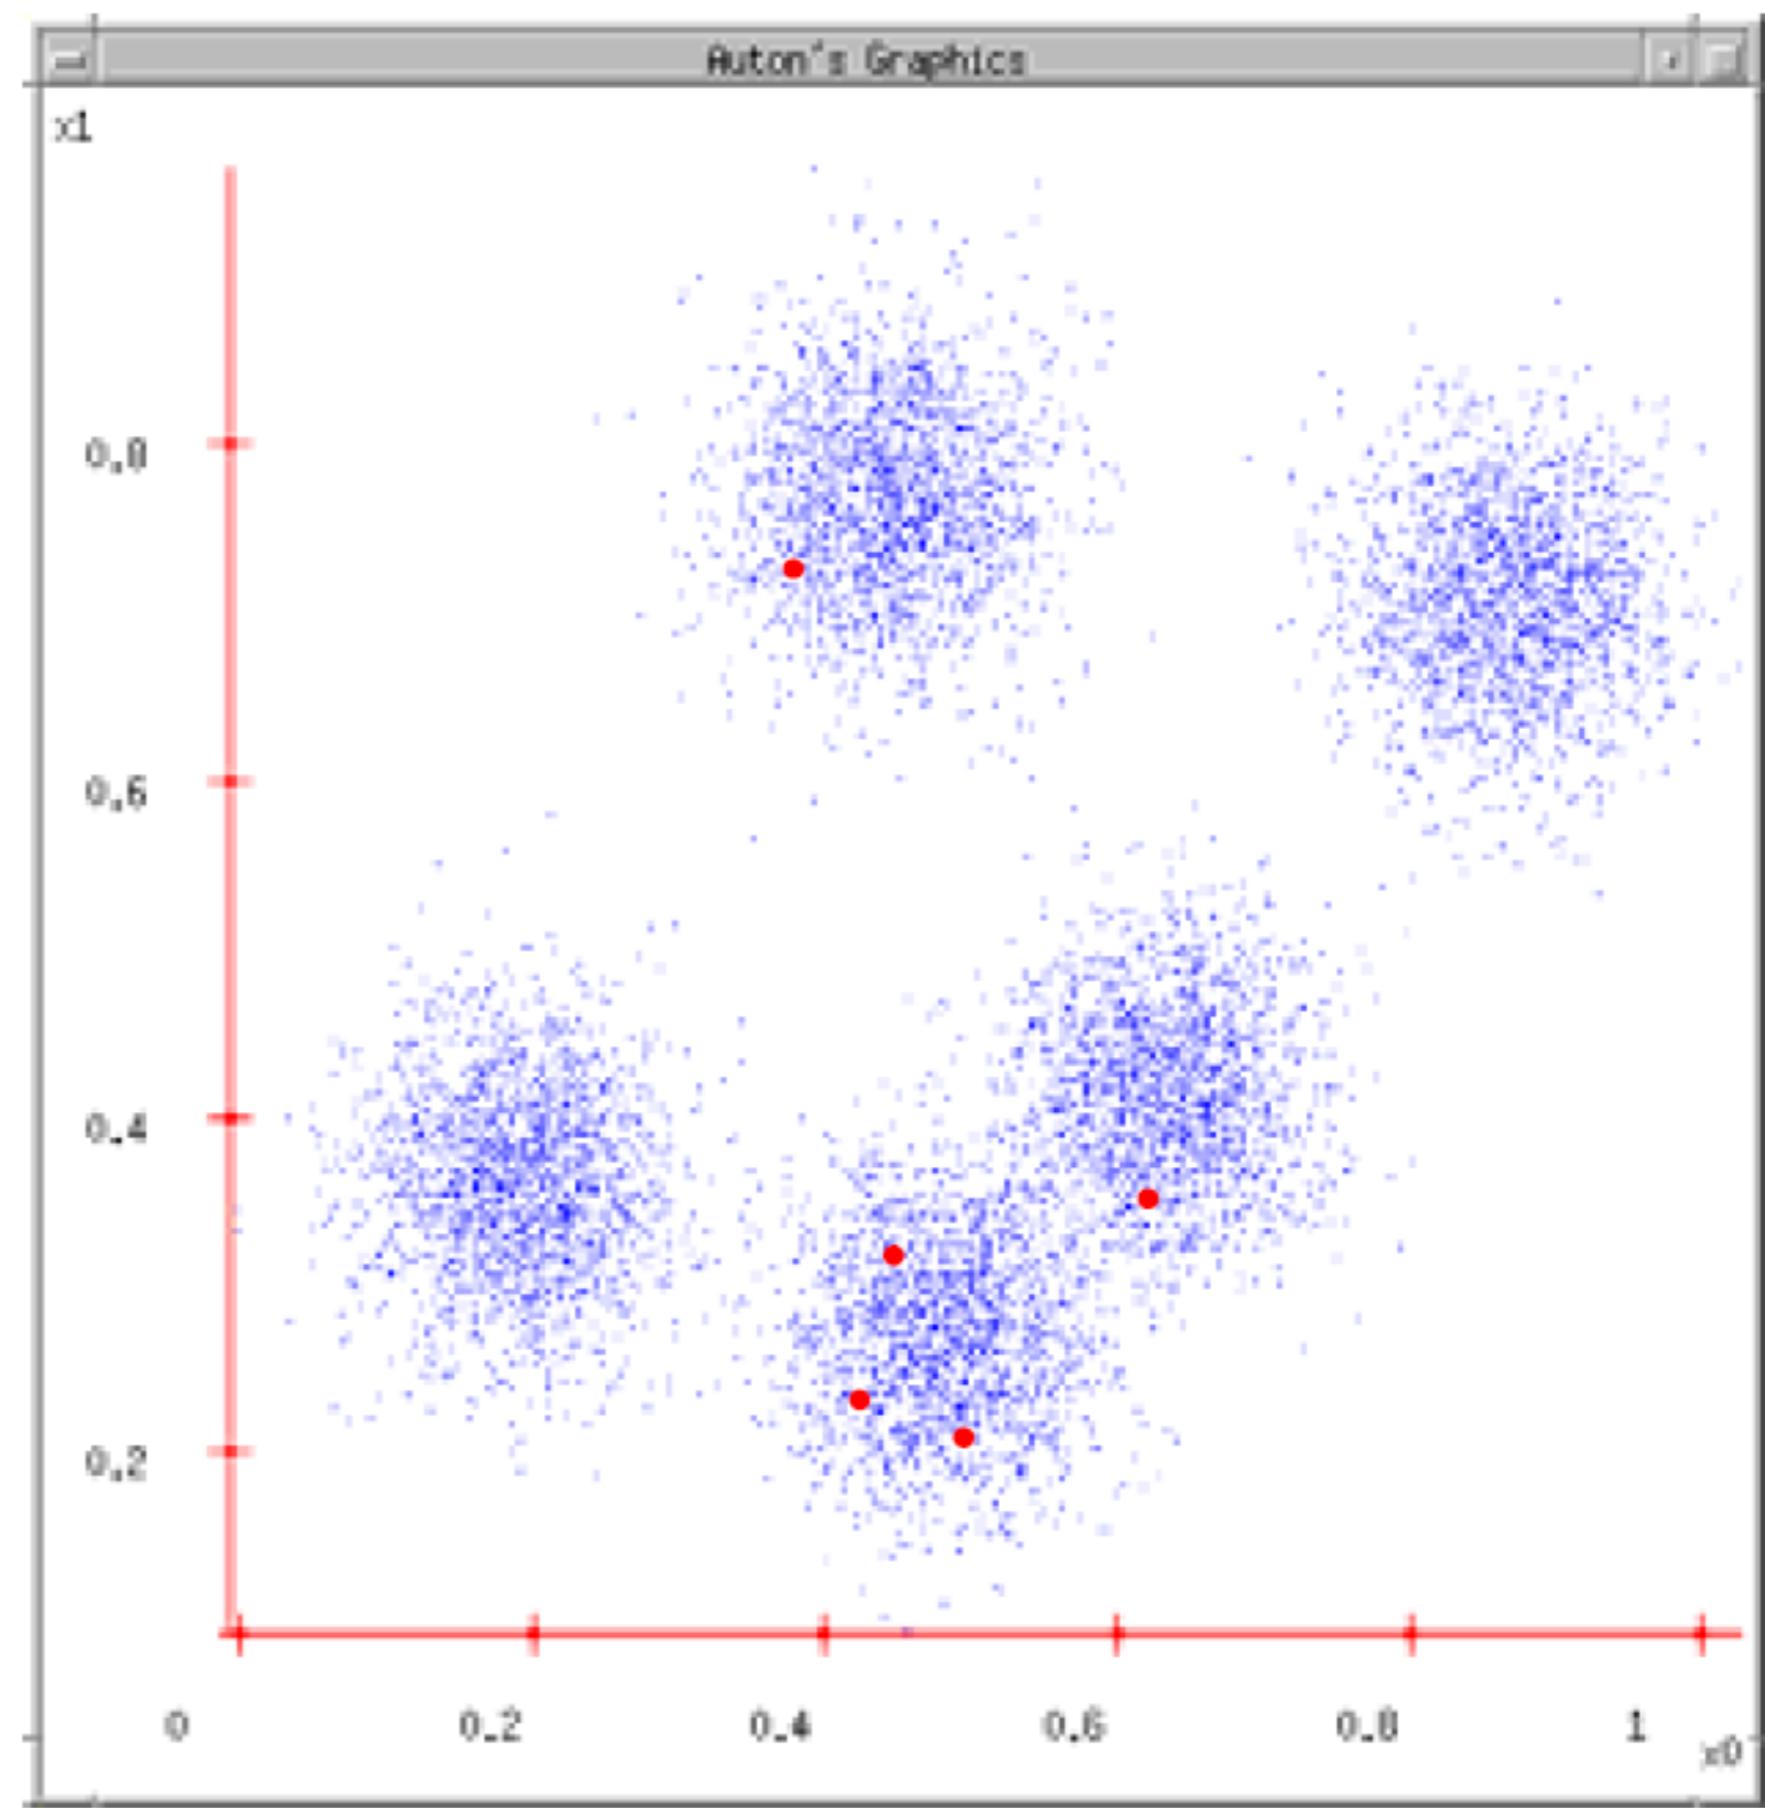
K-means

1. Ask user how many clusters they'd like.
(e.g. $k=5$)



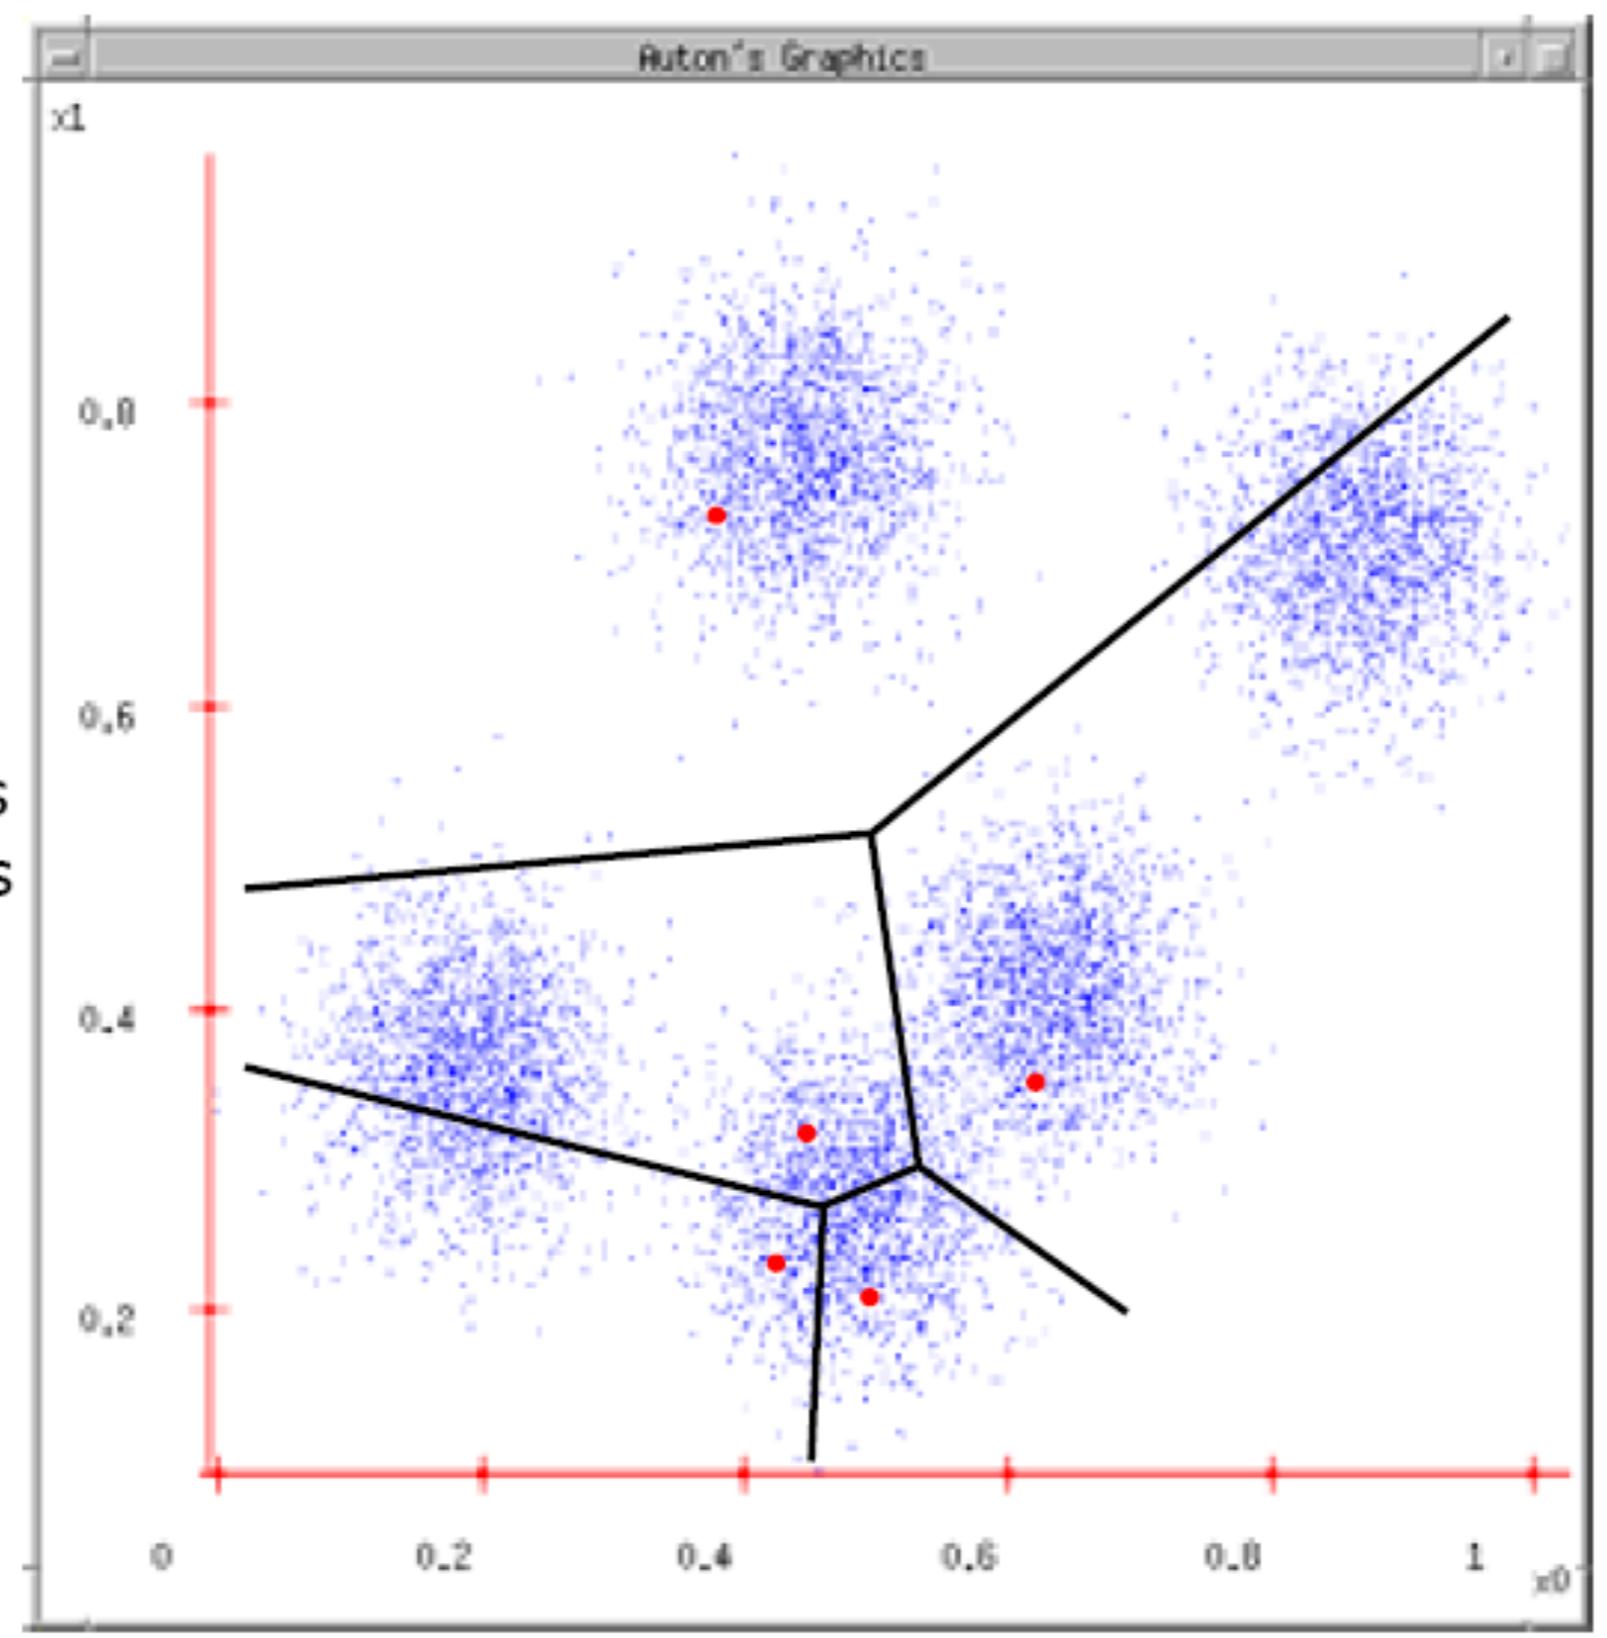
K-means

1. Ask user how many clusters they'd like.
(e.g. k=5)
2. Randomly guess k cluster Center locations



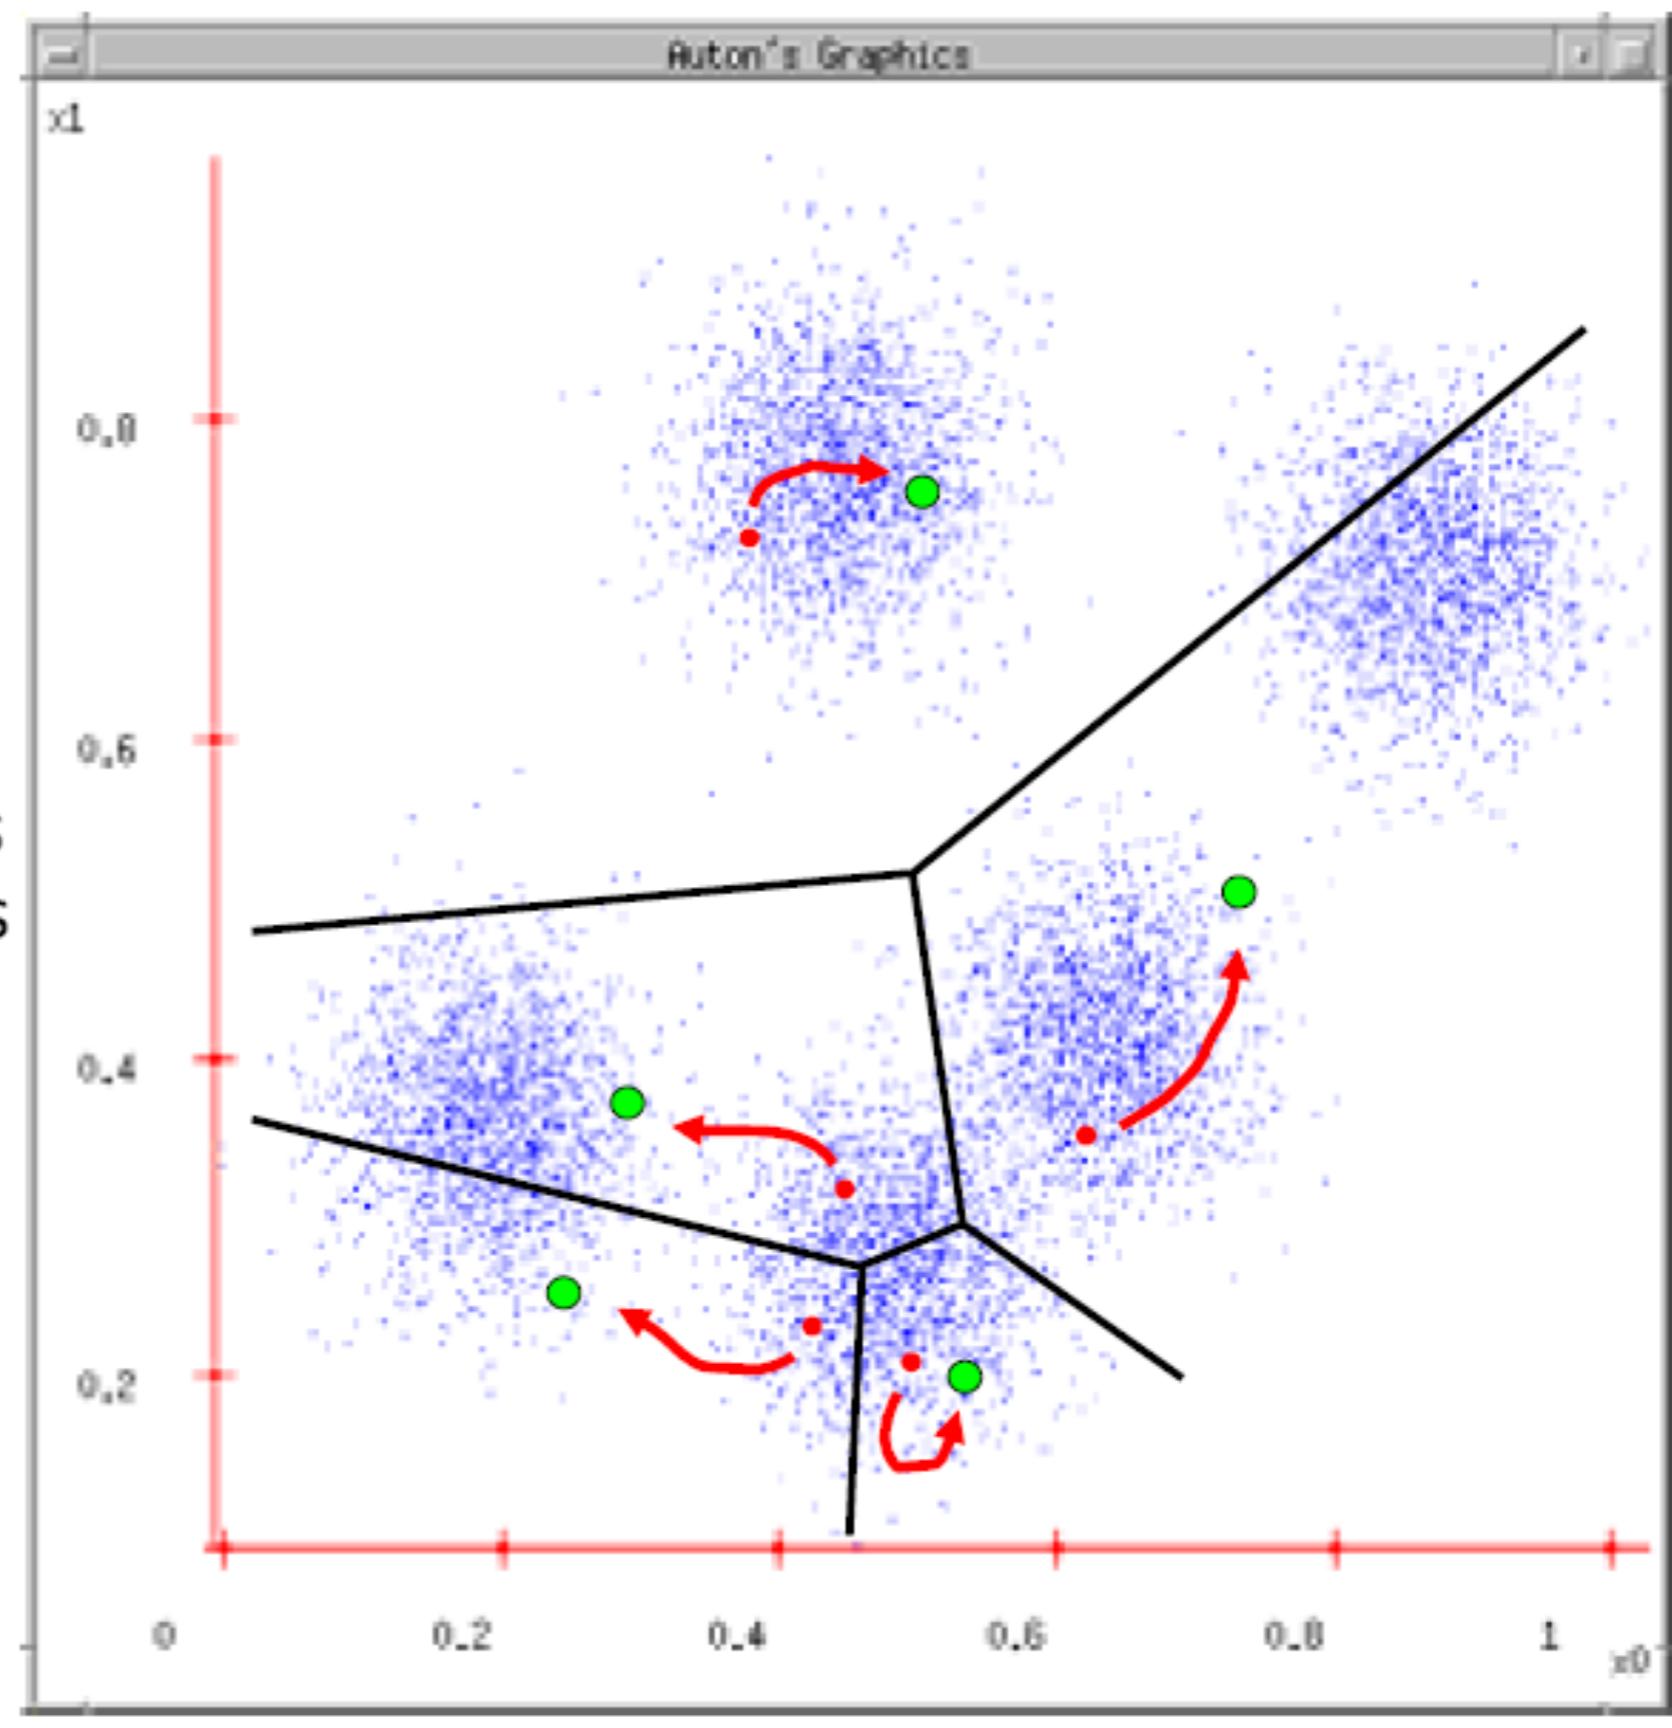
K-means

1. Ask user how many clusters they'd like.
(e.g. k=5)
2. Randomly guess k cluster Center locations
3. Each datapoint finds out which Center it's closest to. (Thus each Center "owns" a set of datapoints)



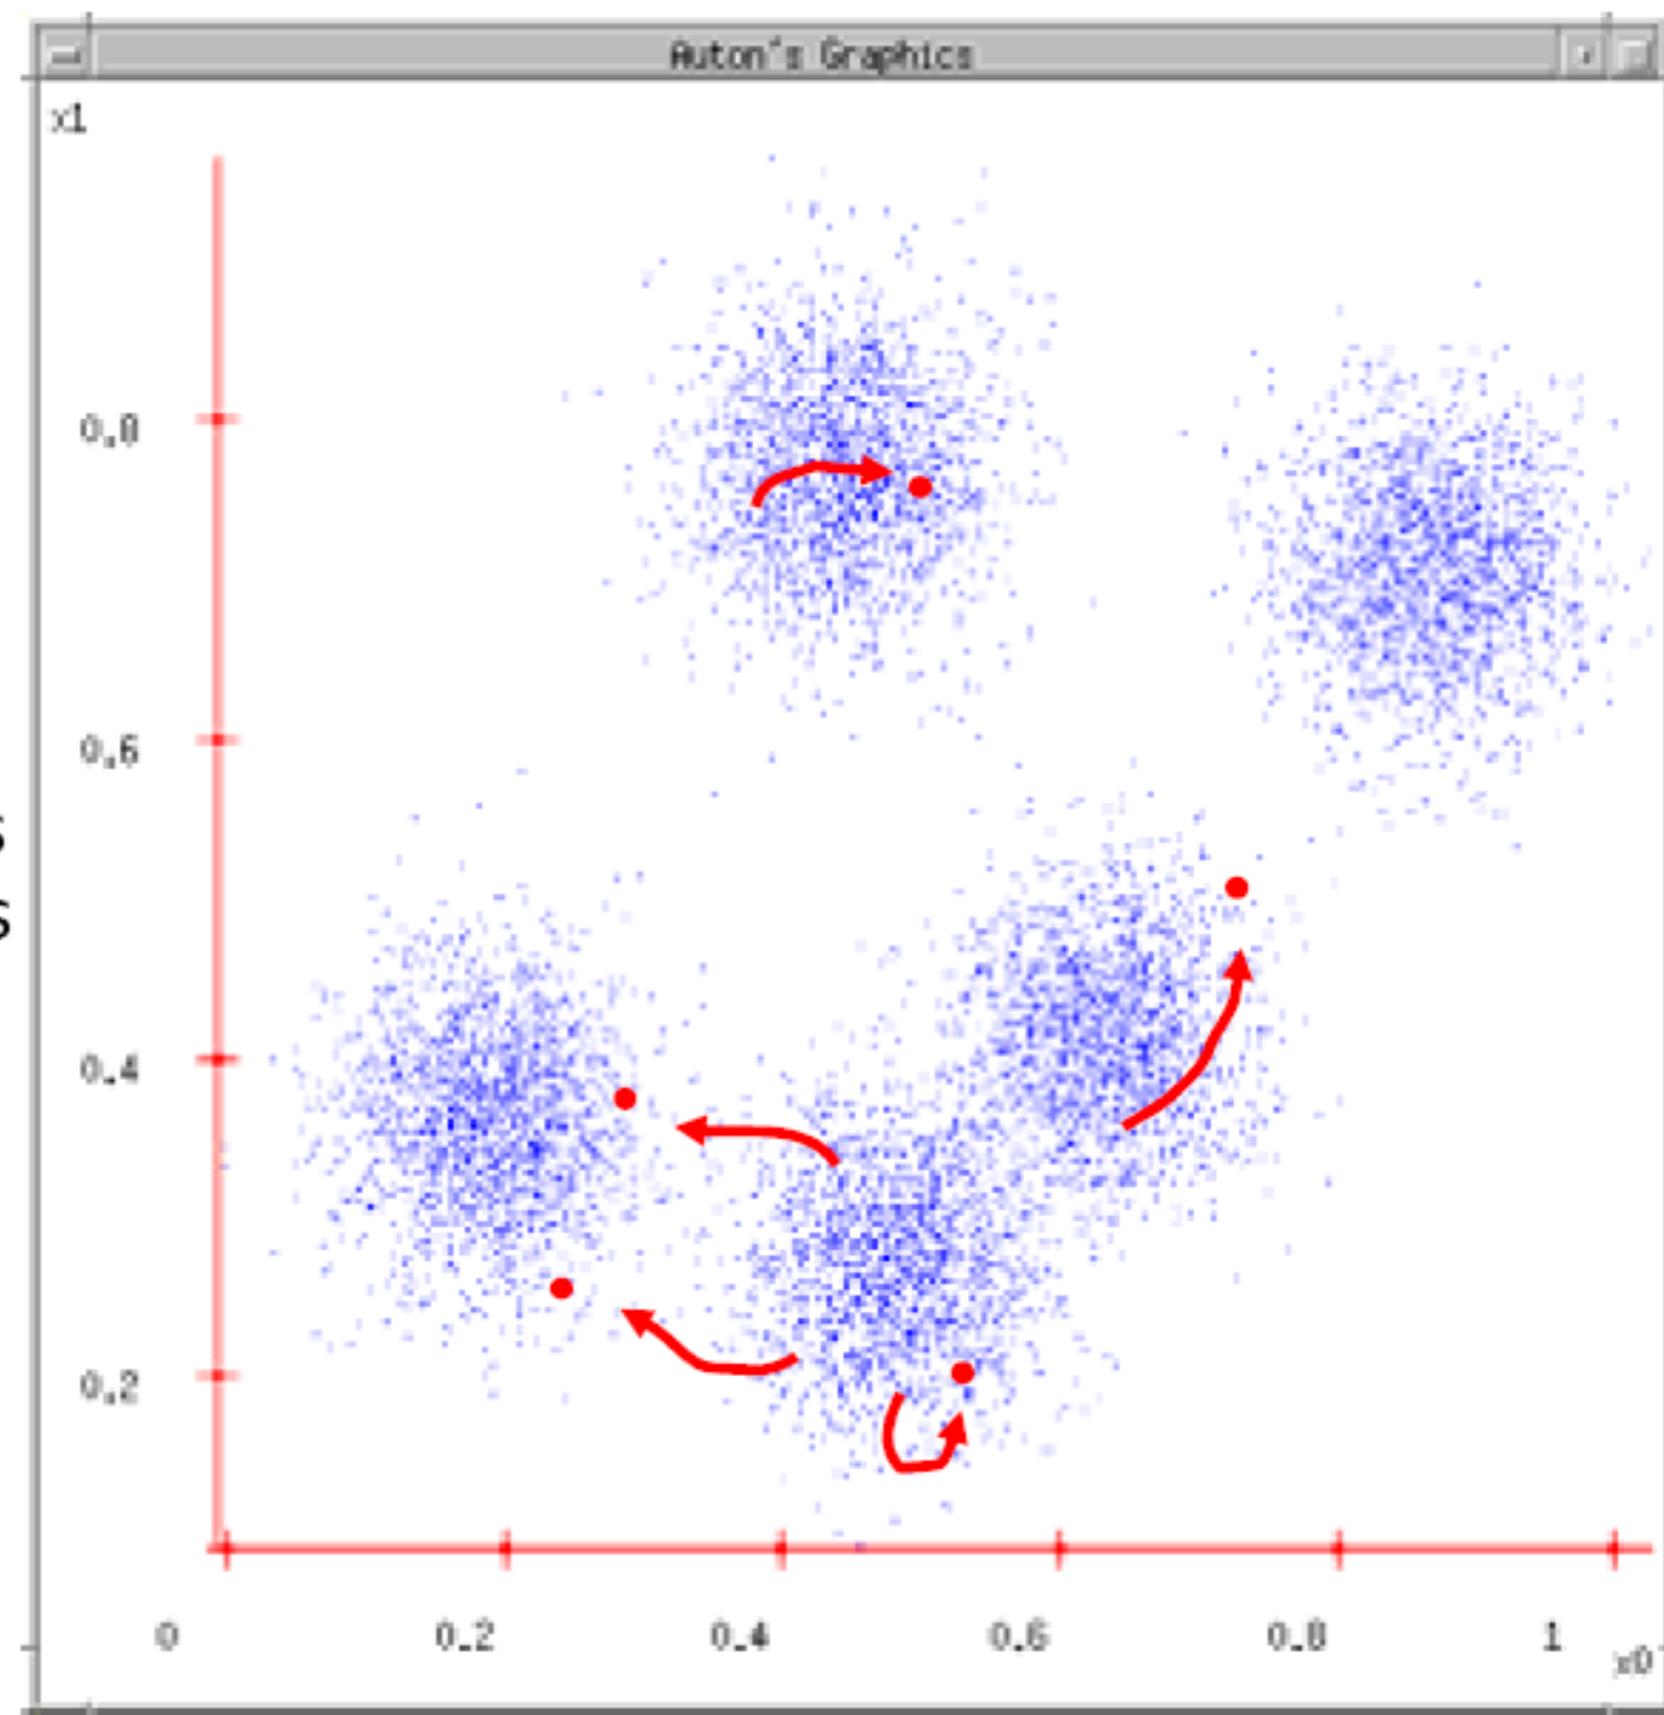
K-means

1. Ask user how many clusters they'd like.
(e.g. $k=5$)
2. Randomly guess k cluster Center locations
3. Each datapoint finds out which Center it's closest to.
4. Each Center finds the centroid of the points it owns



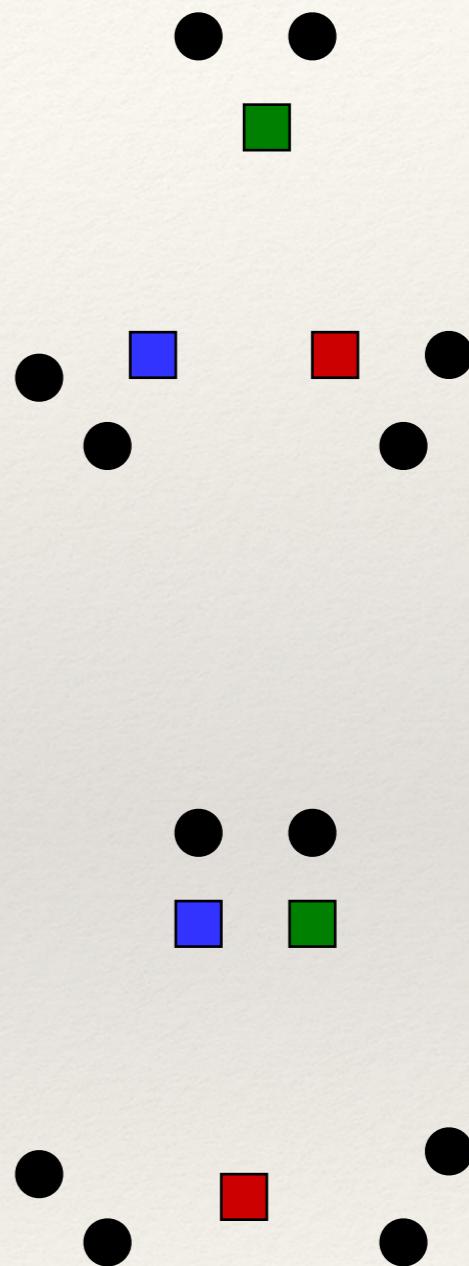
K-means

1. Ask user how many clusters they'd like.
(e.g. k=5)
2. Randomly guess k cluster Center locations
3. Each datapoint finds out which Center it's closest to.
4. Each Center finds the centroid of the points it owns...
5. ...and jumps there
6. ...Repeat until terminated!



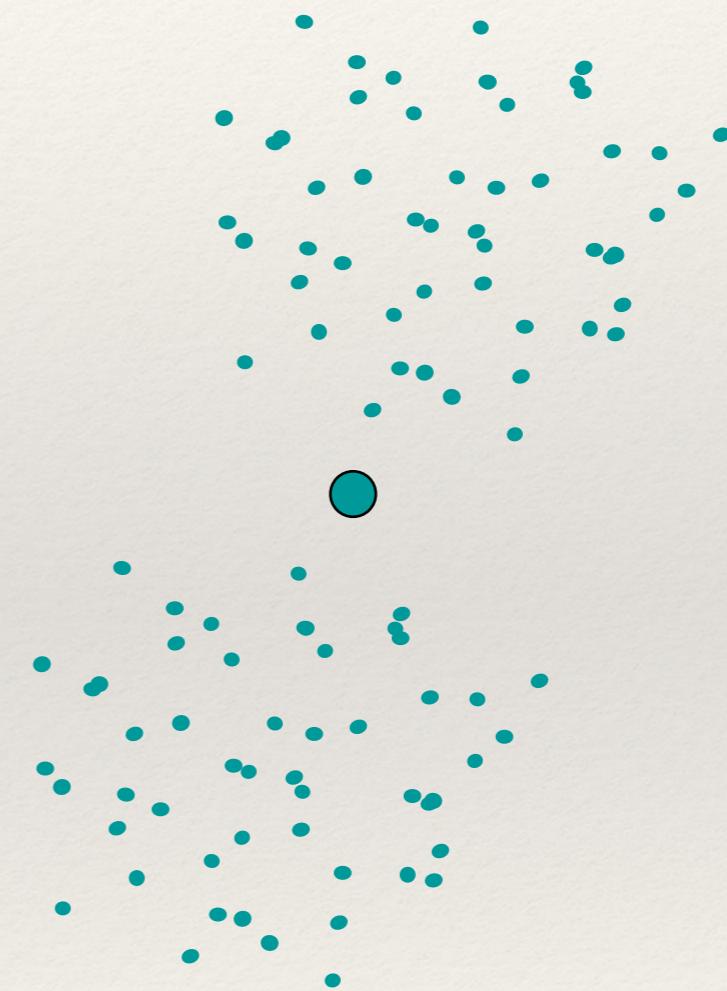
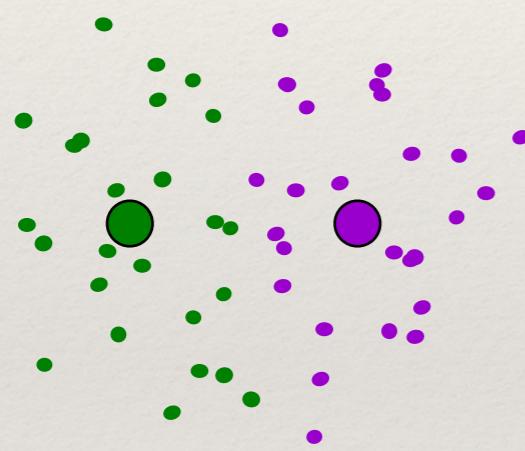
Initialization

- ❖ K-means is non-deterministic
 - ❖ Requires initial means
 - ❖ It does matter what you pick!
 - ❖ What can go wrong?
 - ❖ Various schemes for preventing this kind of thing



K-Means getting stuck

- ❖ A local optimum:



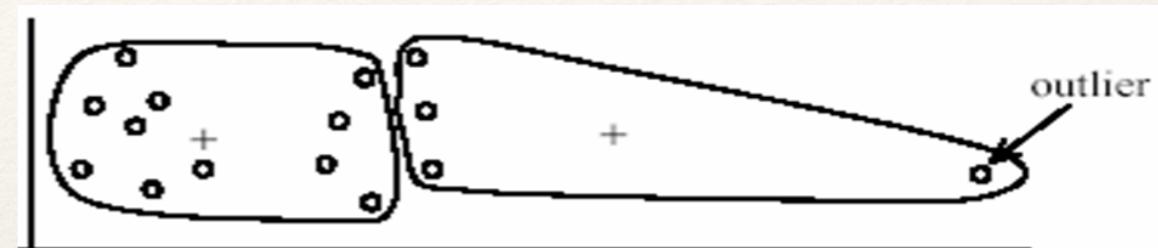
K-means: pros and cons

Pros

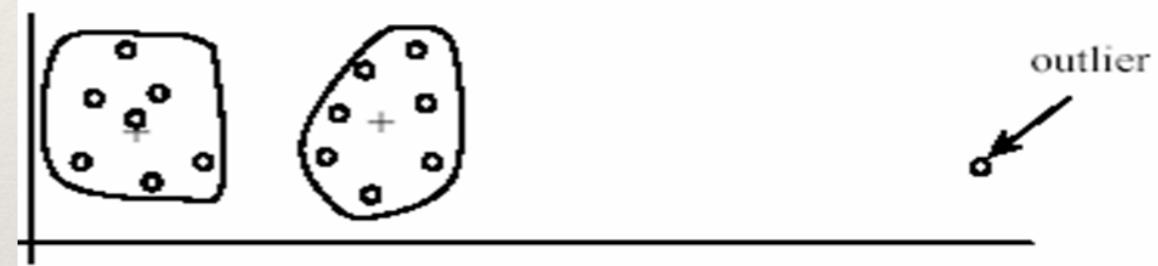
- ❖ Simple, fast to compute
- ❖ Converges to local minimum of within-cluster squared error

Cons/issues

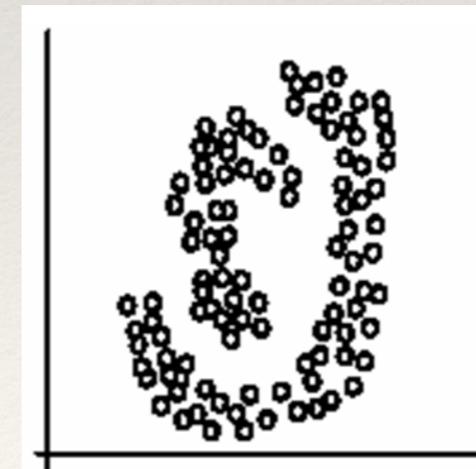
- ❖ Setting k?
- ❖ Sensitive to initial centers
- ❖ Sensitive to outliers
- ❖ Detects spherical clusters
- ❖ Assuming means can be computed



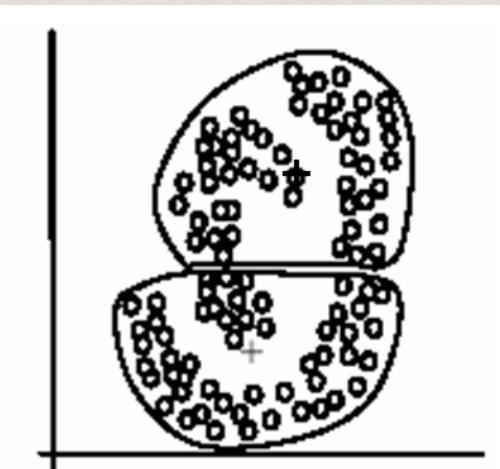
(A): Undesirable clusters



(B): Ideal clusters



(A): Two natural clusters



(B): k -means clusters

K-Means questions

- ❖ Will K-means converge?
 - ❖ To a global optimum?
- ❖ Will it always find the true patterns in the data?
 - ❖ If the patterns are very very clear?
- ❖ Will it find something interesting?
- ❖ How many clusters to pick?
- ❖ Do people ever use it?

Probabilistic clustering

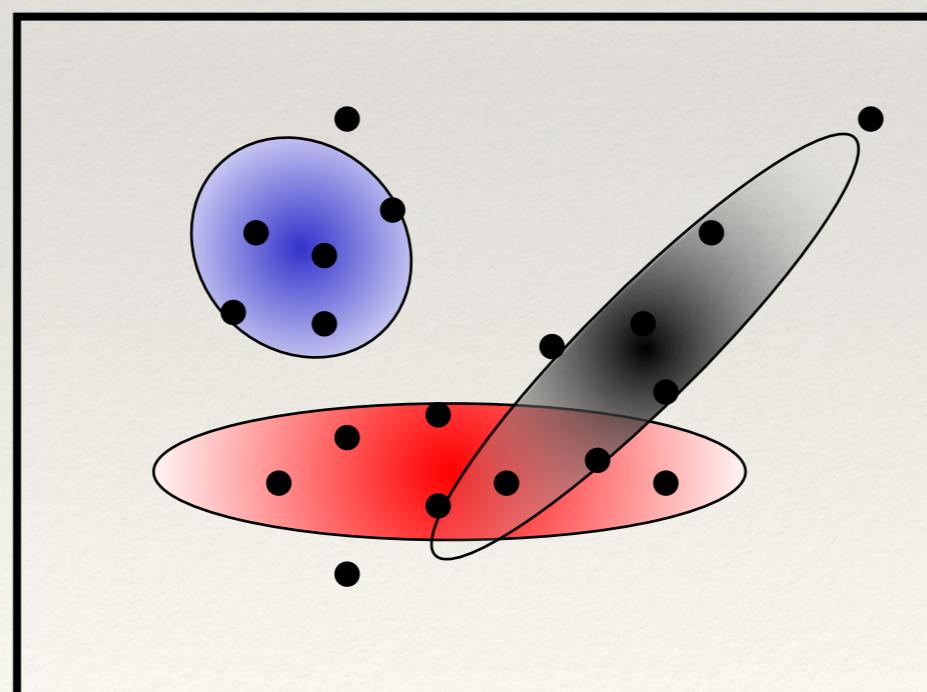
Basic questions

- ❖ what's the probability that a point \mathbf{x} is in cluster m ?
- ❖ what's the shape of each cluster?

K-means doesn't answer these questions

Probabilistic clustering (basic idea)

- ❖ Treat each cluster as a Gaussian density function

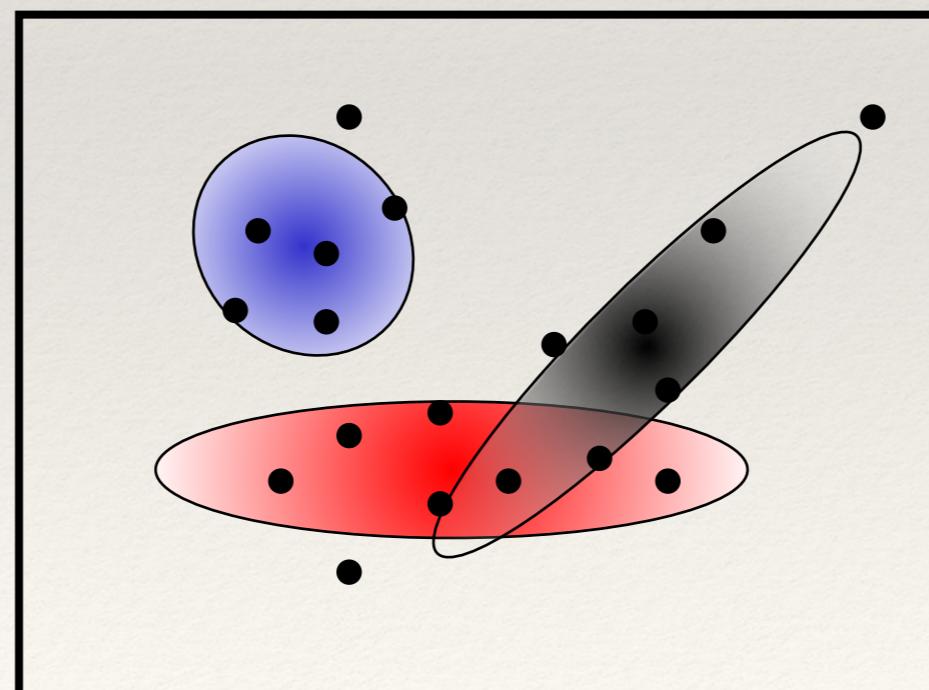


Slide credit: Steve Seitz

Expectation Maximization (EM)

A probabilistic variant of K-means:

- ❖ E step: “soft assignment” of points to clusters
 - ❖ estimate probability that a point is in a cluster
- ❖ M step: update cluster parameters
 - ❖ mean and variance info (covariance matrix)
- ❖ maximizes the likelihood of the points given the clusters



Slide credit: Steve Seitz

Segmentation as clustering

Depending on what we choose as the *feature space*, we can group pixels in different ways.

Grouping pixels based on **intensity** similarity



Feature space: intensity value (1-d)



K=2

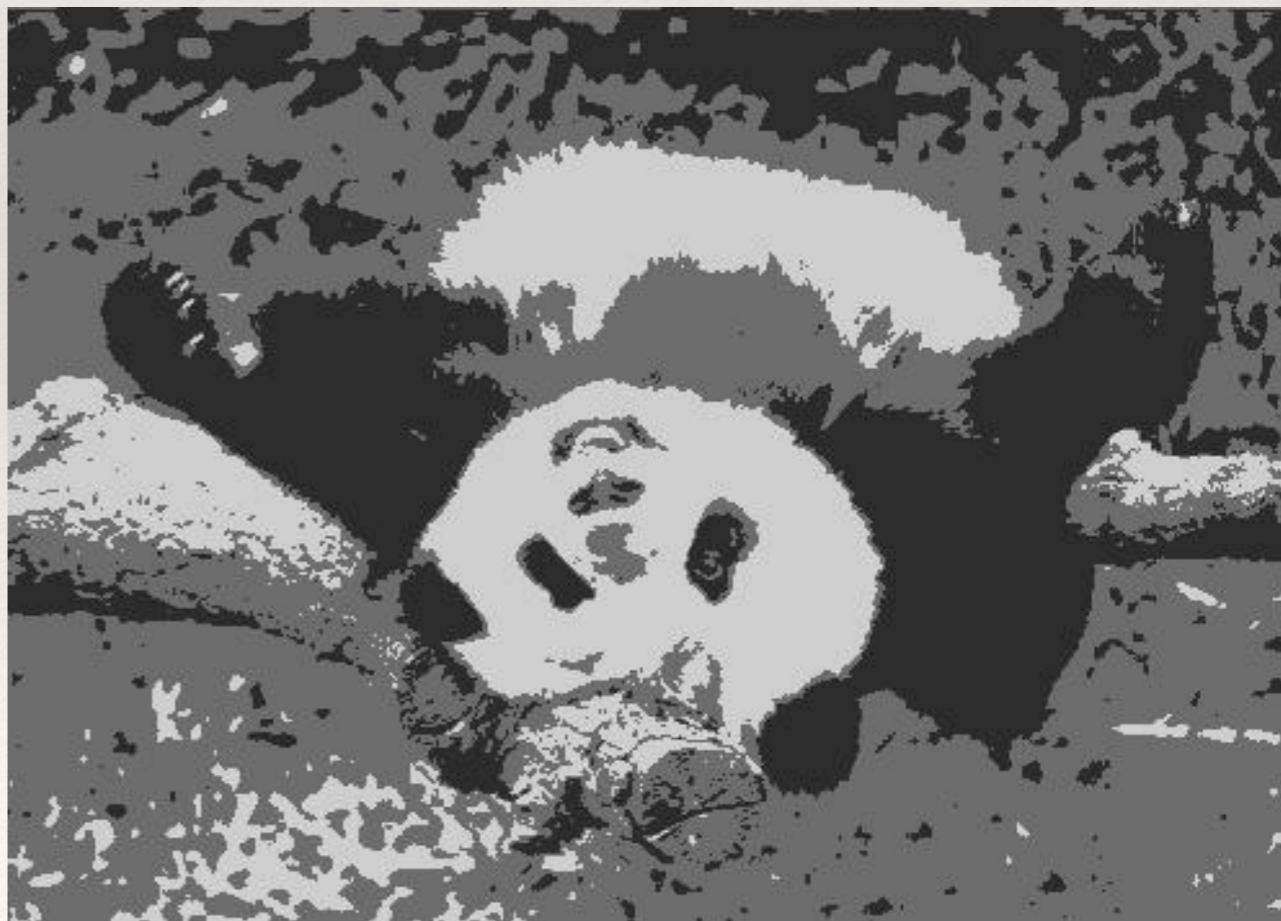


K=3

quantization of the feature space;
segmentation label map

```
img_as_col = double(im(:));
cluster_membs = kmeans(img_as_col, K);

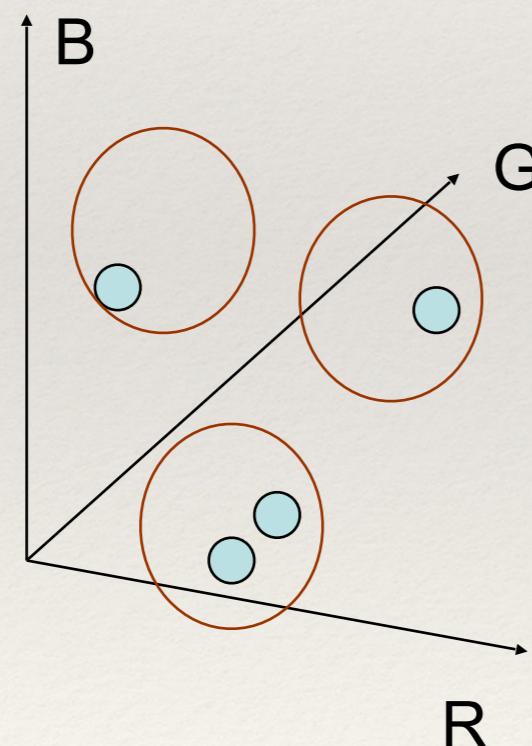
labelim = zeros(size(im));
for i=1:k
    inds = find(cluster_membs==i);
    meanval = mean(img_as_column(inds));
    labelim(inds) = meanval;
end
```



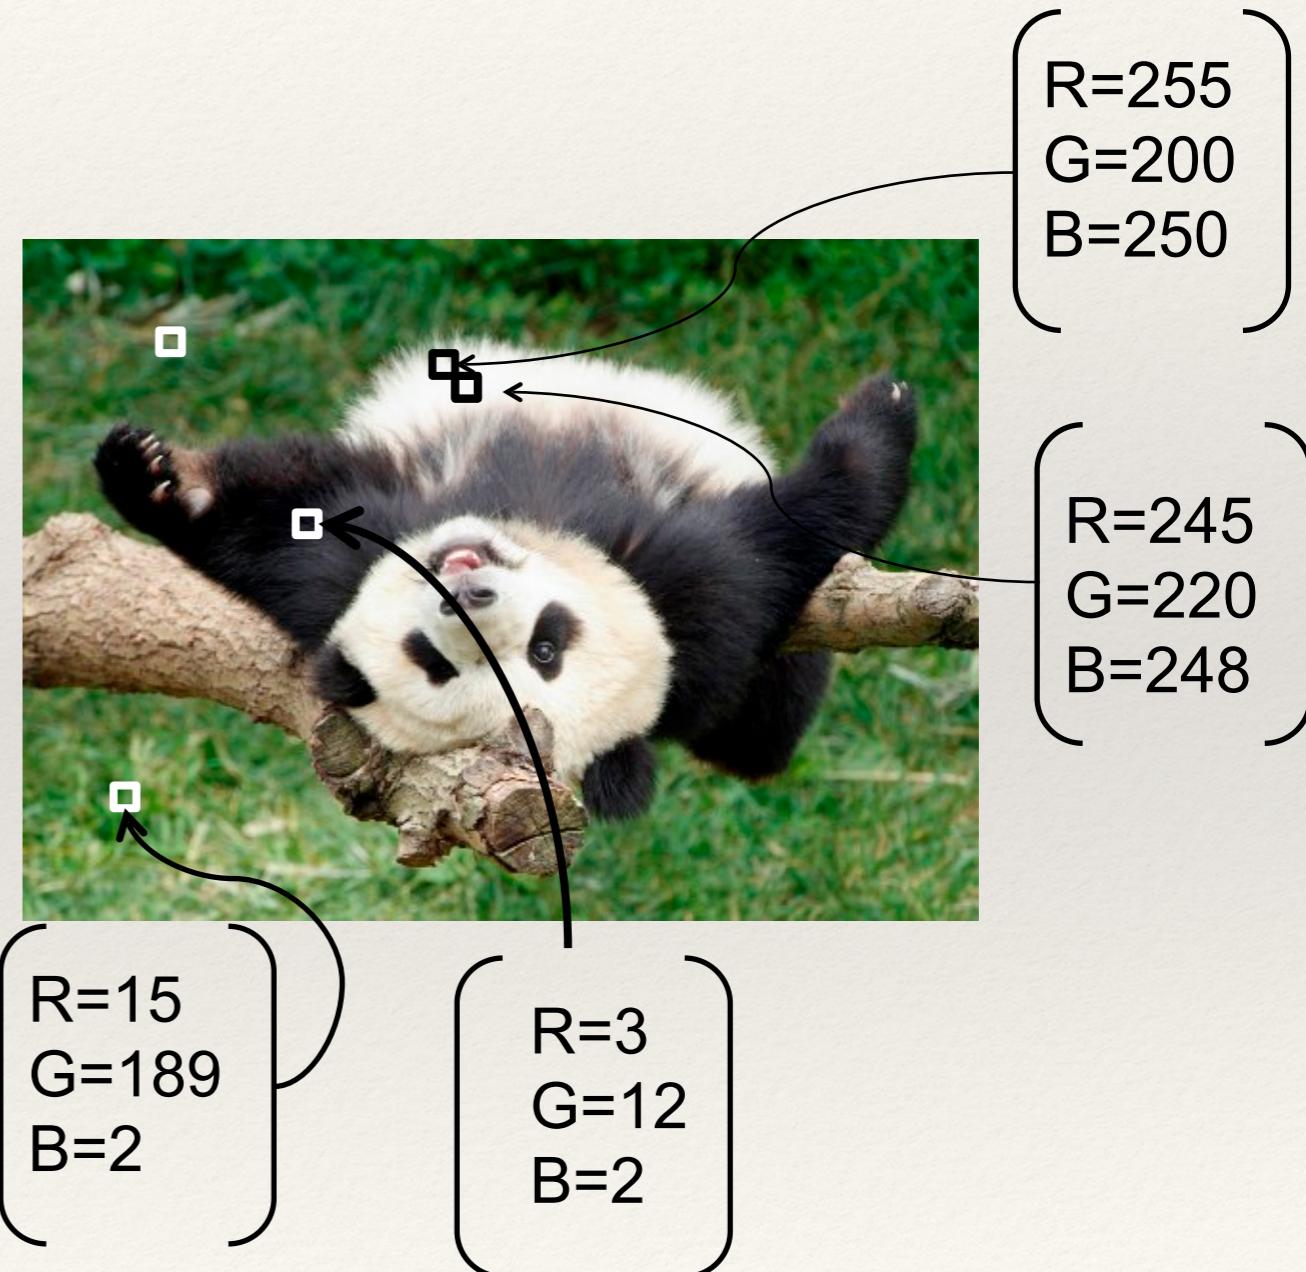
Segmentation as clustering

Depending on what we choose as the *feature space*, we can group pixels in different ways.

Grouping pixels based on **color** similarity



Feature space: color value (3-d)



Slide credit: Kristen Grauman

Segmentation as clustering

Depending on what we choose as the *feature space*, we can group pixels in different ways.

Grouping pixels based
on **intensity** similarity



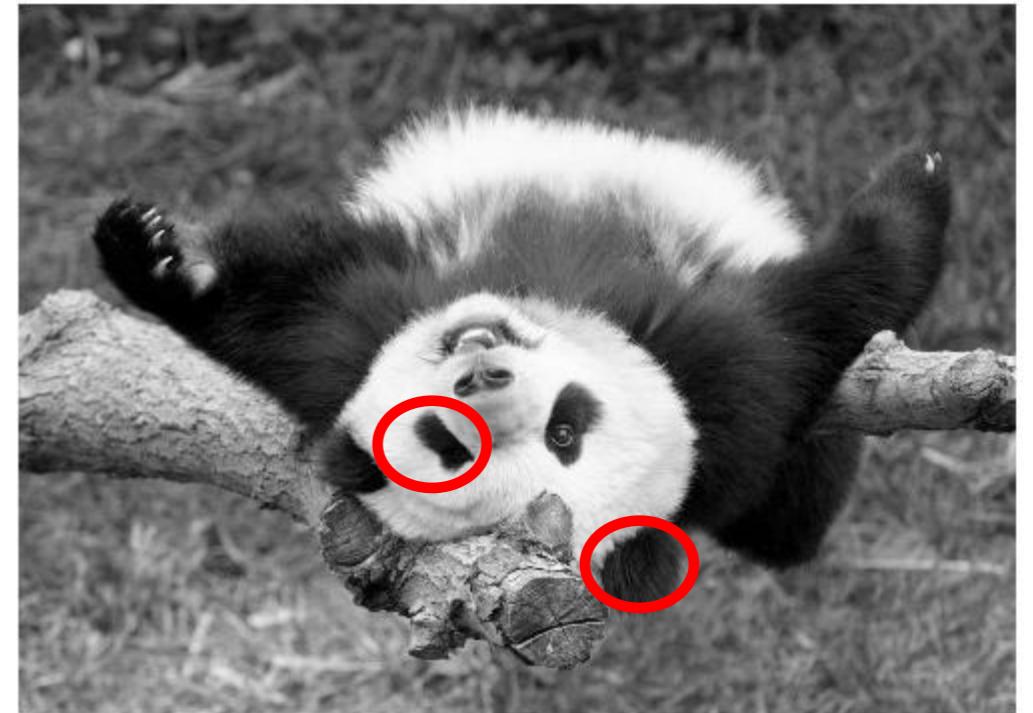
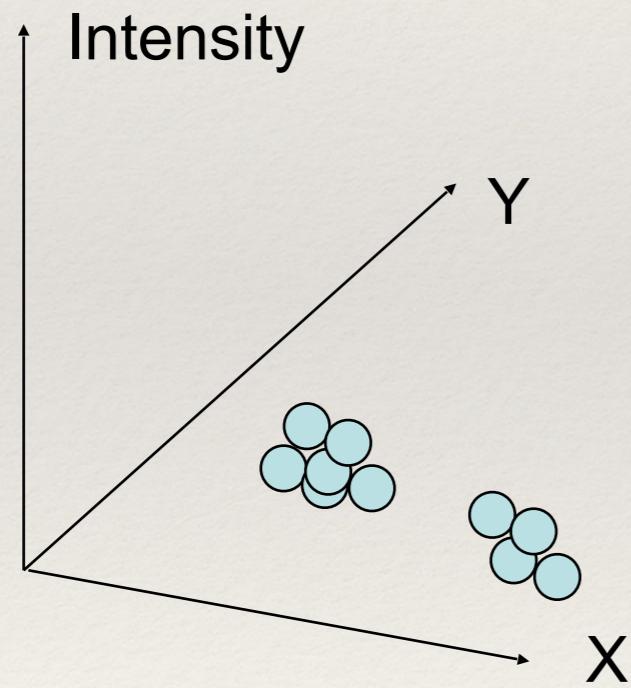
Clusters based on intensity
similarity don't have to be
spatially coherent.



Segmentation as clustering

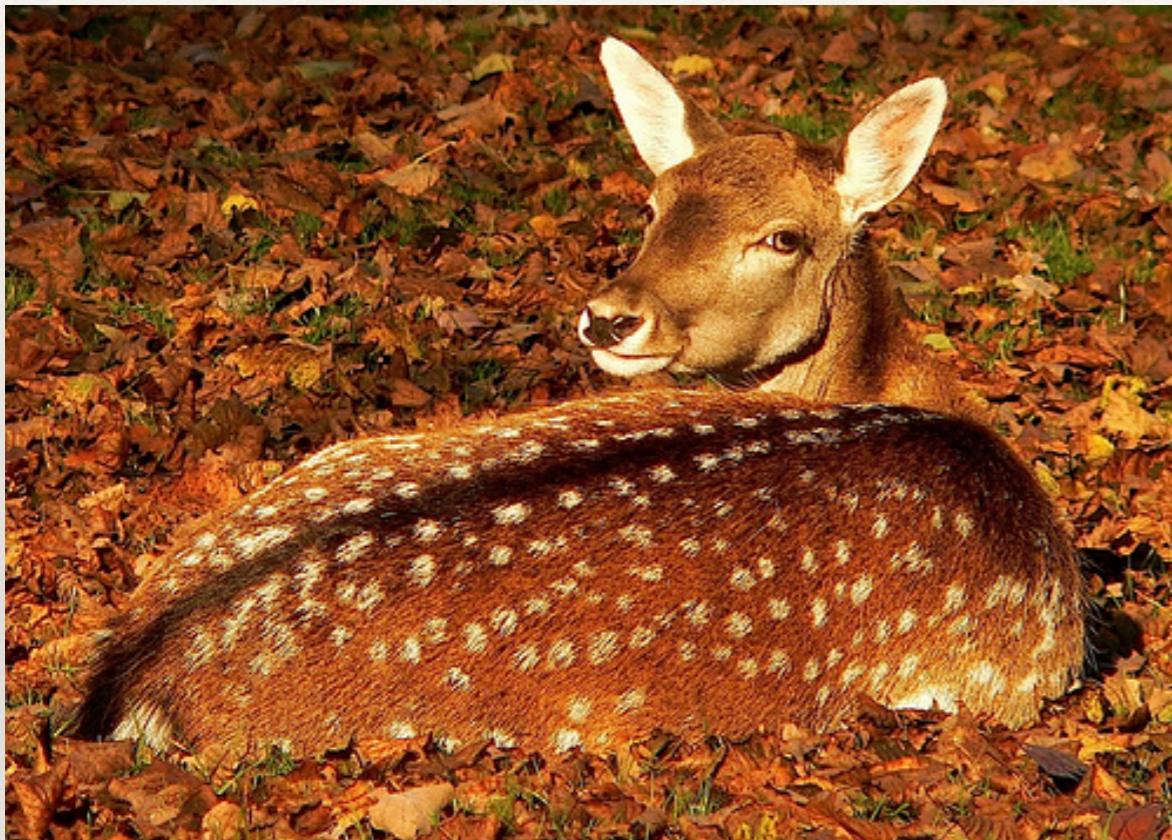
Depending on what we choose as the *feature space*, we can group pixels in different ways.

Grouping pixels based on
intensity+position similarity



Both regions are black, but if we also include **position (x,y)**, then we could group the two into distinct segments; way to encode both similarity & proximity.

Segmentation as clustering

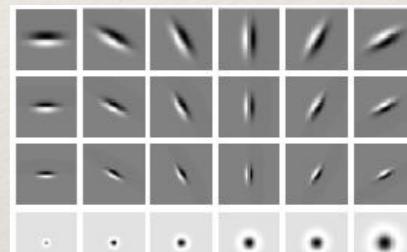
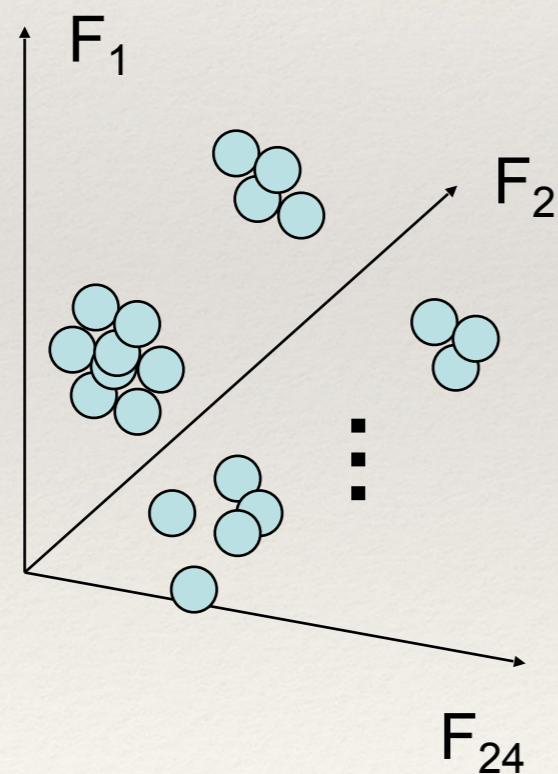


- ❖ Color, brightness, position alone are not enough to distinguish all regions...

Segmentation as clustering

Depending on what we choose as the *feature space*, we can group pixels in different ways.

Grouping pixels based
on **texture** similarity



Filter bank
of 24 filters

Feature space: filter bank responses (e.g., 24-d)

Segmentation with texture features

- Find “textons” by **clustering** vectors of filter bank outputs
- Describe texture in a window based on *texton histogram*

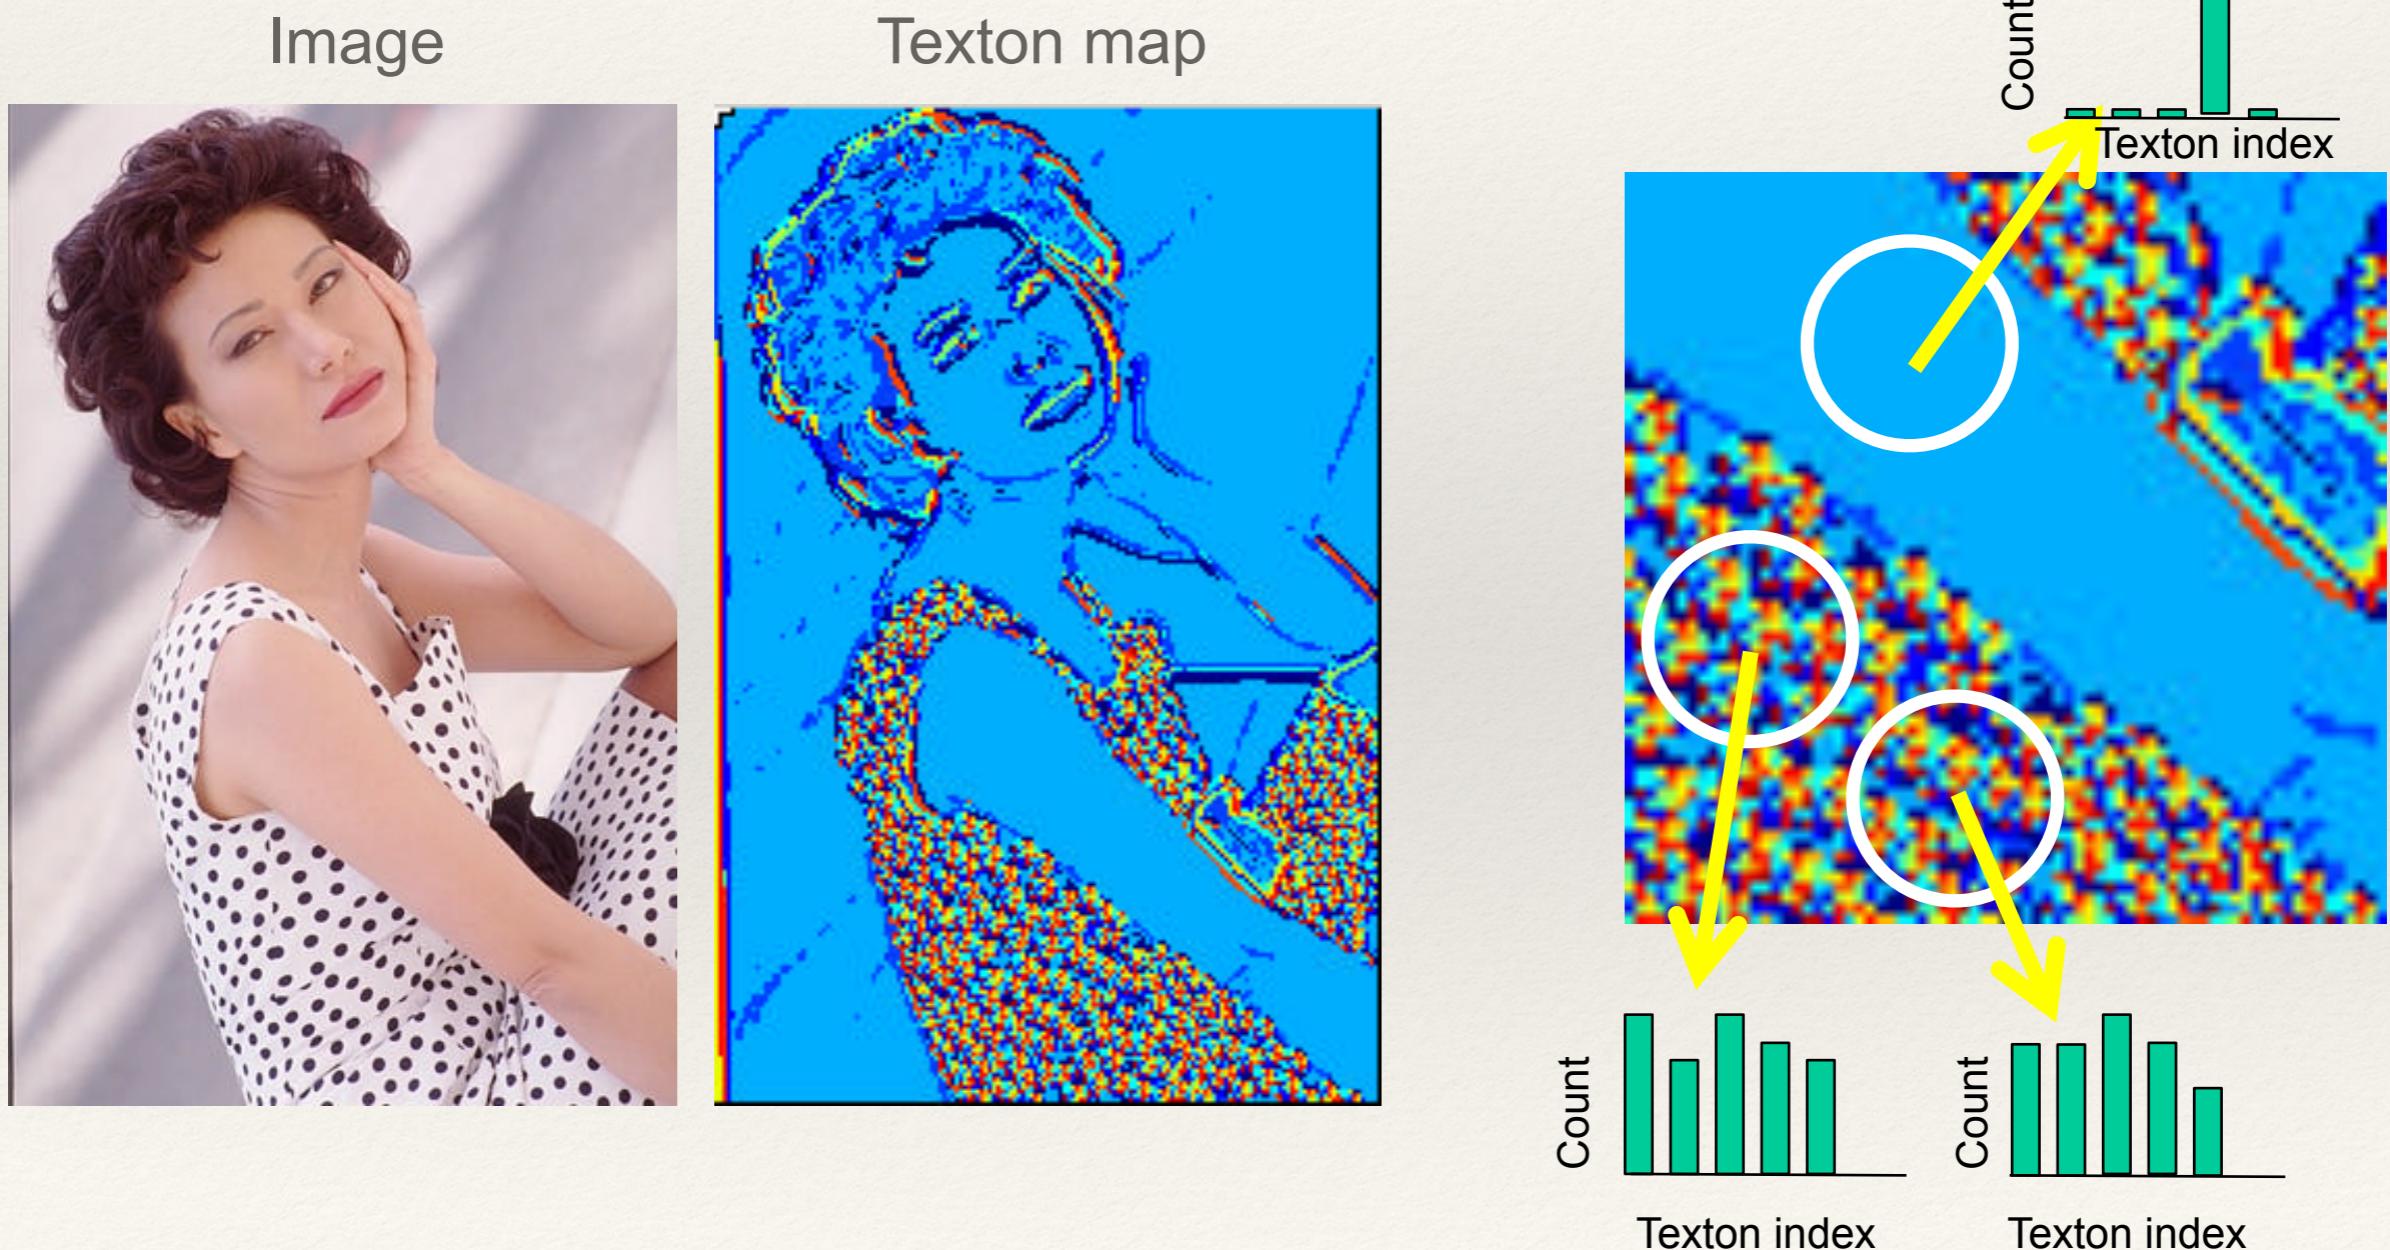
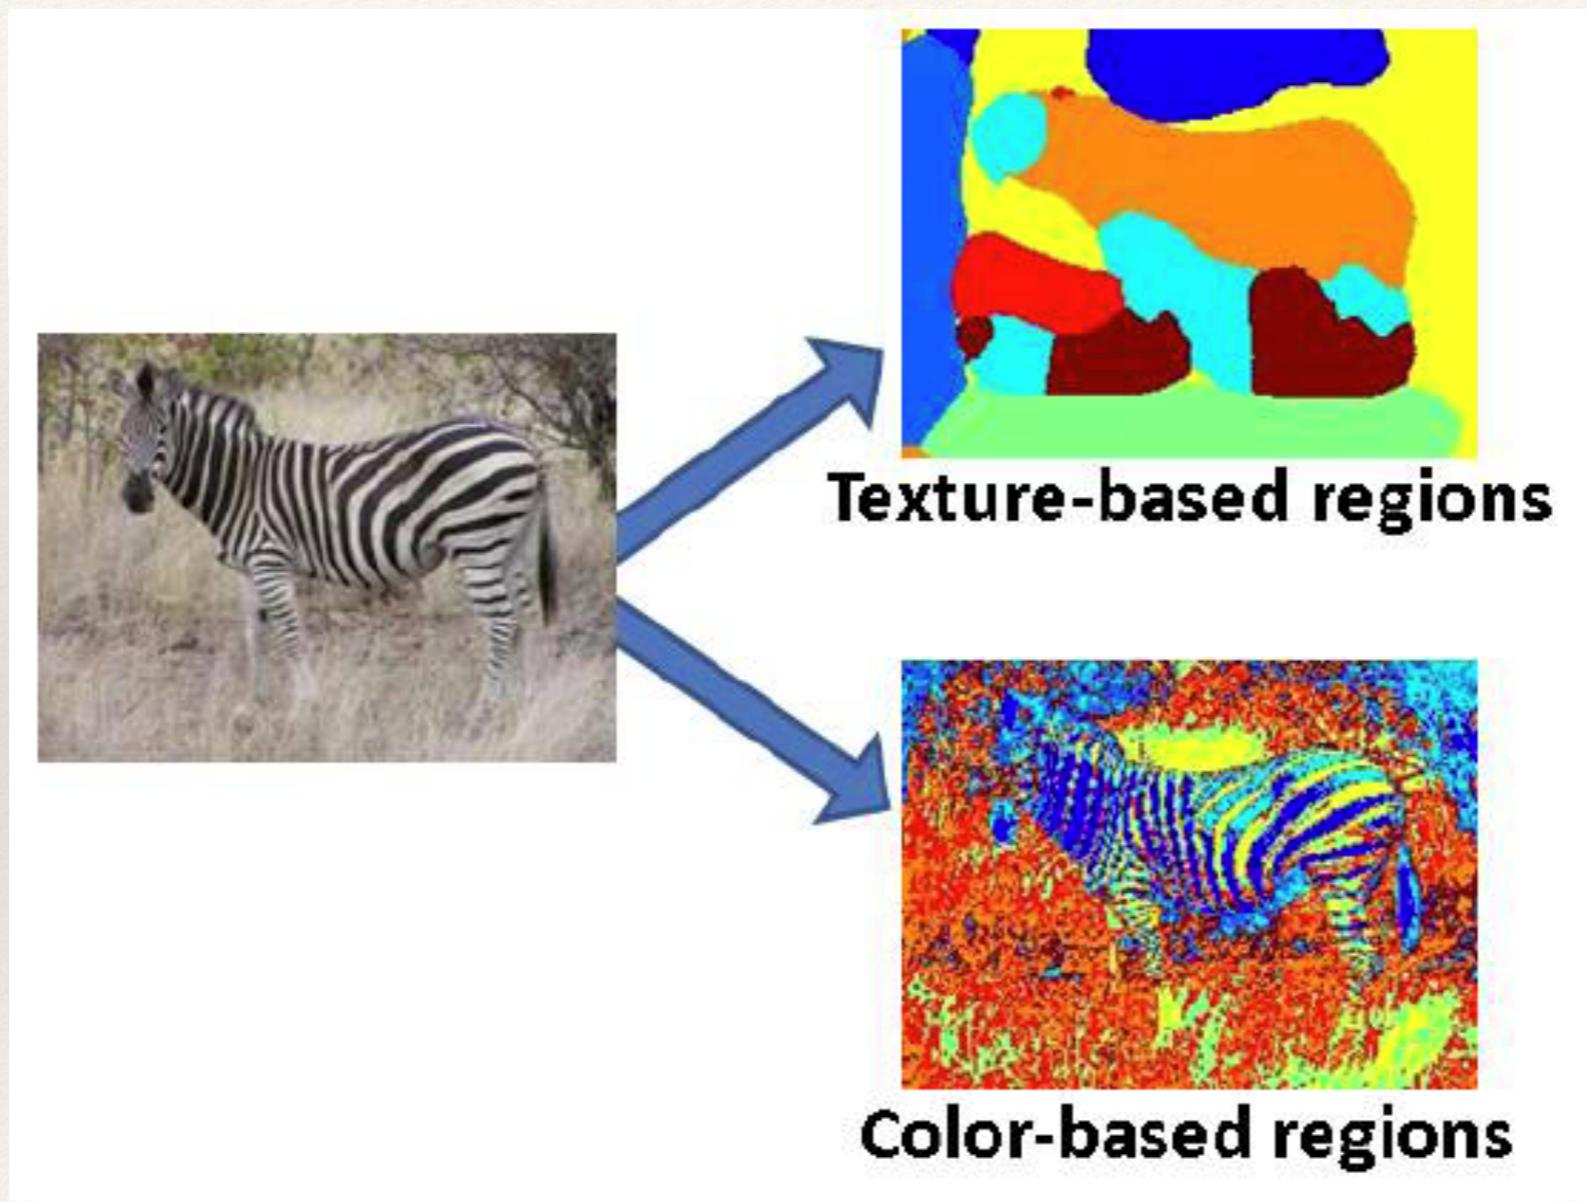


Image segmentation example



Pixel properties vs. neighborhood properties

query



query

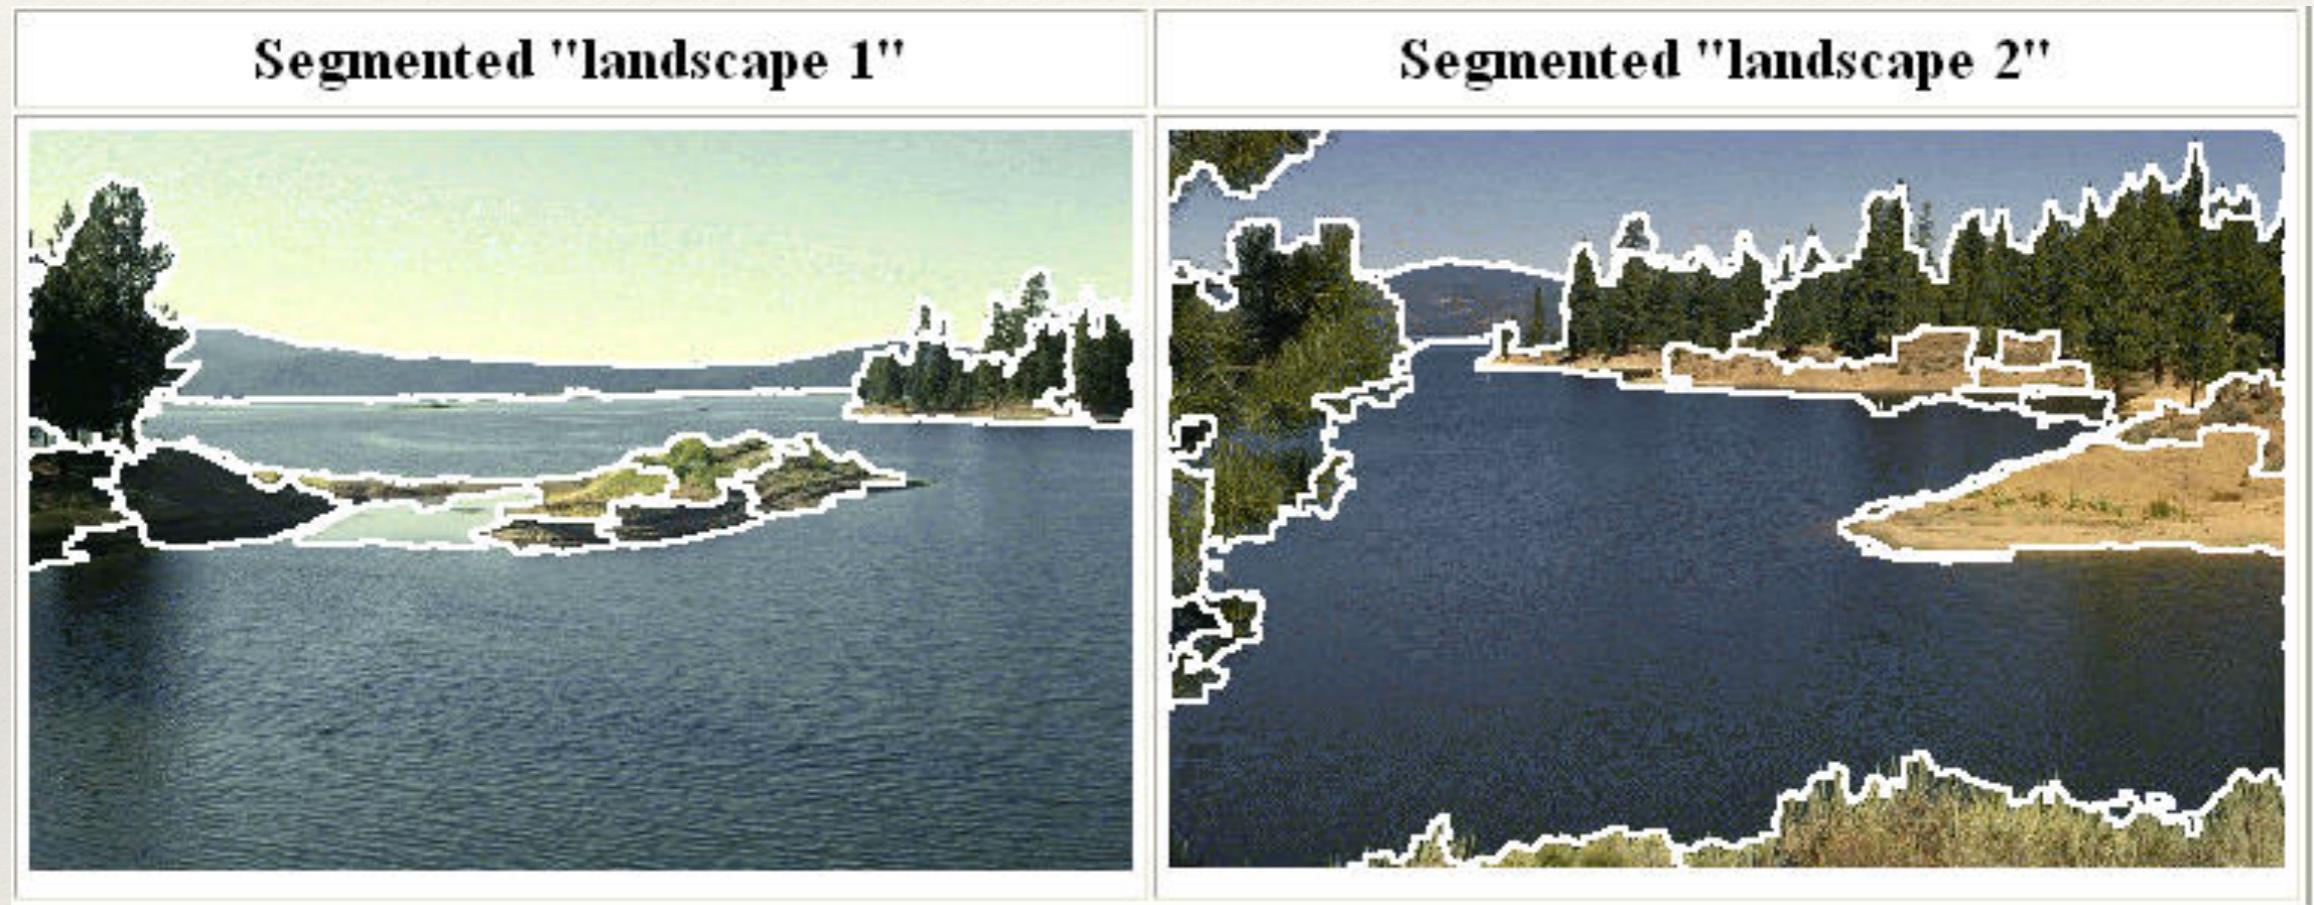


These look very similar in terms of their color distributions (histograms).

How would their *texture* distributions compare?

Mean shift clustering and segmentation

- ❖ An advanced and versatile technique for clustering-based segmentation



<http://www.caip.rutgers.edu/~comanici/MSPAMI/msPamiResults.html>

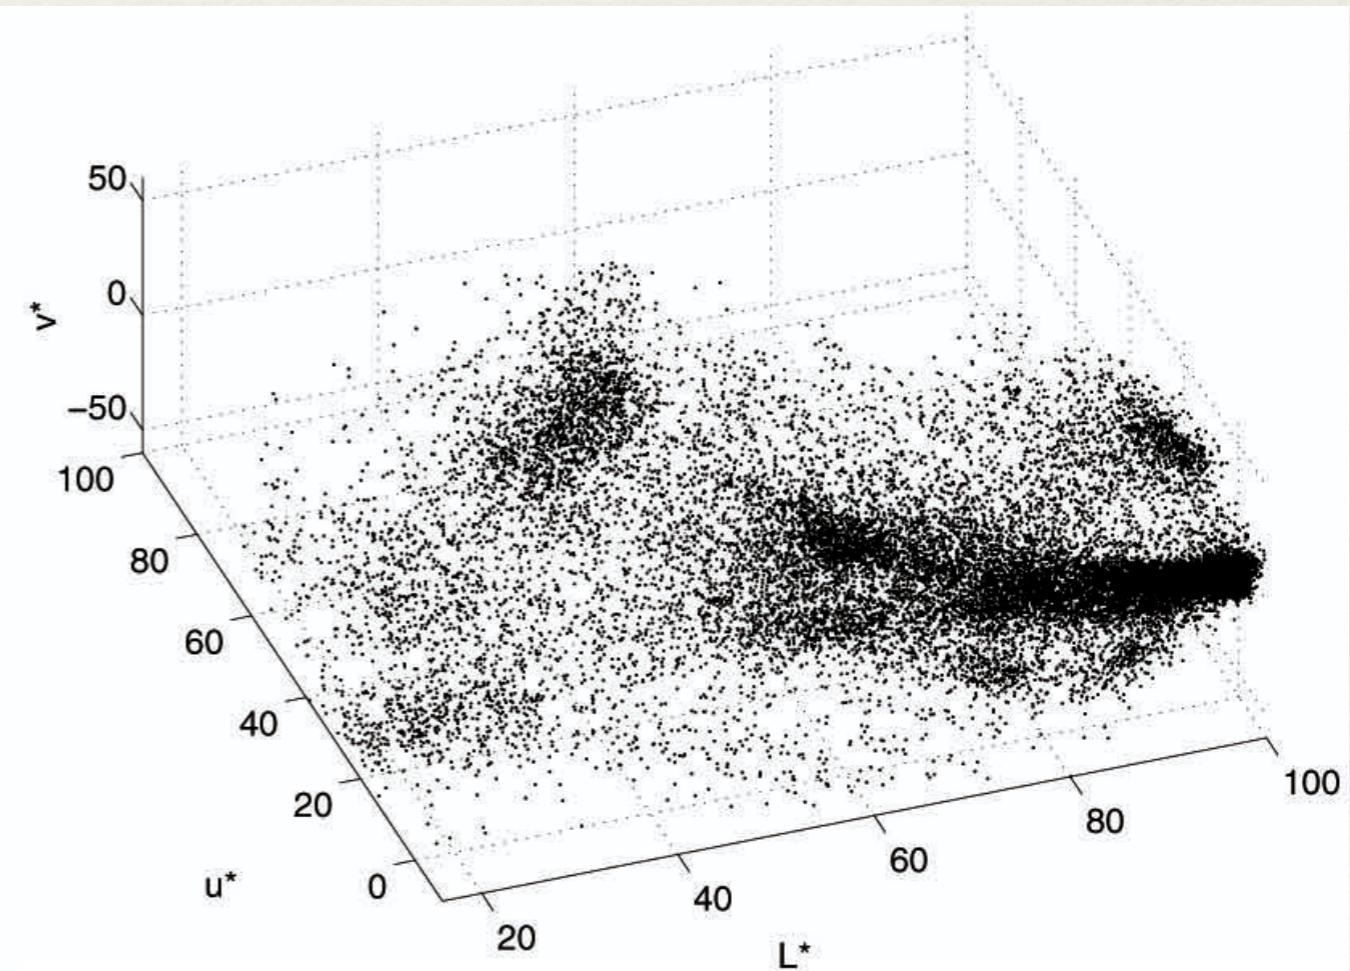
D. Comaniciu and P. Meer, [Mean Shift: A Robust Approach toward Feature Space Analysis](#), PAMI 2002.

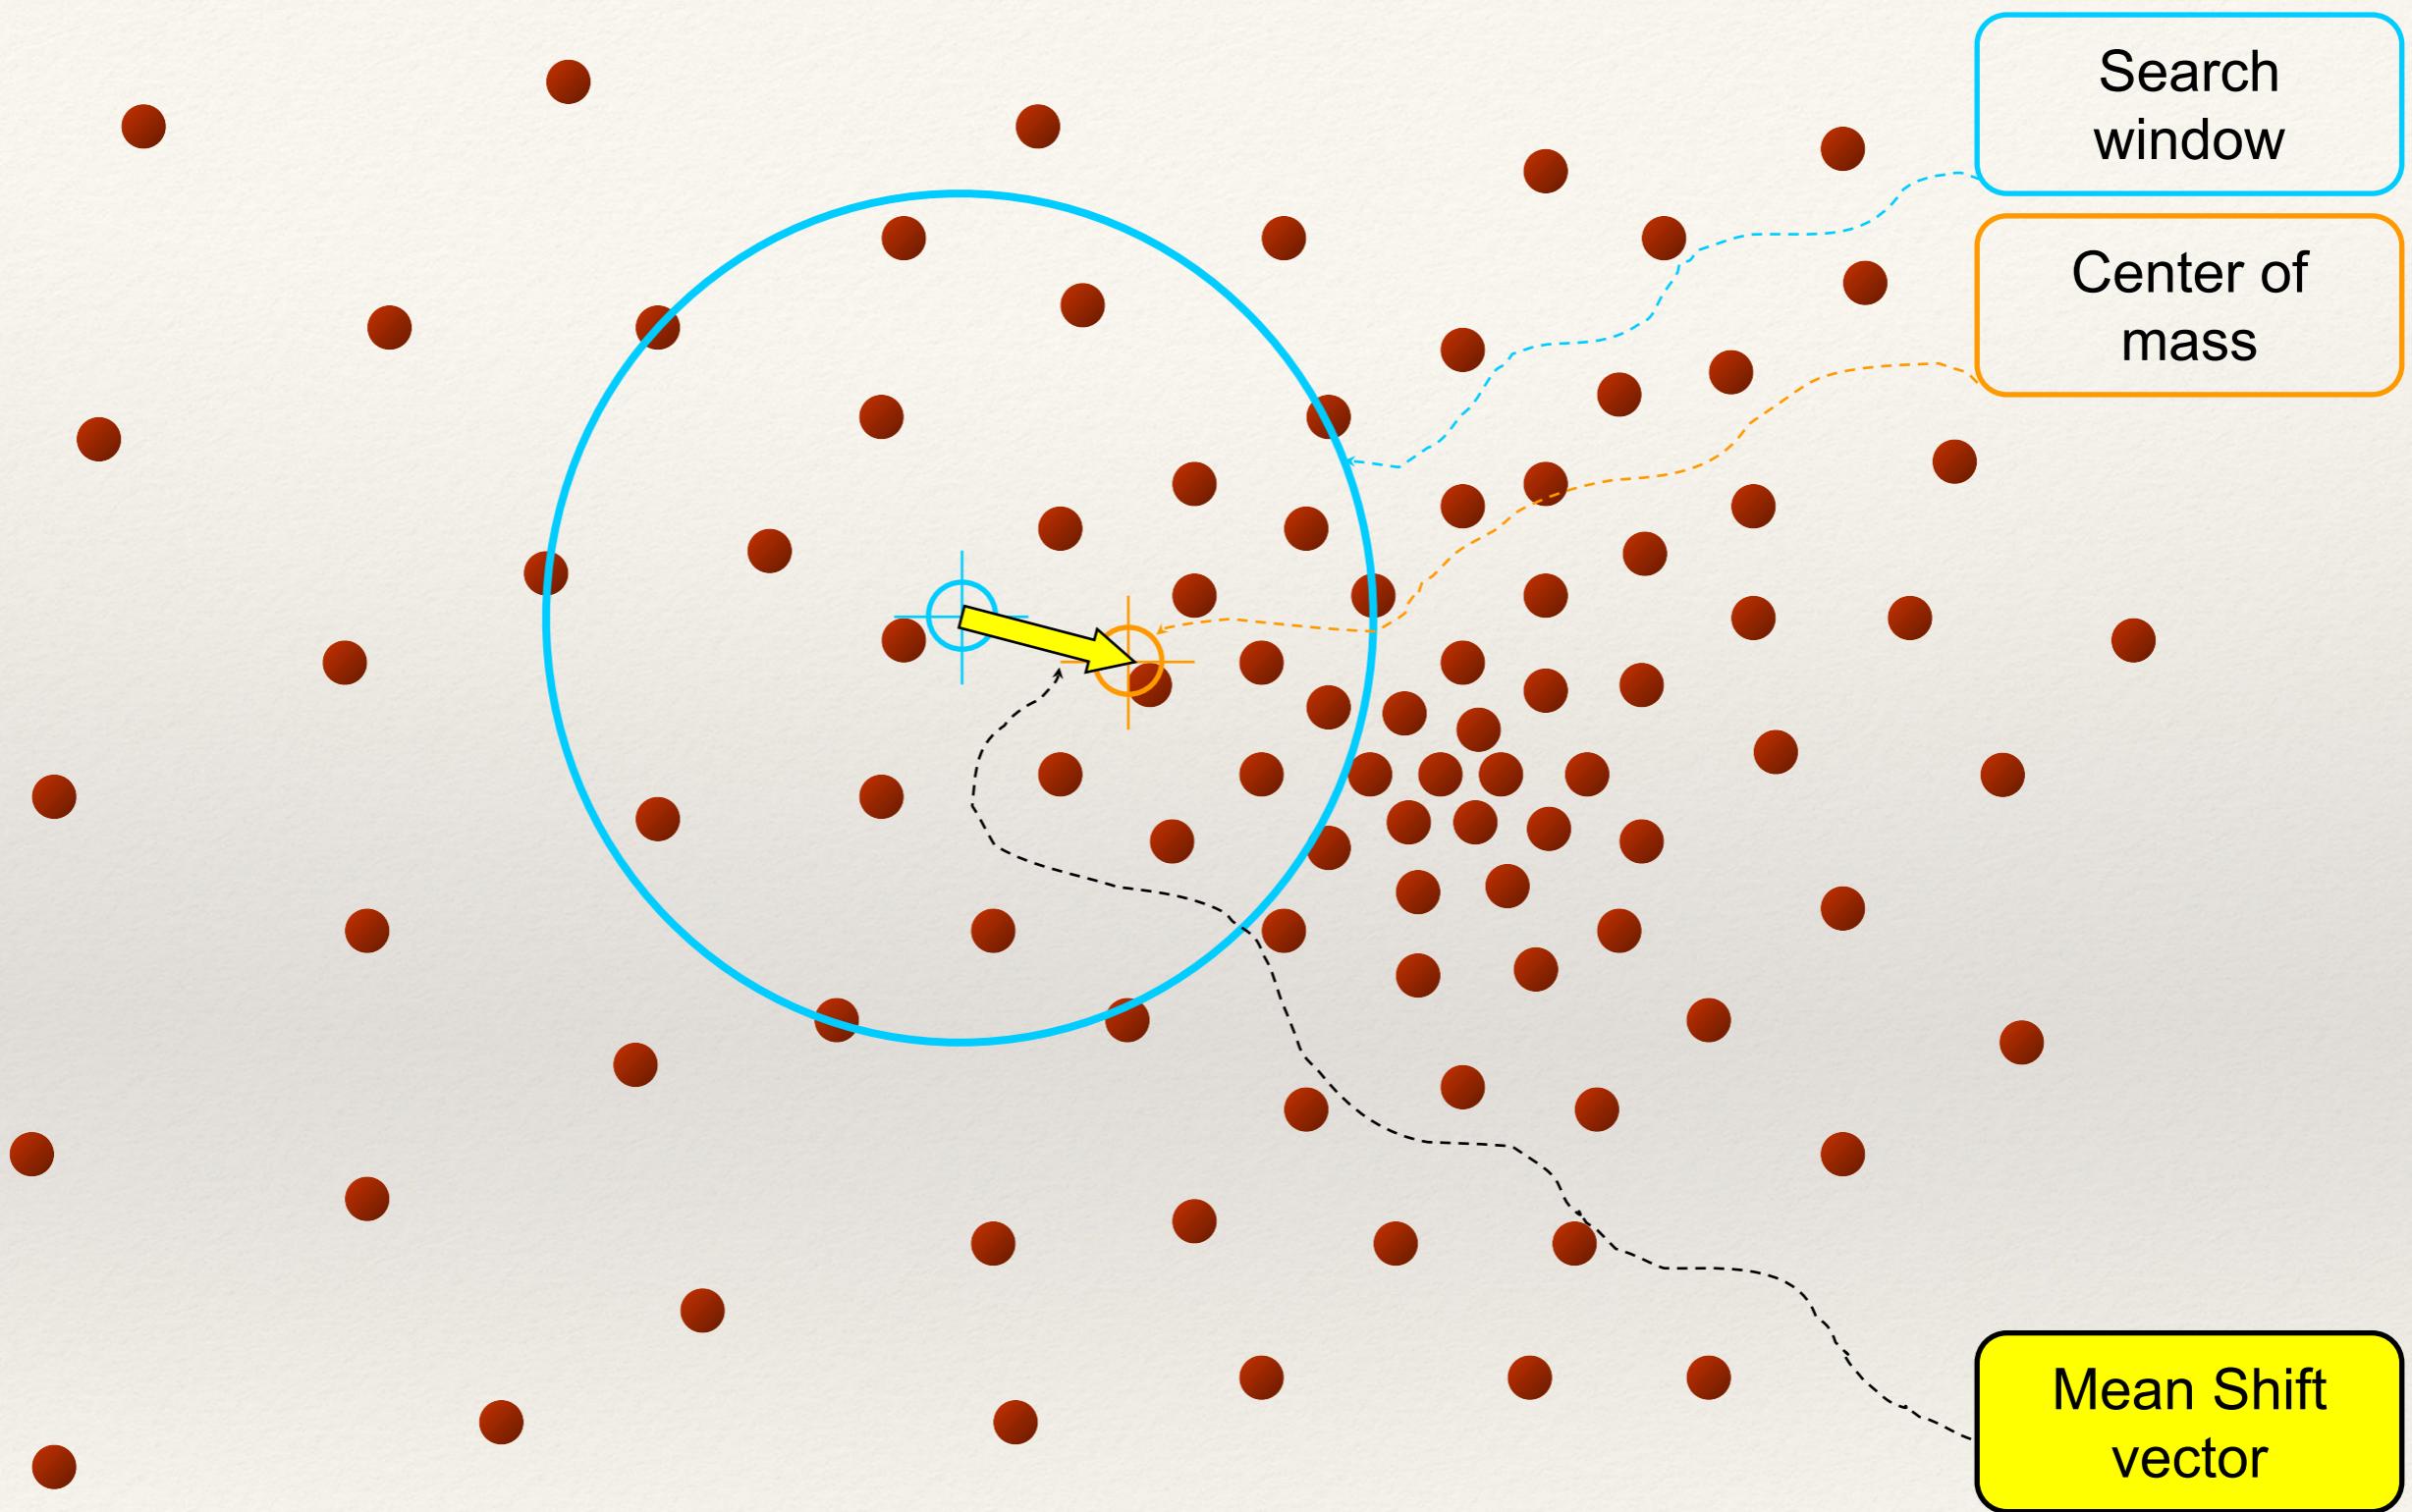
Mean shift algorithm

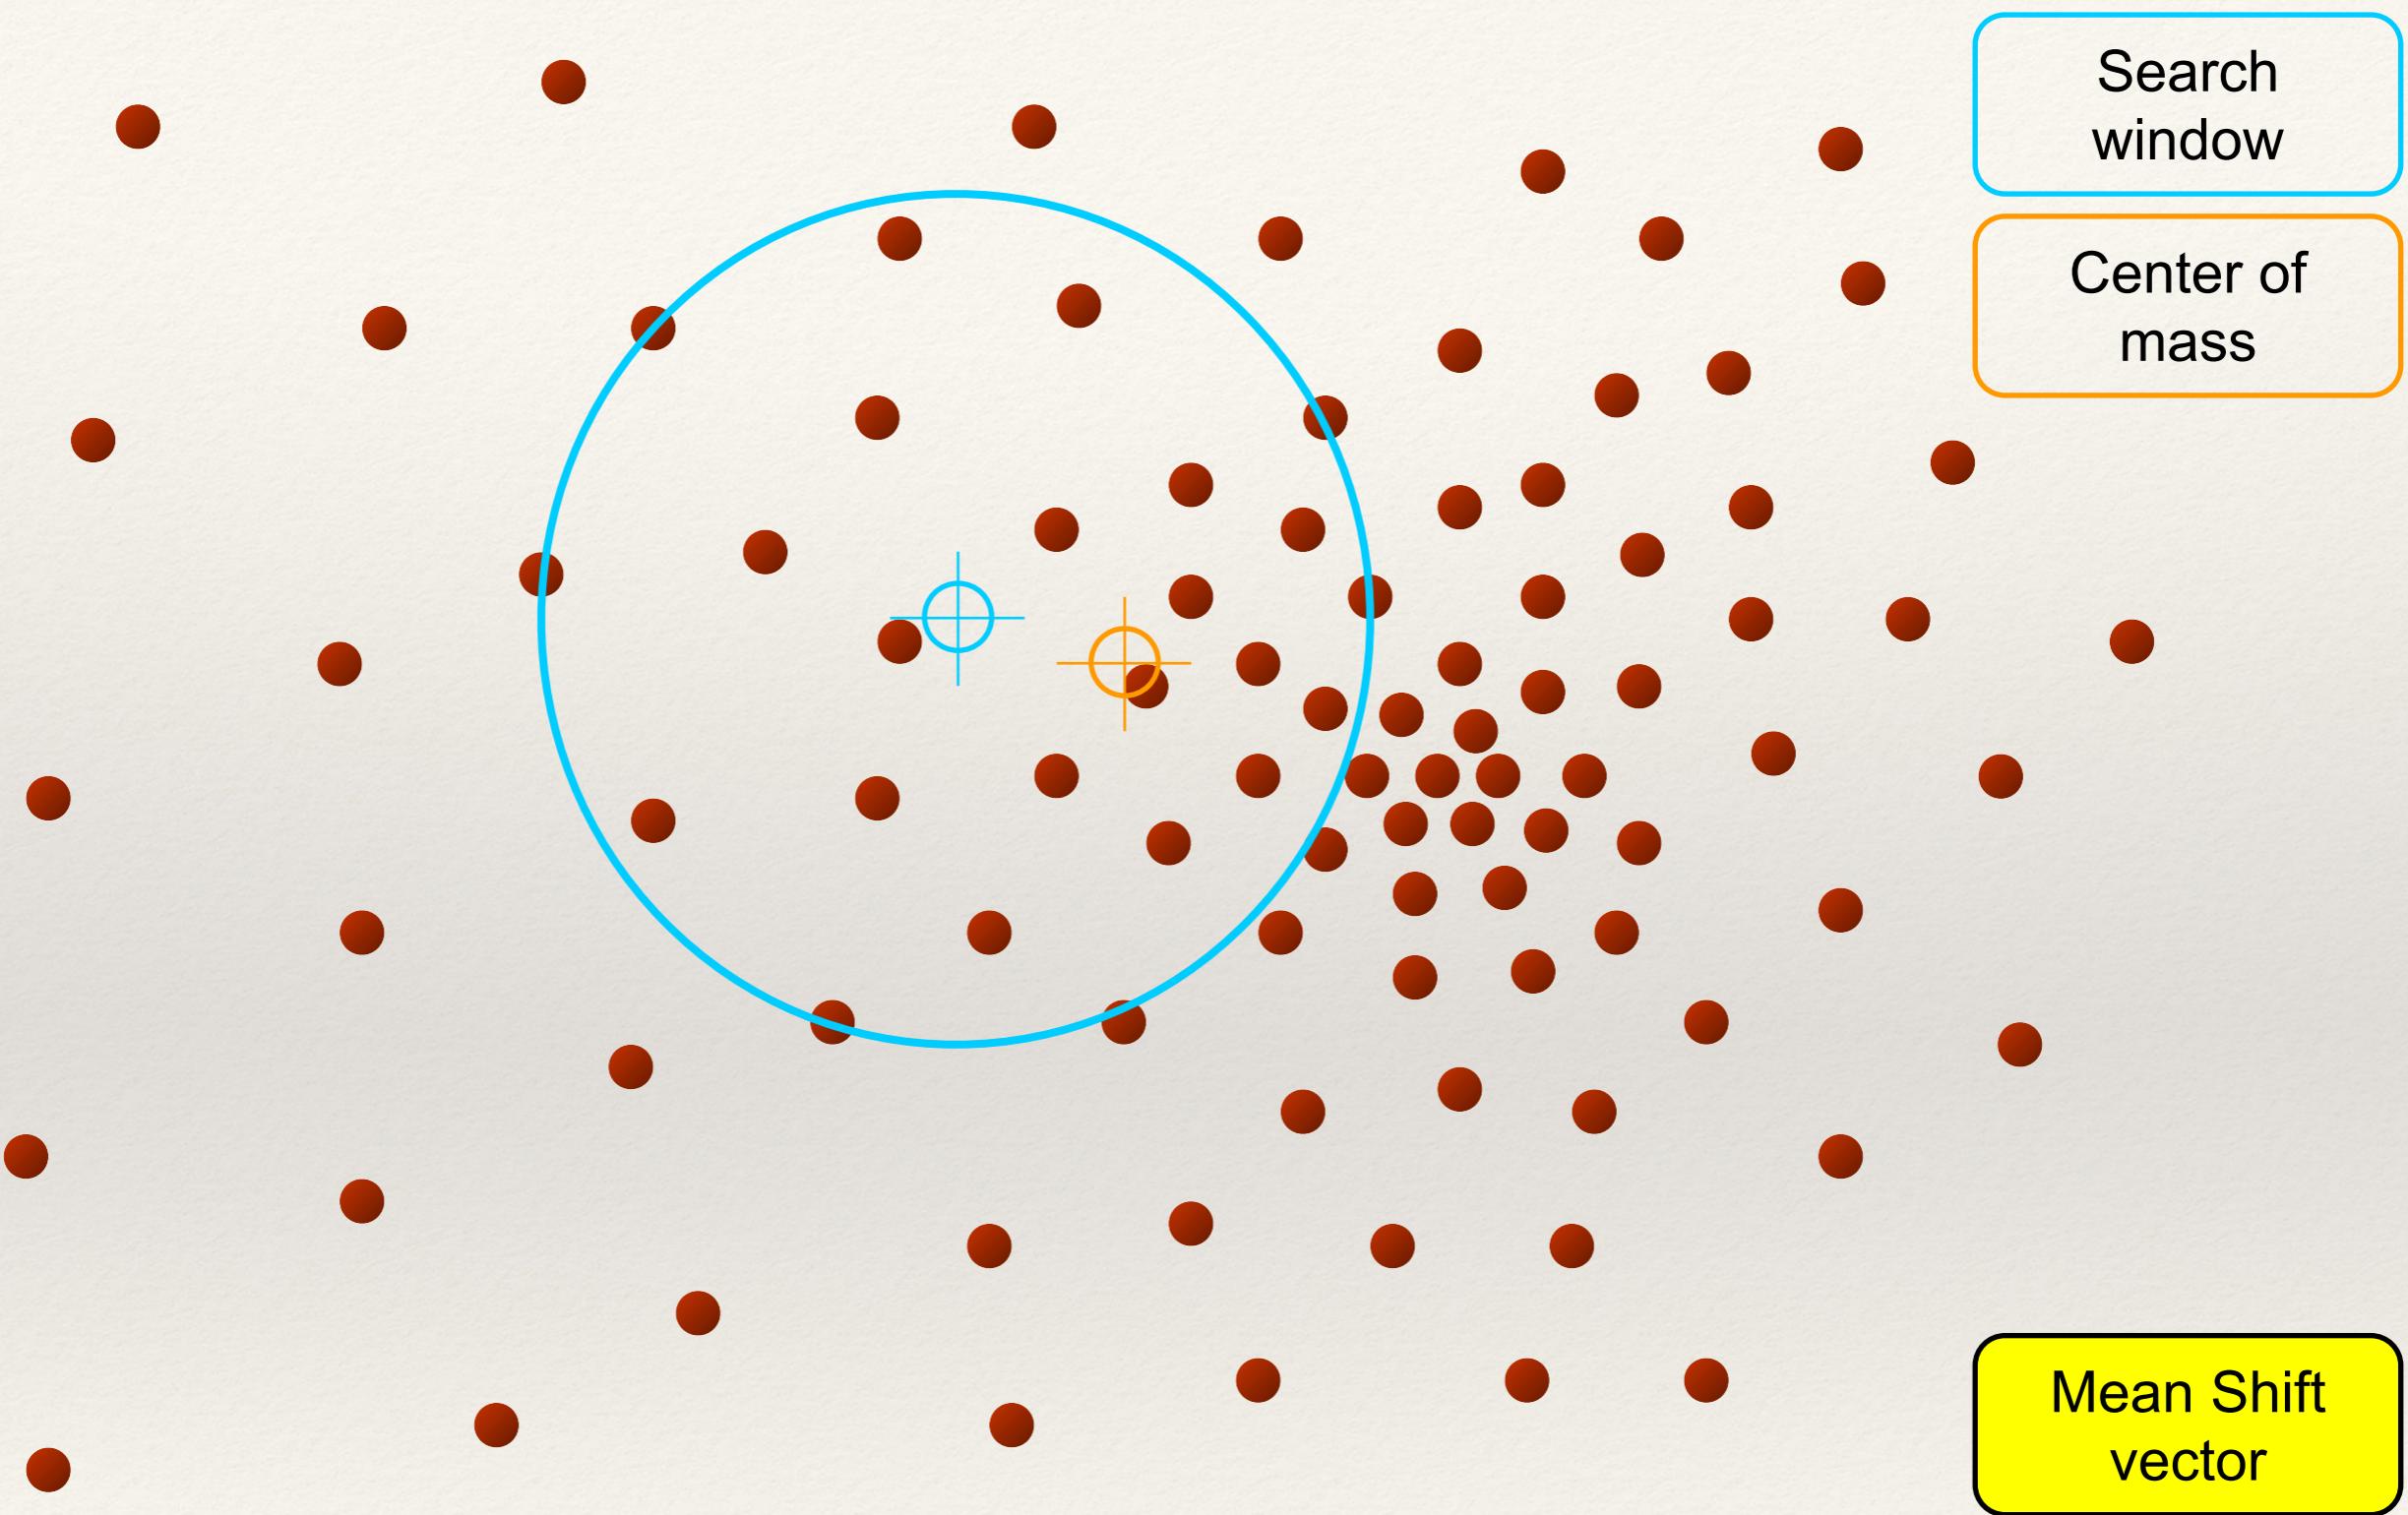
image

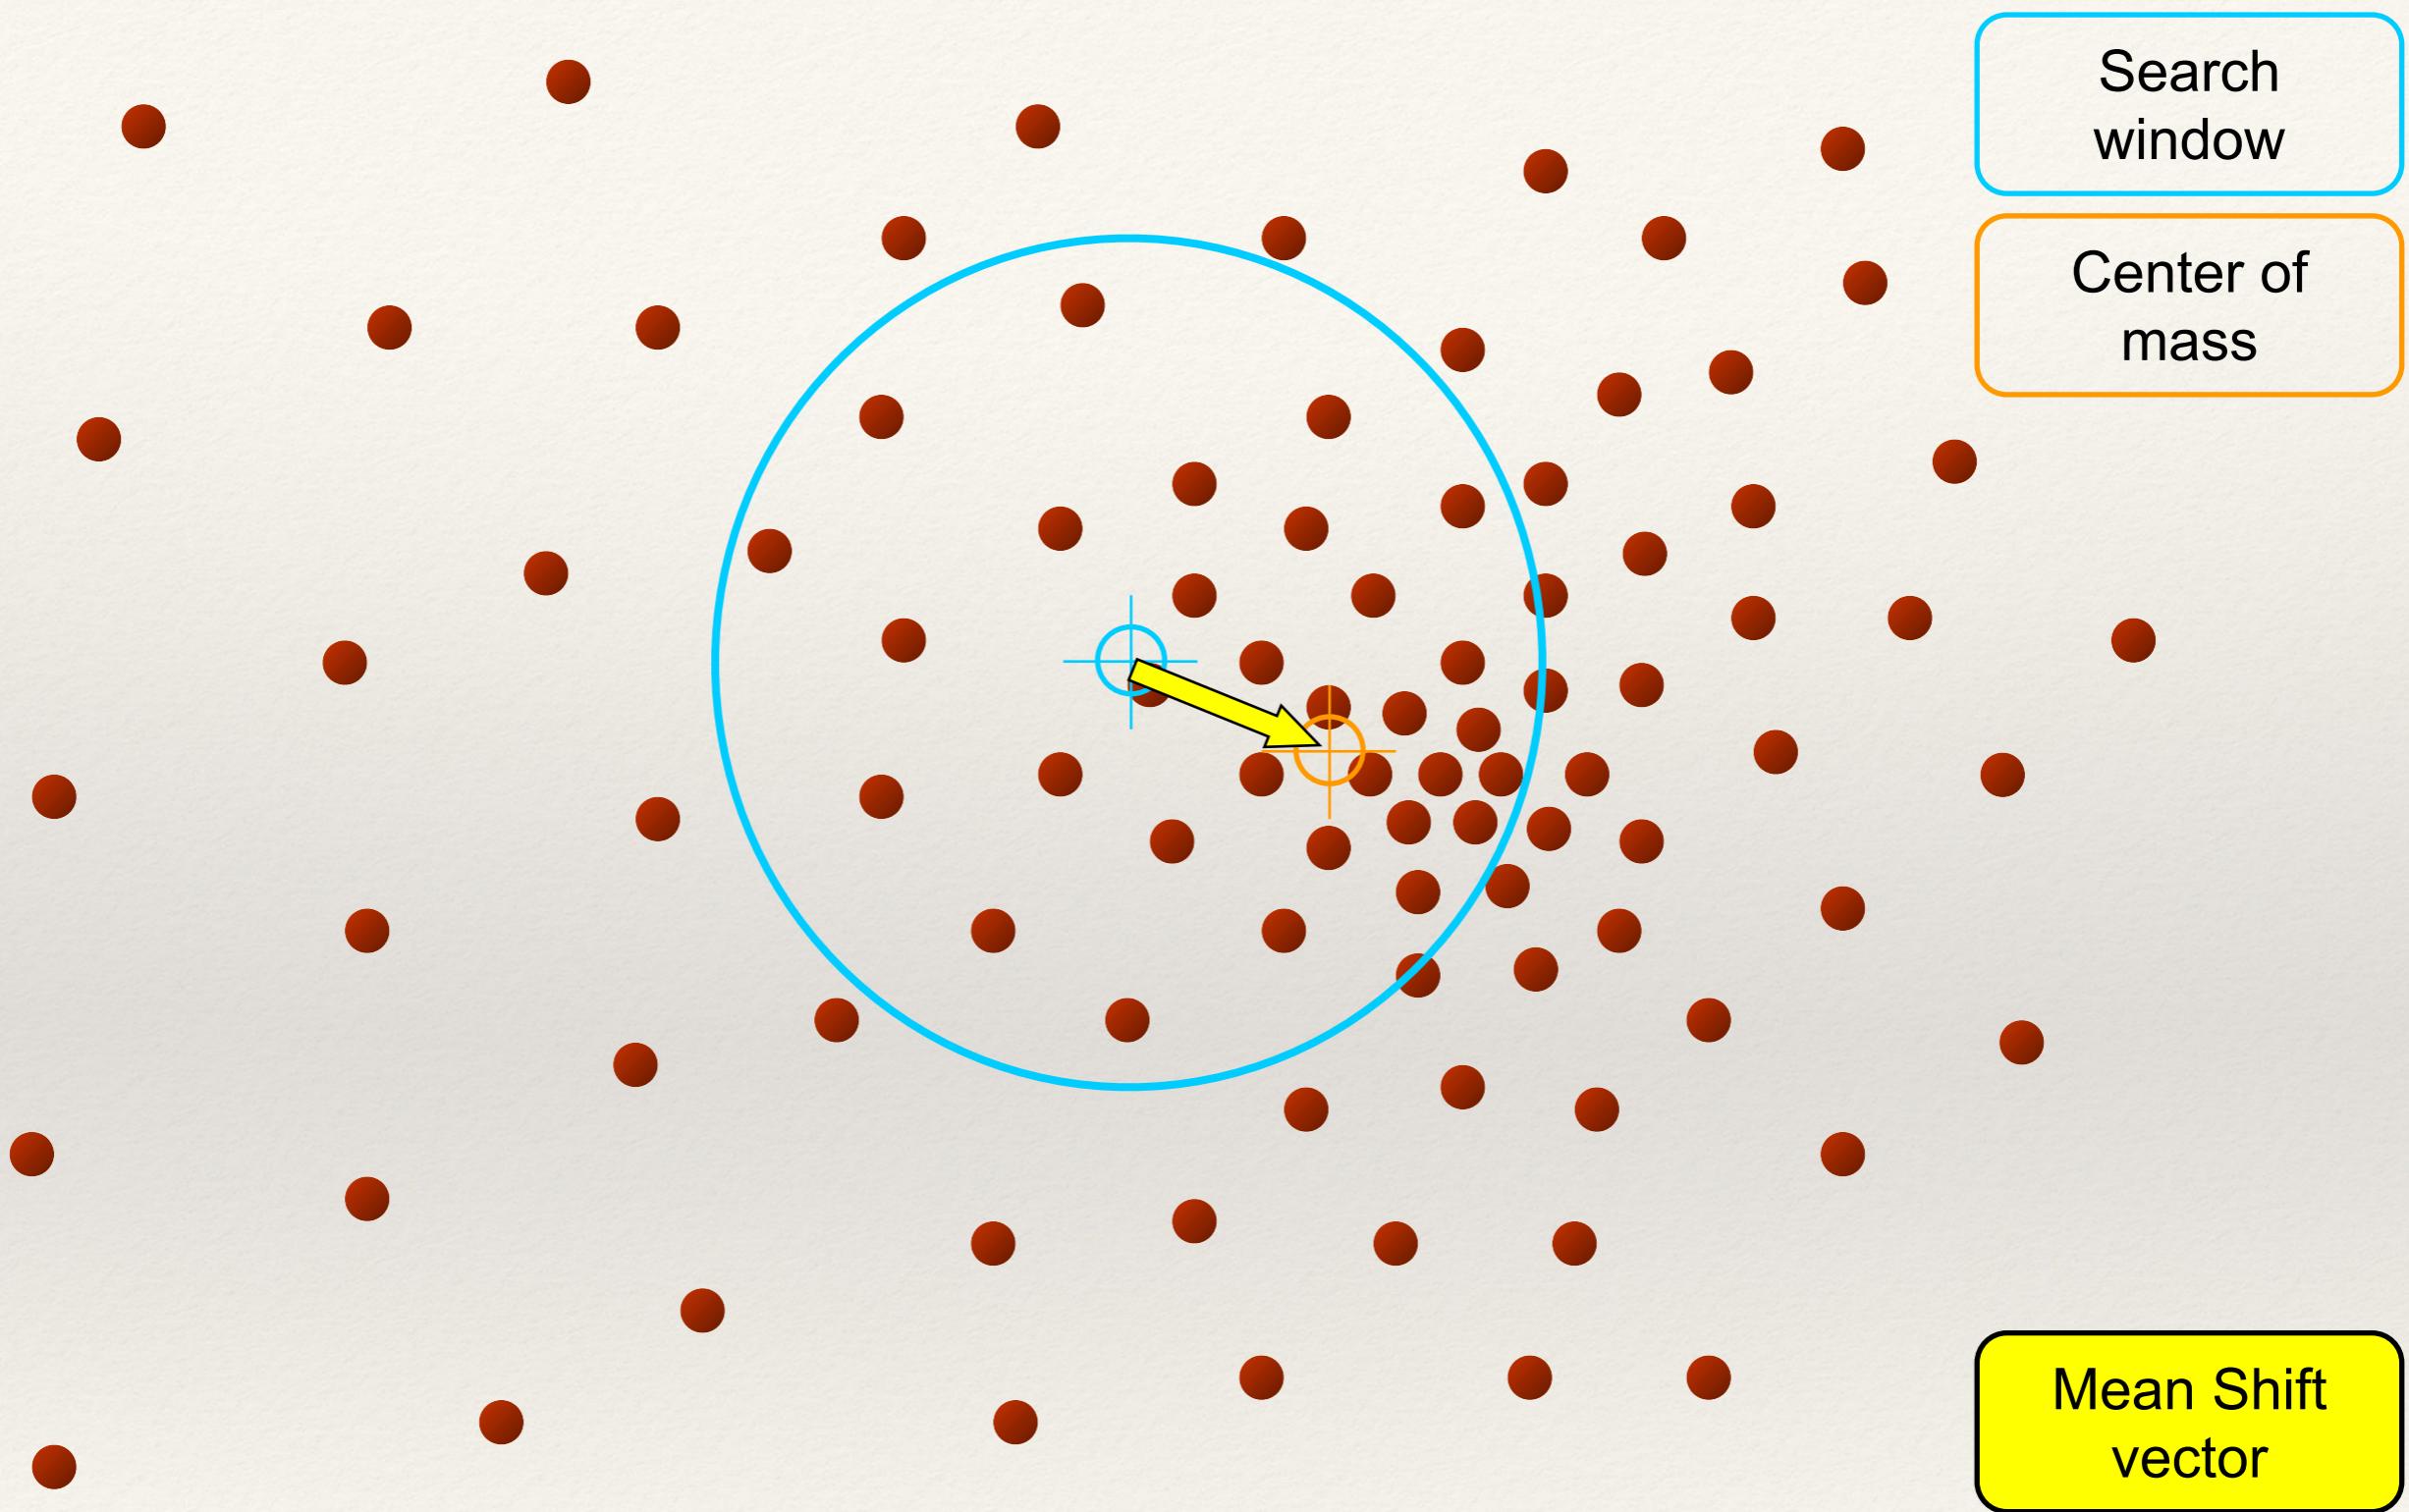


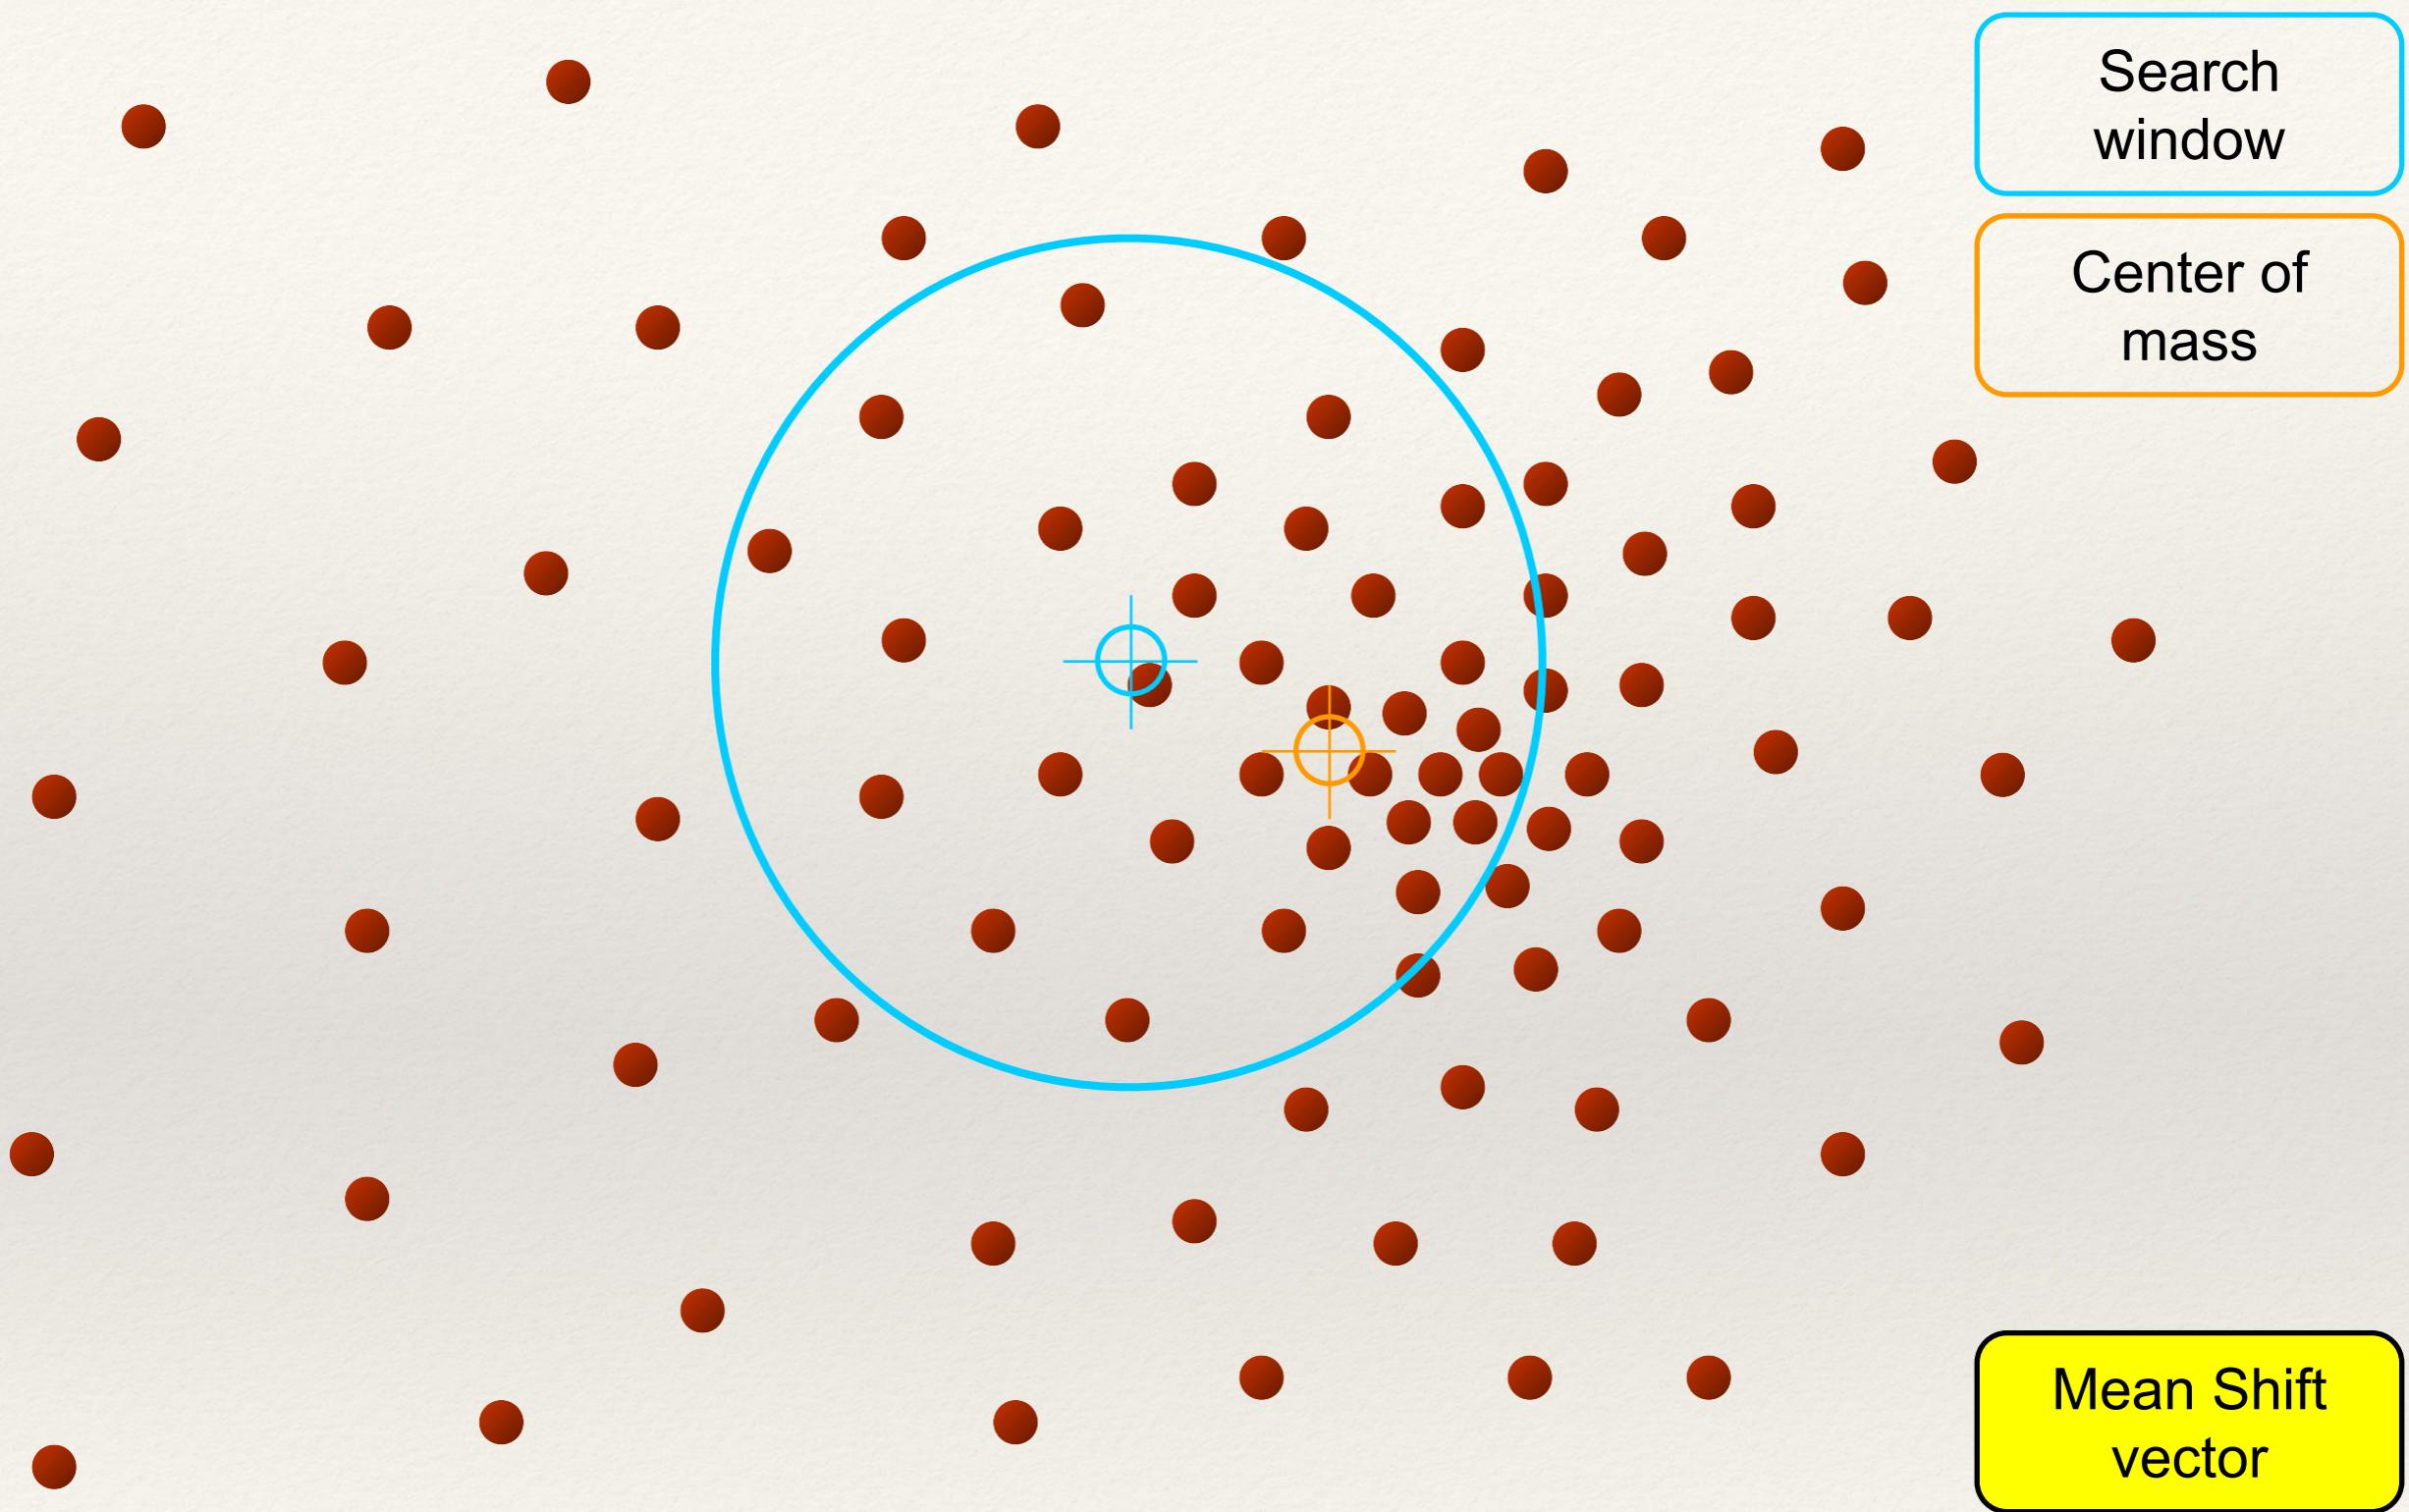
Feature space
($L^*u^*v^*$ color values)

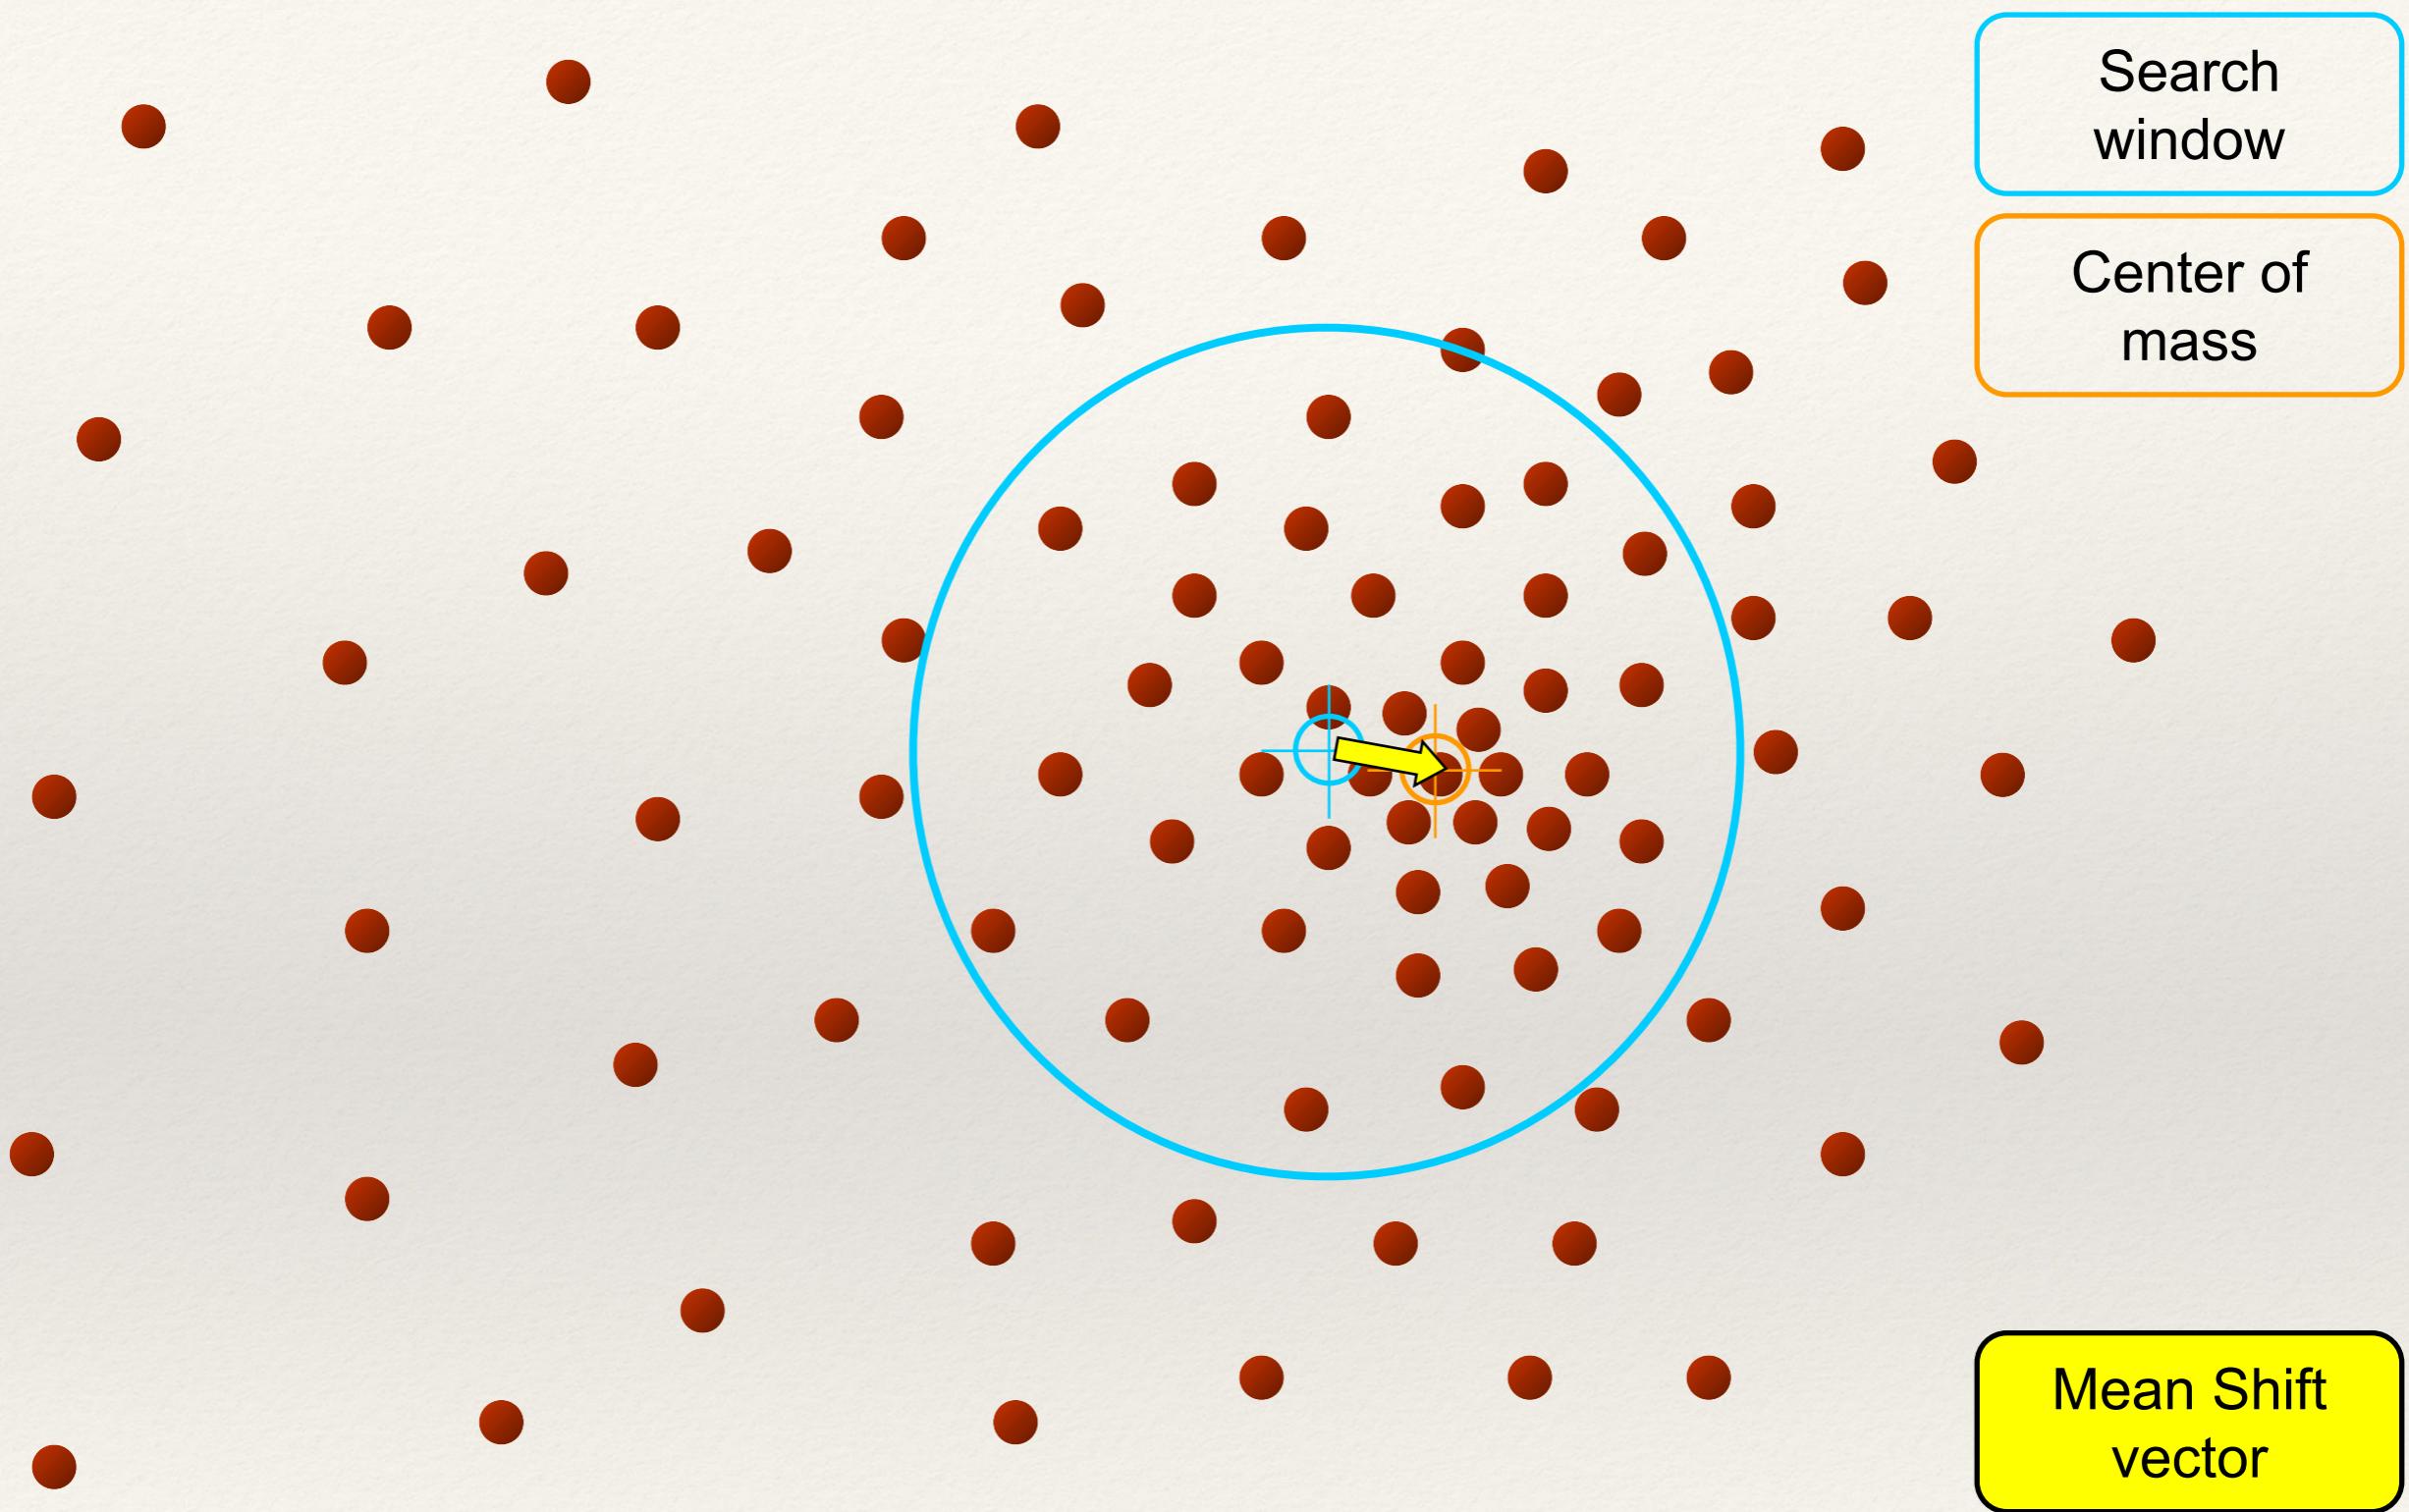


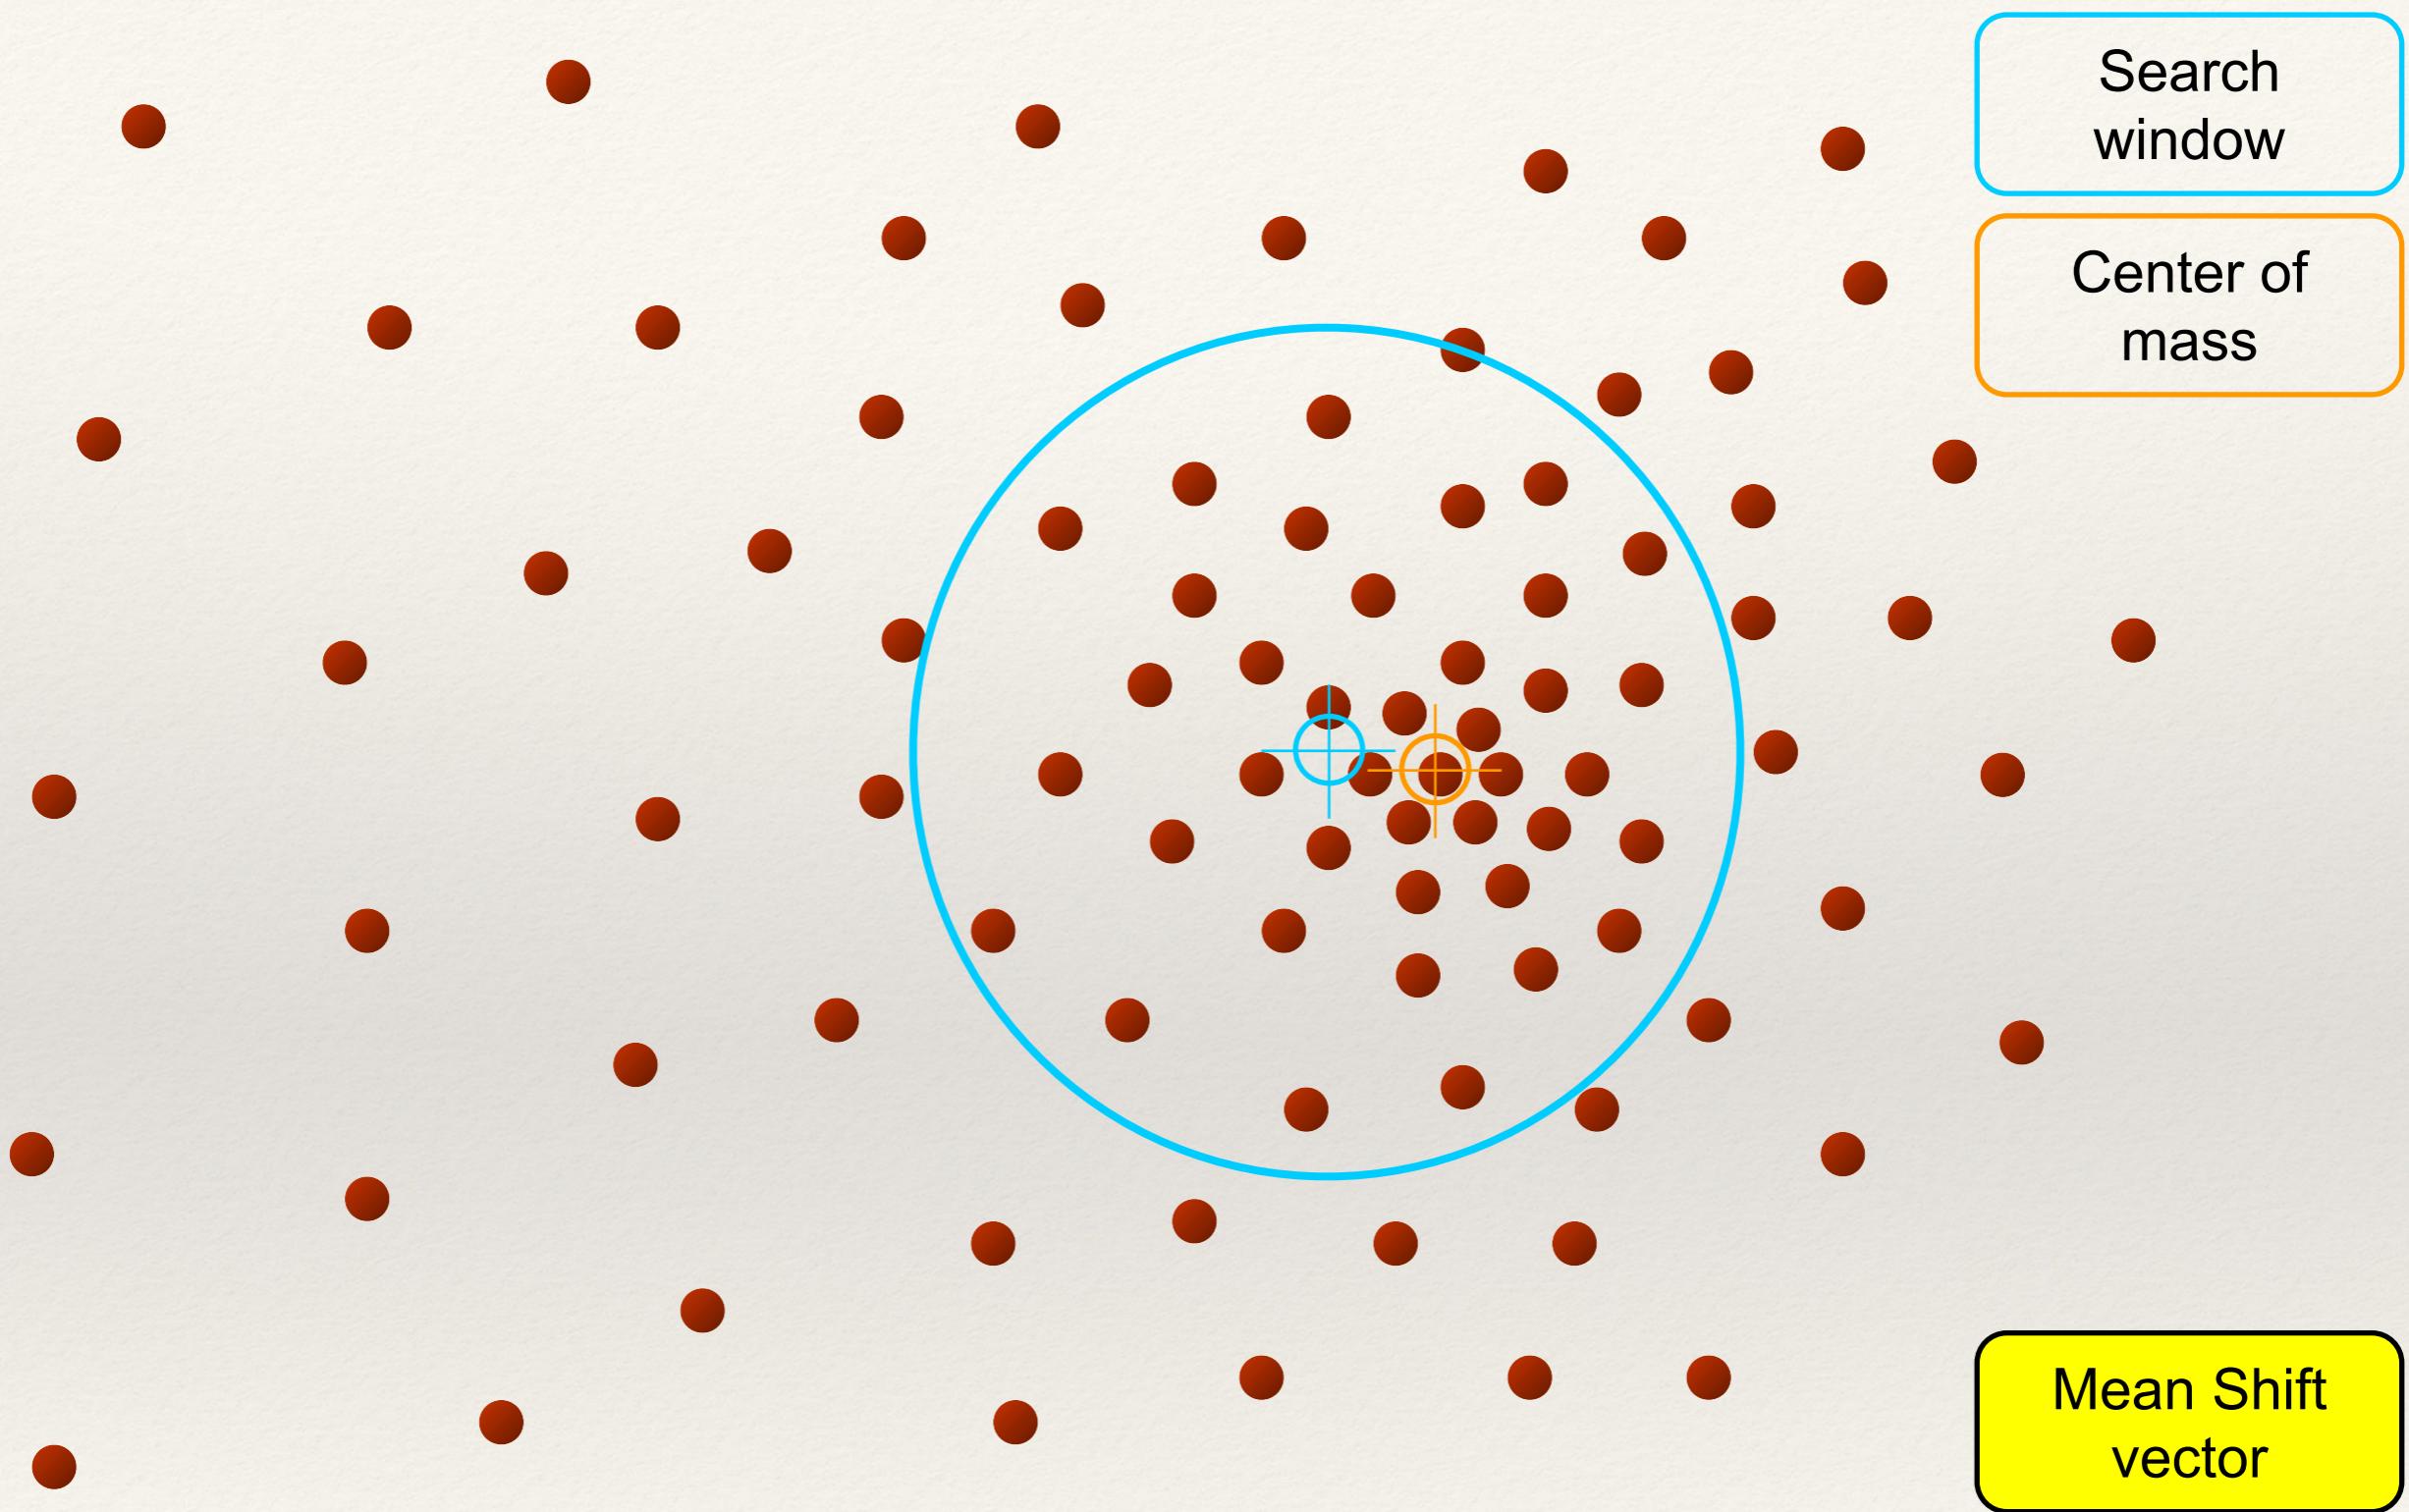


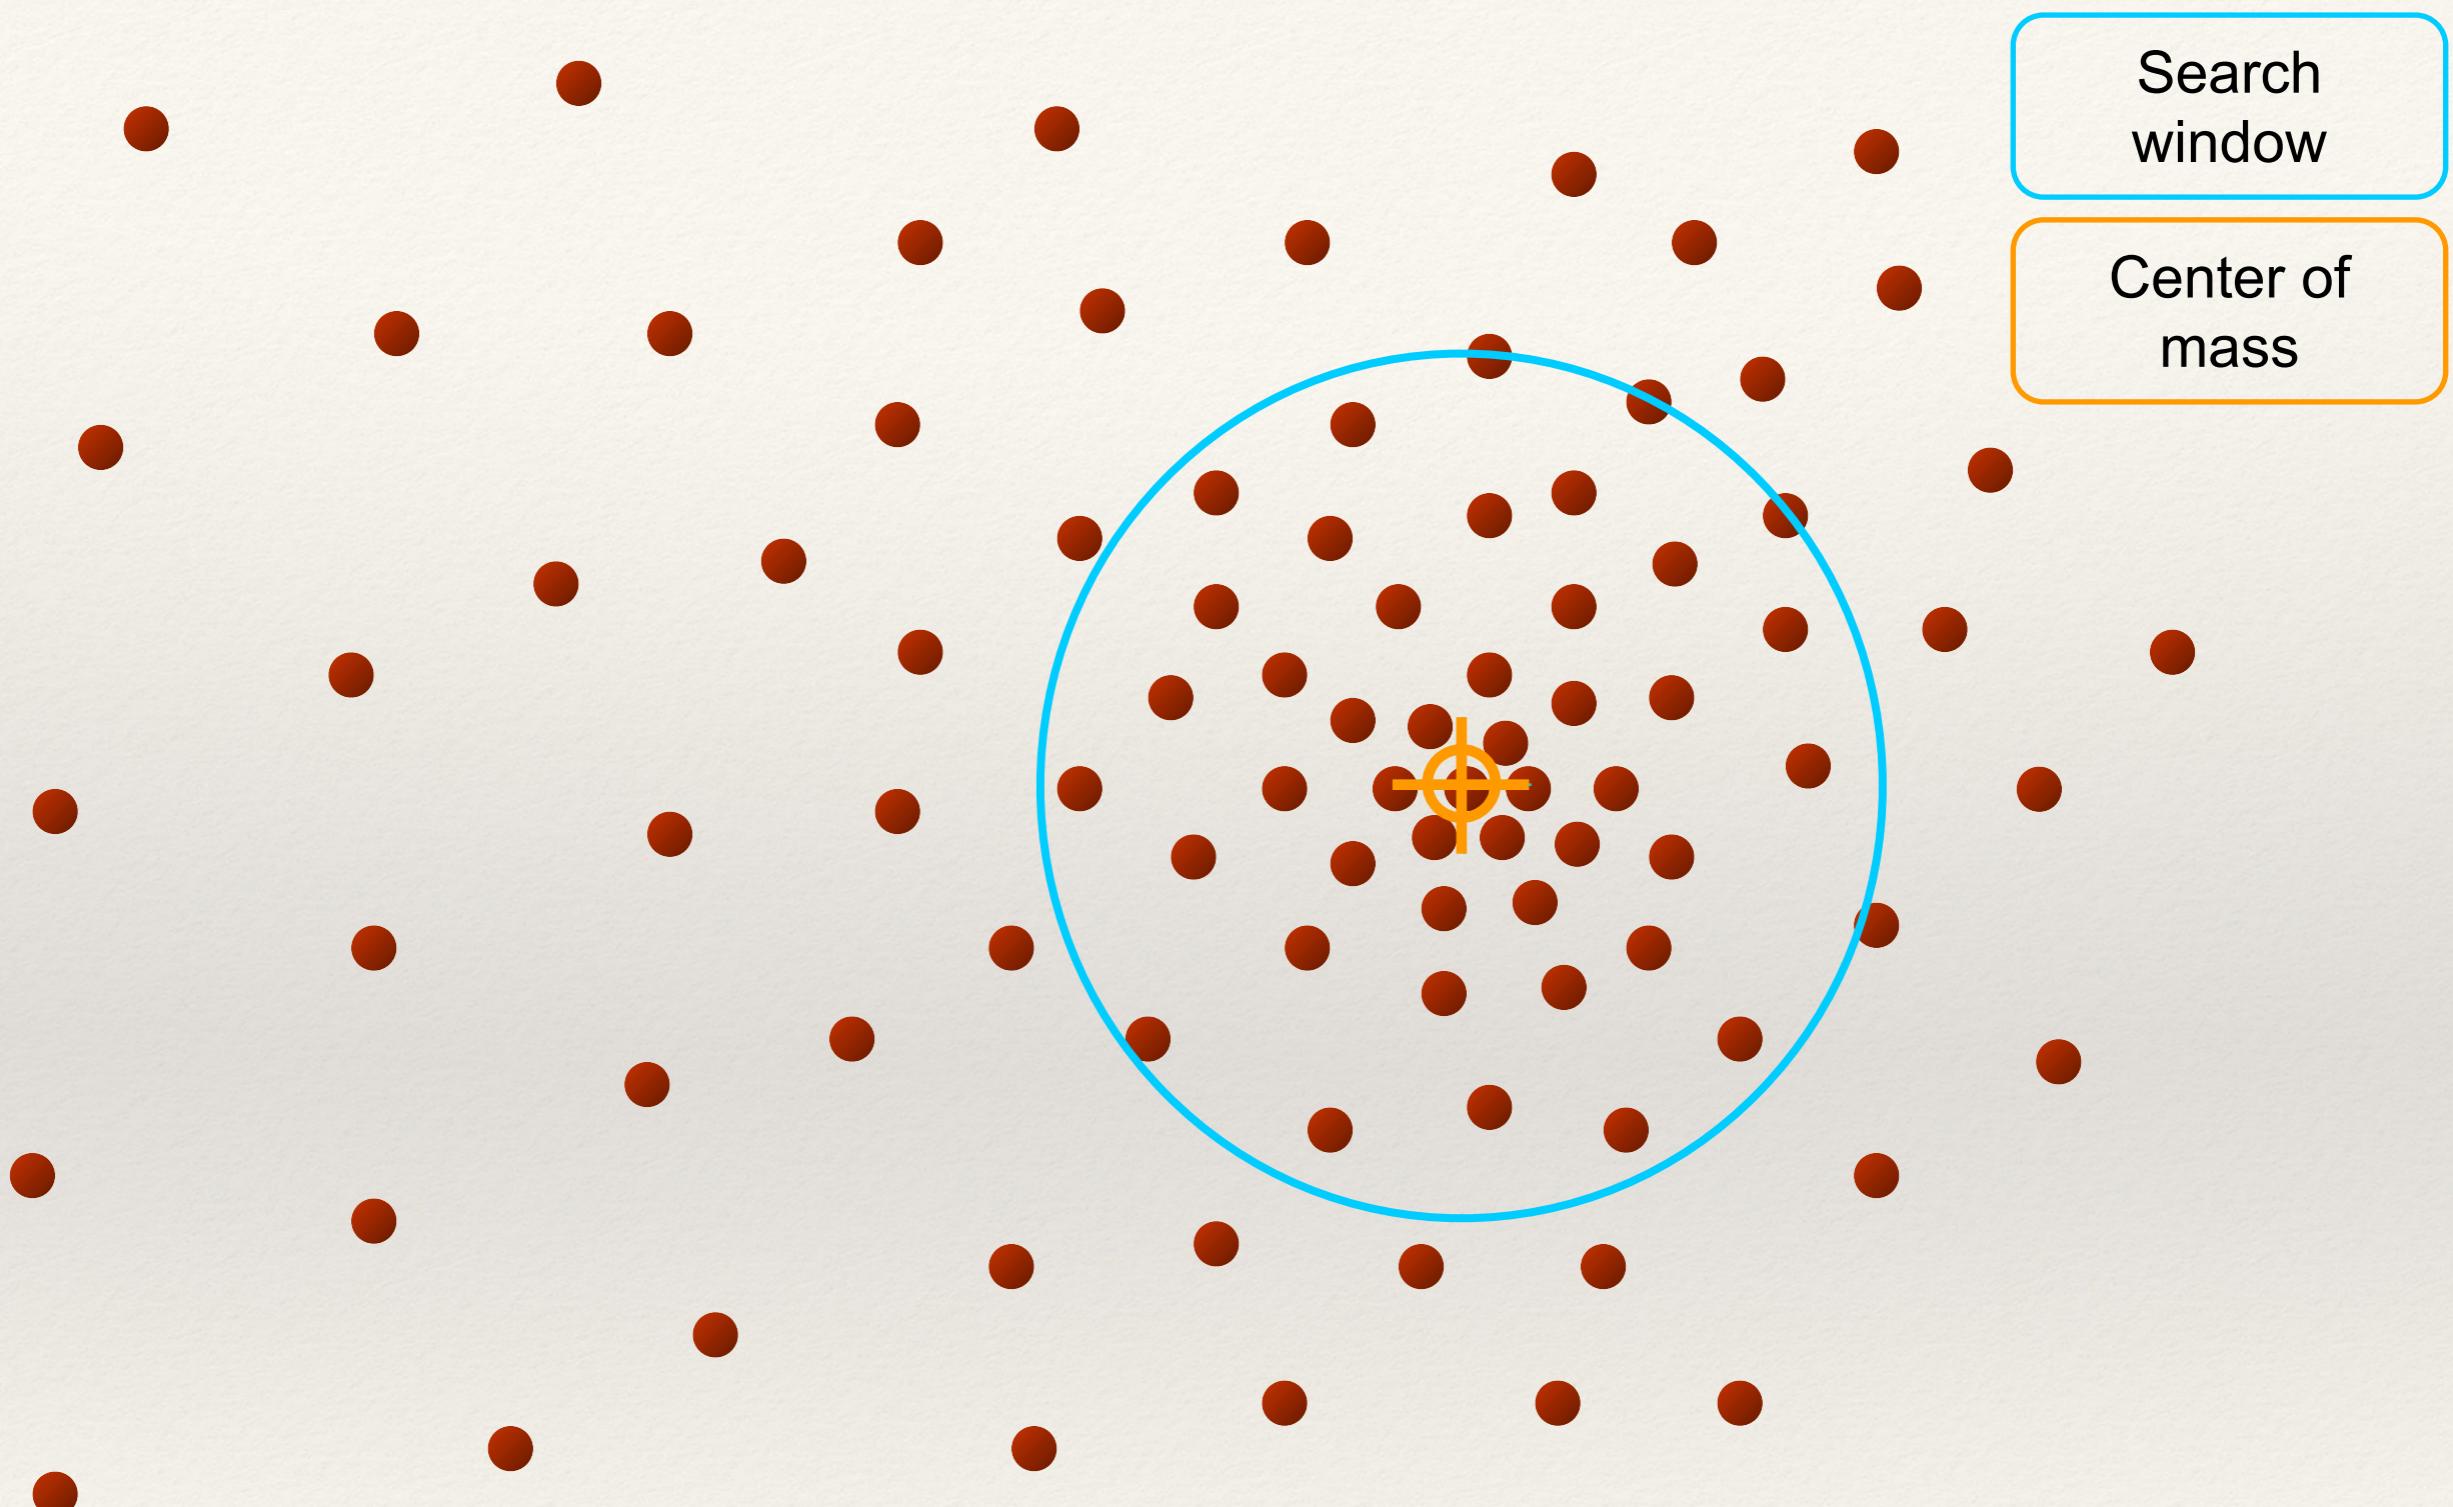






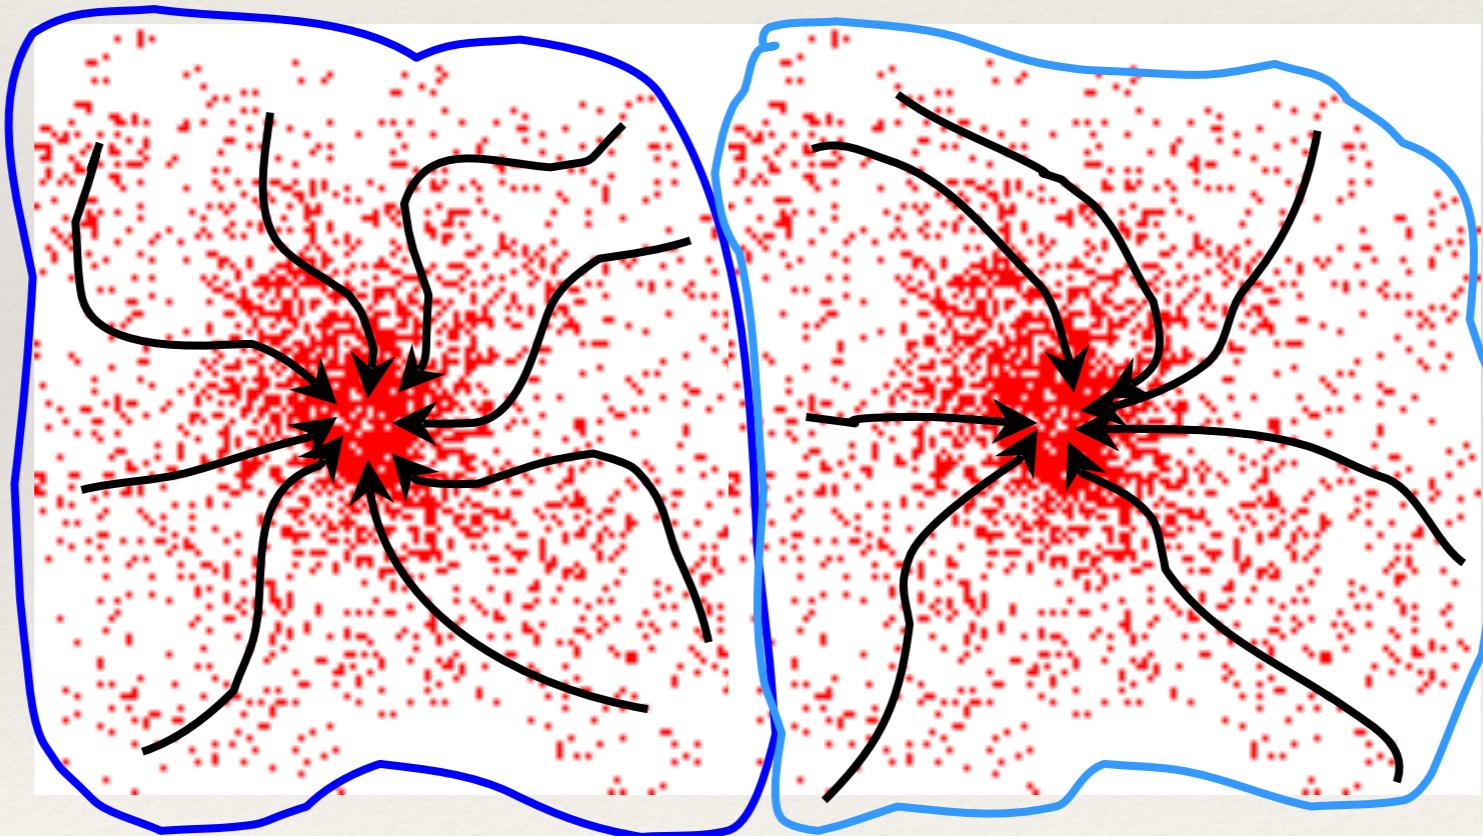






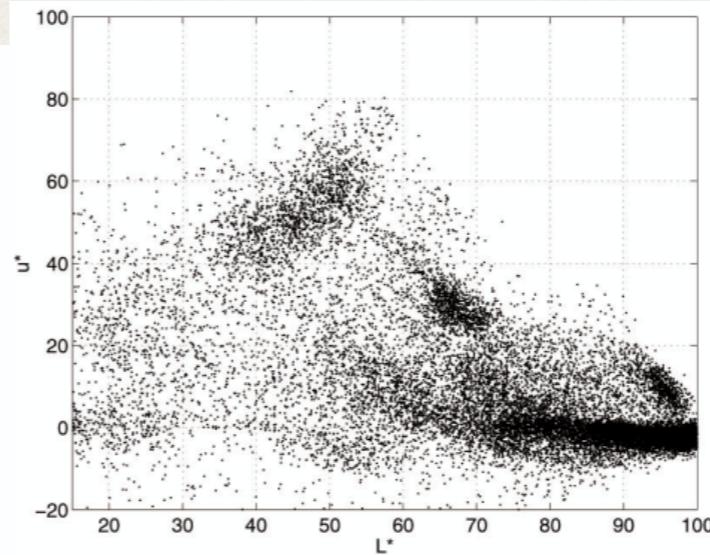
Mean shift clustering

- ❖ Cluster: all data points in the attraction basin of a mode
- ❖ Attraction basin: the region for which all trajectories lead to the same mode

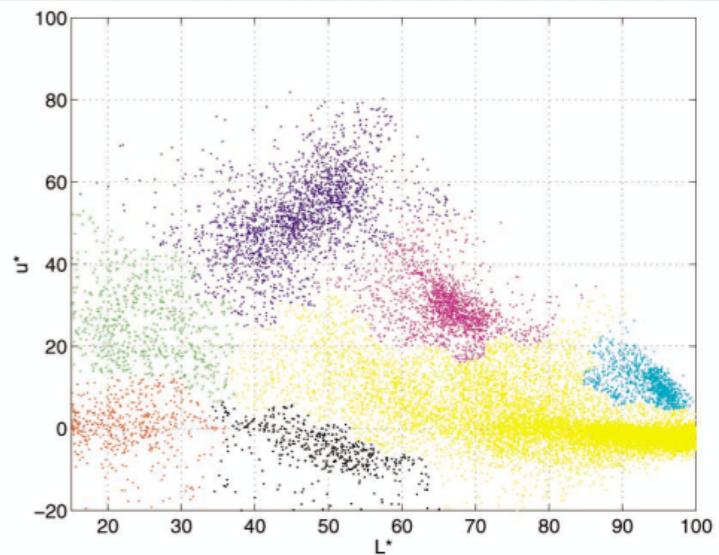


Mean shift clustering/segmentation

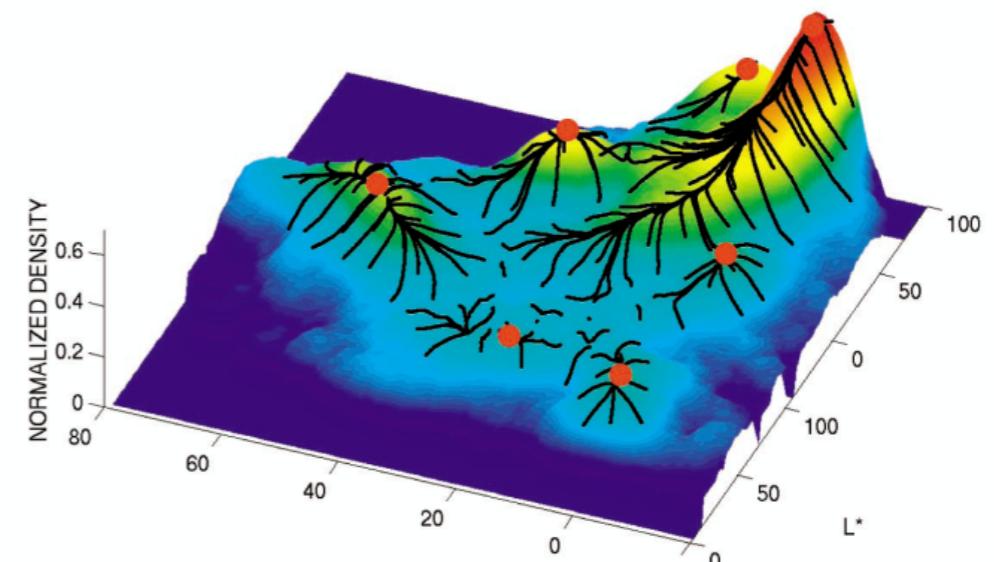
- Find features (color, gradients, texture, etc)
- Initialize windows at individual feature points
- Perform mean shift for each window until convergence
- Merge windows that end up near the same “peak” or mode



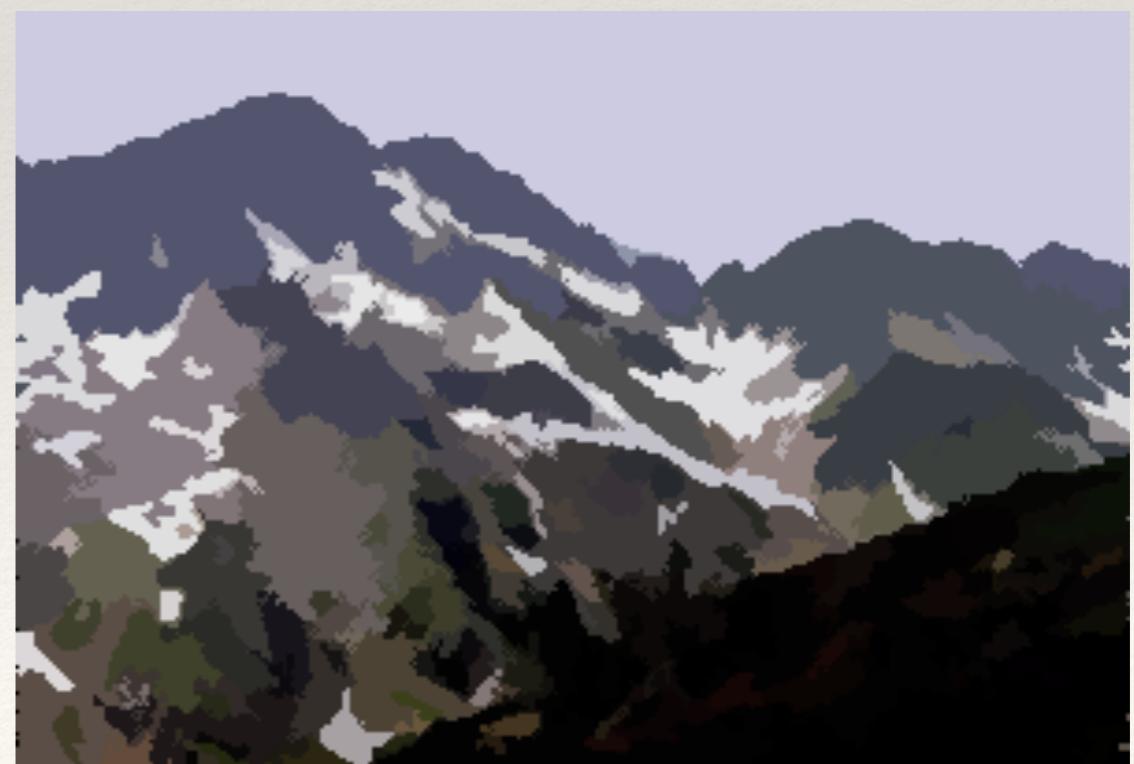
(a)



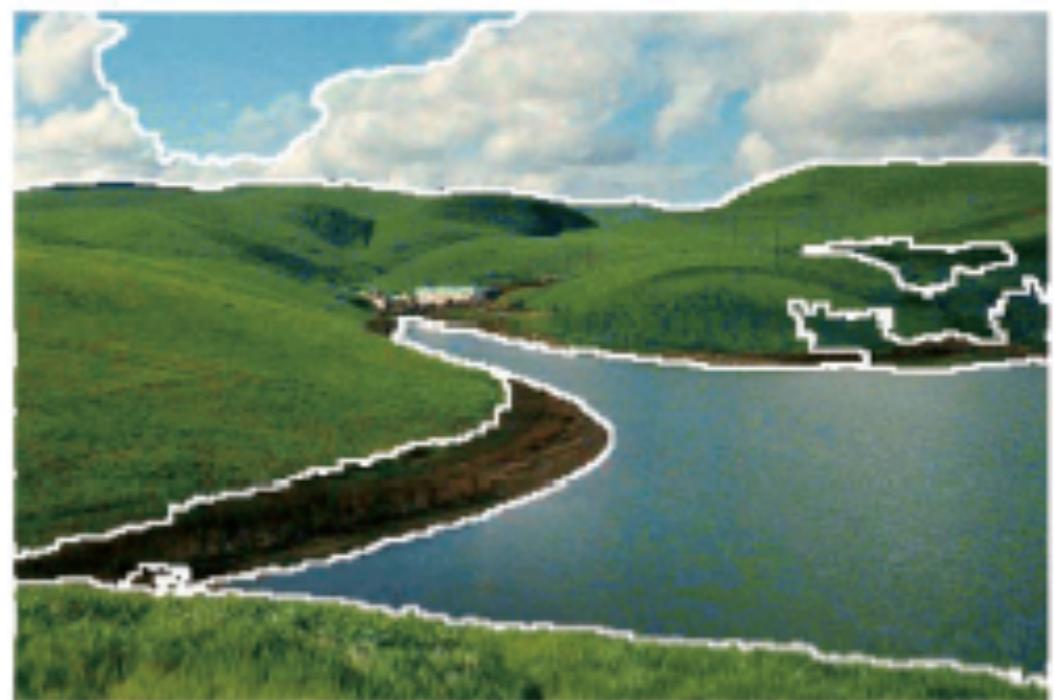
(b)



Mean shift segmentation results



Mean shift segmentation results



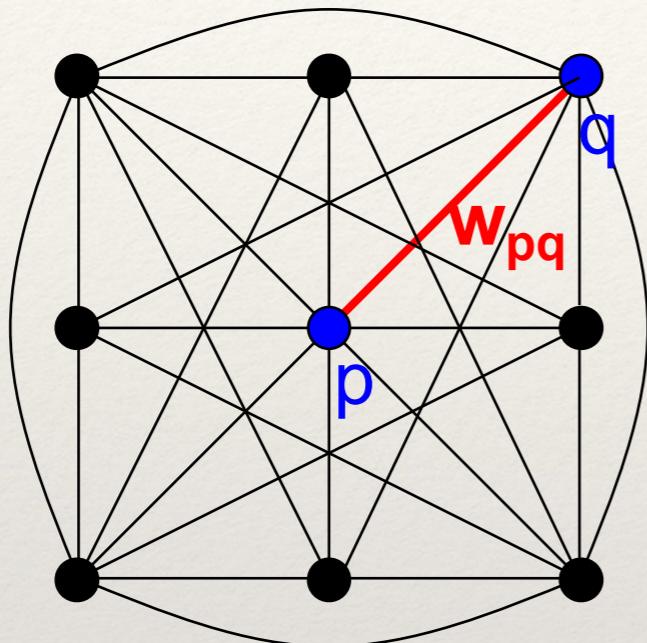
Mean shift segmentation results



Mean shift

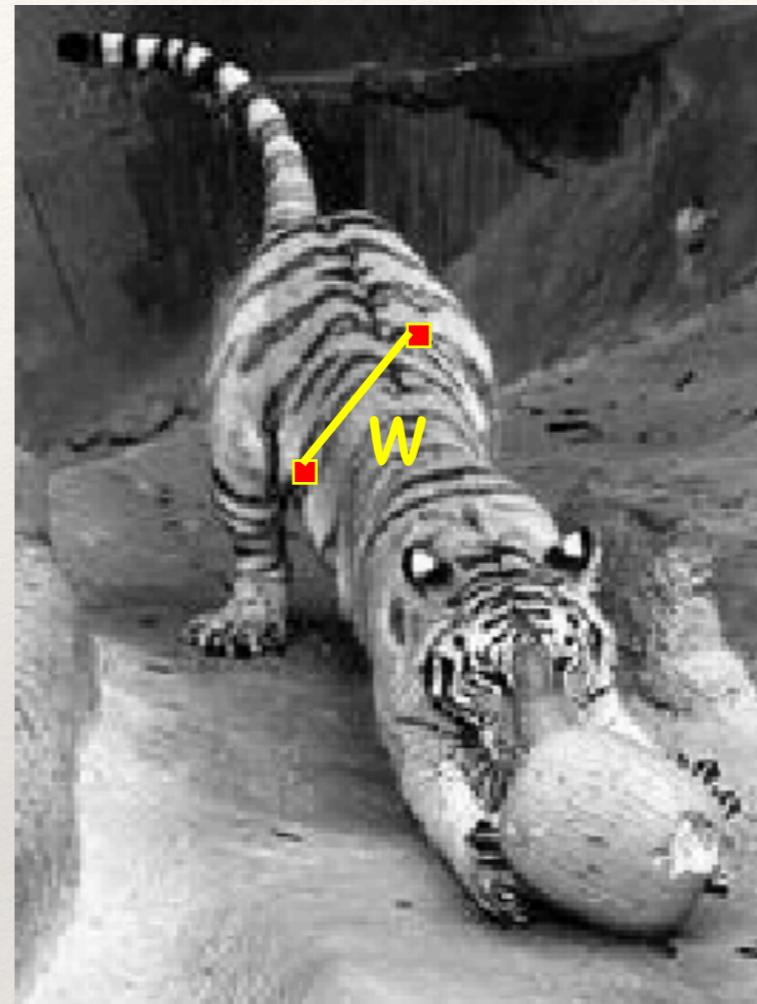
- ❖ Pros:
 - ❖ Does not assume shape on clusters
 - ❖ One parameter choice (window size, aka “bandwidth”)
 - ❖ Generic technique
 - ❖ Find multiple modes
- ❖ Cons:
 - ❖ Selection of window size
 - ❖ Does not scale well with dimension of feature space

Images as graphs



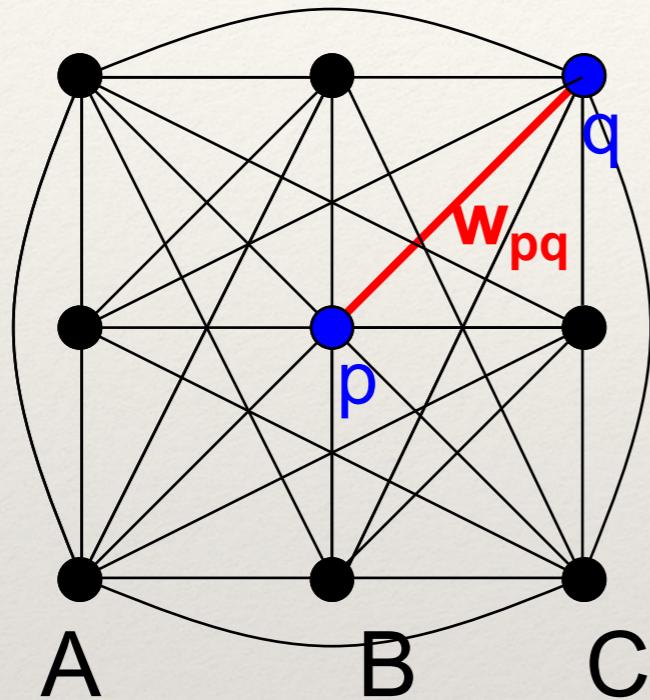
Fully-connected graph

- node (vertex) for every pixel
- link between every pair of pixels, \mathbf{p}, \mathbf{q}
- affinity weight $w_{\mathbf{pq}}$ for each link (edge)
 - $w_{\mathbf{pq}}$ measures *similarity*
 - » similarity is *inversely proportional* to difference (in color and position...)



Source: Steve Seitz

Segmentation by Graph Cuts



Break Graph into Segments

- Want to delete links that cross **between** segments
- Easiest to break links that have low similarity (low weight)
 - similar pixels should be in the same segments
 - dissimilar pixels should be in different segments

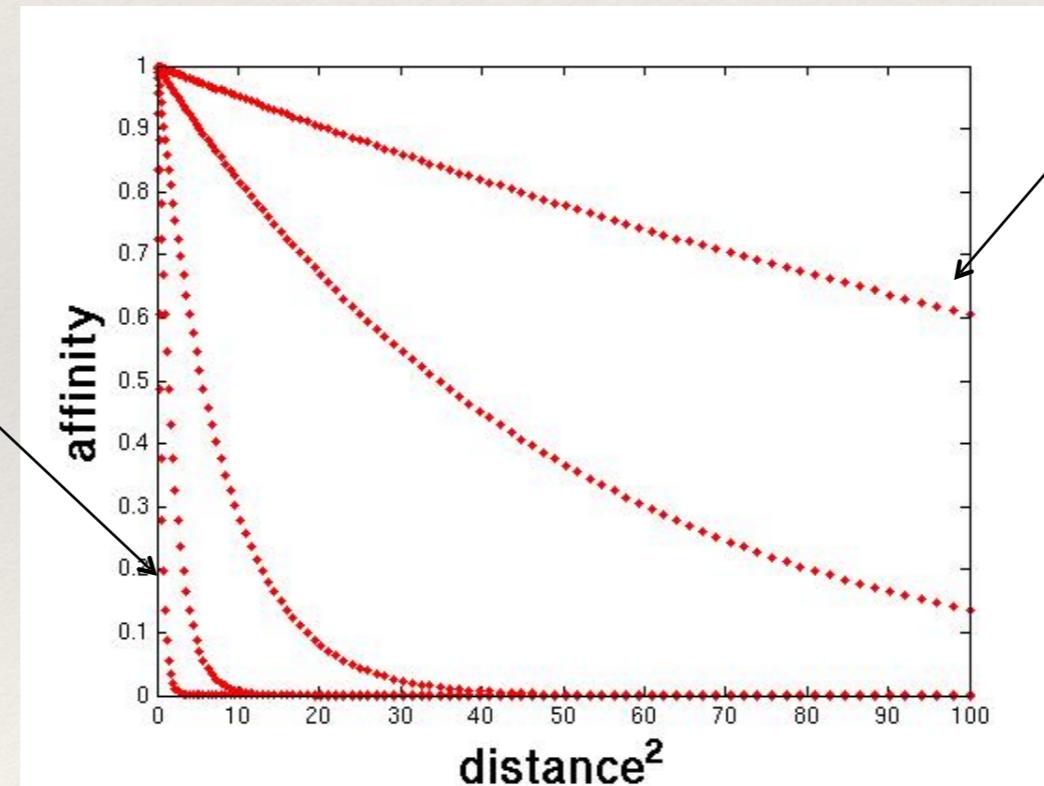


Measuring affinity

- ❖ One possibility:

$$aff(x, y) = \exp\left\{-\left(\frac{1}{2\sigma_d^2}\right)\|x - y\|^2\right\}$$

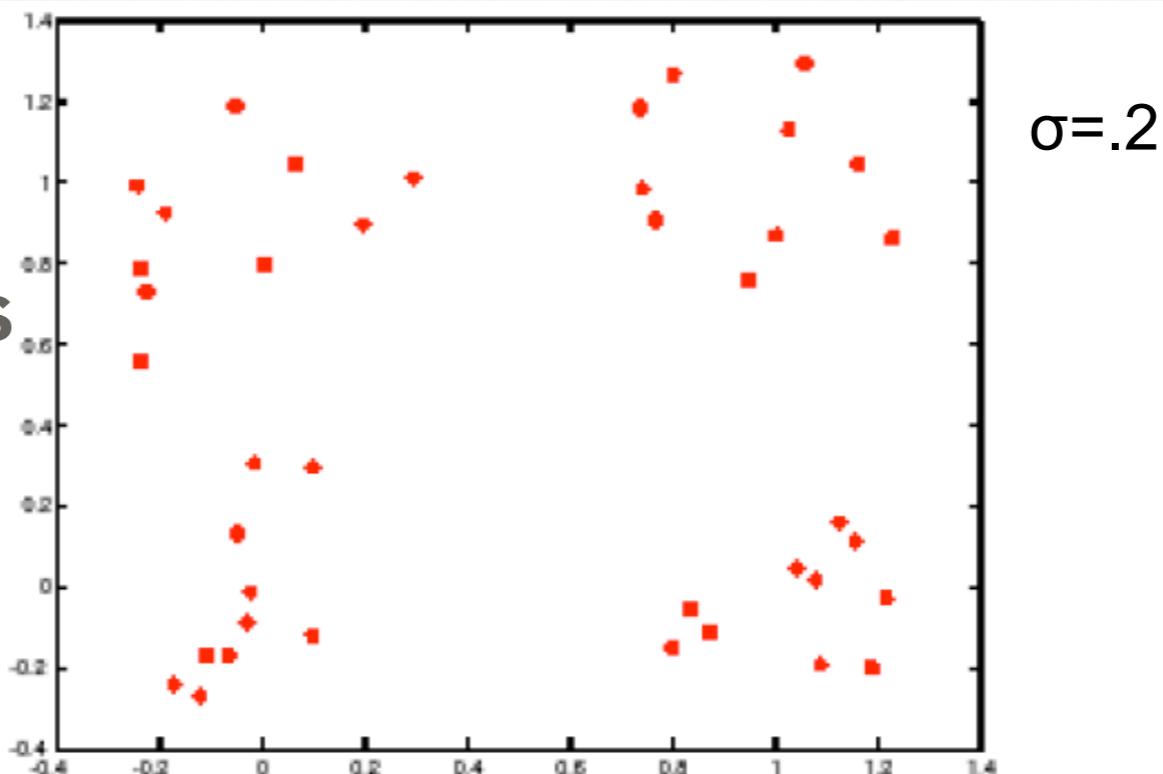
Small sigma: group only nearby points



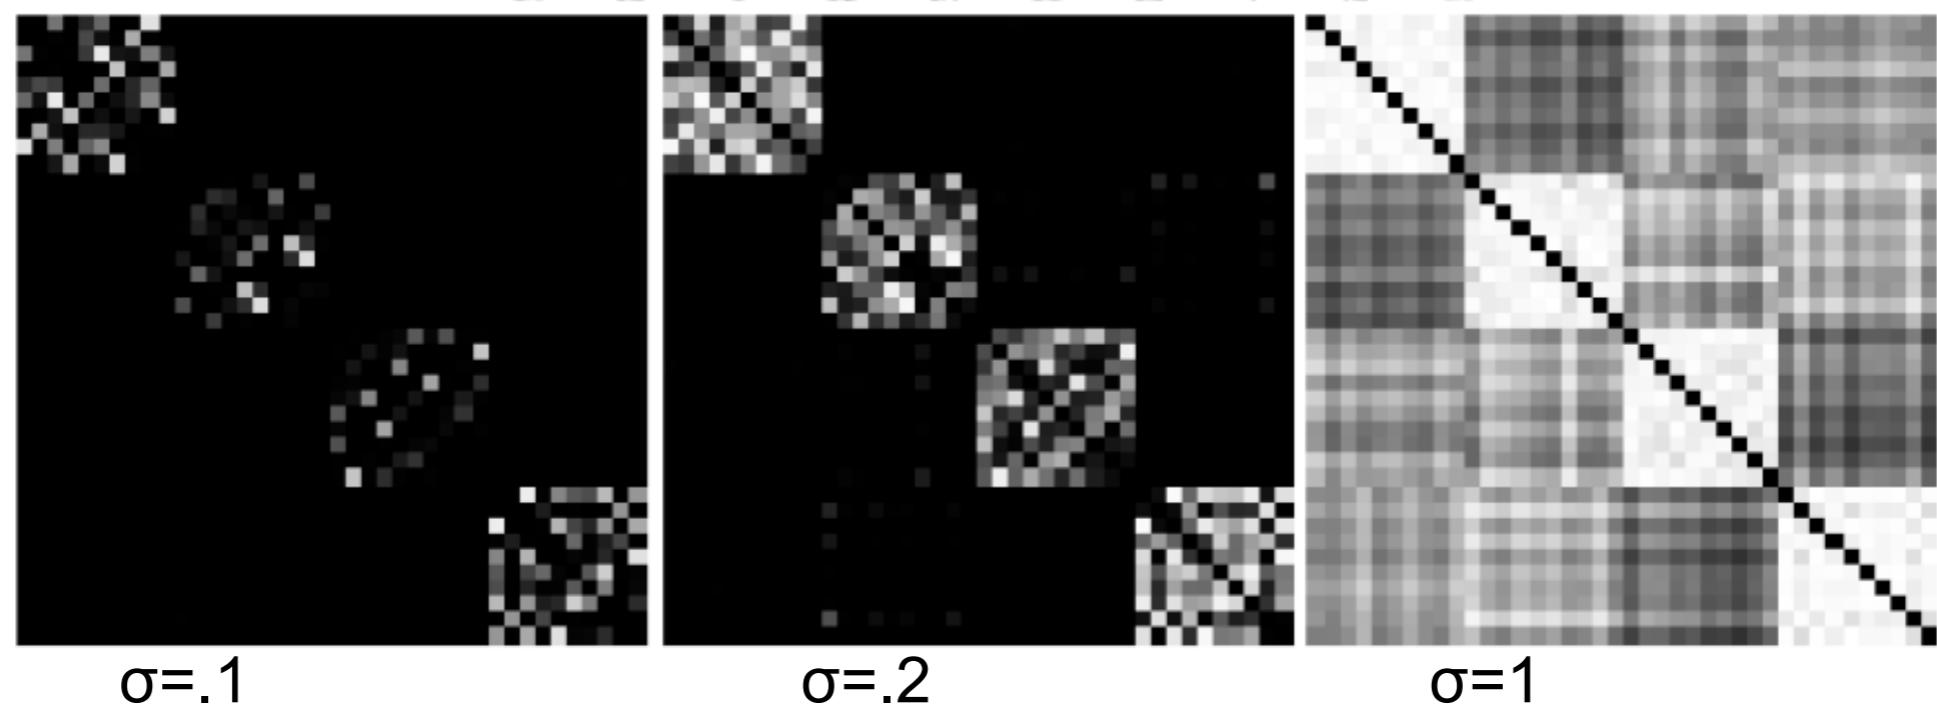
Large sigma: group distant points

Measuring affinity

Data points



Affinity
matrices

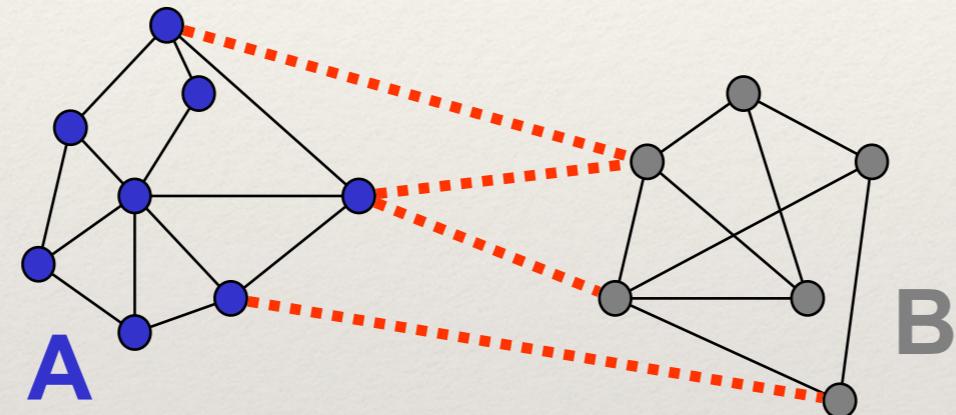


Slide credit: Kristen Grauman

Cuts in a graph: Min cut

Link Cut

- set of links whose removal makes a graph disconnected
- cost of a cut:



$$cut(A, B) = \sum_{p \in A, q \in B} w_{p,q}$$

Find minimum cut

- gives you a segmentation
- fast algorithms exist for doing this

Minimum cut

- Problem with minimum cut:

Weight of cut proportional to number of edges in the cut; tends to produce small, isolated components.

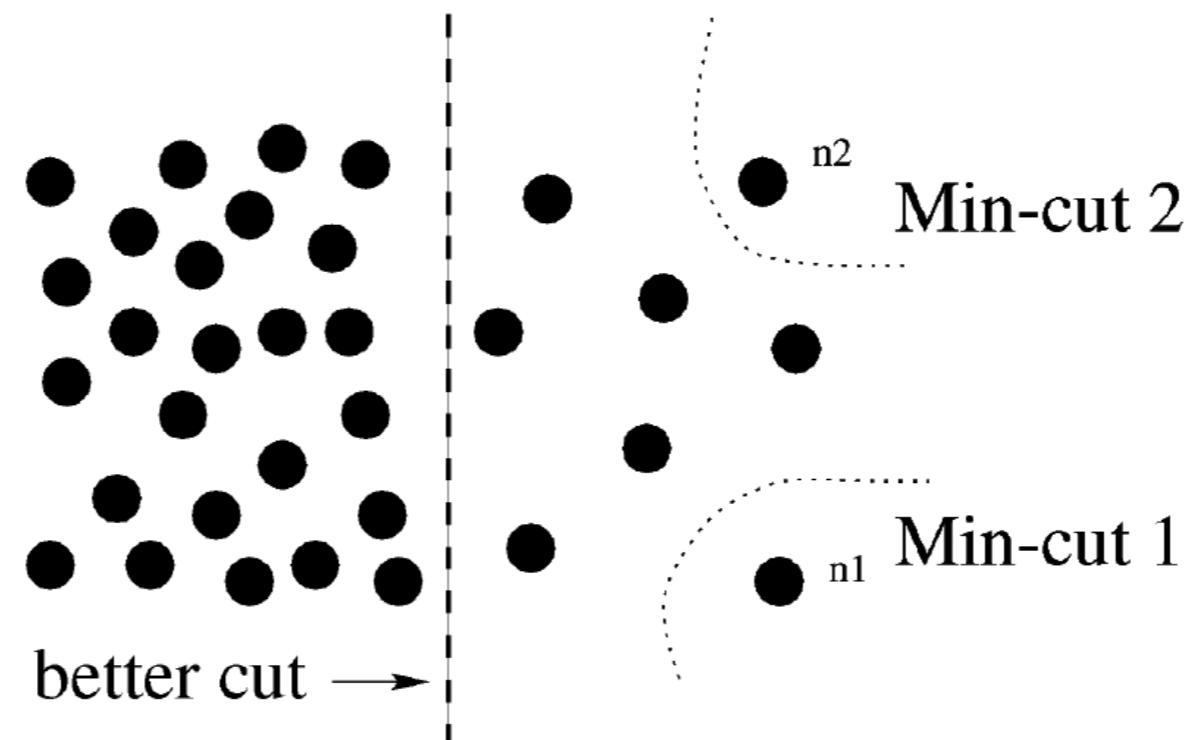
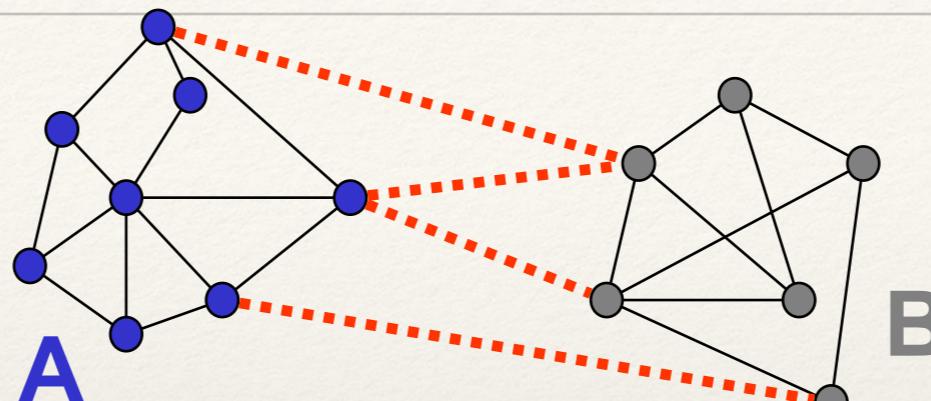


Fig. 1. A case where minimum cut gives a bad partition.

Cuts in a graph: Normalized cut

Normalized Cut



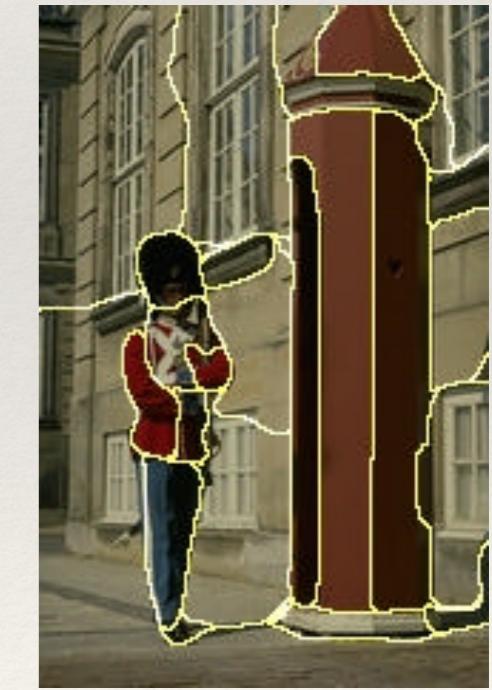
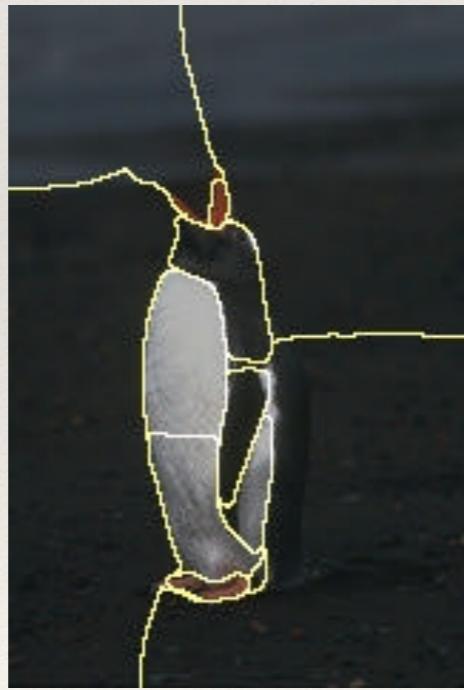
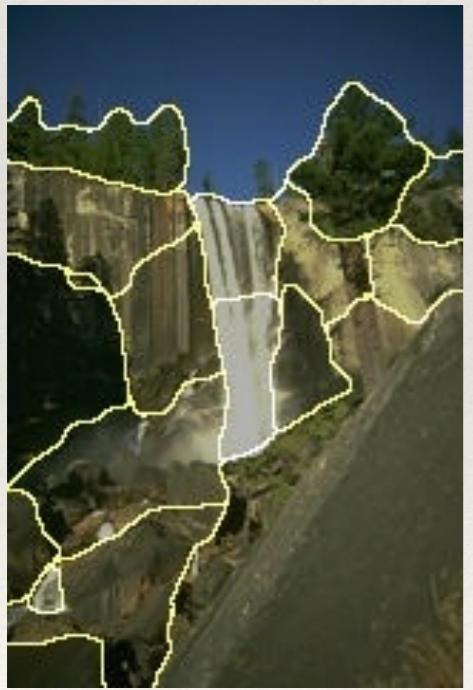
- fix bias of Min Cut by **normalizing** for size of segments:

$$Ncut(A, B) = \frac{cut(A, B)}{assoc(A, V)} + \frac{cut(A, B)}{assoc(B, V)}$$

$assoc(A, V)$ = sum of weights of all edges that touch A

- Ncut value small when we get two clusters with many edges with high weights, and few edges of low weight between them
- Approximate solution for minimizing the Ncut value : generalized eigenvalue problem.





<http://www.cs.berkeley.edu/~fowlkes/BSE/>

Normalized cuts: pros and cons

Pros:

- ❖ Generic framework, flexible to choice of function that computes weights ("affinities") between nodes
- ❖ Does not require model of the data distribution

Cons:

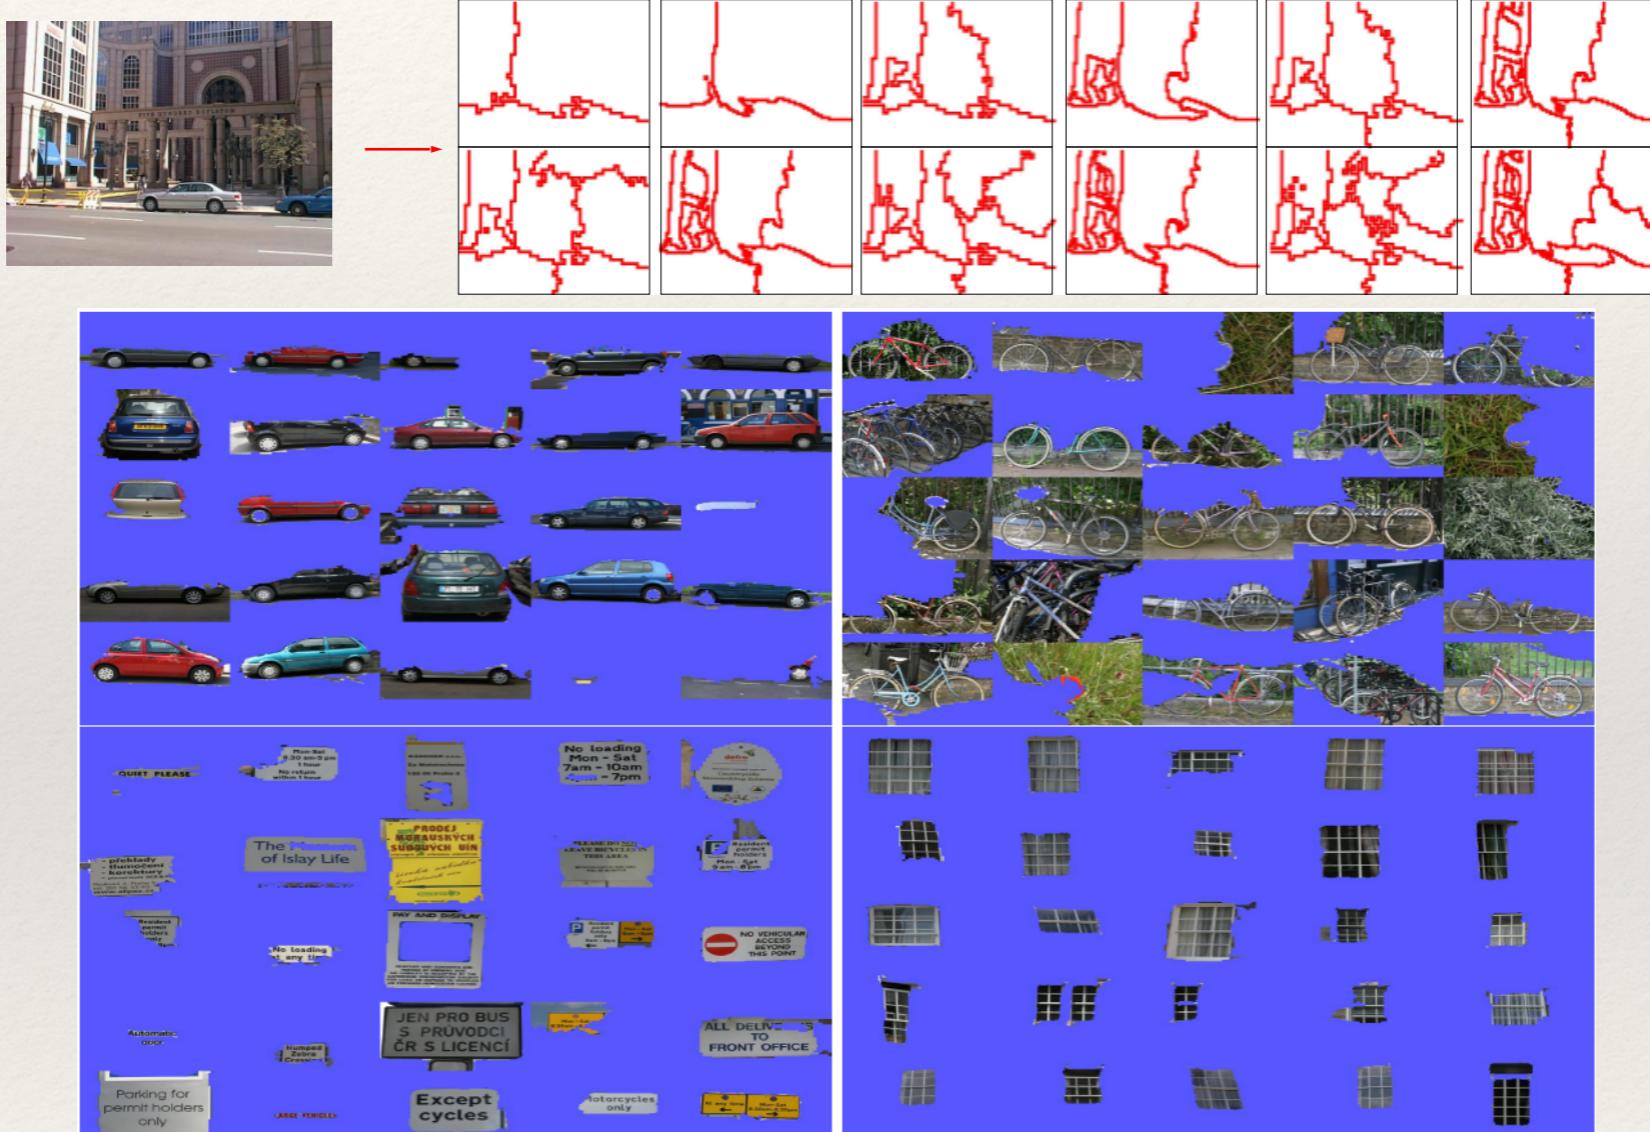
- ❖ Time complexity can be high
 - ❖ Dense, highly connected graphs → many affinity computations
 - ❖ Solving eigenvalue problem
- ❖ Preference for balanced partitions

Summary

- ❖ Segmentation to find object boundaries or mid-level regions, tokens.
- ❖ Bottom-up segmentation via clustering
 - ❖ General choices -- features, affinity functions, and clustering algorithms
- ❖ Grouping also useful for quantization, can create new feature summaries
 - ❖ Texton histograms for texture within local region
- ❖ Example clustering methods
 - ❖ K-means
 - ❖ Mean shift
 - ❖ Graph cut, normalized cuts

Segments as primitives for recognition

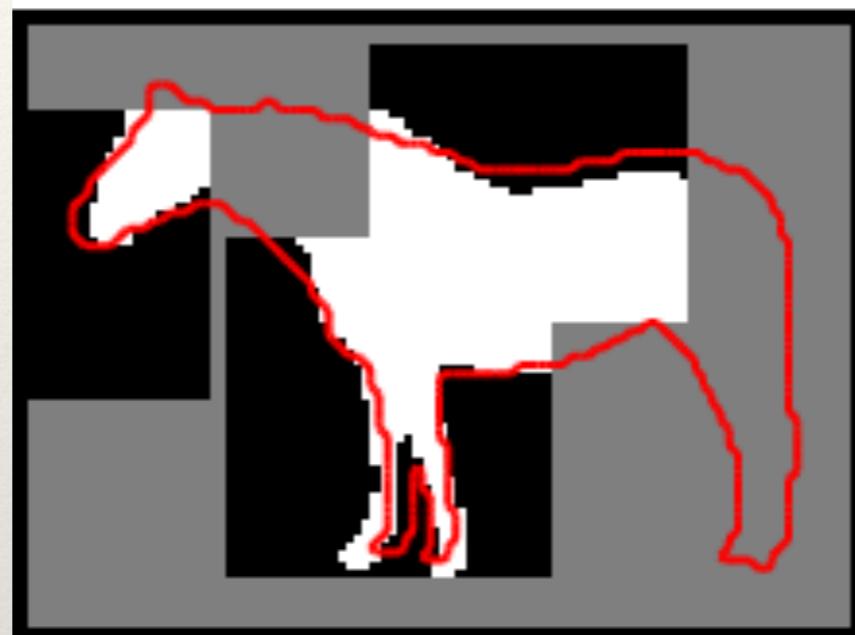
Multiple segmentations



B. Russell et al., ["Using Multiple Segmentations to Discover Objects and their Extent in Image Collections,"](#) CVPR 2006

Slide credit: Lana Lazebnik

Top-down segmentation



- E. Borenstein and S. Ullman, “[Class-specific, top-down segmentation,](#)” ECCV 2002
A. Levin and Y. Weiss, “[Learning to Combine Bottom-Up and Top-Down Segmentation,](#)” ECCV 2006.

Top-down segmentation

Normalized cuts



Top-down segmentation



E. Borenstein and S. Ullman, ["Class-specific, top-down segmentation,"](#) ECCV 2002

A. Levin and Y. Weiss, ["Learning to Combine Bottom-Up and Top-Down Segmentation,"](#) ECCV 2006.

Motion segmentation



Input sequence

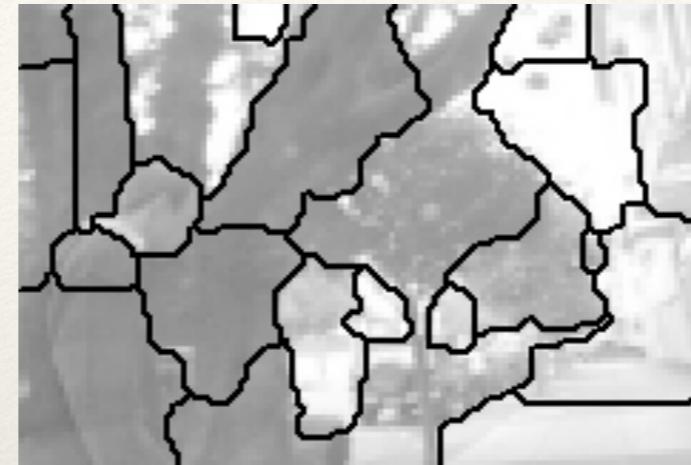


Image Segmentation



Motion Segmentation



Input sequence



Image Segmentation



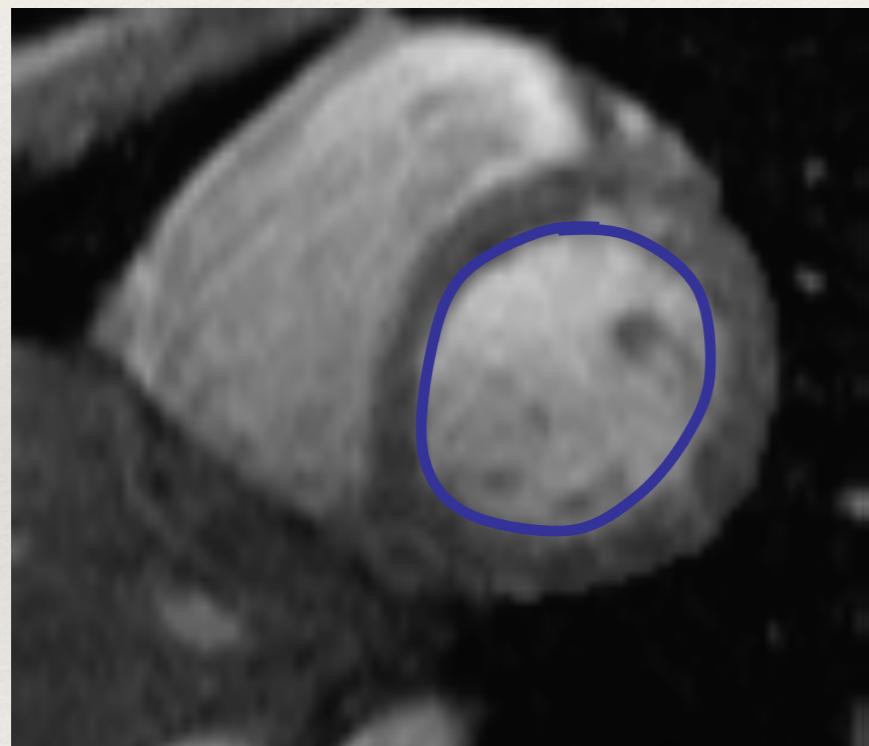
Motion Segmentation

A. Barbu, S.C. Zhu. Generalizing Swendsen-Wang to sampling arbitrary posterior probabilities, *IEEE Trans. PAMI*, August 2005.

Deformable contours

Given: initial contour (model) near desired object

Goal: evolve the contour to fit exact object boundary

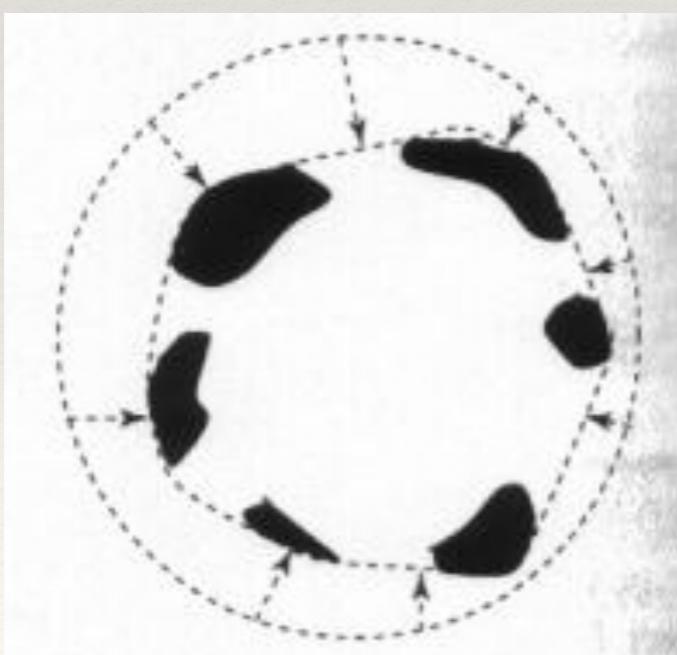
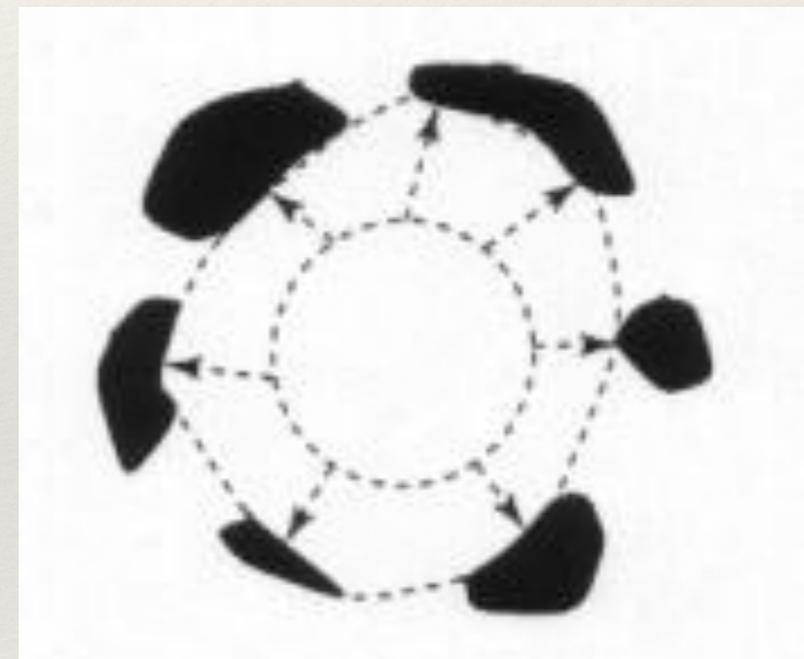
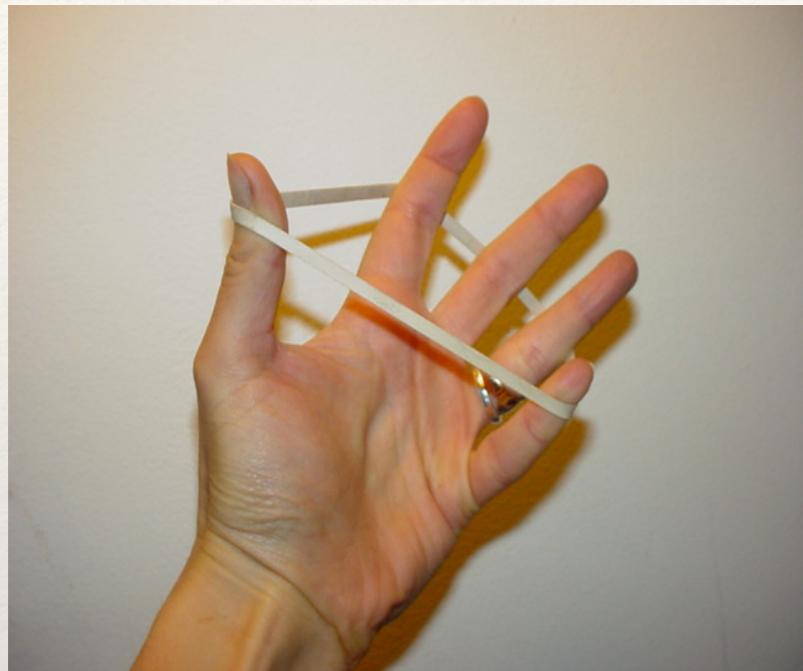


Main idea: elastic band is iteratively adjusted so as to

- be near image positions with high gradients, **and**
- satisfy shape “preferences” or contour priors

a.k.a. active contours, snakes

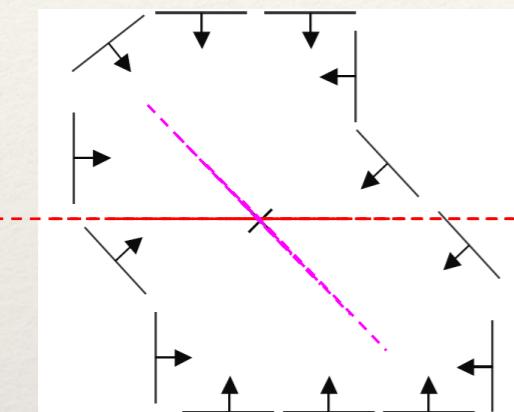
Deformable contours: intuition



Slide credit: Kristen Grauman

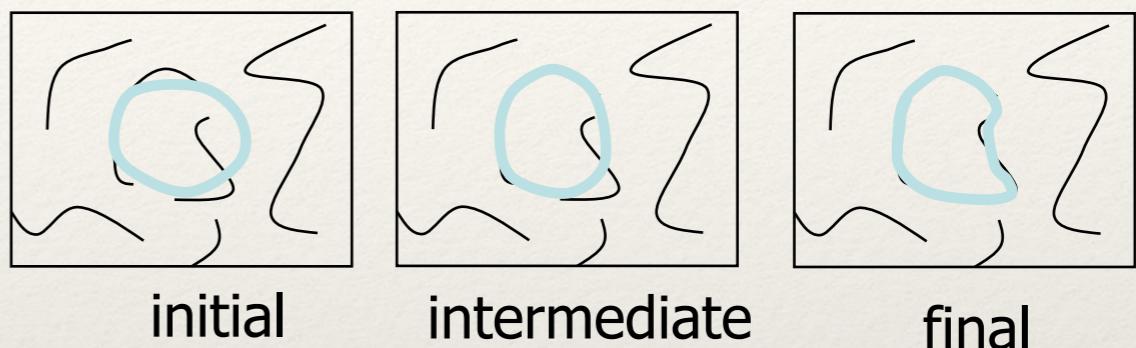
Deformable contours vs. Hough

Like generalized Hough transform, useful for shape fitting; but



Hough

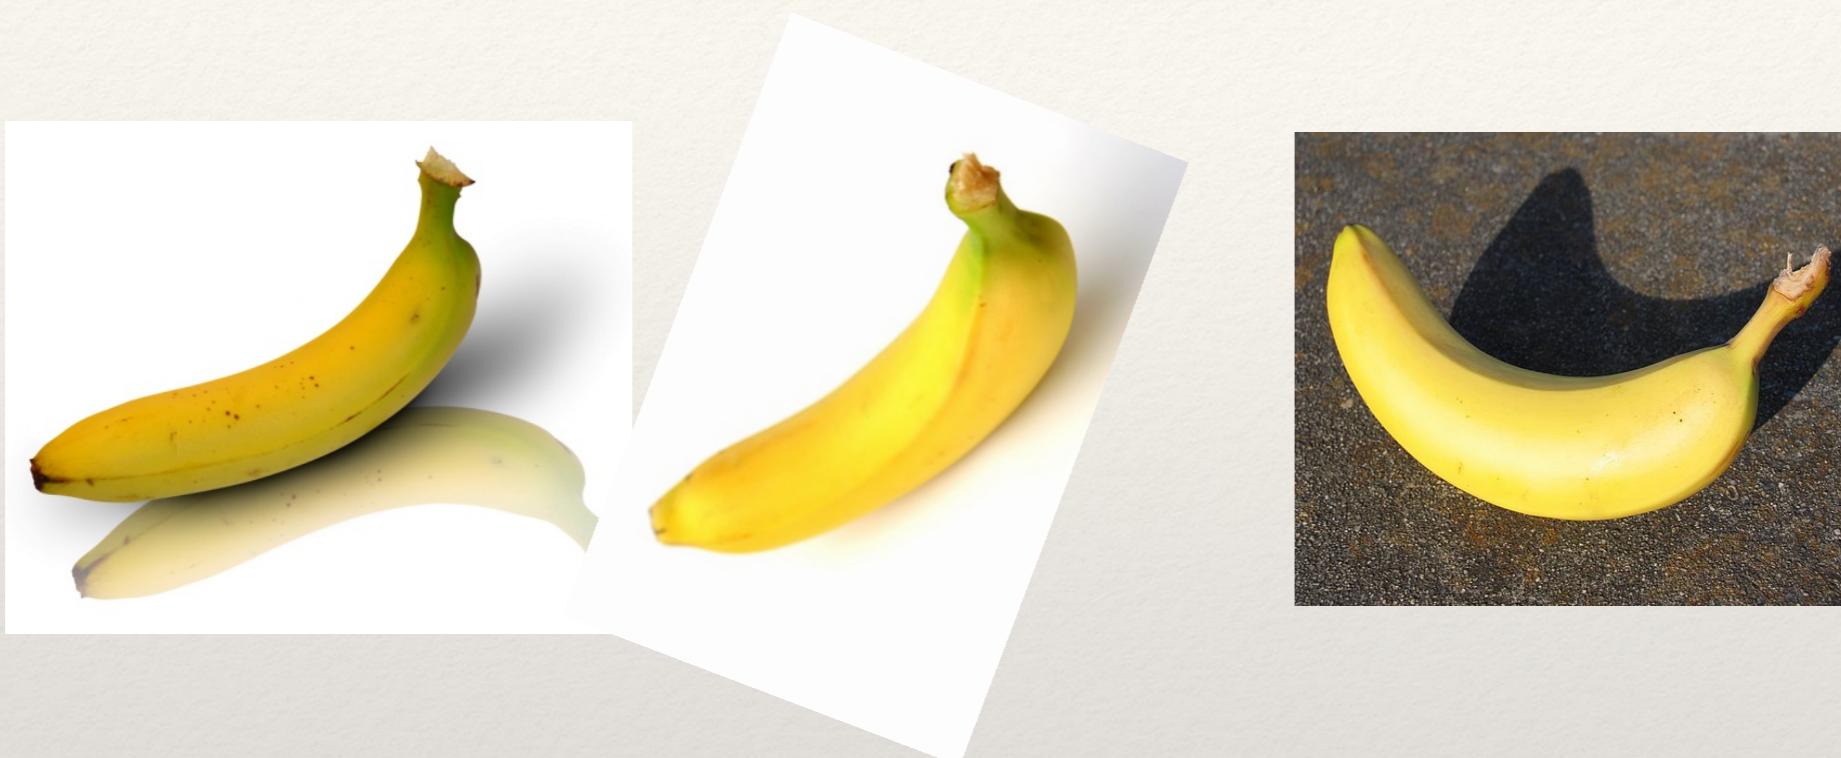
Rigid model shape
Single voting pass can
detect multiple
instances



Deformable contours

Prior on shape types, but shape
iteratively adjusted (*deforms*)
Requires initialization nearby
One optimization “pass” to fit a
single contour

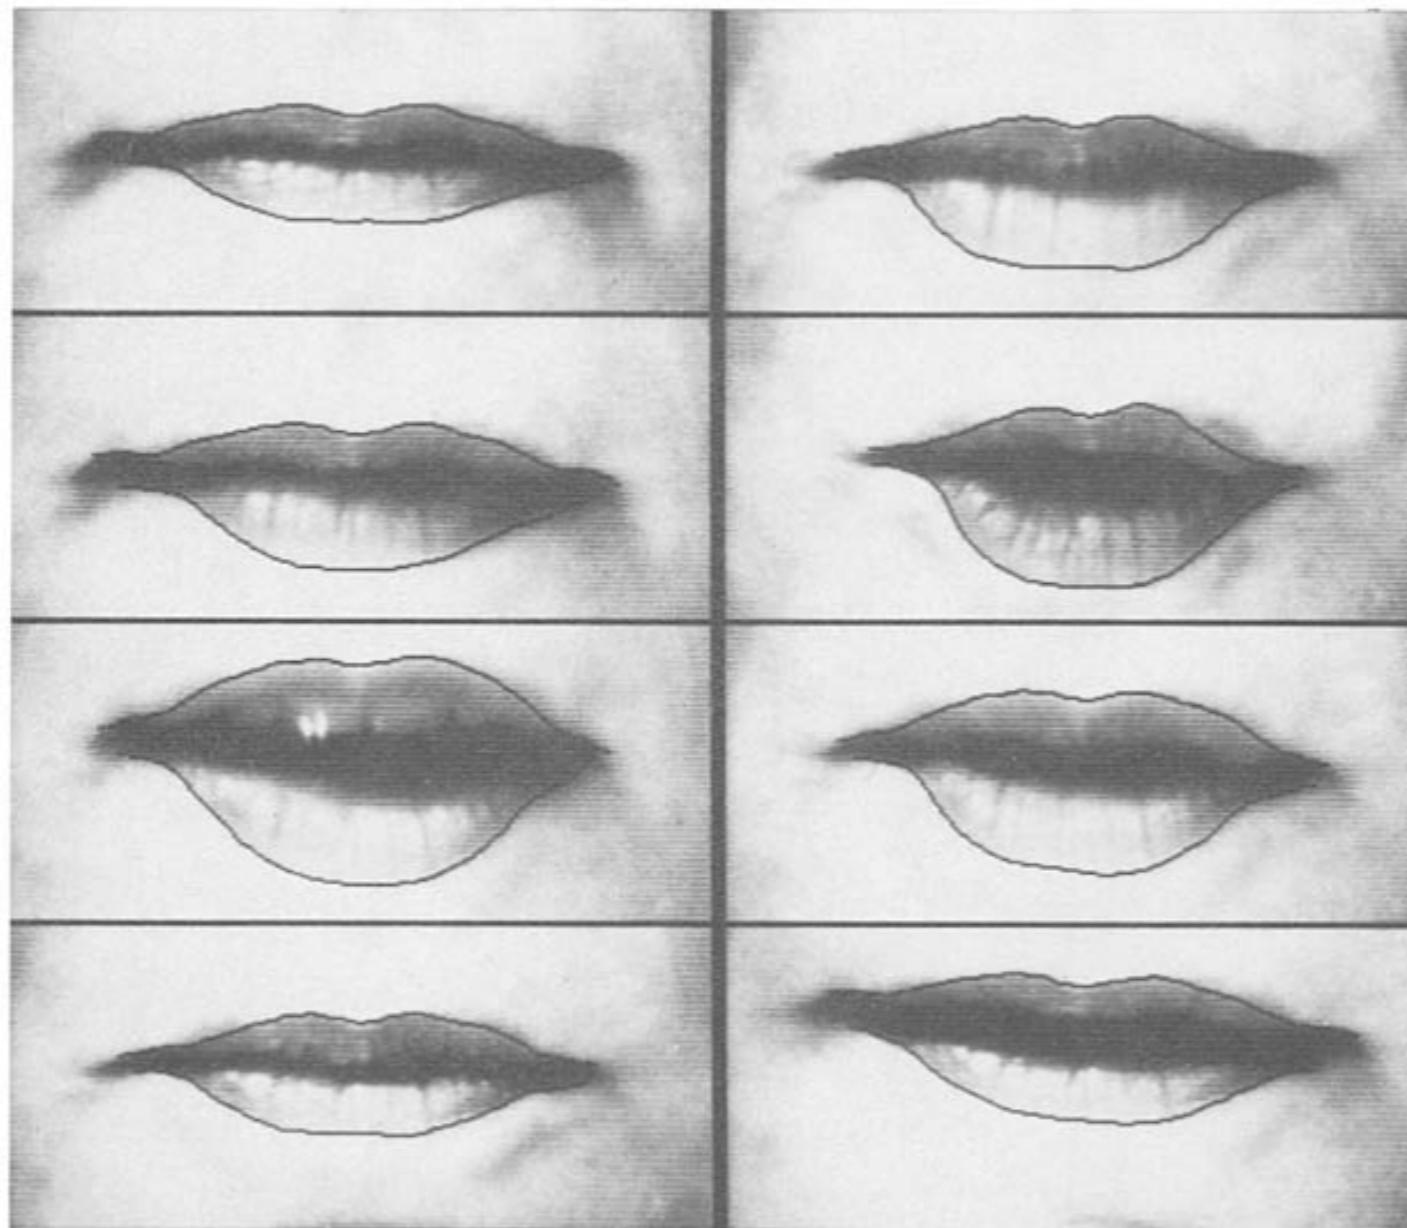
Why do we want to fit deformable shapes?



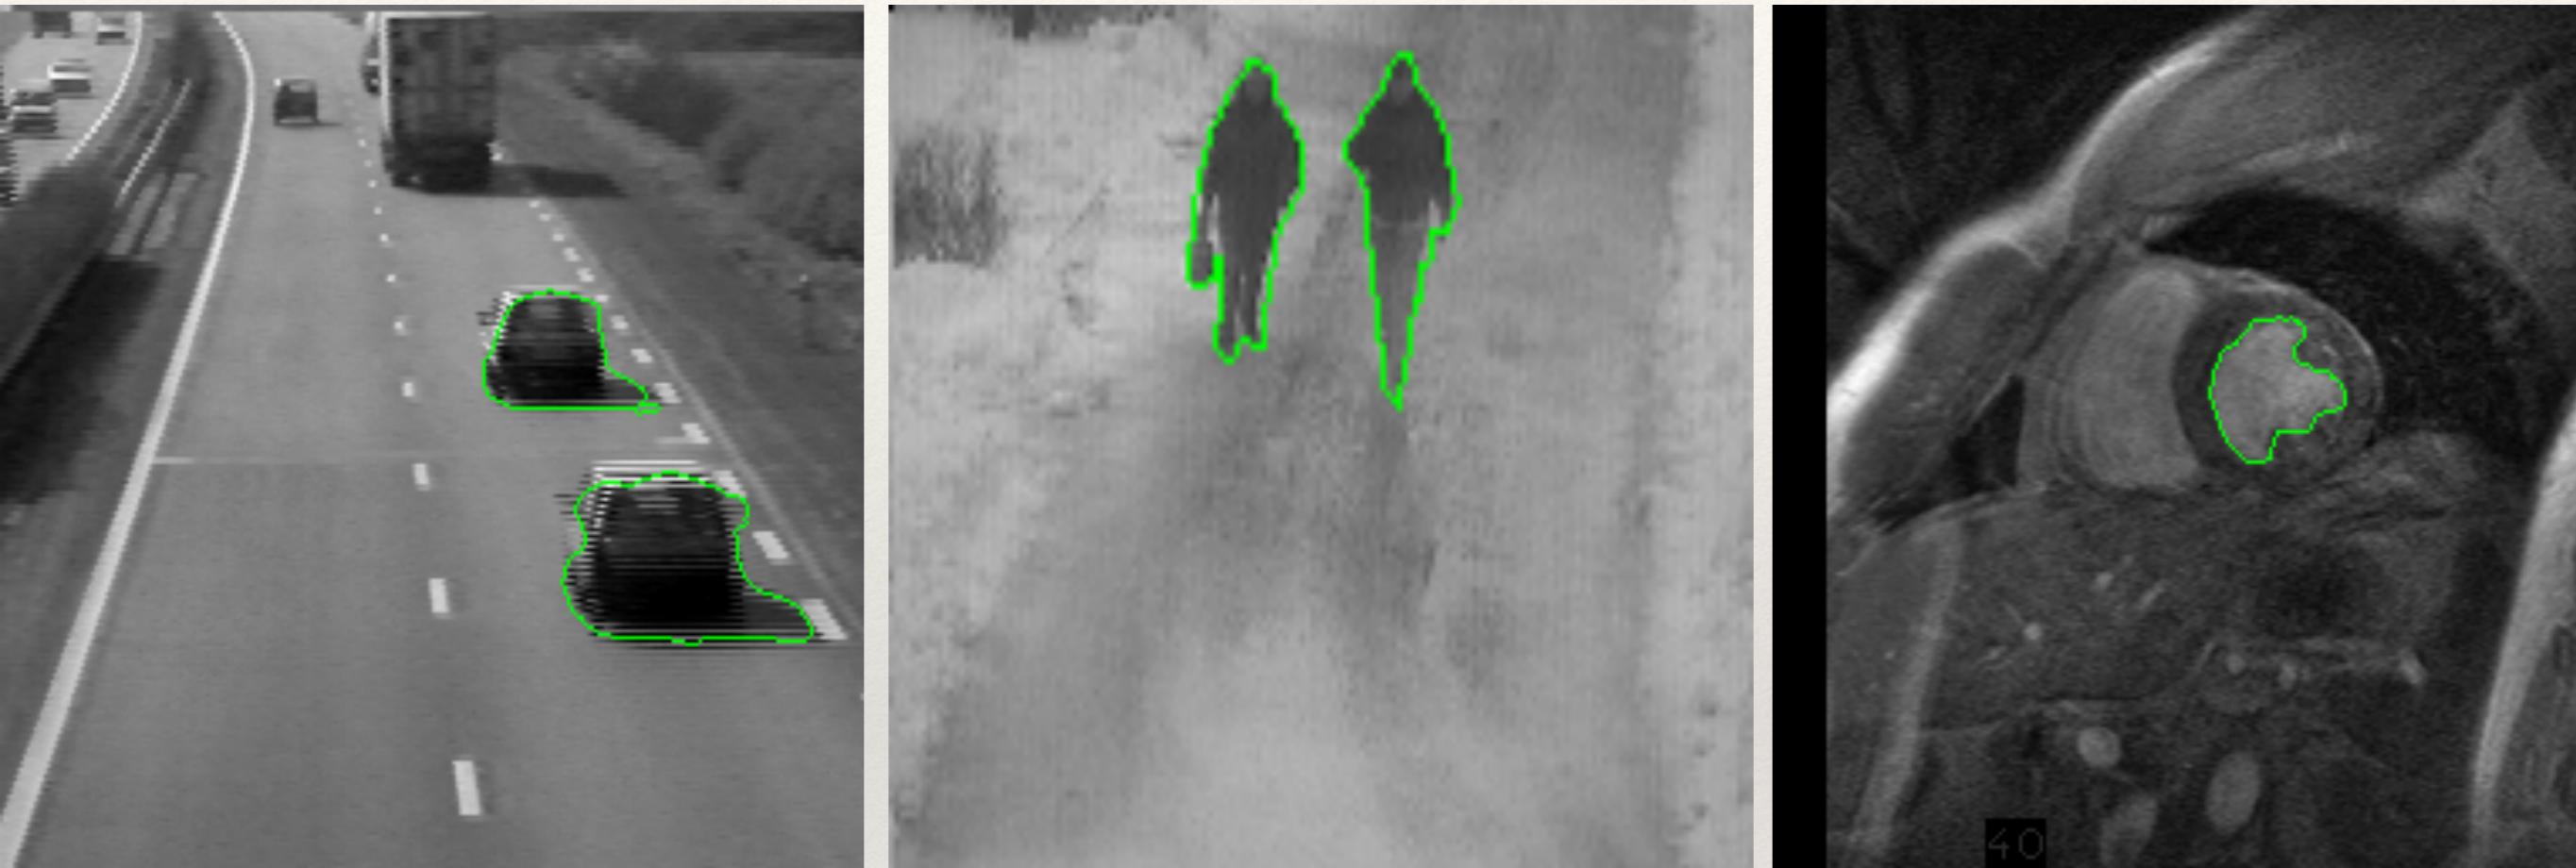
- ❖ Some objects have similar basic form but some variety in the contour shape.

Why do we want to fit deformable shapes?

- Non-rigid, deformable objects can change their shape over time, e.g. lips, hands...



Why do we want to fit deformable shapes?



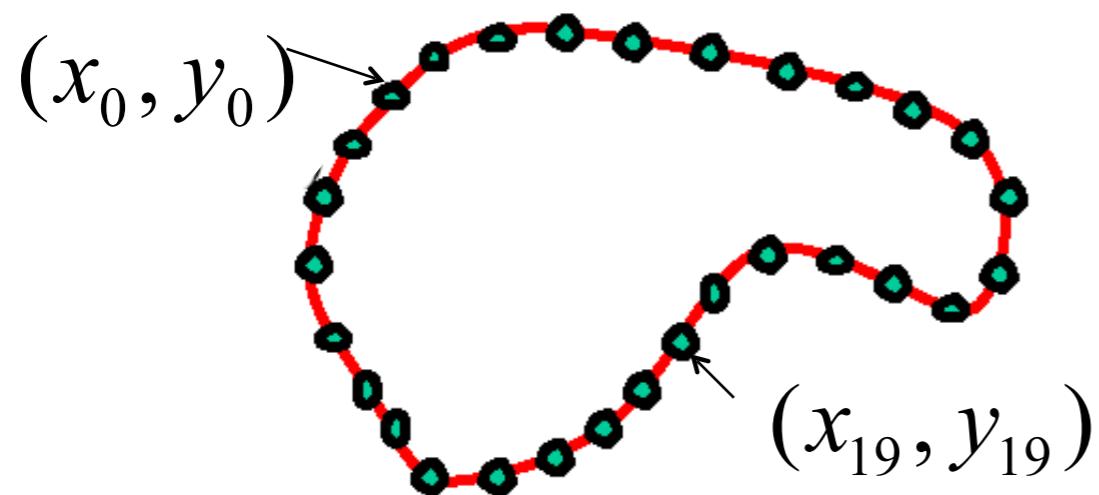
- Non-rigid, deformable objects can change their shape over time.

Aspects we need to consider

- ❖ Representation of the contours
- ❖ Defining the energy functions
 - ❖ External
 - ❖ Internal
- ❖ Minimizing the energy function
- ❖ Extensions:
 - ❖ Tracking
 - ❖ Interactive segmentation

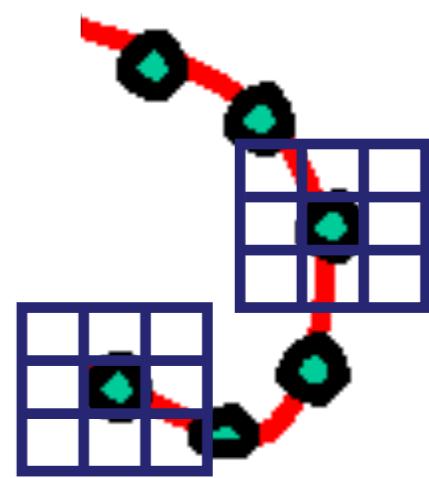
Representation

- ❖ We'll consider a discrete representation of the contour, consisting of a list of 2d point positions ("vertices").



$$\mathbf{v}_i = (x_i, y_i), \quad \text{for } i = 0, 1, \dots, n - 1$$

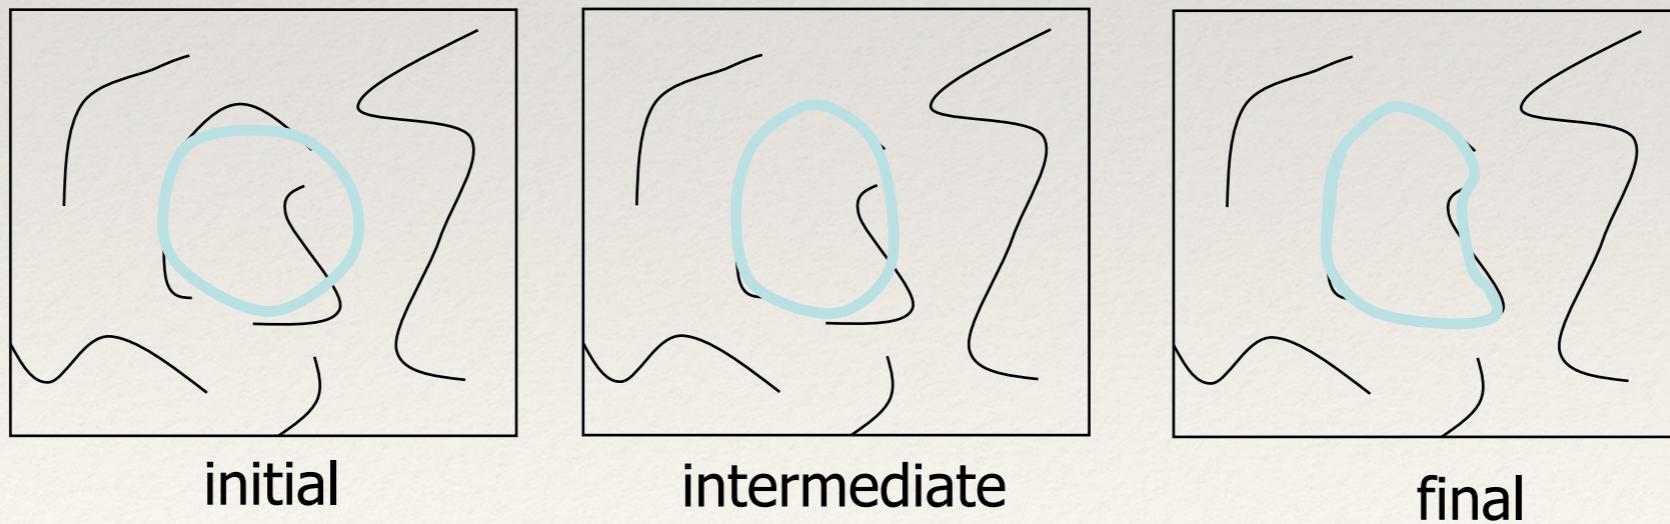
- ❖ At each iteration, we'll have the option to move each vertex to another nearby location ("state").



Fitting deformable contours

How should we adjust the current contour to form the new contour at each iteration?

- Define a cost function (“energy” function) that says how good a candidate configuration is.
- Seek next configuration that minimizes that cost function.



Energy function

The total energy (cost) of the current snake is defined as:

$$E_{total} = E_{internal} + E_{external}$$

Internal energy: encourage *prior* shape preferences: e.g., smoothness, elasticity, particular known shape.

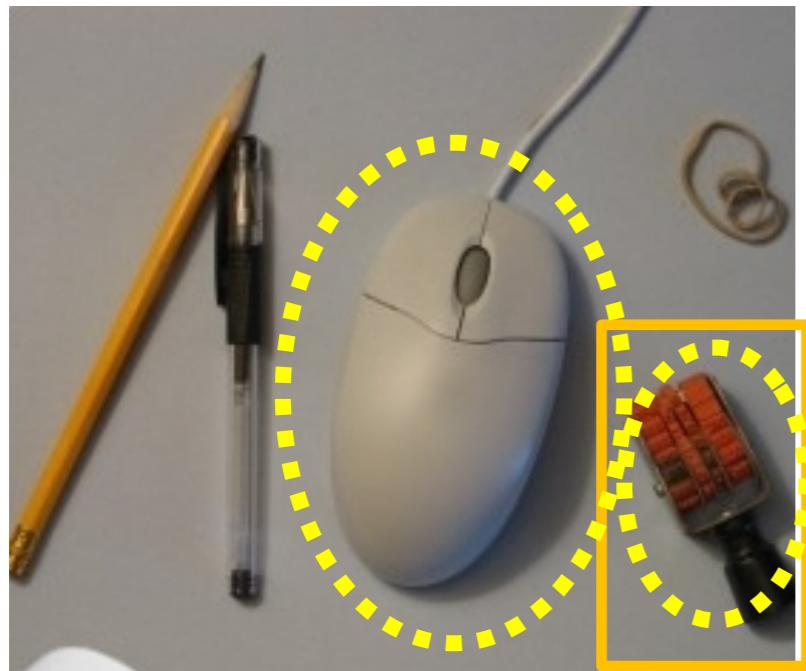
External energy ("image" energy): encourage contour to fit on places where image structures exist, e.g., edges.

A good fit between the current deformable contour and the target shape in the image will yield a **low** value for this cost function.

External energy: intuition

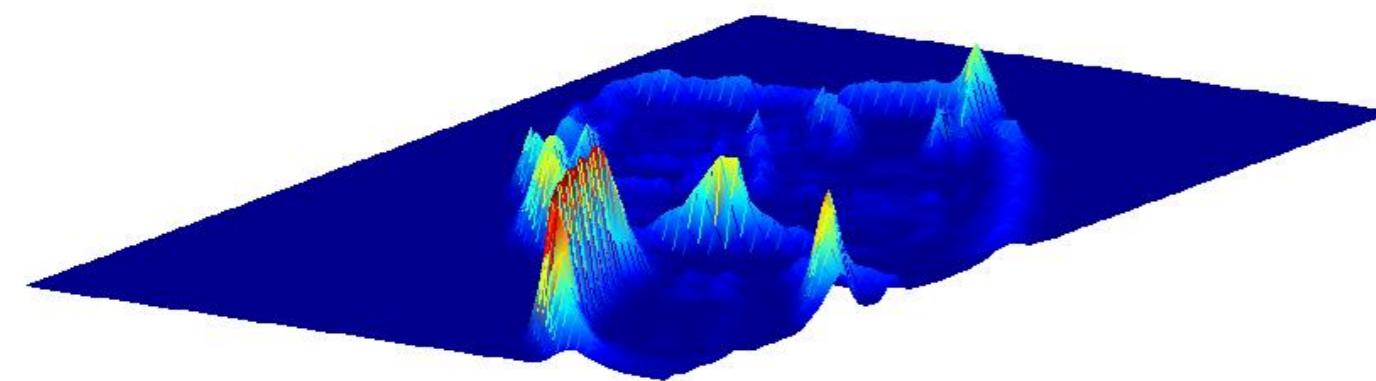
- ❖ Measure how well the curve matches the image data
- ❖ “Attract” the curve toward different image features
 - ❖ Edges, lines, texture gradient, etc.

External image energy



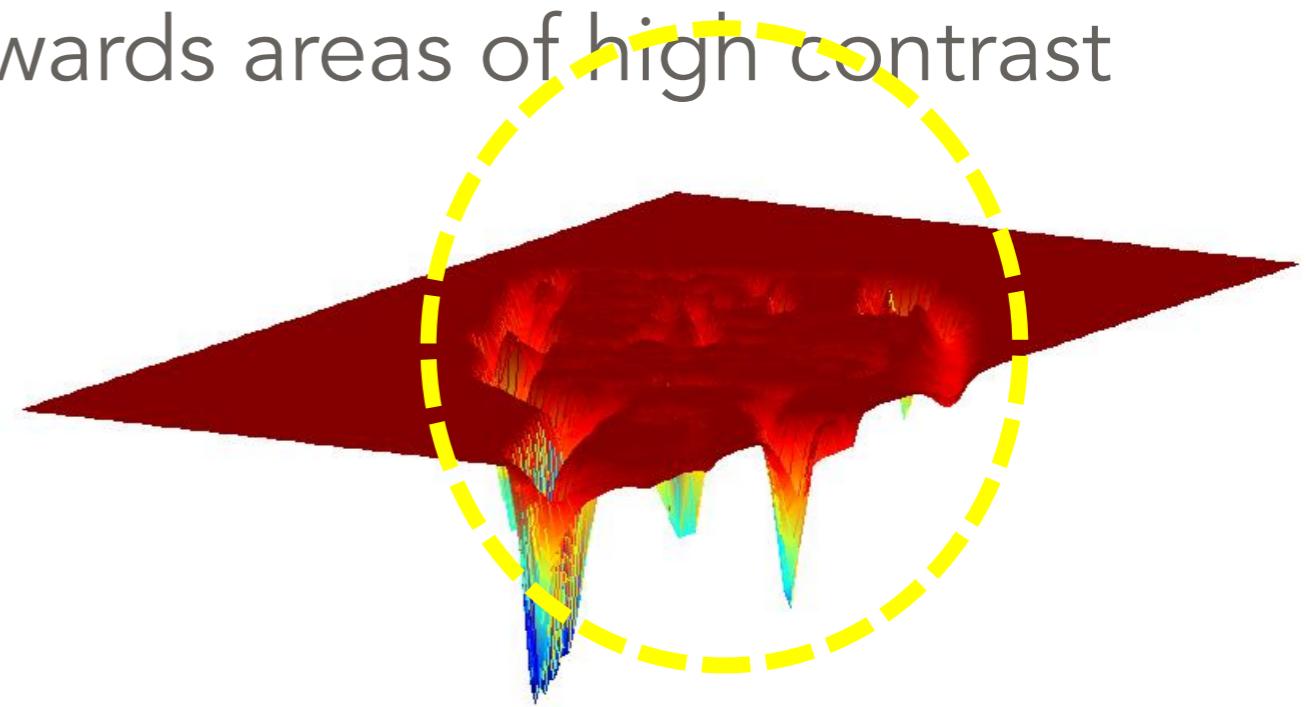
How do edges affect “snap” of rubber band?

Think of external energy from image as gravitational pull towards areas of high contrast



Magnitude of gradient

$$G_x(I)^2 + G_y(I)^2$$

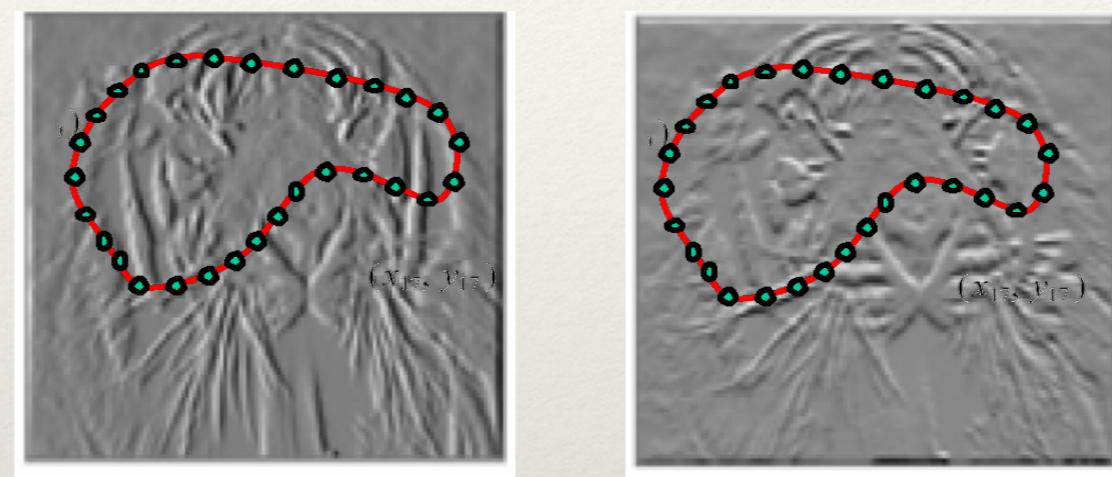


- (Magnitude of gradient)

$$- \left(G_x(I)^2 + G_y(I)^2 \right)$$

External image energy

- Gradient images $G_x(x, y)$ and $G_y(x, y)$



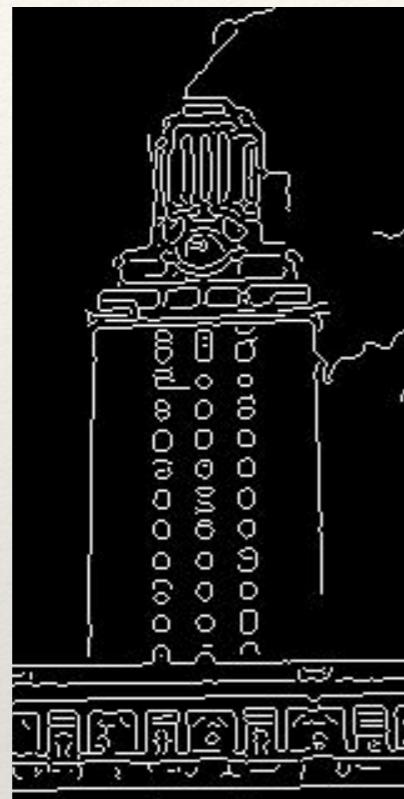
- External energy at a point on the curve is:

$$E_{external}(\mathbf{v}) = -(|G_x(\mathbf{v})|^2 + |G_y(\mathbf{v})|^2)$$

- External energy for the whole curve:

$$E_{external} = - \sum_{i=0}^{n-1} |G_x(x_i, y_i)|^2 + |G_y(x_i, y_i)|^2$$

Internal energy: intuition



What are the underlying boundaries in this fragmented edge image?



And in this one?

Internal energy: intuition

A priori, we want to favor **smooth** shapes, contours with **low curvature**, contours similar to a **known shape**, etc. to balance what is actually observed (i.e., in the gradient image).



Internal energy

For a *continuous* curve, a common internal energy term is the “bending energy”.
At some point $v(s)$ on the curve, this is:

$$E_{internal}(v(s)) = \alpha \left| \frac{dv}{ds} \right|^2 + \beta \left| \frac{d^2v}{d^2s} \right|^2$$

Tension,
Elasticity

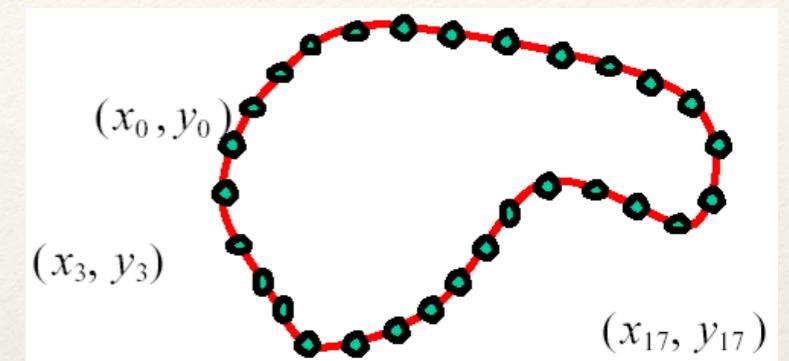
Stiffness,
Curvature



Internal energy

- ❖ For our discrete representation,

$$\mathbf{v}_i = (x_i, y_i) \quad i = 0 \dots n - 1$$



$$\frac{d\mathbf{v}}{ds} \approx \mathbf{v}_{i+1} - \mathbf{v}_i \quad \frac{d^2\mathbf{v}}{ds^2} \approx (\mathbf{v}_{i+1} - \mathbf{v}_i) - (\mathbf{v}_i - \mathbf{v}_{i-1}) = \mathbf{v}_{i+1} - 2\mathbf{v}_i + \mathbf{v}_{i-1}$$

*Note these are derivatives relative to **position**--not spatial image gradients.*

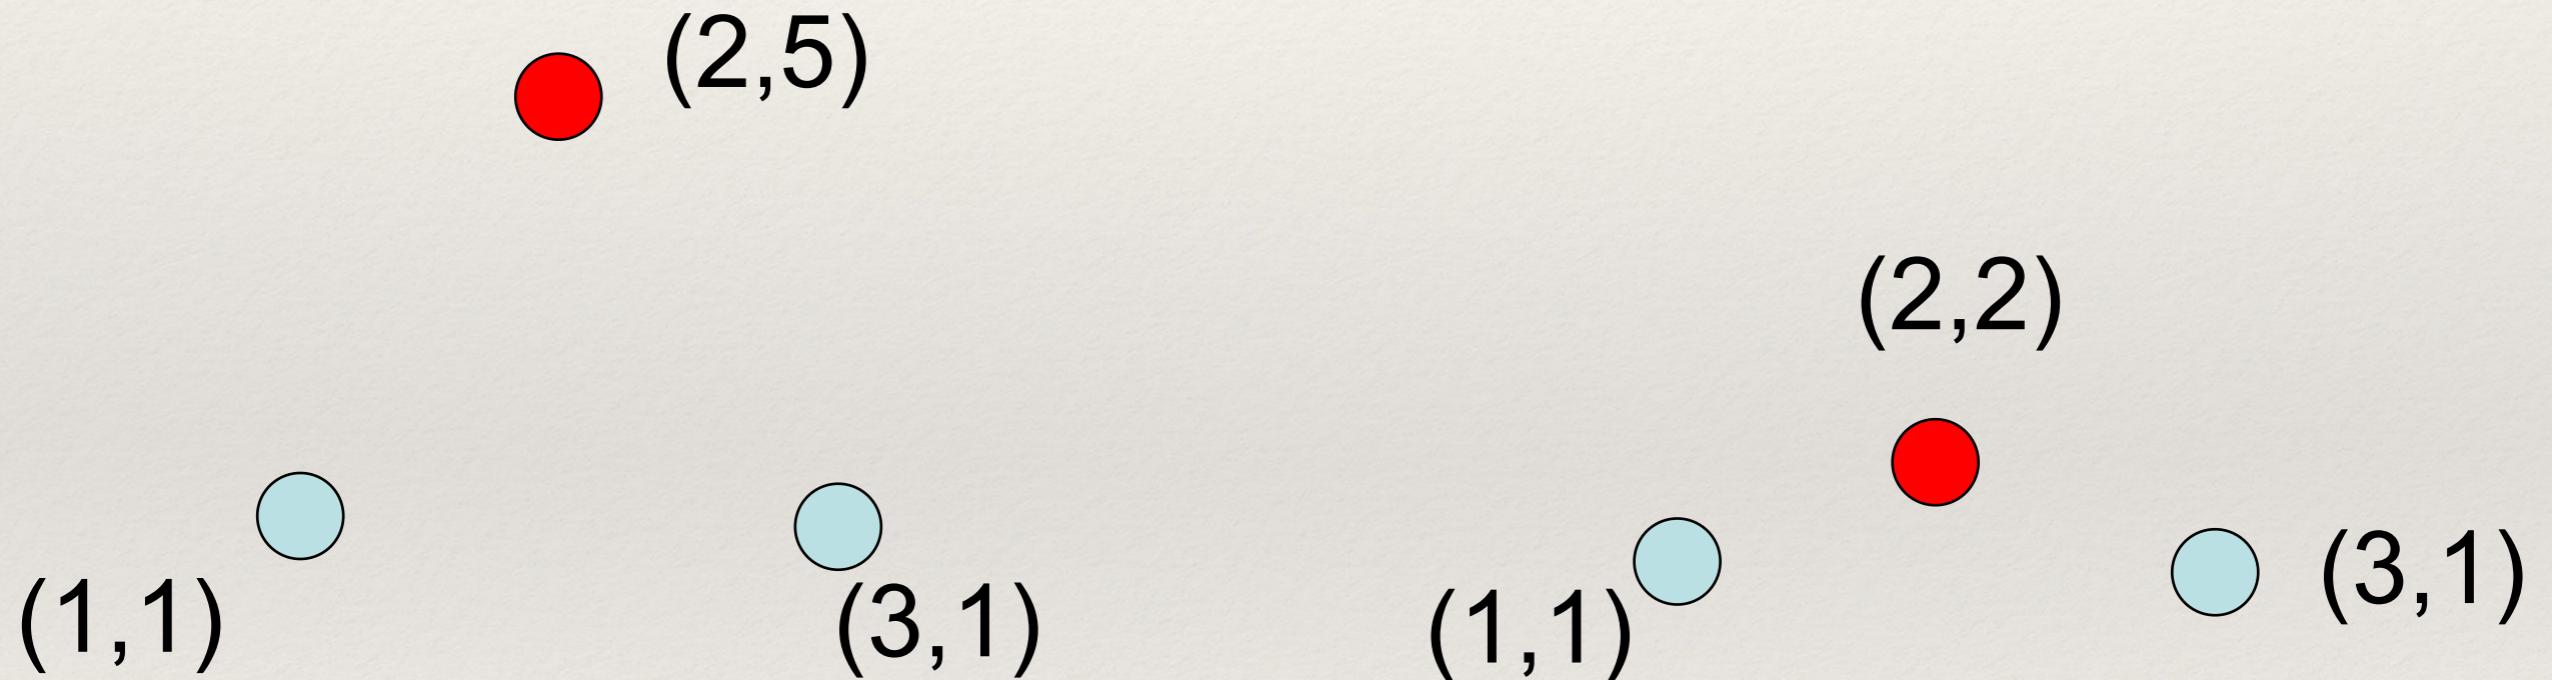
- ❖ Internal energy for the whole curve:

$$E_{internal} = \sum_{i=0}^{n-1} \alpha \|\mathbf{v}_{i+1} - \mathbf{v}_i\|^2 + \beta \|\mathbf{v}_{i+1} - 2\mathbf{v}_i + \mathbf{v}_{i-1}\|^2$$

*Why do these reflect **tension** and **curvature**?*

Example: compare curvature

$$E_{curvature}(v_i) = \left\| v_{i+1} - 2v_i + v_{i-1} \right\|^2 \\ = (x_{i+1} - 2x_i + x_{i-1})^2 + (y_{i+1} - 2y_i + y_{i-1})^2$$

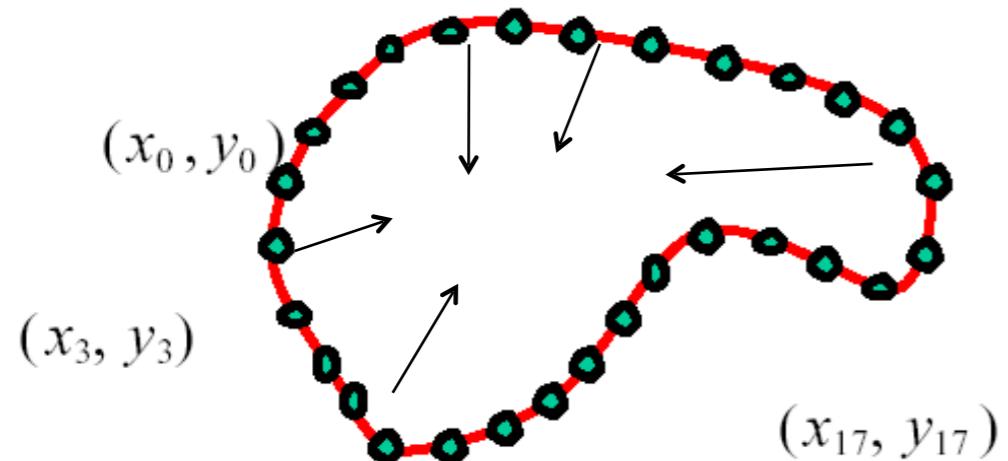


$$(3 - 2(2) + 1)2 + (1 - 2(5) + 1)2 \\ = (-8)2 = 64$$

$$(3 - 2(2) + 1)2 + (1 - 2(2) + 1)2 \\ = (-2)2 = 4$$

Penalizing elasticity

$$\begin{aligned} E_{elastic} &= \sum_{i=0}^{n-1} \alpha \|\mathbf{v}_{i+1} - \mathbf{v}_i\|^2 \\ &= \alpha \cdot \sum_{i=0}^{n-1} (x_{i+1} - x_i)^2 + (y_{i+1} - y_i)^2 \end{aligned}$$



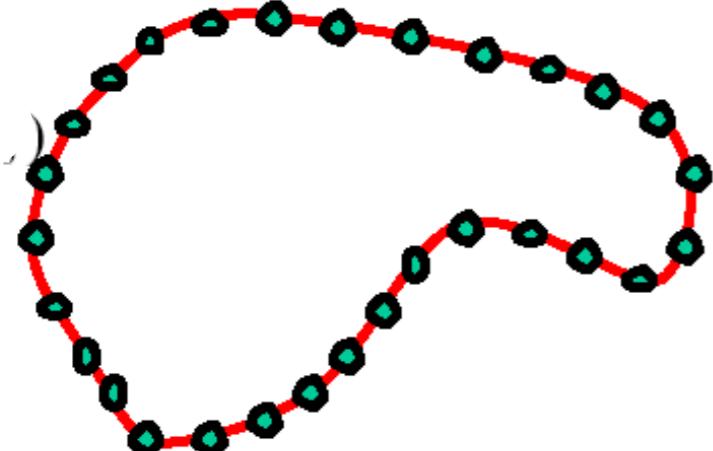
What is the possible problem with this definition?

Penalizing elasticity

$$E_{elastic} = \sum_{i=0}^{n-1} \alpha \|\mathbf{v}_{i+1} - \mathbf{v}_i\|^2$$

Instead:

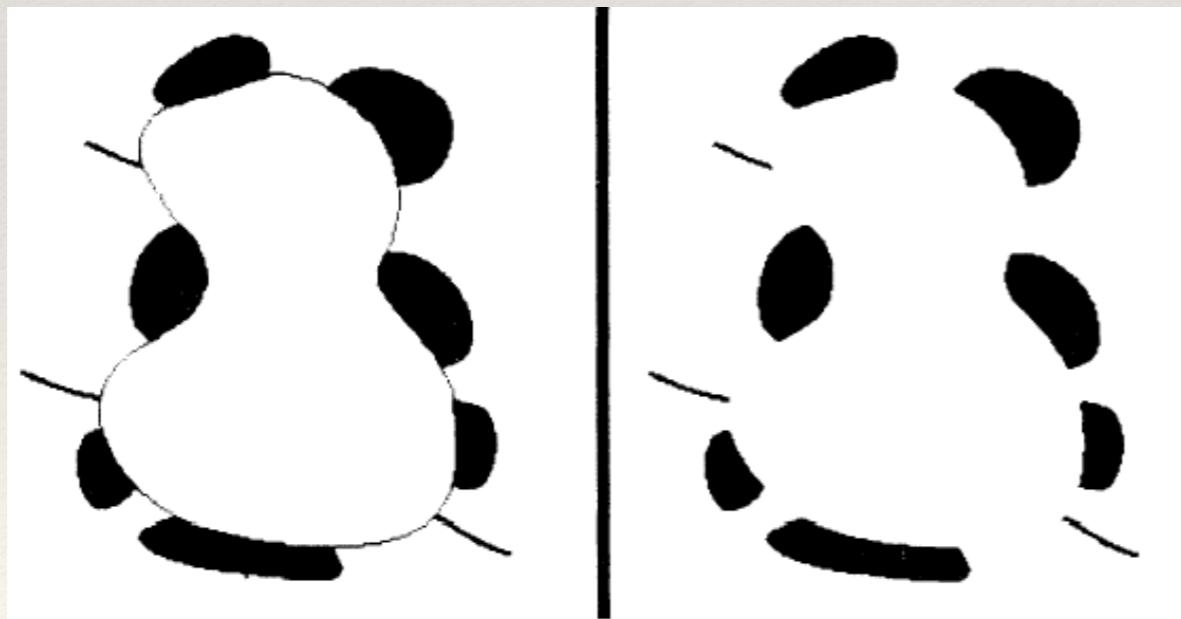
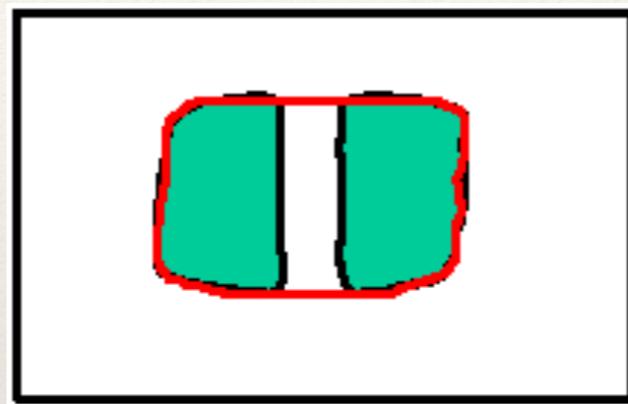
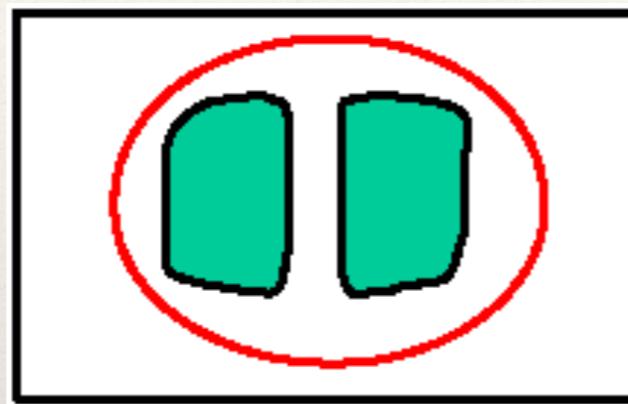
$$= \alpha \cdot \sum_{i=0}^{n-1} \left((x_{i+1} - x_i)^2 + (y_{i+1} - y_i)^2 - \bar{d} \right)^2$$



where d is the average distance between pairs of points – updated at each iteration.

Dealing with missing data

- ❖ The preferences for low-curvature, smoothness help deal with missing data:



Illusory contours found!

[Figure from Kass et al. 1987]

Extending the internal energy: capture shape prior

- ❖ If object is some smooth variation on a known shape, we can use a term that will penalize deviation from that shape:

$$E_{internal} + = \alpha \cdot \sum_{i=0}^{n-1} (\mathbf{v}_i - \hat{\mathbf{v}}_i)^2$$

where $\{\hat{\mathbf{v}}_i\}$ are the points of the known shape.

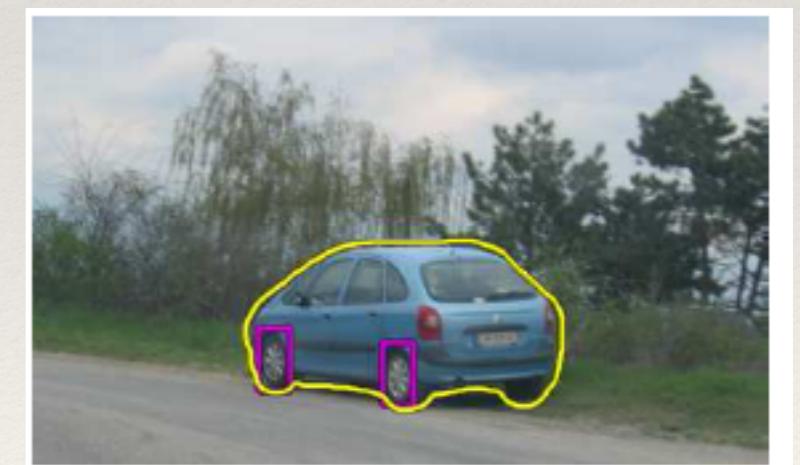
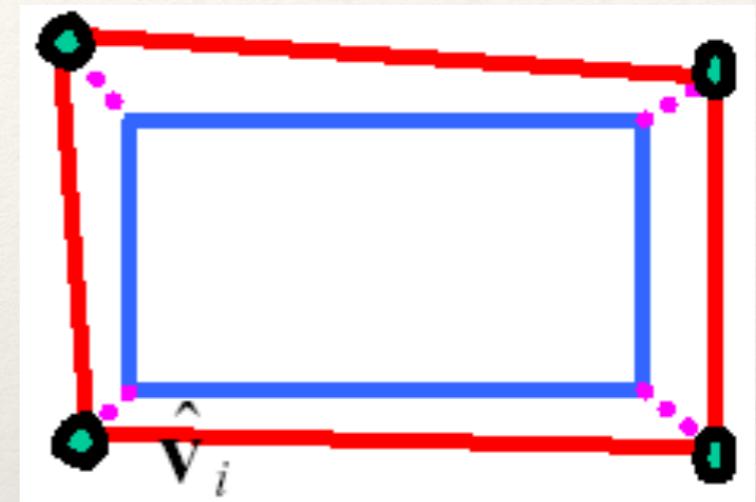


Fig from Y. Boykov

Total energy: function of the weights

$$E_{total} = E_{internal} + \gamma E_{external}$$

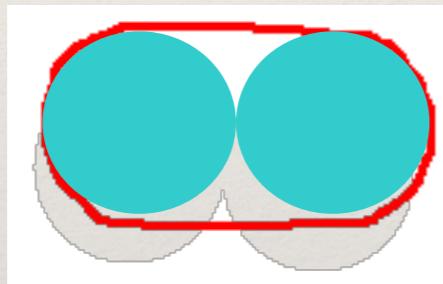
$$E_{external} = - \sum_{i=0}^{n-1} |G_x(x_i, y_i)|^2 + |G_y(x_i, y_i)|^2$$

$$E_{internal} = \sum_{i=0}^{n-1} \alpha (\bar{d} - \|v_{i+1} - v_i\|) + \beta \|v_{i+1} - 2v_i + v_{i-1}\|^2$$

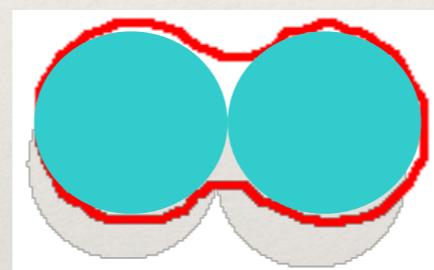
Total energy: function of the weights

e.g., weight controls the penalty for internal elasticity

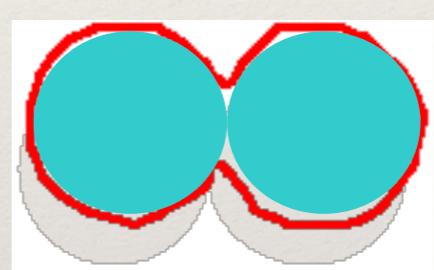
α



large α



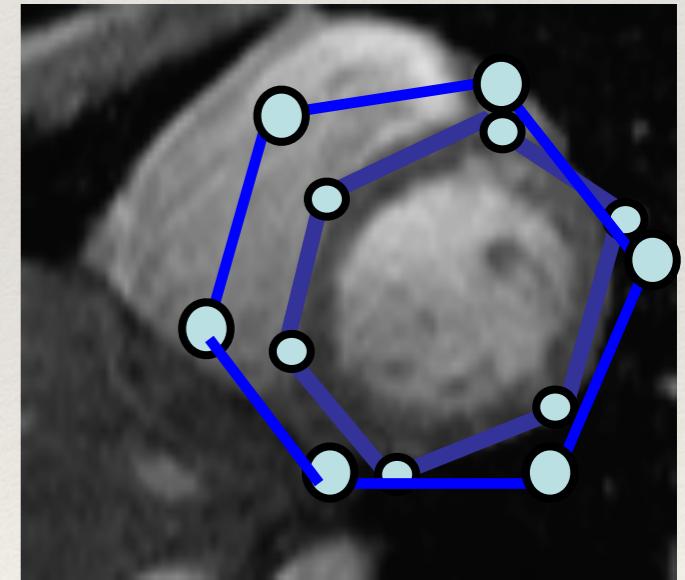
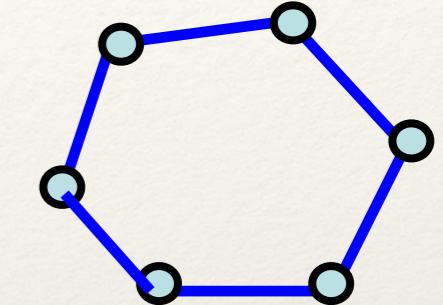
medium α



small α

Recap: deformable contour

- ❖ A simple elastic snake is defined by:
 - ❖ A set of n points,
 - ❖ An internal energy term (tension, bending, plus optional shape prior)
 - ❖ An external energy term (gradient-based)
- ❖ To use to segment an object:
 - ❖ Initialize in the vicinity of the object
 - ❖ Modify the points to minimize the total energy

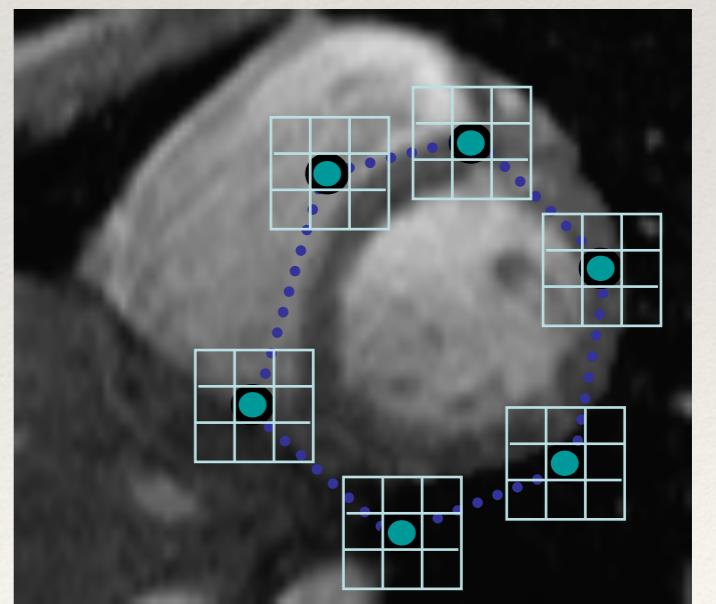


Energy minimization

- ❖ Several algorithms have been proposed to fit deformable contours.
- ❖ We'll look at two:
 - ❖ Greedy search
 - ❖ Dynamic programming (for 2d snakes)

Energy minimization: greedy

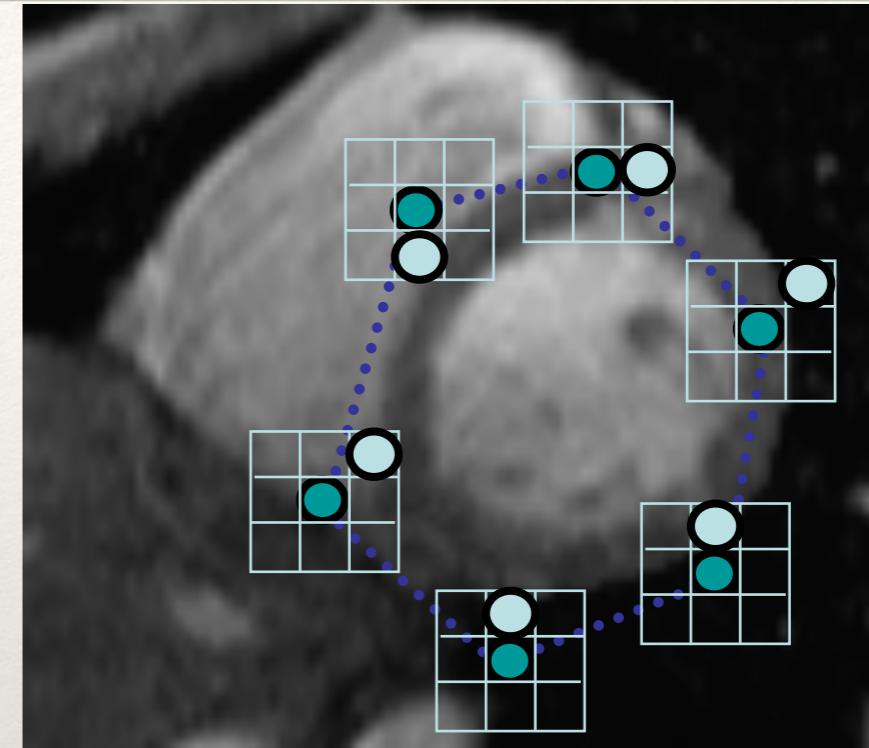
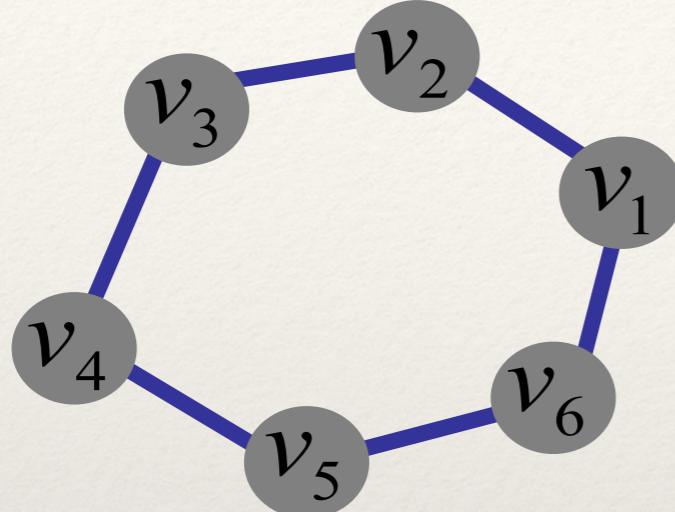
- ❖ For each point, search window around it and move to where energy function is minimal
 - ❖ Typical window size, e.g., 5×5 pixels
- ❖ Stop when predefined number of points have not changed in last iteration, or after max number of iterations
- ❖ Note:
 - ❖ Convergence not guaranteed
 - ❖ Need decent initialization



Energy minimization

- ❖ Several algorithms have been proposed to fit deformable contours.
- ❖ We'll look at two:
 - ❖ Greedy search
 - ❖ Dynamic programming (for 2d snakes)

Energy minimization: dynamic programming



- ❖ With this form of the energy function, we can minimize using dynamic programming, with the *Viterbi* algorithm.
- ❖ Iterate until optimal position* for each point is the center of the box, i.e., the snake is optimal in the local search space constrained by boxes.
- ❖ *optimal within resolution of considered window

Fig from Y. Boykov
[Amini, Weymouth, Jain, 1990]

Energy minimization: dynamic programming

- ❖ Possible because snake energy can be rewritten as a sum of pair-wise interaction potentials:

$$E_{total}(\mathbf{v}_1, \dots, \mathbf{v}_n) = \sum_{i=1}^{n-1} E_i(\mathbf{v}_i, \mathbf{v}_{i+1})$$

- ❖ Or sum of triple-interaction potentials.

$$E_{total}(\mathbf{v}_1, \dots, \mathbf{v}_n) = \sum_{i=1}^{n-1} E_i(\mathbf{v}_{i-1}, \mathbf{v}_i, \mathbf{v}_{i+1})$$

Snake energy: pair-wise interactions

$$\begin{aligned} E_{total}(x_1, \dots, x_n, y_1, \dots, y_n) &= - \sum_{i=1}^{n-1} |G_x(x_i, y_i)|^2 + |G_y(x_i, y_i)|^2 \\ &\quad + \alpha \cdot \sum_{i=1}^{n-1} (x_{i+1} - x_i)^2 + (y_{i+1} - y_i)^2 \end{aligned}$$

Re-writing the above with $v_i = (x_i, y_i)$:

$$E_{total}(v_1, \dots, v_n) = - \sum_{i=1}^{n-1} \|G(v_i)\|^2 + \alpha \cdot \sum_{i=1}^{n-1} \|v_{i+1} - v_i\|^2$$

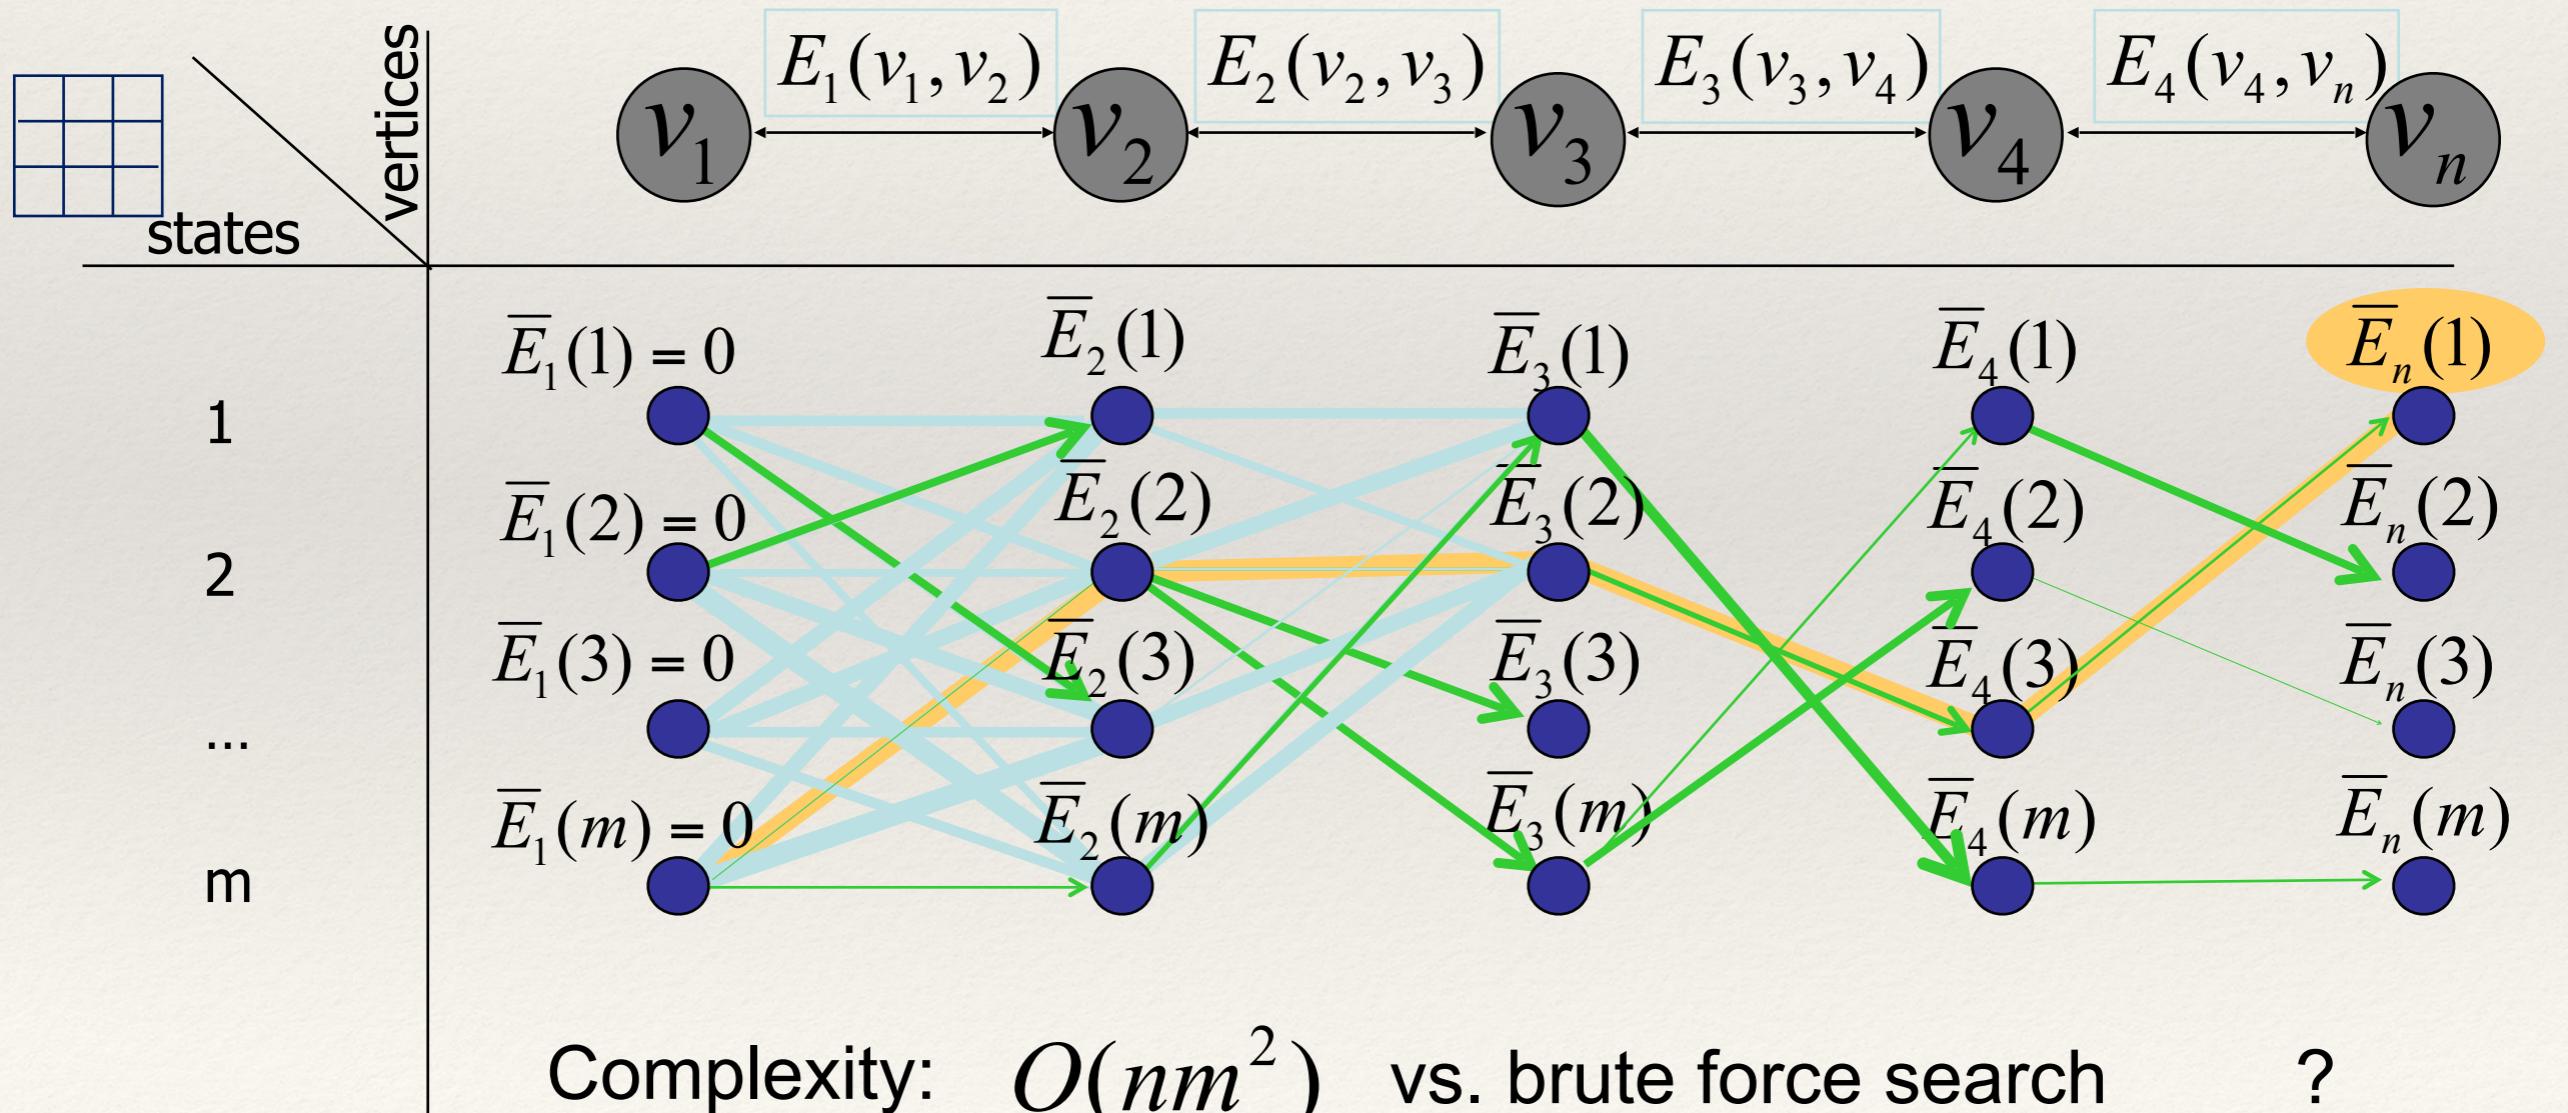
$$E_{total}(v_1, \dots, v_n) = E_1(v_1, v_2) + E_2(v_2, v_3) + \dots + E_{n-1}(v_{n-1}, v_n)$$

where $E_i(v_i, v_{i+1}) = -\|G(v_i)\|^2 + \alpha \|v_{i+1} - v_i\|^2$

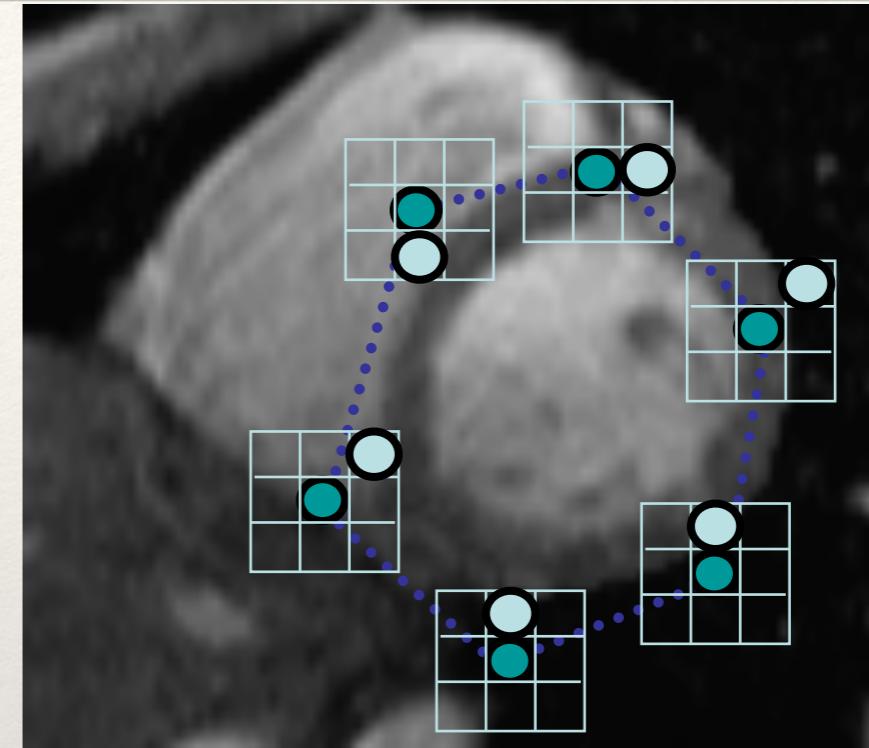
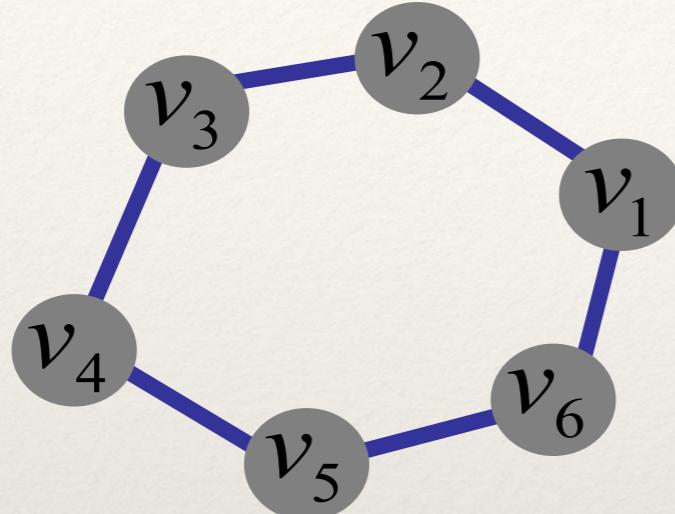
Viterbi algorithm

Main idea: determine optimal position (state) of predecessor, for each possible position of self. Then backtrack from best state for last vertex.

$$E_{total} = E_1(v_1, v_2) + E_2(v_2, v_3) + \dots + E_{n-1}(v_{n-1}, v_n)$$



Energy minimization: dynamic programming



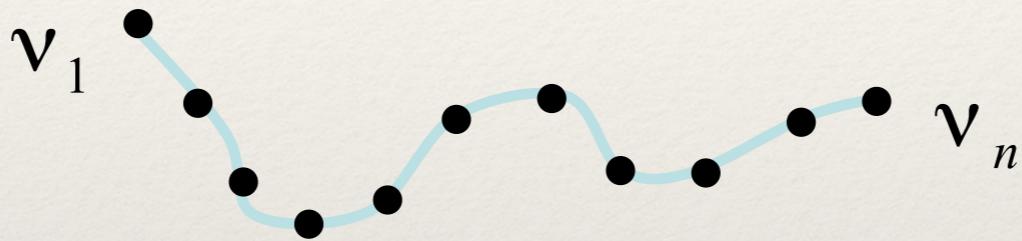
- ❖ With this form of the energy function, we can minimize using dynamic programming, with the *Viterbi* algorithm.
- ❖ Iterate until optimal position for each point is the center of the box, i.e., the snake is optimal in the local search space constrained by boxes.

Fig from Y. Boykov
[Amini, Weymouth, Jain, 1990]

Energy minimization: dynamic programming

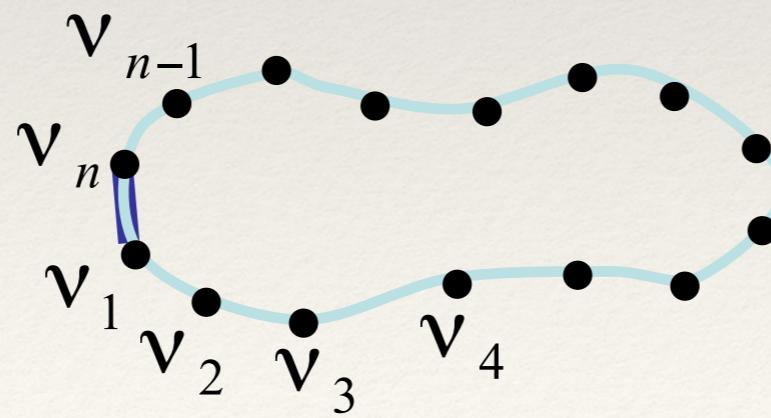
DP can be applied to optimize an open ended snake

$$E_1(v_1, v_2) + E_2(v_2, v_3) + \dots + E_{n-1}(v_{n-1}, v_n)$$



For a closed snake, a “loop” is introduced into the total energy.

$$E_1(v_1, v_2) + E_2(v_2, v_3) + \dots + E_{n-1}(v_{n-1}, v_n) + \boxed{E_n(v_n, v_1)}$$



Work around:

- 1) Fix v_1 and solve for rest .
- 2) Fix an intermediate node at its position found in (1), solve for rest.

Aspects we need to consider

- ❖ Representation of the contours
- ❖ Defining the energy functions
 - ❖ External
 - ❖ Internal
- ❖ Minimizing the energy function
- ❖ Extensions:
 - ❖ Tracking
 - ❖ Interactive segmentation

Tracking via deformable contours

1. Use final contour/model extracted at frame t as an initial solution for frame $t+1$
2. Evolve initial contour to fit exact object boundary at frame $t+1$
3. Repeat, initializing with most recent frame.



Tracking Heart Ventricles
(multiple frames)

Tracking via deformable contours



[Visual Dynamics Group](#), Dept. Engineering Science, University of Oxford.

Applications:

- Traffic monitoring
- Human-computer interaction
- Animation
- Surveillance
- Computer assisted diagnosis in medical imaging

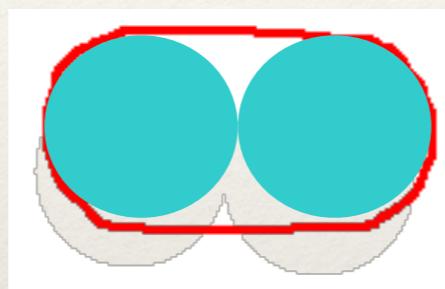
3D active contours

- ❖ Jörgen Ahlberg
- ❖ [http://www.cvl.isy.liu.se/
ScOut/Masters/Papers/
Ex1708.pdf](http://www.cvl.isy.liu.se/ScOut/Masters/Papers/Ex1708.pdf)



Limitations

- ❖ May over-smooth the boundary



- ❖ Cannot follow topological changes of objects



Limitations

- ❖ External energy: snake does not really “see” object boundaries in the image unless it gets very close to it.

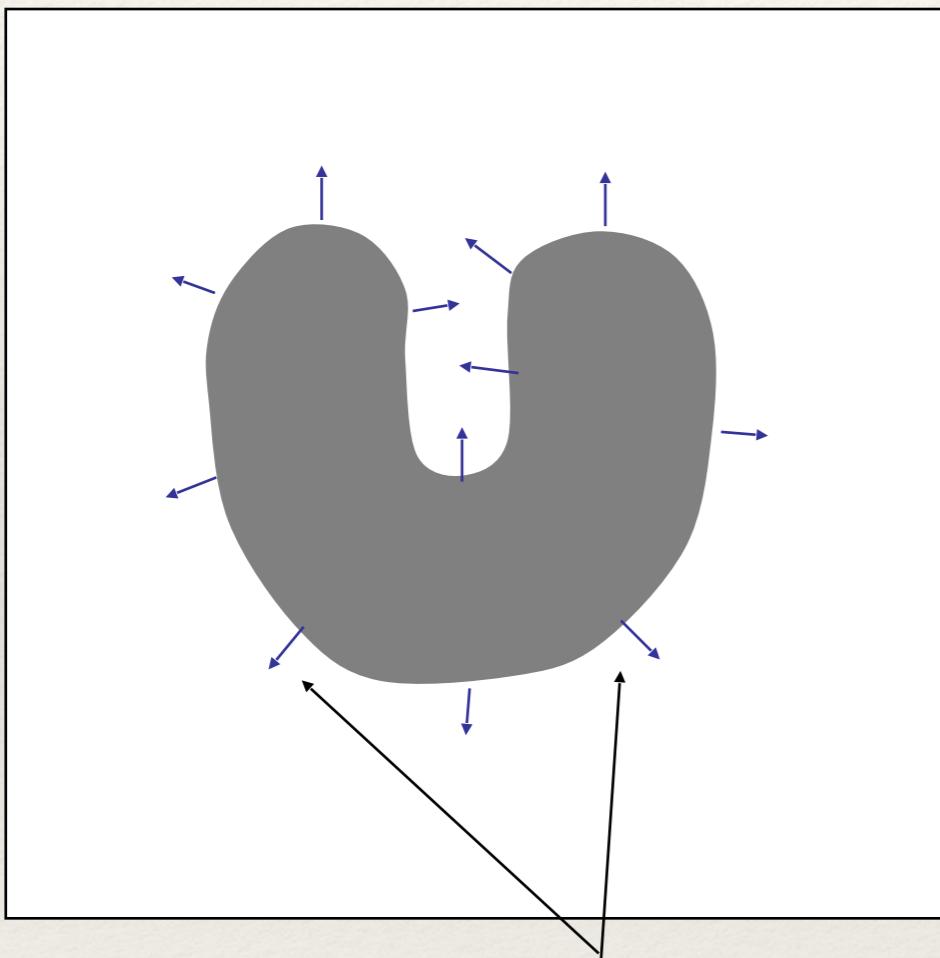
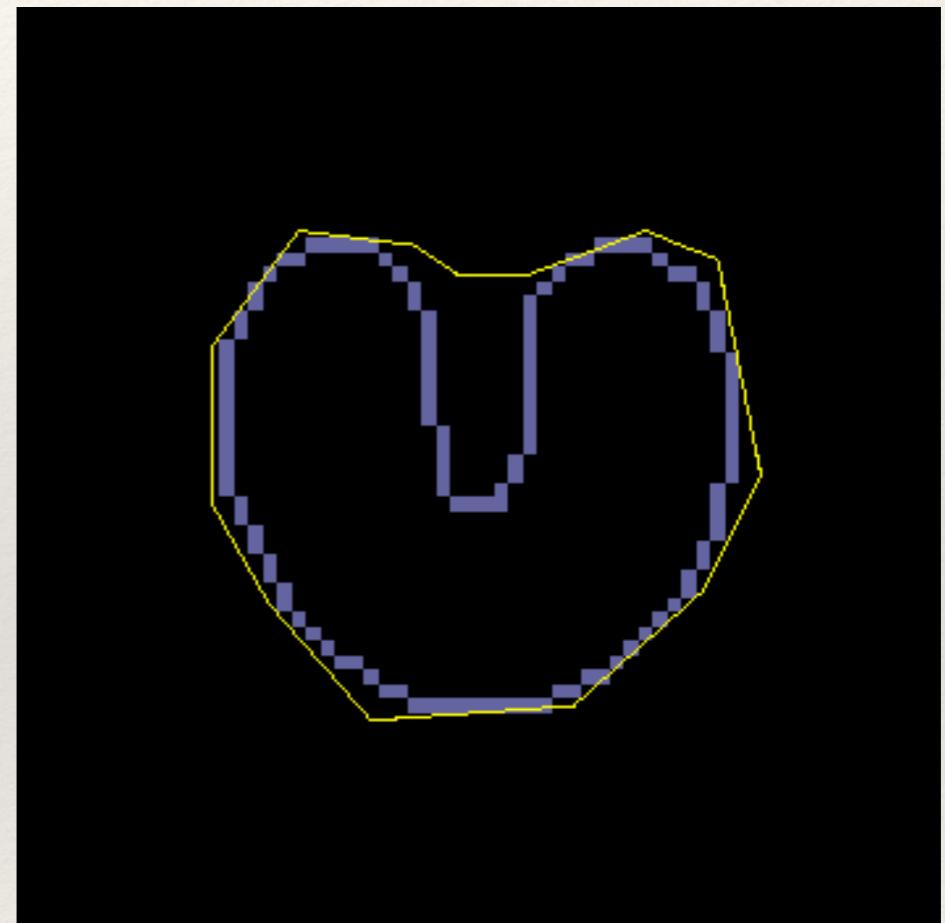


image gradients ∇I
are large only directly on the boundary



Distance transform

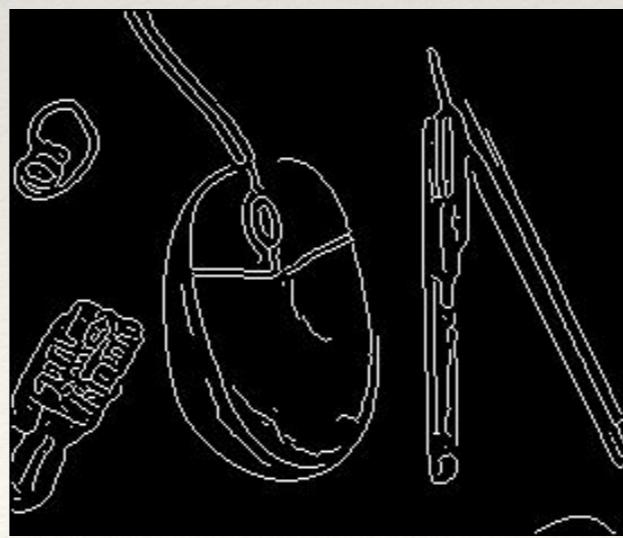
- ❖ External image can instead be taken from the **distance transform** of the edge image.



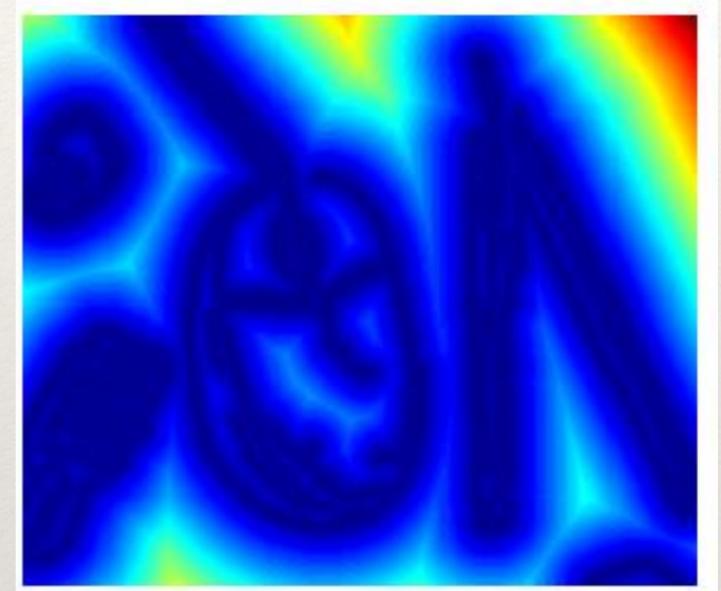
original



-gradient



edges



distance transform

Value at (x,y) tells how far that position is from the nearest edge point (or other binary mage structure)

>> **help bwdist**

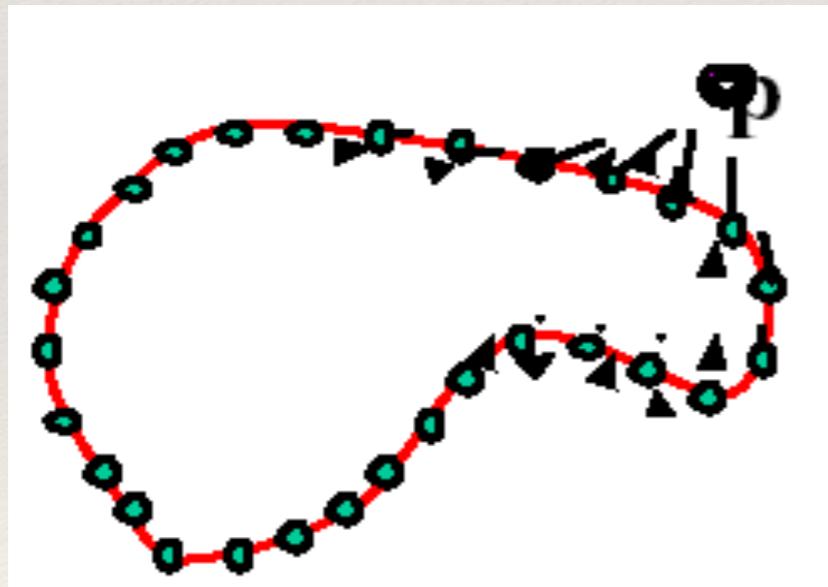
Interactive forces



How can we implement such an *interactive* force with deformable contours?

Interactive forces

- ❖ An energy function can be altered online based on user input – use the cursor to push or pull the initial snake away from a point.
- ❖ Modify external energy term to include:



$$E_{push} = \sum_{i=0}^{n-1} \frac{r^2}{|\mathbf{v}_i - p|^2}$$

Nearby points get pushed hardest

Intelligent scissors

Another form of
interactive
segmentation:

Compute optimal paths
**from every point to the
seed** based on edge-
related costs.

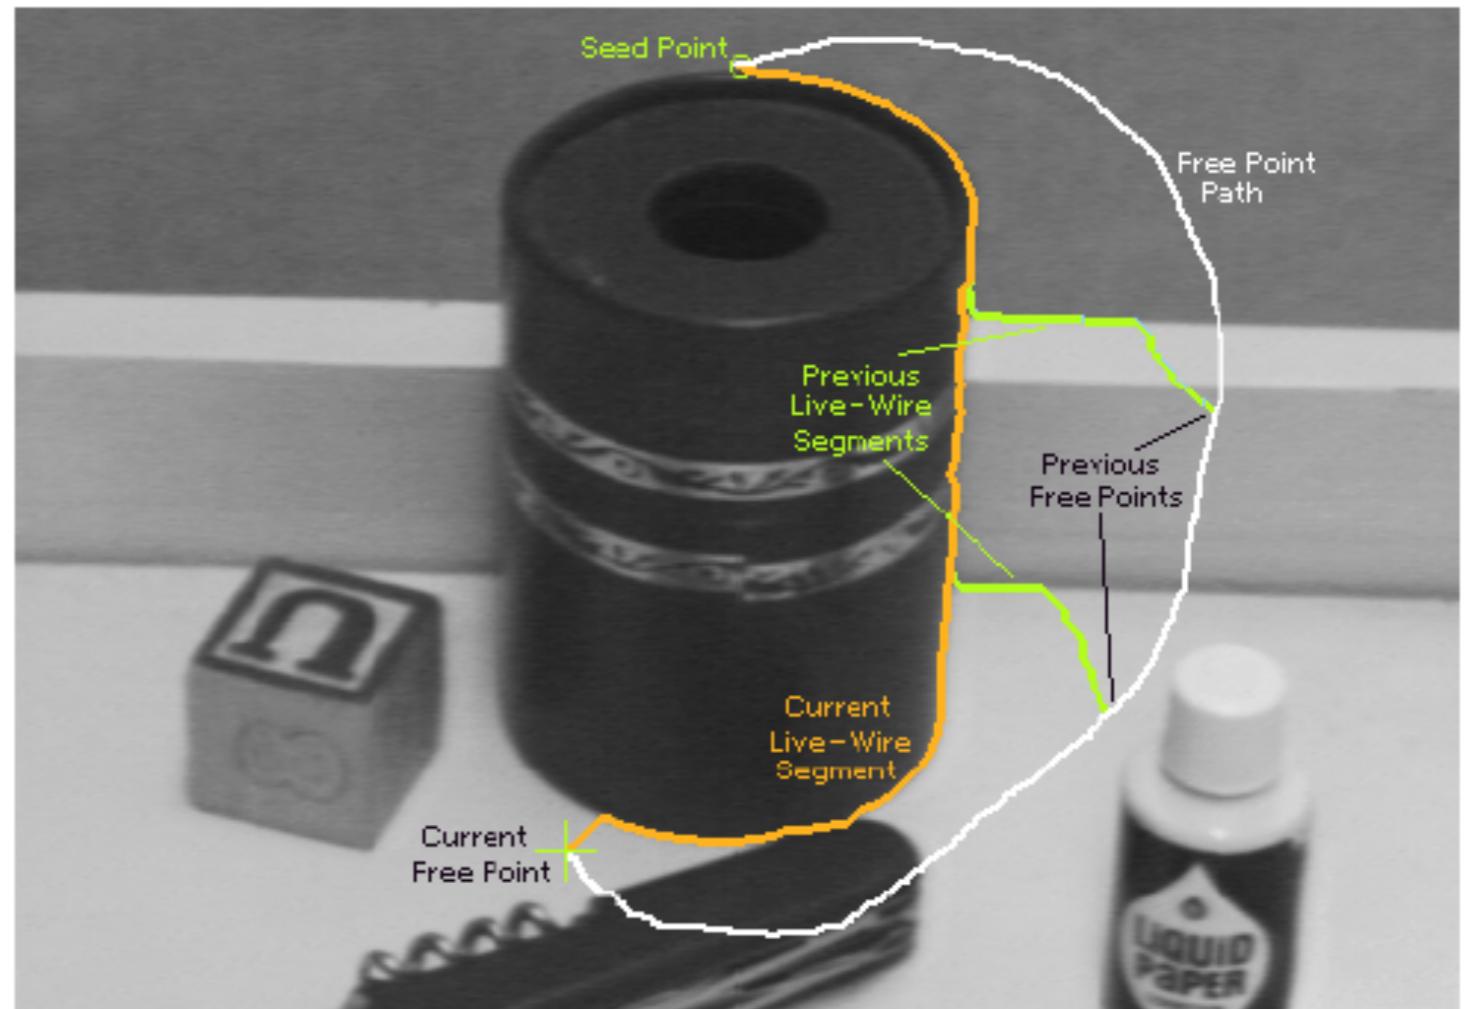
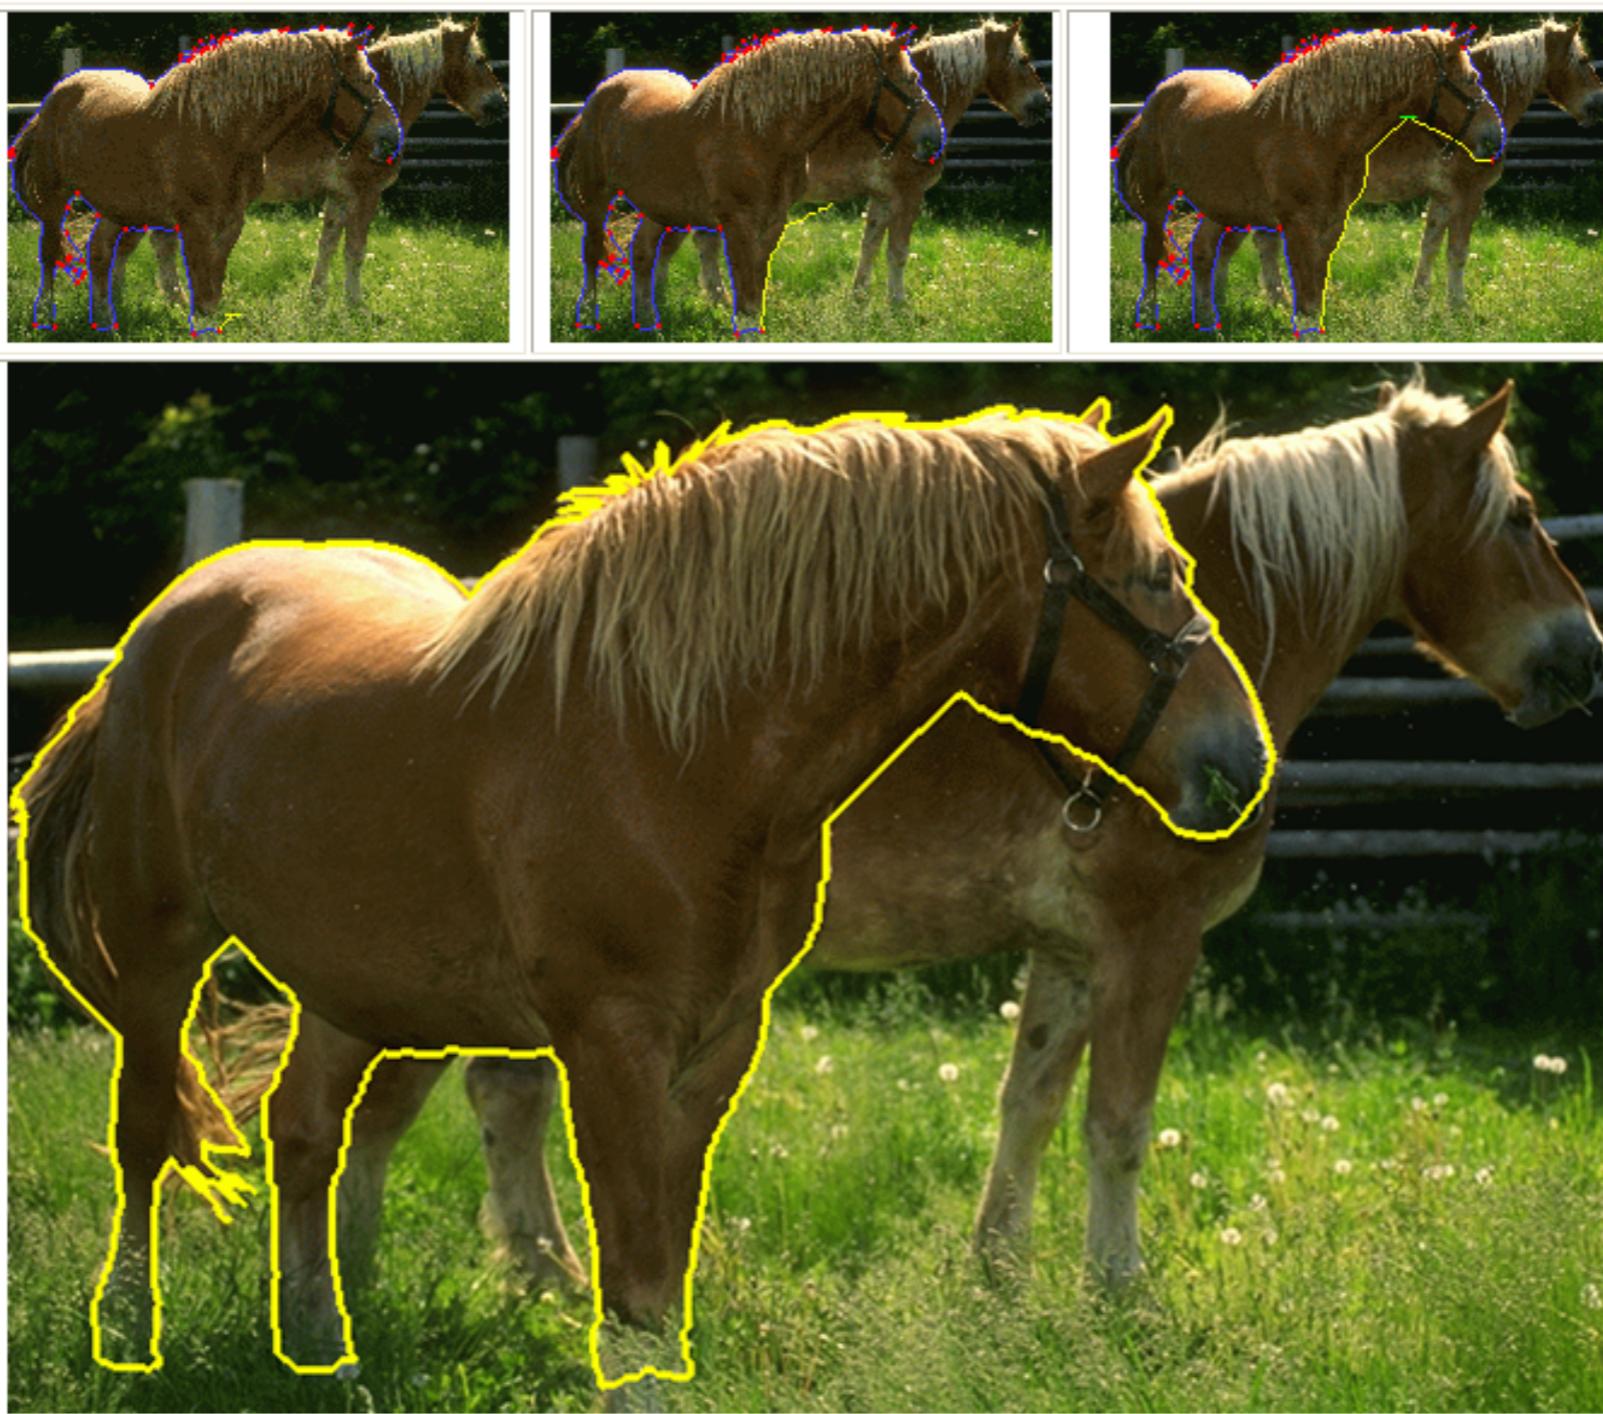


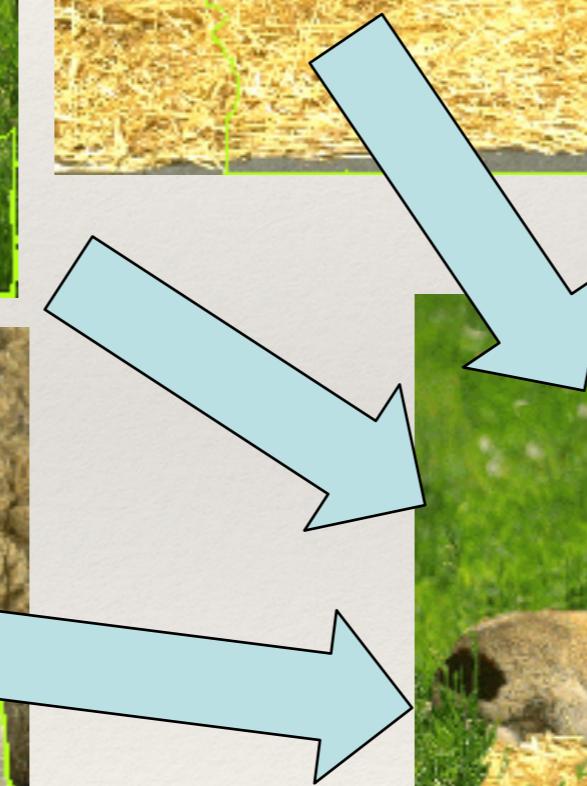
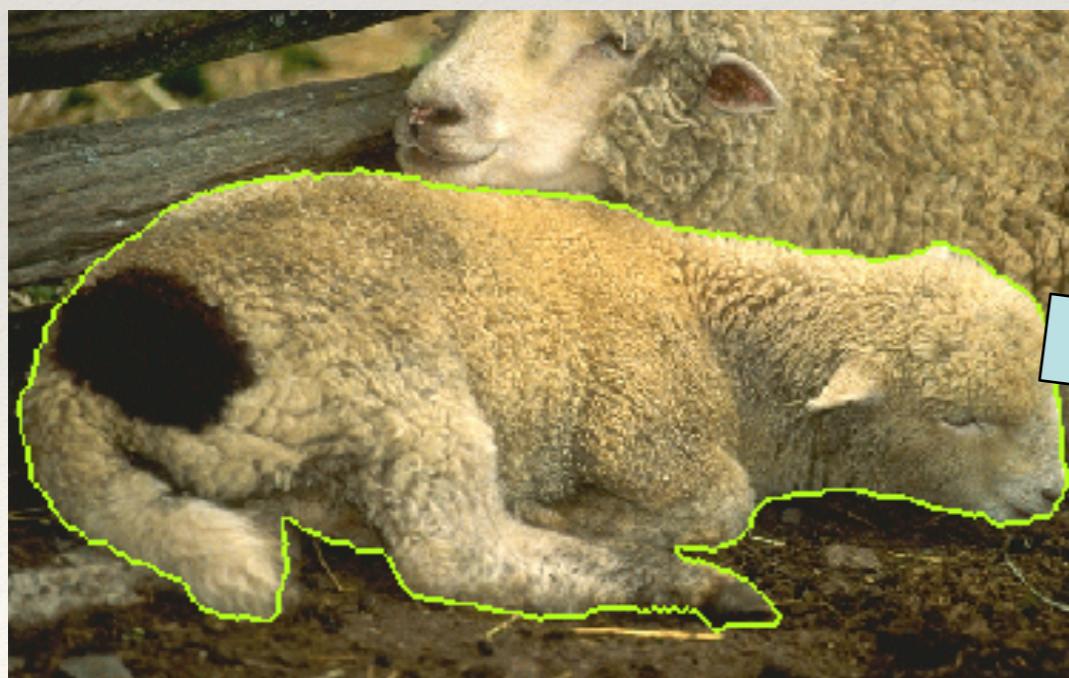
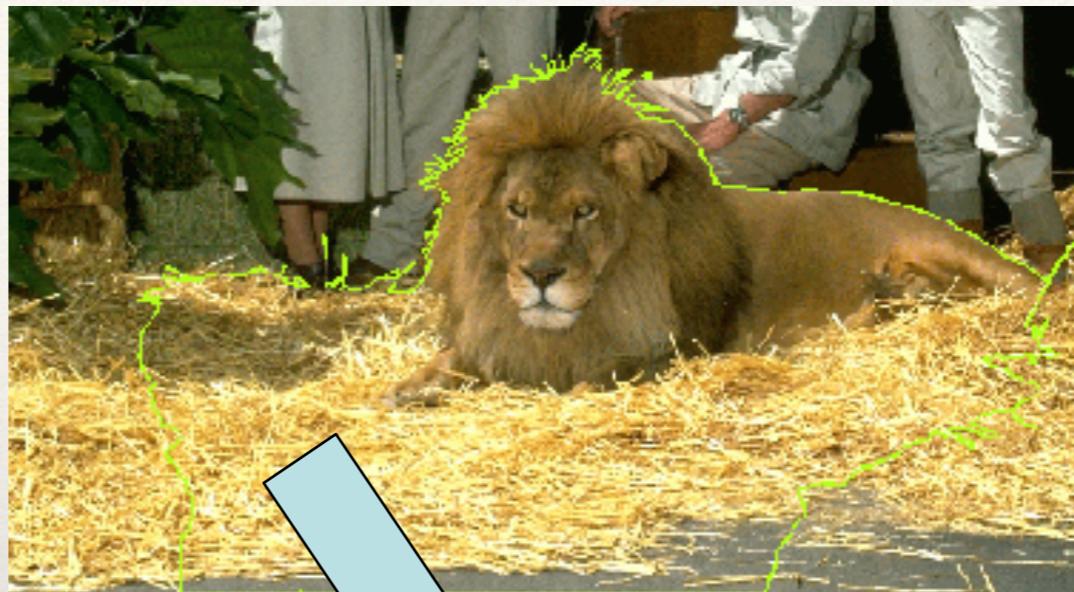
Figure 2: Image demonstrating how the live-wire segment adapts and snaps to an object boundary as the free point moves (via cursor movement). The path of the free point is shown in white. Live-wire segments from previous free point positions (t_0 , t_1 , and t_2) are shown in green.

Intelligent scissors



- <http://rivit.cs.byu.edu/Eric/Eric.html>

Intelligent scissors



Deformable contours: pros and cons

Pros:

- ❖ Useful to track and fit non-rigid shapes
- ❖ Contour remains connected
- ❖ Possible to fill in “subjective” contours
- ❖ Flexibility in how energy function is defined, weighted.

Cons:

- ❖ Must have decent initialization near true boundary, may get stuck in local minimum
- ❖ Parameters of energy function must be set well based on prior information

Summary

- ❖ Deformable shapes and active contours are useful for
 - ❖ Segmentation: fit or “snap” to boundary in image
 - ❖ Tracking: previous frame’s estimate serves to initialize the next
- ❖ Fitting active contours:
 - ❖ Define terms to encourage certain shapes, smoothness, low curvature, push/pulls, ...
 - ❖ Use weights to control relative influence of each component cost
 - ❖ Can optimize 2d snakes with Viterbi algorithm.
- ❖ Image structure (esp. gradients) can act as attraction force for *interactive* segmentation methods.